

LFD5461

Kubernetes Administration and Development, Capital One Custom Course

Version 1.20



Version 1.20

© Copyright the Linux Foundation 2021. All rights reserved.

© Copyright the Linux Foundation 2021. All rights reserved.

The training materials provided or developed by The Linux Foundation in connection with the training services are protected by copyright and other intellectual property rights.

Open source code incorporated herein may have other copyright holders and is used pursuant to the applicable open source license.

The training materials are provided for individual use by participants in the form in which they are provided. They may not be copied, modified, distributed to non-participants or used to provide training to others without the prior written consent of The Linux Foundation.

No part of this publication may be reproduced, photocopied, stored on a retrieval system, or transmitted without express prior written consent.

Published by:

the **Linux Foundation**

<https://www.linuxfoundation.org>

No representations or warranties are made with respect to the contents or use of this material, and any express or implied warranties of merchantability or fitness for any particular purpose or specifically disclaimed.

Although third-party application software packages may be referenced herein, this is for demonstration purposes only and shall not constitute an endorsement of any of these software applications.

Linux is a registered trademark of Linus Torvalds. Other trademarks within this course material are the property of their respective owners.

If there are any questions about proper and fair use of the material herein, please contact info@linuxfoundation.org.

Contents

1	Introduction	1
1.1	Objectives	1
1.2	Who You Are	1
1.3	The Linux Foundation	2
1.4	Linux Foundation Training	3
1.5	Certification Programs and Digital Badging	3
1.6	Preparing Your System	4
1.7	Course Registration	5
1.8	Labs	5
2	Kubernetes Architecture	7
2.1	What Is Kubernetes?	7
2.2	Components of Kubernetes	8
2.3	Challenges	9
2.4	The Borg Heritage	10
2.5	Kubernetes Architecture	10
2.6	Terminology	11
2.7	Master Node	12
2.8	Minion (Worker) Nodes	13
2.9	Pods	14
2.10	Services	14
2.11	Operators	15
2.12	Single IP per Pod	15
2.13	Networking Setup	16
2.14	CNI Network Configuration File	19
2.15	Pod-to-Pod Communication	20
2.16	Cloud Native Computing Foundation	20
2.17	Resource Recommendations	21
2.18	Labs	21
3	Build	39
3.1	Container Options	39
3.2	Containerizing an Application	41
3.3	Creating the Dockerfile	41
3.4	Hosting a Local Repository	41
3.5	Creating a Deployment	42
3.6	Running Commands in a Container	42
3.7	Multi-Container Pod	42
3.8	readinessProbe	43
3.9	livenessProbe	43
3.10	startupProbe	43
3.11	Testing	43
3.12	Labs	44
4	APIs and Access	63
4.1	API Access	64

4.2	Annotations	68
4.3	Working with A Simple Pod	69
4.4	kubectl and API	70
4.5	Swagger and OpenAPI	77
4.6	Labs	79
5	API Objects	85
5.1	API Objects	86
5.2	The v1 Group	87
5.3	API Resources	89
5.4	RBAC APIs	94
5.5	Labs	95
6	Design	105
6.1	Traditional Applications: Considerations	105
6.2	Decoupled Resources	105
6.3	Transience	106
6.4	Flexible Framework	106
6.5	Managing Resource Usage	106
6.6	Using Label Selectors	107
6.7	Multi-Container Pods	108
6.8	Sidecar Container	109
6.9	Adapter Container	109
6.10	Ambassador	109
6.11	Points to Ponder	109
6.12	Jobs	110
6.13	Labs	110
7	Deployment Configuration	123
7.1	Volumes Overview	123
7.2	Introducing Volumes	124
7.3	Volume Spec	125
7.4	Volume Types	126
7.5	Shared Volume Example	126
7.6	Persistent Volumes and Claims	127
7.7	Persistent Volume	128
7.8	Persistent Volume Claim	128
7.9	Dynamic Provisioning	129
7.10	Secrets	130
7.11	Using Secrets via Environment Variables	130
7.12	Mounting Secrets as Volumes	131
7.13	Portable Data with ConfigMaps	131
7.14	Using ConfigMaps	132
7.15	Deployment Configuration Status	132
7.16	Scaling and Rolling Updates	133
7.17	Deployment Rollbacks	134
7.18	Labs	134
8	Custom Resource Definition	155
8.1	Overview	156
8.2	Custom Resource Definitions	157
8.3	Aggregated APIs	161
8.4	Labs	162
9	Scheduling	165
9.1	Overview	166
9.2	Scheduler Settings	167
9.3	Policies	169

9.4	Affinity Rules	172
9.5	Taints and Tolerations	177
9.6	Labs	180
10	Security	187
10.1	Security Overview	187
10.2	Accessing the API	187
10.3	Authentication	188
10.4	Authorization	189
10.5	ABAC	189
10.6	RBAC	190
10.7	RBAC Process Overview	190
10.8	Admission Controller	190
10.9	Security Contexts	191
10.10	Pod Security Policies	191
10.11	Network Security Policies	192
10.12	Network Security Policy Example	192
10.13	Default Policy Example	193
10.14	Labs	194
11	Exposing Applications	213
11.1	Service Types	213
11.2	Services Diagram	214
11.3	Service Update Pattern	215
11.4	Accessing an Application with a Service	215
11.5	Service without a Selector	215
11.6	ClusterIP	216
11.7	NodePort	216
11.8	LoadBalancer	216
11.9	ExternalName	216
11.10	Ingress Resource	217
11.11	Ingress Controller	217
11.12	Service Mesh	217
11.13	Labs	218
12	Troubleshooting	231
12.1	Troubleshooting Overview	231
12.2	Basic Troubleshooting Steps	231
12.3	Ongoing (Constant) Change	232
12.4	Basic Troubleshooting Flow: Pods	232
12.5	Basic Troubleshooting Flow: Node and Security	234
12.6	Basic Troubleshooting Flow: Agents	234
12.7	Monitoring	234
12.8	Logging Tools	235
12.9	Monitoring Applications	235
12.10	System and Agent Logs	236
12.11	Conformance Testing	237
12.12	More Resource	237
12.13	Labs	238
13	Logging and Troubleshooting - II	247
13.1	Overview	248
13.2	Troubleshooting Flow	249
13.3	Basic Start Sequence	251
13.4	Monitoring	252
13.5	Plugins	253
13.6	Logging	256
13.7	Troubleshooting Resources	257

13.8	Labs	258
14	Helm	267
14.1	Overview	268
14.2	Helm	269
14.3	Using Helm	273
14.4	Labs	276
15	Closing and Evaluation Survey	281
15.1	Evaluation Survey	281

List of Figures

1.1	Linux Foundation Digital Badges	4
2.1	Kubernetes Logo	7
2.2	Legacy vs Cloud Architecture	9
2.3	Kubernetes Lineage	10
2.4	Kubernetes Architecture	11
2.5	Single IP, multiple volumes	16
2.6	Service Relationships	17
2.7	Pod and Services Example	18
2.8	Another Cluster Example	19
2.9	Cloud Native Computing Foundation (CNCF)	20
2.10	Core Concepts Domain	36
2.11	Exam Handbook Guidelines and Tips	36
3.1	Observability Domain	61
4.1	Swagger Screenshot	77
6.1	Multiple Domain	122
7.1	Multiple volumes in a pod	125
7.2	Deployment Related Domain Topics	153
10.1	Accessing the API	188
10.2	Domain	210
11.1	Traffic from Cluster IP to Pod	214
11.2	Ingress Controller	217
11.3	Istio Service Mesh	218
11.4	Accessing the API	229
11.5	Service Domain Topic	229
12.1	Grafana Dashboard Example	235
12.2	Jaeger Tracing Example	236
12.3	Troubleshooting Domain Topic	245
13.1	External Access via Browser	263
13.2	External Access via Browser	264
13.3	External Access via Browser	265
14.1	Basic Helm and Tiller Flow	269
15.1	Course Survey	281

Chapter 1

Introduction



We lay out our objectives and describe the **Linux Foundation** and our target audience.

1.1	Objectives	1
1.2	Who You Are	1
1.3	The Linux Foundation	2
1.4	Linux Foundation Training	3
1.5	Certification Programs and Digital Badging	3
1.6	Preparing Your System	4
1.7	Course Registration	5
1.8	Labs	5

1.1 Objectives

- By the end of this course, you should be able to:
 - Containerize and deploy a new **Python** script.
 - Configure the deployment with ConfigMaps, Secrets, and SecurityContexts.
 - Understand multi-container Pod design.
 - Configure probes for pod health.
 - Update and roll back an application.
 - Implement services and NetworkPolicies.
 - Use PersistentVolumeClaims for state persistence.

1.2 Who You Are

- This course is for developers looking to learn how to deploy, configure, and test their containerized applications on a multi-node **Kubernetes** cluster.
- For a successful learning experience, basic **Linux** command line and file editing skills are required.

- Familiarity with using a programming language (such as **Python**, **Node.js**, **Go**) and Cloud Native application concepts and architectures is helpful.
- Our free **LFS158x: Introduction to Kubernetes MOOC** on [edX.org](https://www.edx.org/course/introduction-to-kubernetes) is a useful preparation for this course, and can be found at: <https://www.edx.org/course/introduction-to-kubernetes>.
- You might be doing this just for fun, but more likely, this task is part of your job. The purpose here is to ease your path and perhaps shorten the amount of time it takes to reach the level of competence required.
- How much you get out of this and how bug-free, efficient, and optimized your code will be, depends on how good a programmer you were before you started with the present material. There is no intent or time here to teach you the elements of good programming or the **Linux** operating system design in great detail.

1.3 The Linux Foundation

- The **Linux Foundation** is the non-profit consortium dedicated to fostering the growth of **Linux** and many other **Open Source Software (OSS)** projects and communities.
- The **Linux Foundation** supports the creation of sustainable **OSS** ecosystems by providing:
 - Financial and intellectual resources and services
 - Training
 - Events
- Founded in 2000, it sponsors the work of **Linux** creator **Linus Torvalds** and is supported by leading technology companies and developers from around the world.
- Today it has grown well beyond that mission to support a host of other **OSS** projects and communities, with the list of hosted foundations and organizations growing daily.
- The **Linux Foundation** provides a neutral, trusted hub for developers to code, manage, and scale open technology projects. Founded in 2000, The **Linux Foundation** is supported by more than 1,000 members and is the world's leading home for collaboration on open source software, open standards, open data and open hardware. The **Linux Foundation's** methodology focuses on leveraging best practices and addressing the needs of contributors, users and solution providers to create sustainable models for open collaboration.
- The **Linux Foundation** hosts **Linux**, the world's largest and most pervasive open source software project in history. It is also home to **Linux** creator Linus Torvalds and lead maintainer Greg Kroah-Hartman. The success of **Linux** has catalyzed growth in the open source community, demonstrating the commercial efficacy of open source and inspiring countless new projects across all industries and levels of the technology stack.
- As a result, the **Linux Foundation** today hosts far more than **Linux**; it is the umbrella for many critical open source projects that power corporations today, spanning virtually all industry sectors. Some of the technologies we focus on include big data and analytics, networking, embedded systems and IoT, web tools, cloud computing, edge computing, automotive, security, blockchain, and many more.

Over 85,000 open source technologists and leaders worldwide gather at **Linux Foundation** events annually to share ideas, learn and collaborate. **Linux Foundation** events are the meeting place of choice for open source maintainers, developers, architects, infrastructure managers, and sysadmins and technologists leading open source program offices, and other critical leadership functions.

These events are the best place to gain visibility within the open source community quickly and advance open source development work by forming connections with the people evaluating and creating the next generation of technology. They provide a forum to share and gain knowledge, help organizations identify software trends early to inform future technology investments, connect employers with talent, and showcase technologies and services to influential open source professionals, media, and analysts around the globe.

The **Linux Foundation** hosts an ever increasing number of events each year. Some are held in multiple locations yearly, such as North America, Europe and Asia. Here is a partial list:

- Open Source Summit
 - Embedded Linux Conference
 - Open Networking & Edge Summit
 - KubeCon + CloudNativeCon
 - Automotive Linux Summit
 - KVM Forum
 - Linux Storage Filesystem and Memory Management Summit
 - Linux Security Summit
 - Linux Kernel Maintainer Summit
 - The Linux Foundation Member Summit
 - Open Compliance Summit
 - And many more.
- The **Linux Foundation** also provides services to the **Linux** development community, including an open source developer travel fund, legal support and other administrative assistance. Through its work groups, members and developers can collaborate on key technical areas which can include everything from Cloud Computing to Printing in **Linux**. Finally, users can access **Linux** training through its technical training program. (<https://training.linuxfoundation.org>)
 - The **Linux Foundation** website (<https://www.linuxfoundation.org>) contains fuller descriptions of the ongoing work it is involved in.

1.4 Linux Foundation Training

- The **Linux Foundation** offers several types of training:
 - Physical Classroom (often On-Site)
 - Online Virtual Classroom
 - Individual Self-Paced E-learning over the Internet
 - Events-Based
- To get more information about specific courses offered by the **Linux Foundation**:
 - **Linux Programming & Development Training:**
<https://training.linuxfoundation.org/linux-courses/development-training>
 - **Enterprise IT & Linux System Administration Courses:**
<https://training.linuxfoundation.org/linux-courses/system-administration-training>
 - **Open Source Compliance Courses:**
<https://training.linuxfoundation.org/linux-courses/open-source-compliance-courses>

1.5 Certification Programs and Digital Badging

Certification Programs

- The **Linux Foundation** offers a two-level **Certification Program**:



LFCS (Linux Foundation Certified Sysadmin)



LFCE (Linux Foundation Certified Engineer)

- Full details about this program can be found at <https://training.linuxfoundation.org/certification>. This information includes a thorough description of the **Domains** and **Competencies** covered by each exam.
- Besides the **LFCS** and **LFCE** exams, the **Linux Foundation** is currently associated with other open source certification programs involving its collaborative projects, for which it has created (or is currently creating) the exams. These include:

- Certified Kubernetes Administrator (CKA)
- Certified Kubernetes Application Developer (CKAD)
- Certified Kubernetes Application Security Specialist (CKAS)
- Linux Foundation Certified IT Associate (LFCA)
- FinOps Certified Practitioner (FOCP)
- Certified Hyperledger Fabric Administrator (CHFA)
- Certified Hyperledger Sawtooth Administrator (CHSA)
- OpenJS Node.js Services Developer (JSNSD)
- OpenJS Node.js Application Developer (JSNAD)
- Cloud Foundry Certified Developer (CFCDD)
- Certified ONAP Professional (COP)

- For additional information, including technical requirements and other logistics, see <https://training.linuxfoundation.org>.

Linux Foundation Digital Badges



Figure 1.1: Linux Foundation Digital Badges

- Digital Badges communicate abilities and credentials
- Can be used in email signatures, digital resumes, social media sites (**LinkedIn**, **Facebook**, **Twitter**, etc)
- Contain verified metadata describing qualifications and process that earned them.
- Available to students who successfully complete **Linux Foundation** courses and certifications
- Details at <https://training.linuxfoundation.org/badges/>

How it Works

1. You will receive an email notifying you to claim your badge at our partner **Credly's Acclaim** platform website.
2. Click the link in that email.
3. Create an account on the **Acclaim** platform site and confirm your email.
4. Claim your badge.
5. Start sharing.

1.6 Preparing Your System

- You will need a computer installed with a current **Linux** distribution, with the important developer tools (for compiling, etc.) properly deployed.
- Once you have installed your favorite OS, go to <https://training.linuxfoundation.org/cm/prep/> and download a copy of the `ready-for.sh` script, either by downloading through your browser, or using **wget**:

```
$ wget https://training.linuxfoundation.org/cm/prep/ready-for.sh
$ chmod 755 ready-for.sh
$ ./ready-for.sh LFD5461
```

- If there are missing dependencies, you can do:

```
$ ./ready-for.sh --install LFD5461
```

1.7 Course Registration

- Please go to: <https://training.linuxfoundation.org/surveys/registration/>, and using the unique **key** your instructor provides you, register for this class.
- This will be used to generate your **Certificate of Completion** at the end of the class. Please be sure to check the spelling of your name and use correct capitalization as this is how your name will appear on the certificate.
- At the end of the class you will be asked to submit an evaluation survey, at: <https://training.linuxfoundation.org/surveys/evaluation/>, using the same key and email address you used for registration.
- Thank you!

1.8 Labs

Exercise 1.1: Obtaining Some System Information

- Open up a command line terminal window; there are a number of ways to do this as we shall discuss later. If you do not see it under Applications->Accessories or Applications->System Tools, do Alt-F2 and type in **gnome-terminal**.
- The **uname** utility will provide some basics. If you type

```
$ uname --help
```

you will get:

```
$ uname --help
```

```
1 Usage: uname [OPTION]...
2 Print certain system information.  With no OPTION, same as -s.
3
4 -a, --all                print all information, in the following order,
5                          except omit -p and -i if unknown:
6 -s, --kernel-name        print the kernel name
7 -n, --nodename           print the network node hostname
8 -r, --kernel-release     print the kernel release
9 -v, --kernel-version     print the kernel version
10 -m, --machine            print the machine hardware name
11 -p, --processor          print the processor type or "unknown"
```

```
12 -i, --hardware-platform  print the hardware platform or "unknown"
13 -o, --operating-system   print the operating system
14 --help                  display this help and exit
15 --version               output version information and exit
```

Try some of the options.

- Other basic commands you might try are:

```
$ df
$ free
$ less /proc/meminfo
$ more /proc/cpuinfo
$ cat /proc/version
```

Chapter 2

Kubernetes Architecture



In this chapter we are going to cover a little bit of the history of Kubernetes. This helps to understand the why things are where they are. Then we'll talk about master nodes and the main agents that work to make Kubernetes what it is. We'll also talk about minion or worker nodes. These are nodes that use API calls back to the master node to understand what the configuration should be, and report back status. We'll also talk about CNI network plugin configuration, which continues to be the future of Kubernetes networking. And we'll cover some common terms so as you learn more about Kubernetes, you'll understand major components and how they work with the other components in our decoupled and transient environment.

2.1	What Is Kubernetes?	7
2.2	Components of Kubernetes	8
2.3	Challenges	9
2.4	The Borg Heritage	10
2.5	Kubernetes Architecture	10
2.6	Terminology	11
2.7	Master Node	12
2.8	Minion (Worker) Nodes	13
2.9	Pods	14
2.10	Services	14
2.11	Operators	15
2.12	Single IP per Pod	15
2.13	Networking Setup	16
2.14	CNI Network Configuration File	19
2.15	Pod-to-Pod Communication	20
2.16	Cloud Native Computing Foundation	20
2.17	Resource Recommendations	21
2.18	Labs	21

2.1 What Is Kubernetes?

- Running a container on a laptop is relatively simple. Deploying and connecting containers across multiple hosts, scaling them, deploying applications without downtime, and service discovery among several aspects can be complex.
- Kubernetes addresses those challenges from the start with a set of primitives and a powerful open and extensible API. The ability to add new objects and operators allows easy customization for various production needs. The ability to add multiple schedulers and multiple API servers adds to this customization.

- According to the <https://kubernetes.io> website, **Kubernetes** is:
“an open-source system for automating deployment, scaling, and management of containerized applications”
- A key aspect of Kubernetes is that it builds on 15 years of experience at Google in a project called **Borg**.
- Google’s infrastructure started reaching high scale before virtual machines became pervasive in the datacenter, and containers provided a fine-grained solution for packing clusters efficiently. Efficiency in using clusters and managing distributed applications has been at the core of Google challenges.



Figure 2.1: Kubernetes Logo

In Greek, κυβερνήτης means the Helmsman, or pilot of the ship. Keeping with the maritime theme of Docker containers, Kubernetes is the pilot of a ship of containers. Due to the difficulty in pronouncing the name, many will use a nickname, **K8s**, as Kubernetes has eight letters between K and S. The nickname is said like **Kate’s**.

- Kubernetes can be an integral part of **Continuous Integration/Continuous Delivery (CI/CD)**, as it offers many of the necessary components.
 - **Continuous Integration** - A consistent way to build and test software. Deploying new packages with code written each day, or every hour, instead of quarterly. Tools like **Helm** and **Jenkins** are often part of this with Kubernetes.
 - **Continuous Delivery** - An automated way to test and deploy software into various environments. Kubernetes handles the life cycle of containers and connection of infrastructure resources to make rolling updates and rollbacks easy, among other deployment schemes.
 - There are several options and possible configurations when building a CI/CD pipeline. Tools such as Jenkins, Spinnaker, GitHub, GitLab, and Helm among others may be part of your particular pipeline. To learn more about CI/CD take a look at **DevOps and SRE Fundamentals: Implementing Continuous Delivery (LFS261)** at <https://training.linuxfoundation.org/training/devops-and-sre-fundamentals-implementing-continuous-delivery-lfs261/>

2.2 Components of Kubernetes

- Deploying containers and using Kubernetes may require a change in the development and the system administration approach to deploying applications. In a traditional environment, an application (such as a web server) would be a monolithic application placed on a dedicated server. As the web traffic increases, the application would be tuned, and perhaps moved to bigger and bigger hardware. After a couple of years, a lot of customization may have been done in order to meet the current web traffic needs.
- Instead of using a large server, Kubernetes approaches the same issue by deploying a large number of small web servers, or **microservices**. The server and client sides of the application expect that there are one or more microservices, called **replicas**, available to respond to a request. It is also important that clients expect server processes to be terminated and eventually be replaced, leading to a **transient** server deployment. Instead of a large tightly-coupled Apache web server with several **httpd** daemons responding to page requests, there would be many **nginx** servers, each responding without knowledge of each other.
- The transient nature of smaller services also allows for **decoupling**. Each aspect of the traditional application is replaced with a dedicated, but transient, microservice or agent. To join these agents, or their replacements together, we use

services and API calls. A service ties traffic from one agent to another (for example, a frontend web server to a backend database) and handles new IP or other information, should either one die and be replaced.

- Developers new to Kubernetes sometimes assume it is another virtual-machine manager, similar to what they have been using for decades and continue to develop applications in the same way as prior to using Kubernetes. **This is a mistake.** The decoupled, transient, microservice architecture is not the same. Most legacy applications will need to be rewritten to optimally run in a cloud. In this diagram we see the legacy deployment strategy on the left with a monolithic applications deployed to nodes. On the right we see an example of the same functionality, on the same hardware, using multiple microservices.

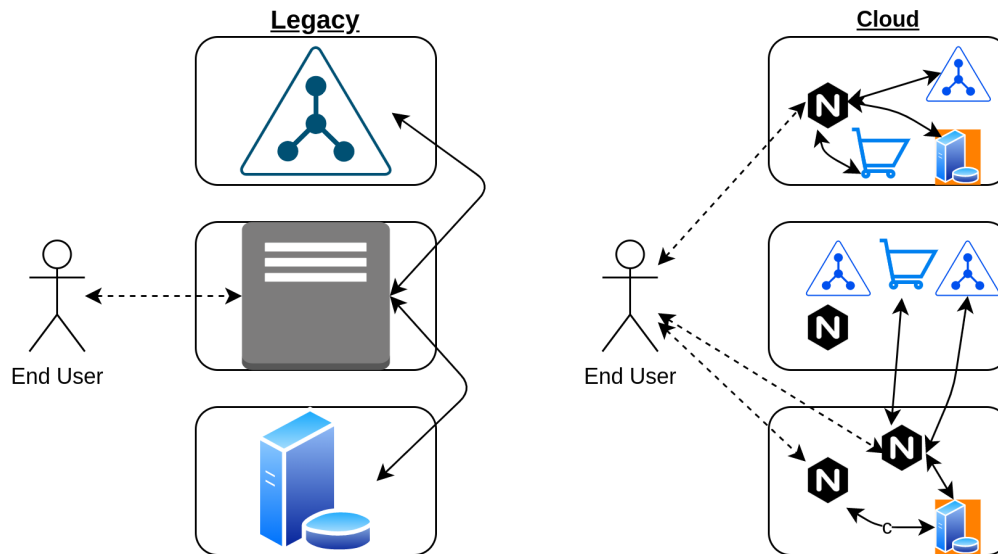


Figure 2.2: Legacy vs Cloud Architecture

- Communication to, as well as internally, between components is API call-driven, which allows for flexibility. Configuration information is stored in a JSON format, but is most often written in YAML. Kubernetes agents convert the YAML to JSON prior to persistence to the database.

Kubernetes is written in Go Language, a portable language which is like a hybridization between C++, Python, and Java. Some claim it incorporates the best (while some claim the worst) parts of each.

2.3 Challenges

- Containers have seen a huge rejuvenation. They provide a great way to package, ship, and run applications - that is the **Docker** motto, now provided by many tools.
- The developer experience has been boosted tremendously thanks to containers. Containers, and **Docker** specifically, have empowered developers with ease of building container images, simplicity of sharing images via registries, and providing a powerful user experience to manage containers. Now several tools such as **Buildah**, **Podman**, **cri-o**, **containerd**, **frakti**, and others allow for easy container creation and management.
- However, managing containers at scale and architecting a distributed application based on microservices' principles is still challenging.
- You first need a continuous integration pipeline to build your container images, test them, and verify them. Then, you need a cluster of machines acting as your base infrastructure on which to run your containers. You also need a system to launch your containers, and watch over them when things fail and self-heal. You must be able to perform rolling updates and rollbacks, and eventually tear down the resource when no longer needed.

- All of these actions require flexible, scalable, and easy-to-manage network and storage. As containers are launched on any worker node, the network must join the resource to other containers, while still keeping the traffic secure from others. We also need a storage structure which provides and keeps or recycles storage in a seamless manner.



Please Note

Would users notice if you ran **Chaos Monkey**, which terminates any container randomly? If so you may have more work making containers and applications more decoupled and transient.

2.4 The Borg Heritage

- What primarily distinguishes Kubernetes from other systems is its heritage. Kubernetes is inspired by Borg - the internal system used by Google to manage its applications (e.g. Gmail, Apps, GCE).
- With Google pouring the valuable lessons they learned from writing and operating Borg for over 15 years into Kubernetes, this makes Kubernetes a safe choice when having to decide on what system to use to manage containers. While a powerful tool, part of the current growth in Kubernetes is making it easier to work with and handle workloads not found in a Google data center.

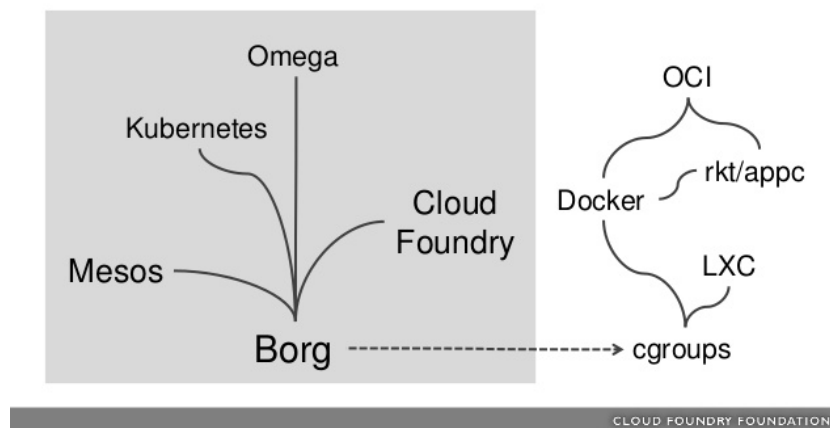


Figure 2.3: Kubernetes Lineage

(The above figure was obtained from

<https://www.slideshare.net/chipchilders/cloud-foundry-the-platform-for-forging-cloud-native-applications>.)

- To learn more about the ideas behind Kubernetes, you can read the **Large-scale cluster management at Google with Borg paper** at <https://ai.google/research/pubs/pub43438>.
- Borg has inspired current data center systems, as well as the underlying technologies used in container runtime today. Google contributed cgroups to the Linux kernel in 2007; it limits the resources used by collection of processes. Both cgroups and Linux namespaces are at the heart of containers today, including Docker.
- Mesos was inspired by discussions with Google when Borg was still a secret. Indeed, Mesos builds a multi-level scheduler, which aims to better use a data center cluster.
- The Cloud Foundry Foundation embraces the 12 factor application principles at <https://12factor.net/>.
- These principles provide great guidance to build web applications that can scale easily, can be deployed in the cloud, and whose build is automated. Borg and Kubernetes address these principles as well.

2.5 Kubernetes Architecture

- To quickly demystify Kubernetes, let's have a look at a Kubernetes architecture graphic, which shows a high-level architecture diagram of the system components.

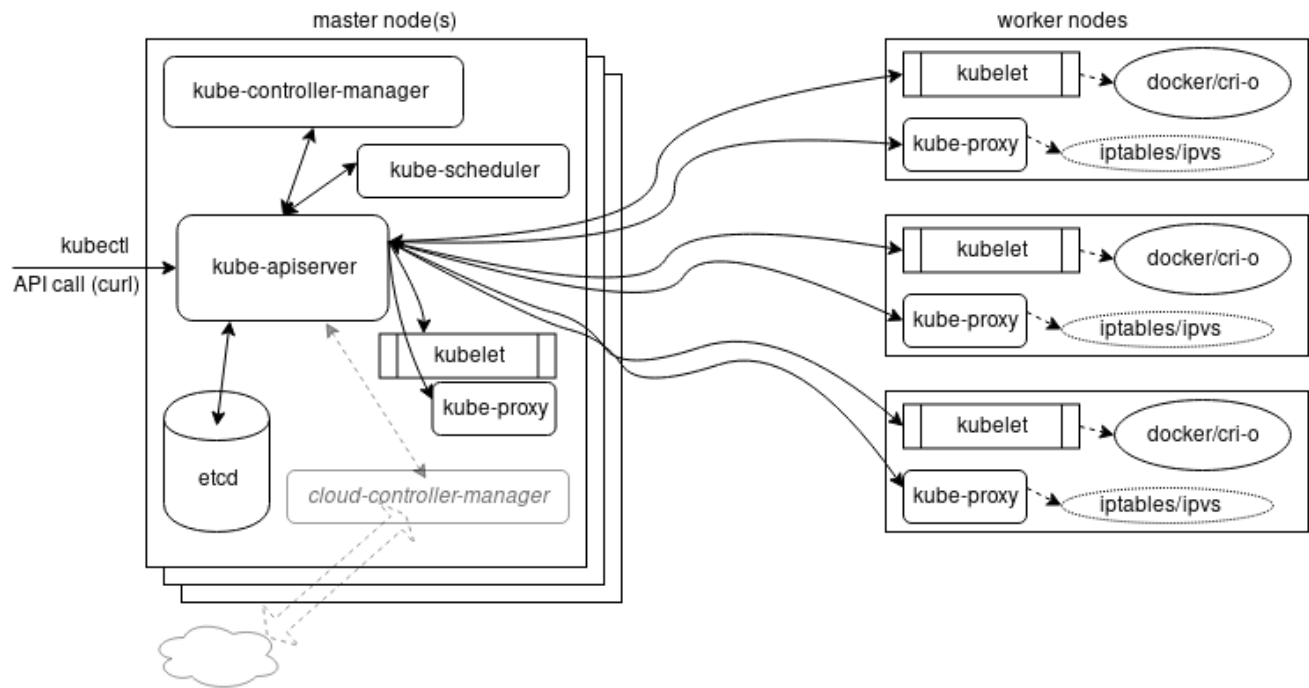


Figure 2.4: Kubernetes Architecture

- In its simplest form, Kubernetes is made of one or more central managers (aka master) and worker nodes (we will see in a follow-on chapter how you can actually run everything on a single node for testing purposes). The manager runs an API server, a scheduler, various operators and a datastore to keep the state of the cluster, container settings, and the networking configuration.
- Kubernetes exposes an API via the API server: you can communicate with the API using a local client called **kubectl** or you can write your own client. The **kube-scheduler** sees the API requests for running a new container and finds a suitable node to run that container. Each node in the cluster runs two containers: **kubelet** and **kube-proxy**. The **kubelet** container receives spec information for container configuration, downloads and manages any necessary resources and works with the container engine on the local node to ensure the container runs or is restarted upon failure. The **kube-proxy** container creates and manages local firewall rules and networking configuration to expose containers on the network.

2.6 Terminology

- We have learned that Kubernetes is an orchestration system to deploy and manage containers. Containers are not managed individually; instead, they are part of a larger object called a **Pod**. A Pod consists of one or more containers which share an IP address, access to storage and namespace. Typically, one container in a Pod runs an application, while other containers support the primary application.
- Orchestration is managed through a series of watch-loops, also known as **operators** or controllers. Each operator interrogates the **kube-apiserver** for a particular object state, modifying the object until the declared state matches the current state. The default, newest, and feature-filled operator for containers is a **Deployment**. A Deployment deploys and manages a different operator called a **ReplicaSet**. A replicaSet is an operator which deploys multiple pods, each with the same spec information. These are called replicas. The Kubernetes architecture is made up of many operators such as **Jobs** and **CronJobs** to handle single or recurring tasks, or custom resource definitions and purpose built operators.

- There are several other API objects which can be used to deploy pods, other than ensuring a certain number of replicas is running somewhere. A **DaemonSet** will ensure that a single pod is deployed on every **node**. These are often used for logging, metrics, and security Pods. A **StatefulSet** can be used to deploy Pods in a particular order, such that following pods are only deployed if previous pods report a ready status. This is useful for legacy applications which are not cloud-friendly.
- To easily manage thousands of Pods across hundreds of nodes can be a difficult. To make management easier, we can use **labels**, arbitrary strings which become part of the object metadata. These are selectors which can then be used when checking or changing the state of objects without having to know individual names or UIDs. Nodes can have **taints**, an arbitrary string in the node metadata, to inform the scheduler on Pod assignments used along with **toleration** in Pod metadata, indicating it should be scheduled on a node with the particular taint.
- There is also space in metadata for **annotations** which remain with the object but cannot be used as a selector, but could be leveraged by other objects or Pods.
- While using lots of smaller, commodity hardware could allow every user to have their very own cluster, often multiple users and teams share access to one or more clusters. This is referred to as **multi-tenancy**. Some form of isolation is necessary in a multi-tenant cluster, using a combination of the following, which we introduce here but will learn more about in the future:
 - **namespace** - A segregation of resources, upon which resource quotas and permissions can be applied. Kubernetes objects may be created in a namespace or cluster-scoped. Users can be limited by the object verbs allowed per namespace. Also the `LimitRange` admission controller constrains resource usage in that namespace. Two objects cannot have the same `Name`: value in the same namespace.
 - **context** - A combination of user, cluster name and namespace. A convenient way to switch between combinations of permissions and restrictions. For example you may have a development cluster and a production cluster, or may be part of both the operations and architecture namespaces. This information is referenced from `~/.kube/config`.
 - **Resource Limits** - A way to limit the amount of resources consumed by a pod, or to request a minimum amount of resources reserved, but not necessarily consumed, by a pod. Limits can also be set per-namespaces, which have priority over those in the `PodSpec`.
 - **Pod Security Policies** - A policy to limit the ability of pods to elevate permissions or modify the node upon which they are scheduled. This wide-ranging limitation may prevent a pod from operating properly. The use of PSPs may be replaced by `Open Policy Agent` in the future.
 - **Network Policies** - The ability to have an inside-the-cluster firewall. Ingress and Egress traffic can be limited according to namespaces and labels as well as typical network traffic characteristics.

2.7 Master Node

- The Kubernetes master runs various server and manager processes for the cluster. Among the components of the master node are the **kube-apiserver**, the **kube-scheduler**, and the **etcd** database. As the software has matured, new components have been created to handle dedicated needs, such as the **cloud-controller-manager**; it handles tasks once handled by the **kube-controller-manager** to interact with other tools, such as **Rancher** or **Digital Ocean** for third-party cluster management and reporting.
- There are several add-ons which have become essential to a typical production cluster, such as DNS services. Others are third-party solutions where Kubernetes has not yet developed a local component, such as cluster-level logging and resource monitoring.

kube-apiserver

- The **kube-apiserver** is central to the operation of the Kubernetes cluster. All calls, both internal and external traffic, are handled via this agent. All actions are accepted and validated by this agent, and it is the only agent which connects to the **etcd** database. As a result, it acts as a master process for the entire cluster, and acts as a front end of the cluster's shared state. Each API call goes through three steps: authentication, authorization, and several admission controllers.

kube-scheduler

- The **kube-scheduler** uses an algorithm to determine which node will host a Pod of containers. The scheduler will try to view available resources (such as available CPU) to, and then assign the Pod based on availability and success. The scheduler uses pod-count by default, but complex configuration is often done if cluster-wide metrics are collected.
- There are several ways you can affect the algorithm, or a custom scheduler could be used simultaneously instead. A Pod can also be assigned bind to a particular node in the pod spec, though the Pod may remain in a pending if the node or other declared resource is unavailable.
- One of the first configurations referenced during creation is if the Pod can be deployed within the current quota restrictions. If so, then the taints and tolerations, and labels of the Pods are used along with those of the nodes to determine the proper placement. Some is done as an admission controller in the kube-apiserver, the rest is done by the chosen scheduler. You can find more details about the scheduler can be found on GitHub at <https://github.com/kubernetes/kubernetes/blob/master/pkg/scheduler/scheduler.go>.

EtcD Database

- The state of the cluster, networking, and other persistent information is kept in an **etcd** database, or, more accurately, a **b+tree key-value** store. Rather than finding and changing an entry, values are always appended to the end. Previous copies of the data are then marked for future removal by a compaction process. It works with **curl** and other HTTP libraries, and provides reliable watch queries.
- Simultaneous requests to update a particular value all travel via the **kube-apiserver**, which then passes along the request to **etcd** in a series. The first request would update the database. The second request would no longer have the same version number as found in the object, in which case the **kube-apiserver** would reply with an error 409 to the requester. There is no logic past that response on the server side, meaning the client needs to expect this and act upon the denial to update.
- There is a **master** database along with possible **followers**. They communicate with each other on an ongoing basis to determine which will be **master**, and determine another in the event of failure. While very fast and potentially durable, there have been some hiccups with some features like whole cluster upgrades. The **kubeadm** cluster creation tool allows easy deployment of a multi-master cluster with stacked **etcd** or an external database cluster.

Other Agents

- The **kube-controller-manager** is a core control loop daemon which interacts with the **kube-apiserver** to determine the state of the cluster. If the state does not match, the manager will contact the necessary controller to match the desired state. There are several controllers in use, such as endpoints, namespace, and replication. The full list has expanded as Kubernetes has matured.
- Remaining in **beta** as of v1.16, the **cloud-controller-manager** interacts with agents outside of the cloud. It handles tasks once handled by **kube-controller-manager**. This allows faster changes without altering the core Kubernetes control process. Each **kubelet** must use the `--cloud-provider-external` settings passed to the binary.

2.8 Minion (Worker) Nodes

- All worker nodes run the **kubelet** and **kube-proxy**, as well as the container engine, such as **Docker** or **cri-o** <https://cri-o.io>. Other management daemons are deployed to watch these agents or provide services not yet included with Kubernetes.
- The **kubelet** interacts with the underlying Docker Engine also installed on all the nodes, and makes sure that the containers that need to run are actually running. The **kube-proxy** is in charge of managing the network connectivity to the containers. It does so through the use of **iptables** entries. It also has the **userspace** mode, in which it monitors Services and Endpoints using a random high-number port to proxy traffic. Use of **ipvs** can be enabled, with the expectation it becoming the default, replacing **iptables**.

- Kubernetes does not have cluster-wide logging yet. Instead, another CNCF project is used, called **Fluentd** (<https://www.fluentd.org/>). When implemented, it provides a unified logging layer for the cluster, which filters, buffers, and routes messages.
- Cluster-wide metrics is not quite fully mature, so **Prometheus** (<https://Prometheus.io>) is also often deployed to gather metrics from nodes and perhaps some applications.

Kubelet

- The **kubelet** agent is the heavy lifter for changes and configuration on worker nodes. It accepts the API calls for Pod specifications (a **PodSpec** is a JSON or YAML file that describes a pod). It will work to configure the local node until the specification has been met.
- Should a Pod require access to storage, Secrets or ConfigMaps, the **kubelet** will ensure access or creation. It also sends back status to the **kube-apiserver** for eventual persistence.

2.9 Pods

- The whole point of Kubernetes is to orchestrate the life cycle of a container. We do not interact with particular containers. Instead, the smallest unit we can work with is a **Pod**. Some would say a pod of whales or peas-in-a-pod. Due to shared resources, the design of a Pod typically follows a one-process-per-container architecture.
- Containers in a Pod are started in parallel by default. As a result, there is no way to determine which container becomes available first inside a pod. The use of `initContainers` can be used to ensure some containers are ready before others in a pod. To support a single process running in a container, you may need logging, a proxy, or special adapter. These tasks are often handled by other containers in the same pod.
- There is only one IP address per Pod, with most network plugins. (**HPE Labs** made a plugin which allows more than one IP per pod) As a result if there is more than one container, they must share the IP. To communicate with each other they can use IPC, the loopback interface, or a shared filesystem.
- While Pods are often deployed with one application container in each, a common reason to have multiple containers in a Pod is for logging. You may find the term **sidecar** for a container dedicated to performing a helper task, like handling logs and responding to requests, as the primary application container may have this ability.
- Pods, and other objects, can be created in several ways. They can be created by using a generator, which historically has changed with each release:

```
kubectl run newpod --image=nginx --generator=run-pod/v1
```

Or they can be created and deleted using a properly formatted JSON or YAML file:

```
kubectl create -f newpod.yaml
```

```
kubectl delete -f newpod.yaml
```

Other objects will be created by operators/watch-loops to ensure the spec and current status are the same.

2.10 Services

- With every object and agent decoupled we need a flexible and scalable operator which connects resources together and will reconnect, should something die and a replacement is spawned. Each Service is a microservice handling a particular bit of traffic, such as a single **NodePort** or a **LoadBalancer** to distribute inbound requests among many Pods.
- A **Service** also handles access policies for inbound requests, useful for resource control, as well as for security.
- A service, as well as **kubectl**, uses a **selector** in order to know which objects to connect. There are two selectors currently supported:

- **equality-based** - Filters by label keys and their values. Three operators can be used such as `=`, `==`, and `!=`. If multiple values or keys are used all must be included for a match.
- **set-based** - Filters according to a set of values. The operators are `in`, `notin`, and `exists`. For example the use of

```
status notin (dev, test, maint)
```

would select resources with the key of `status` which did not have a value of `dev`, `test`, nor `maint`.

2.11 Operators

- An important concept for orchestration is the use of operators. These are also known as watch-loops and controllers. These query the current state, compare it against the spec, and execute code based on how they differ. Various operators ship with Kubernetes, and you can create your own, as well. A simplified view of an operator is an agent, or **Informers**, and a downstream store. Using a `DeltaFIFO` queue, the source and downstream are compared. A loop process receives an `obj` or object, which is an array of deltas from the FIFO queue. As long as the delta is not of the type `Deleted`, the logic of the operator is used to create or modify some object until it matches the specification.
- The **Informers** which uses the API server as a source requests the state of an object via an API call. The data is cached to minimize API server transactions. A similar agent is the **SharedInformers**; objects are often used by multiple other objects. It creates a shared cache of the state for multiple requests.
- A **Workqueue** uses a key to hand out tasks to various workers. The standard **Go** work queues of rate limiting, delayed, and time queue are typically used.
- The `endpoints`, `namespace`, and `serviceaccounts` operators each manage the eponymous resources for Pods.

2.12 Single IP per Pod

- A pod represents a group of co-located containers with some associated data volumes. All containers in a pod share the same network **namespace**.
- The diagram below shows a pod with two containers, `MainApp` and `Logger`, and two data volumes, made available under two mount points. Containers `MainApp` and `Logger` share the network namespace of a third container, known as the **pause container**. The pause container is used to get an IP address, then all the containers in the pod will use its network namespace. You won't see this container from the Kubernetes perspective, but you would by running **sudo docker ps**. The volumes are shown for completeness and will be discussed later.




Figure 2.5: Single IP, multiple volumes

- To communicate with each other, Pods can use the loopback interface, write to files on a common filesystem, or via inter-process communication (IPC). As a result co-locating applications in the same pod may have issues. There is a network plugin which will allow more than one IP address, but so far has only been used within HPE labs.
- Support for dual-stack, IPv4 and IPv6, continues to increase with each release. For example, in a recent release **kube-proxy** iptables supports both stacks simultaneously.

2.13 Networking Setup

- Getting all the previous components running is a common task for system administrators who are accustomed to configuration management. But, to get a fully functional Kubernetes cluster, the network will need to be set up properly, as well.
- A detailed explanation about the Kubernetes networking model can be seen on the Cluster Networking page in the Kubernetes documentation at <https://kubernetes.io/docs/concepts/cluster-administration/networking>
- If you have experience deploying virtual machines (VMs) based on IaaS solutions, this will sound familiar. The only caveat is that, in Kubernetes, the lowest compute unit is not a **container**, but what we call a **pod**.
- A pod is a group of co-located containers that share the same IP address. From a networking perspective, a pod can be seen as a virtual machine of physical hosts. The network needs to assign IP addresses to pods, and needs to provide traffic routes between all pods on any nodes.

- The three main networking challenges to solve in a container orchestration system are:
 - Coupled container-to-container communications (solved by the pod concept)
 - Pod-to-pod communications
 - External-to-pod communications
- Kubernetes expects the network configuration to enable pod-to-pod communications to be available; it will not do it for you.
- Pods are assigned an IP address prior to application containers being started. The **service** object is used to connect pods within the network using **ClusterIP** addresses, from outside of the cluster using **NodePort** addresses, and using a load balancer if configured with a **LoadBalancer** service.
- A **ClusterIP** is used for traffic within the cluster. A **NodePort** first creates a **ClusterIP** then associates a port of the node to that new **ClusterIP**. If you create a **LoadBalancer** service it will first create a **ClusterIP**, then a **NodePort** and then make an asynchronous request for an external load balancer. If one is not configured to respond the **EXTERNAL-IP** will remain in pending state for the life of the service.

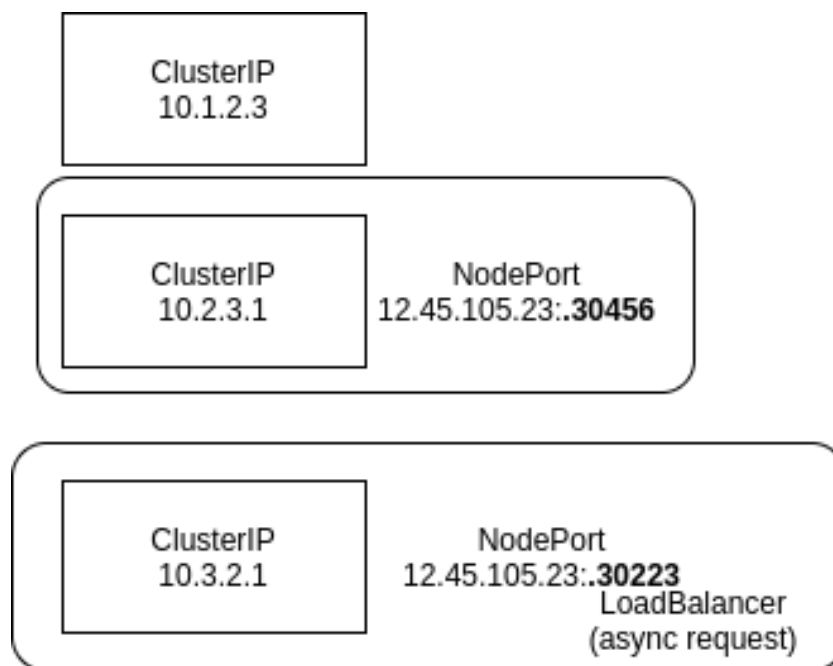


Figure 2.6: Service Relationships

- An **Ingress Controller** or a service mesh like **Istio** can also be used to connect traffic to a pod. This image shows a multi-container pod, two services with one for internal traffic only, and an ingress controller. The sidecar container, acting as a logger, is shown writing out storage, just to show a more complete pod. The pause container, which is only used to retrieve the namespace and IP address is also shown.

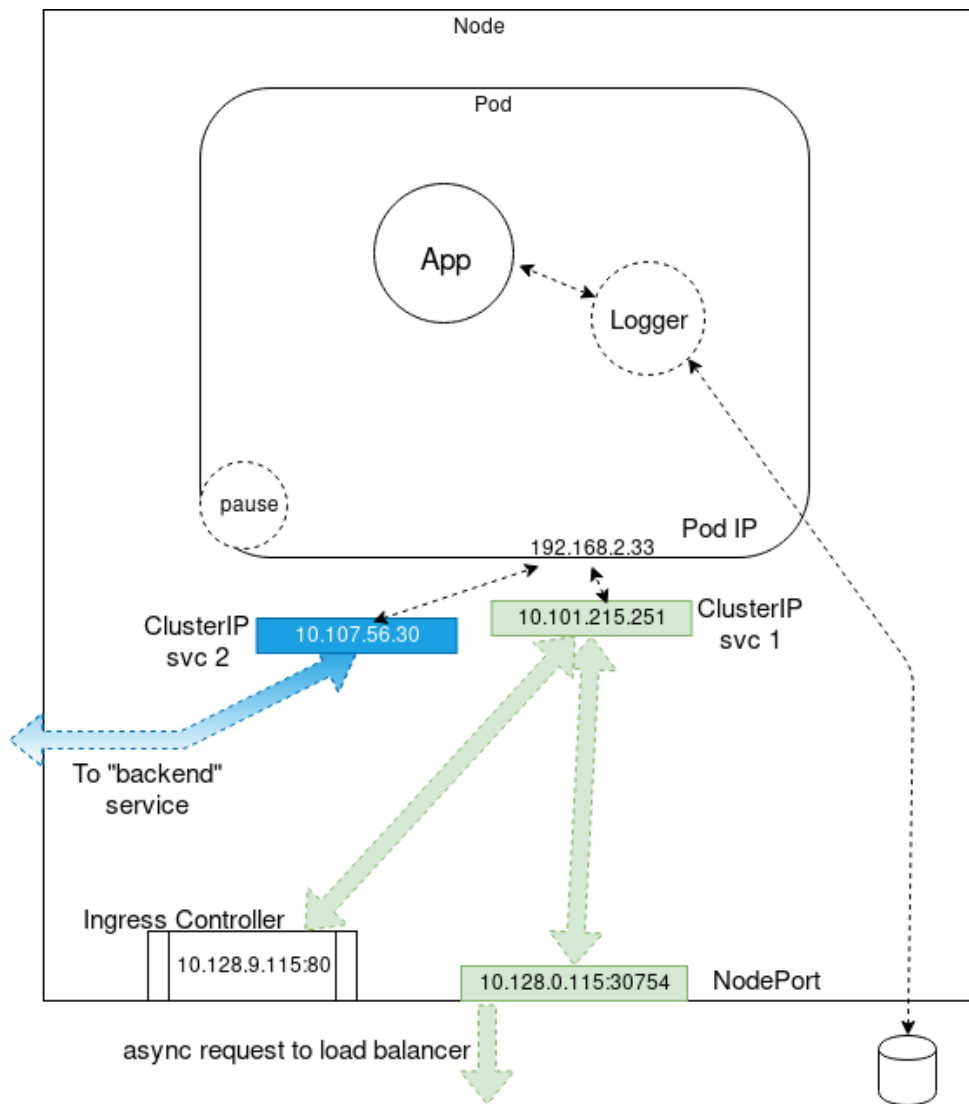


Figure 2.7: Pod and Services Example

- Another possible view of a cluster with multiple pods and services. This graphic shows the **Calico** pod running on each node and communicating with the BIRD protocol. There are also three **ClusterIP** services and on **LoadBalancer** service trying to show how the front end may communicate with other pods. The pods could have been on any worker, and are shown on one only as an example. Note this graphic is not an expansion of the previous graphic.

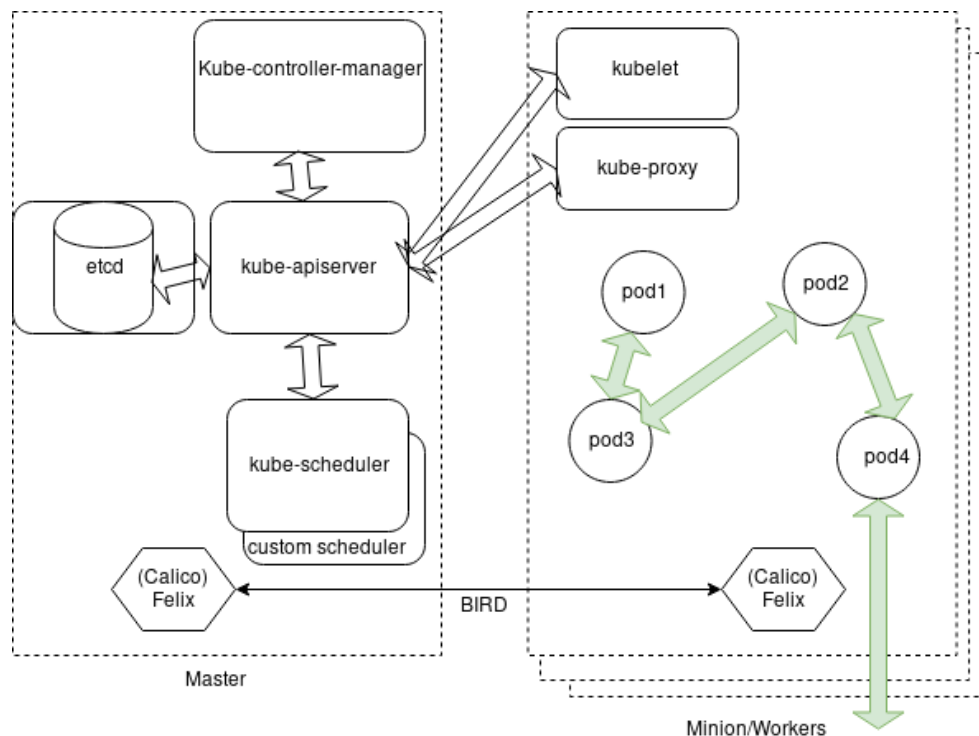


Figure 2.8: Another Cluster Example

- Tim Hockin, one of the lead Kubernetes developers, has created a very useful slide deck to understand the Kubernetes networking An Illustrated Guide to Kubernetes Networking at <https://speakerdeck.com/thockin/illustrated-guide-to-kubernetes-networking>

2.14 CNI Network Configuration File

- To provide container networking, Kubernetes is standardizing on the Container Network Interface (CNI) specification at <https://github.com/containernetworking/cni>. As of v1.6.0, **kubeadm** (the Kubernetes cluster bootstrapping tool) uses CNI as the default network interface mechanism.
- CNI is an emerging specification with associated libraries to write plugins that configure container networking and remove allocated resources when the container is deleted. Its aim is to provide a common interface between the various networking solutions and container runtimes. As the CNI specification is language-agnostic, there are many plugins from Amazon ECS, to SR-IOV, to Cloud Foundry, and more.
- With CNI, you can write a network configuration file:

```
{
  "cniVersion": "0.2.0",
  "name": "mynet",
  "type": "bridge",
  "bridge": "cni0",
  "isGateway": true,
  "ipMasq": true,
  "ipam": {
    "type": "host-local",
    "subnet": "10.22.0.0/16",
    "routes": [
      { "dst": "0.0.0.0/0" }
    ]
  }
}
```

- This configuration defines a standard Linux bridge named `cni0`, which will give out IP addresses in the subnet `10.22.0.0/16`. The **bridge** plugin will configure the network interfaces in the correct namespaces to define the container network properly. The main [README](#) of the CNI GitHub repository (<https://github.com/containernetworking/cni>) has more information.

2.15 Pod-to-Pod Communication

- While a CNI plugin can be used to configure the network of a pod and provide a single IP per pod, CNI does not help you with pod-to-pod communication across nodes.
- The early requirement from Kubernetes was the following:
 - All pods can communicate with each other across nodes.
 - All nodes can communicate with all pods.
 - No Network Address Translation (NAT).
- Basically, all IPs involved (nodes and pods) are routable without NAT. This can be achieved at the physical network infrastructure if you have access to it (e.g. GKE). Or, this can be achieved with a software defined overlay with solutions like:
 - **Weave**
<https://www.weave.works/oss/net/>
 - **Flannel**
<https://github.com/coreos/flannel#flannel>
 - **Calico**
<https://www.projectcalico.org/>
 - **Romana**
<https://romana.io/>
- Most network plugins now support the use of `Network Policies` which act as an internal firewall, limiting ingress and egress traffic.
- See the cluster networking documentation page (<https://kubernetes.io/docs/concepts/cluster-administration/networking/>) or the list of networking add-ons (<https://kubernetes.io/docs/concepts/cluster-administration/addons/>) for a more complete list.

2.16 Cloud Native Computing Foundation



Figure 2.9: Cloud Native Computing Foundation (CNCF)

- Kubernetes is an open source software with an Apache license. **Google** donated Kubernetes to a newly formed collaborative project within the **Linux Foundation** in July 2015, when Kubernetes reached the v1.0 release. This project is known as the **Cloud Native Computing Foundation (CNCF)**.
- CNCF is not just about Kubernetes, it serves as the governing body for open source software that solves specific issues faced by cloud native applications (i.e. applications that are written specifically for a cloud environment).
- CNCF has many corporate members that collaborate, such as Cisco, the Cloud Foundry Foundation, AT&T, Box, Goldman Sachs, and many others.

**Please Note**

Since CNCF now owns the Kubernetes copyright, contributors to the source need to sign a contributor license agreement (**CLA**) with CNCF, just like any contributor to an Apache-licensed project signs a CLA with the Apache Software Foundation.

2.17 Resource Recommendations

- If you want to go beyond this general introduction to Kubernetes, here are a few things we recommend:
 - Read the Borg paper.
<https://ai.google/research/pubs/pub43438>
 - Listen to John Wilkes talking about Borg and Kubernetes.
<https://www.gcppodcast.com/post/episode-46-borg-and-k8s-with-john-wilkes/>
 - Add the Kubernetes community hangout to your calendar, and attend at least once.
<https://github.com/kubernetes/community>
 - Join the community on **Slack** and go in the #kubernetes-users channel.
<https://github.com/kubernetes/community>

2.18 Labs

Exercise 2.1: Overview and Preliminaries

We will create a two-node **Ubuntu 18.04** cluster. Using two nodes allows an understanding of some issues and configurations found in a production environment. Currently 2 vCPU and 8G of memory allows for quick labs. Other Linux distributions should work in a very similar manner, but have not been tested.

**Very Important**

Regardless of the platform used (**VirtualBox**, **VMWare**, **AWS**, **GCE** or even bare metal) please remember that security software like **SELinux**, **AppArmor**, and firewall configurations can prevent the labs from working. While not something to do in production consider disabling the firewall and security software.

GCE requires a new VPC to be created and a rule allowing all traffic to be included. The use of **wireshark** can be a helpful place to start with troubleshooting network and connectivity issues if you're unable to open all ports.

The **kubeadm** utility currently requires that swap be turned off on every node. The **swapoff -a** command will do this until the next reboot, with various methods to disable swap persistently. Cloud providers typically deploy instances with swap disabled.

Download shell scripts and YAML files

To assist with setting up your cluster please download the tarball of shell scripts and YAML files. The **k8sMaster.sh** and **k8sSecond.sh** scripts deploy a Kubernetes cluster using **kubeadm** and use Project Calico for networking. Should the file not be found you can always use a browser to investigate the parent directory.

```
$ wget https://training.linuxfoundation.org/cm/LFD5461/LFD5461_V1.20_SOLUTIONS.tar.xz \
    --user=LFtraining --password=Penguin2014
```

```
$ tar -xvf LFD5461_V1.20_SOLUTIONS.tar.xz
```

(Note: depending on your software, if you are cutting and pasting the above instructions, the underscores may disappear and be replaced by spaces, so you may have to edit the command line by hand!)

Exercise 2.2: Deploy a New Cluster

Deploy a Master Node using Kubeadm

1. Log into your nodes using **PuTTY** or using **SSH** from a terminal window. Unless the instructor tells you otherwise the user name to use will be **student**. You may need to change the permissions on the pem or ppk file as shown in the following commands. Your file and node IP address will probably be different.

```
localTerm:~$ chmod 400 LFD459.pem
localTerm:~$ ssh -i LFD459.pem student@WW.XX.YY.ZZ
```

```
1 student@master:~$
```

2. Review the script to install and begin the configuration of the master kubernetes server. You may need to change the **find** command search directory which uses tilde for your home directory depending on how and where you downloaded the tarball.

A **find** command is shown if you want to locate and copy to the current directory instead of creating the file. Mark the command for reference as it may not be shown for future commands.

```
student@master:~$ find $HOME -name <YAML File>
student@master:~$ cp LFD5461/<Some Path>/<YAML File> .
```

```
student@master:~$ find $HOME -name k8sMaster.sh
```

```
student@master:~$ more LFD5461/SOLUTIONS/s_02/k8sMaster.sh
```

SH

k8sMaster.sh

```
#!/bin/bash -x
## TxS 01-2021
## v1.20.1 CKAD
echo "This script is written to work with Ubuntu 18.04"
sleep 3
echo
echo "Disable swap until next reboot"
echo
sudo swapoff -a

echo "Update the local node"
sudo apt-get update && sudo apt-get upgrade -y
echo
echo "Install Docker"
sleep 3

sudo apt-get install -y docker.io
echo
echo "Install kubeadm, kubelet, and kubect1"
sleep 3

sudo sh -c
↪ "echo 'deb http://apt.kubernetes.io/ kubernetes-xenial main' >> /etc/apt/sources.list.d/kubernetes.list"

sudo sh -c "curl -s https://packages.cloud.google.com/apt/doc/apt-key.gpg | apt-key add -"

sudo apt-get update

sudo apt-get install -y kubeadm=1.20.1-00 kubelet=1.20.1-00 kubect1=1.20.1-00
```

SH

```

sudo apt-mark hold kubelet kubeadm kubectl

echo
echo "Installed - now to get Calico Project network plugin"

## If you are going to use a different plugin you'll want
## to use a different IP address, found in that plugins
## readme file.

sleep 3

sudo kubeadm init --kubernetes-version 1.20.1 --pod-network-cidr 192.168.0.0/16

sleep 5

echo "Running the steps explained at the end of the init output for you"

mkdir -p $HOME/.kube

sleep 2

sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config

sleep 2

sudo chown $(id -u):$(id -g) $HOME/.kube/config

echo "Apply Calico network plugin from ProjectCalico.org"
echo "If you see an error they may have updated the yaml file"
echo "Use a browser, navigate to the site and find the updated file"

kubectl apply -f https://docs.projectcalico.org/manifests/calico.yaml

echo
echo
sleep 3
echo "You should see this node in the output below"
echo "It can take up to a minute for node to show Ready status"
echo
kubectl get node
echo
echo
echo "Script finished. Move to the next step"

```

- Run the script as an argument to the **bash** shell. You will need the `kubeadm join` command shown near the end of the output when you add the worker/minion node in a future step. Use the **tee** command to save the output of the script, in case you cannot scroll back to find the `kubeadm join` in the script output. Please note the following is one command and then its output.

Using **Ubuntu 18** you will be asked questions during the installation. Allow restarts and use the local, installed software if asked during the update, usually option 2.

Copy files to your home directory first.

```
student@master:~$ cp LFD5461/SOLUTIONS/s_02/k8sMaster.sh .
```

```
student@master:~$ bash k8sMaster.sh | tee $HOME/master.out
```

```

1 <output_omitted>
2

```

```

3 Your Kubernetes master has initialized successfully!
4
5     To start using your cluster, you need to run the
6     following as a regular user:
7
8     mkdir -p $HOME/.kube
9     sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
10    sudo chown $(id -u):$(id -g) $HOME/.kube/config
11
12    You should now deploy a pod network to the cluster.
13    Run kubectl apply -f [podnetwork].yaml with one
14    of the options listed at:
15    https://kubernetes.io/docs/concepts/cluster-administration/addons/
16
17    You can now join any number of machines by running the
18    following on each node as root:
19
20    kubeadm join 10.128.0.3:6443 --token 69rdjq.2x2012j9ncexy37b
21    --discovery-token-ca-cert-hash
22
23    sha256:72143e996ef78301191b9a42184124416aebcf0c7f363adf9208f9fa599079bd
24
25    <output_omitted>
26
27    NAME                STATUS    ROLES                  AGE      VERSION
28    master              NotReady control-plane,master   19s     v1.20.1
29
30
31 Script finished. Move to the next step

```

Deploy a Minion Node

4. Open a separate terminal into your **second node**. Having both terminal sessions allows you to monitor the status of the cluster while adding the second node. Change the color or other characteristic of the second terminal to make it visually distinct from the first. This will keep you from running commands on the incorrect instance, which probably won't work.

Use the previous **wget** command to download the tarball to the second node. Extract the files with **tar** as before. Find and copy the `k8sSecond.sh` file to student's home directory then view it. You should see the same early steps as on the master system.

```
student@worker:~$ more k8sSecond.sh
```

SH

k8sSecond.sh

```

#!/bin/bash -x
## TxS 01-2021
## CKAD for 1.20.1
##
echo " This script is written to work with Ubuntu 18.04"
echo
sleep 3
echo " Disable swap until next reboot"
echo
sudo swapoff -a

echo " Update the local node"
sleep 2
sudo apt-get update && sudo apt-get upgrade -y
echo

```


SH

```

sleep 2

echo "  Install Docker"
sleep 3
sudo apt-get install -y docker.io

echo
echo "  Install kubeadm, kubelet, and kubectl"
sleep 2
sudo sh -c
↪ "echo 'deb http://apt.kubernetes.io/ kubernetes-xenial main' >> /etc/apt/sources.list.d/kubernetes.list"

sudo sh -c "curl -s https://packages.cloud.google.com/apt/doc/apt-key.gpg | apt-key add -"

sudo apt-get update

sudo apt-get install -y kubeadm=1.20.1-00 kubelet=1.20.1-00 kubectl=1.20.1-00

sudo apt-mark hold kubelet kubeadm kubectl
echo
echo "  Script finished. You now need the kubeadm join command"
echo "  from the output on the master node"
echo

```

5. Run the script on the **second node**. Again please note you may have questions during the update. Allow daemons to restart and use the local installed version, usually option 2.

```
student@worker:~$ bash k8sSecond.sh
```

```
1 <output_omitted>
```

6. When the script is done the minion node is ready to join the cluster. The `kubeadm join` statement can be found near the end of the `kubeadm init` output on the master node. It should also be in the file `master.out` as well. Your nodes will use a different IP address and hashes than the example below. You'll need to pre-pend **sudo** to run the script copied from the master node. Also note that some non-Linux operating systems and tools insert extra characters when multi-line samples are copied and pasted. Copying one line at a time solves this issue.

```

student@worker:~$ sudo kubeadm join --token 118c3e.83b49999dc5dc034 \
10.128.0.3:6443 --discovery-token-ca-cert-hash \
sha256:40aa946e3f53e38271bae24723866f56c86d77efb49aedeb8a70cc189bfe2e1d

```

```
1 <output_omitted>
```

Configure the Master Node

7. Return to the master node. Install a text editor. While the lab uses **vim**, any text editor such as **emacs** or **nano** will work. Be aware that Windows editors may have issues with special characters. Also install the **bash-completion** package, if not already installed. Use the locally installed version of a package if asked.

```
student@master:~$ sudo apt-get install bash-completion vim -y
```

```
1 <output_omitted>
```

8. We will configure command line completion and verify both nodes have been added to the cluster. The first command will configure completion in the current shell. The second command will ensure future shells have completion. You may need to exit the shell and log back in for command completion to work without error.

```
student@master:~$ source <(kubectl completion bash)
```

```
student@master:~$ echo "source <(kubectl completion bash)" >> $HOME/.bashrc
```

9. Verify that both nodes are part of the cluster. And show a Ready state.

```
student@master:~$ kubectl get node
```

	NAME	STATUS	ROLES	AGE	VERSION
1	master	Ready	control-plane,master	4m11s	v1.20.1
2	worker	Ready	<none>	61s	v1.20.1

10. We will use the **kubectl** command for the majority of work with Kubernetes. Review the help output to become familiar with commands options and arguments.

```
student@master:~$ kubectl --help
```

```
1 kubectl controls the Kubernetes cluster manager.
2
3 Find more information at:
4   https://kubernetes.io/docs/reference/kubectl/overview/
5
6 Basic Commands (Beginner):
7   create          Create a resource from a file or from stdin.
8   expose          Take a replication controller, service,
9   deployment or pod and expose it as a new Kubernetes Service
10  run             Run a particular image on the cluster
11  set             Set specific features on objects
12
13 Basic Commands (Intermediate):
14 <output_omitted>
```

11. With more than 40 arguments, you can explore each also using the `--help` option. Take a closer look at a few, starting with `taint` for example.

```
student@master:~$ kubectl taint --help
```

```
1 Update the taints on one or more nodes.
2
3 * A taint consists of a key, value, and effect. As an argument
4   here, it is expressed as key=value:effect.
5 * The key must begin with a letter or number, and may contain
6   letters, numbers, hyphens, dots, and underscores, up to
7   253 characters.
8 * Optionally, the key can begin with a DNS subdomain prefix
9   and a single '/',
10 like example.com/my-app
11 <output_omitted>
```

12. By default the master node will not allow general containers to be deployed for security reasons. This is via a taint. Only containers which tolerate this taint will be scheduled on this node. As we only have two nodes in our cluster we will remove the taint, allowing containers to be deployed on both nodes. This is not typically done in a production environment for security and resource contention reasons. The following command will remove the taint from all nodes, so you should see one success and one `not found` error. The worker/minion node does not have the taint to begin with. Note the **minus sign** at the end of the command, which removes the preceding value.

```
student@master:~$ kubectl describe nodes | grep -i Taint
```

```
1 Taints:          node-role.kubernetes.io/master:NoSchedule
2 Taints:          <node>
```

```
student@master:~$ kubectl taint nodes --all node-role.kubernetes.io/master-
```

```
1 node/master untainted
2 error: taint "node-role.kubernetes.io/master:" not found
```

13. Check that both nodes are without a Taint. If they both are without taint the nodes should now show as Ready. It may take a minute or two for all infrastructure pods to enter Ready state, such that the nodes will show a Ready state.

```
student@master:~$ kubectl describe nodes | grep -i taint
```

```
1 Taints:          <none>
2 Taints:          <none>
```

```
student@master:~$ kubectl get nodes
```

```
1 NAME          STATUS    ROLES    AGE    VERSION
2 master        Ready    master   6m1s   v1.20.1
3 worker        Ready    <none>   5m31s   v1.20.1
```

Exercise 2.3: Create a Basic Pod

1. The smallest unit we directly control with Kubernetes is the pod. We will create a pod by creating a minimal YAML file. First we will get a list of current API objects and their APIGROUP. If value is not shown it may not exist, as with SHORTNAMES. Note that pods does not declare an APIGROUP. At the moment this indicates it is part of the stable v1 group.

```
student@master:~$ kubectl api-resources
```

```
1 NAME          SHORTNAMES  APIVERSION  NAMESPACE  KIND
2 bindings              v1          true        Binding
3 componentstatuses  cs          v1          false       ComponentStatus
4 configmaps         cm          v1          true        ConfigMap
5 endpoints          ep          v1          true        Endpoints
6 .....
7 pods              po          v1          true        Pod
8 ....
```

2. From the output we see most are v1 which is used to denote a stable object. With that information we will add the other three required sections for pods such as `metadata`, with a `name`, and `spec` which declares which container image to use and a `name` for the container. We will create an eight line YAML file. White space and indentation matters. Don't use **Tabs**. There is a `basic.yaml` file available in the tarball, as well as `basic-later.yaml` which shows what the file will become and can be helpful for figuring out indentation.

YAML

`basic.yaml`

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: basicpod
5 spec:
6   containers:
7   - name: webcont
8     image: nginx
```

3. Create the new pod using the recently created YAML file.

```
student@master:~$ kubectl create -f basic.yaml
```

```
1 pod/basicpod created
```

4. Make sure the pod has been created then use the **describe** sub-command to view the details. Among other values in the output you should be about to find the image and the container name.

```
student@master:~$ kubectl get pod
```

```
1 NAME      READY   STATUS    RESTARTS   AGE
2 basicpod  1/1     Running   0           23s
```

```
student@master:~$ kubectl describe pod basicpod
```

```
1 Name:          basicpod
2 Namespace:     default
3 Priority:       0
4 <output_omitted>
```

5. Shut down the pod and verify it is no longer running.

```
student@master:~$ kubectl delete pod basicpod
```

```
1 pod "basicpod" deleted
```

```
student@master:~$ kubectl get pod
```

```
1 No resources found in default namespace.
```

6. We will now configure the pod to expose port 80. This configuration does not interact with the container to determine what port to open. We have to know what port the process inside the container is using, in this case port 80 as a web server. Add two lines to the end of the file. Line up the indentation with the `image` declaration.

```
student@master:~$ vim basic.yaml
```

YAML

basic.yaml

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: basicpod
5 spec:
6   containers:
7   - name: webcont
8     image: nginx
9     ports:                                #<--Add this and following line
10   - containerPort: 80
```

7. Create the pod and verify it is running. Use the `-o wide` option to see the internal IP assigned to the pod, as well as `NOMINATED NODE`, which is used by the scheduler and `READINESS GATES` which show if experimental features are enabled. Using **curl** and the pods IP address you should get the default nginx welcome web page.

```
student@master:~$ kubectl create -f basic.yaml
```

```
1 pod/basicpod created
```

```
student@master:~$ kubectl get pod -o wide
```

```

1 NAME      READY STATUS  RESTARTS AGE   IP           NODE
2 NOMINATED NODE  READINESS GATES
3 basicpod 1/1    Running 0      9s    192.168.1.3  master
4 <none>      <none>

```

```
student@master:~$ curl http://192.168.1.3
```

```

1 <!DOCTYPE html>
2 <html>
3 <head>
4 <title>Welcome to nginx!</title>
5
6 <output_omitted>

```

```
student@master:~$ kubectl delete pod basicpod
```

```
1 pod "basicpod" deleted
```

8. We will now create a simple service to expose the pod to other nodes and pods in the cluster. The service YAML will have the same four sections as a pod, but different spec configuration and the addition of a selector.

```
student@master:~$ vim basicservice.yaml
```

YAML

basicservice.yaml

```

1 apiVersion: v1
2 kind: Service
3 metadata:
4   name: basicservice
5 spec:
6   selector:
7     type: webserver
8   ports:
9   - protocol: TCP
10     port: 80

```

9. We will also add a label to the pod and a selector to the service so it knows which object to communicate with.

```
student@master:~$ vim basic.yaml
```

YAML

basic.yaml

```

1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: basicpod
5   labels:                                #<-- Add this line
6     type: webserver                      #<-- and this line which matches selector
7 spec:
8   ....

```

10. Create the new pod and service. Verify both have been created.

```
student@master:~$ kubectl create -f basic.yaml
```

```
1 pod/basicpod created
```

```
student@master:~$ kubectl create -f basicservice.yaml
```

```
1 service/basicservice created
```

```
student@master:~$ kubectl get pod
```

```
1 NAME      READY   STATUS    RESTARTS   AGE
2 basicpod   1/1     Running   0           110s
```

```
student@master:~$ kubectl get svc
```

```
1 NAME          TYPE          CLUSTER-IP    EXTERNAL-IP  PORT(S)    AGE
2 basicservice  ClusterIP     10.96.112.50  <none>       80/TCP     14s
3 kubernetes    ClusterIP     10.96.0.1     <none>       443/TCP    4h
```

11. Test access to the web server using the CLUSTER-IP for the basicservice.

```
student@master:~$ curl http://10.96.112.50
```

```
1 <!DOCTYPE html>
2 <html>
3 <head>
4 <title>Welcome to nginx!</title>
5
6 <output_omitted>
```

12. We will now expose the service to outside the cluster as well. Delete the service, edit the file and add a type declaration.

```
student@master:~$ kubectl delete svc basicservice
```

```
1 service "basicservice" deleted
```

```
student@master:~$ vim basicservice.yaml
```

YAML

basicservice.yaml

```
1 apiVersion: v1
2 kind: Service
3 metadata:
4   name: basicservice
5 spec:
6   selector:
7     type: webserver
8   type: NodePort      #<--Add this line
9   ports:
10  - protocol: TCP
11    port: 80
```

13. Create the service again. Note there is a different TYPE and CLUSTER-IP and also a high-numbered port.

```
student@master:~$ kubectl create -f basicservice.yaml
```

```
1 service/basicservice created
```

```
student@master:~$ kubectl get svc
```

	NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
1	basicservice	NodePort	10.100.139.155	<none>	80:31514/TCP	3s
3	kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	47h

14. Using the public IP address of the node and the high port you should be able to test access to the webserver. In the example below the public IP is 35.238.3.83, yours will be different. The high port will also probably be different. Note that testing from within a GCE or AWS node will not work. Use a local to you terminal or web browser to test.

```
local$ curl http://35.238.3.83:31514
```

```
1 <!DOCTYPE html>
2 <html>
3 <head>
4 <title>Welcome to nginx!</title>
5 <style>
6 <output_omitted>
```

✎ Exercise 2.4: Multi-Container Pods

Using a single container per pod allows for the most granularity and decoupling. There are still some reasons to deploy multiple containers, sometimes called *composite containers*, in a single pod. The secondary containers can handle logging or enhance the primary, the *sidecar* concept, or acting as a proxy to the outside, the *ambassador* concept, or modifying data to meet an external format such as an *adapter*. All three concepts are secondary containers to perform a function the primary container does not.

1. We will add a second container to the pod to handle logging. Without going into details of how to use **fluentd** we will add a logging container to the existing pod from its own repository. The second container would act as a *sidecar*. At this state we will just add the second container and verify it is running. In the **Deployment Configuration** chapter we will continue to work on this pod by adding persistent storage and configure **fluentd** via a `configMap`.

Edit the YAML file and add a **fluentd** container. The dash should line up with the previous container dash. At this point a name and image should be enough to start the second container.

```
student@master:~$ vim basic.yaml
```

YAML

basic.yaml

```
1 ....
2 containers:
3   - name: webcont
4     image: nginx
5     ports:
6   - containerPort: 80
7   - name: fdlogger
8     image: fluent/fluentd
```

2. Delete and create the pod again. The commands can be typed on a single line, separated by a semicolon. This time you should see 2/2 under the READY column. You should also find information on the **fluentd** container inside of the `kubectl describe` output.

```
student@master:~$ kubectl delete pod basicpod ; kubectl create -f basic.yaml
```

```
1 pod "basicpod" deleted
2 pod/basicpod created
```

```
student@master:~$ kubectl get pod
```

```
1 NAME          READY   STATUS    RESTARTS   AGE
2 basicpod      2/2     Running   0           2m8s
```

```
student@master:~$ kubectl describe pod basicpod
```

```
1 Name:          basicpod
2 Namespace:     default
3 Priority:      0
4 Node:         master/10.128.0.11
5
6 ....
7
8 fdlogger:
9   Container ID:  docker://f0649457217f00175ce9aec35022d0b238b9b....
10  Image:         fluent/fluentd
11
12 ....
```

3. For now shut down the pod. We will use it again in a future exercise.

```
student@master:~$ kubectl delete pod basicpod
```

```
1 pod "basicpod" deleted
```

Exercise 2.5: Create a Simple Deployment

Creating a pod does not take advantage of orchestration abilities of Kubernetes. We will now create a Deployment which gives us scalability, reliability, and updates.

1. Now run a containerized webserver **nginx**. Use **kubectl create** to create a simple, single replica deployment running the nginx web server. It will create a single pod as we did previously but with new controllers to ensure it runs as well as other features.

```
student@master:~$ kubectl create deployment firstpod --image=nginx
```

```
1 deployment.apps/firstpod created
```

2. Verify the new deployment exists and the desired number of pods matches the current number. Using a comma, you can request two resource types at once. The **Tab** key can be helpful. Type enough of the word to be unique and press the **Tab** key, it should complete the word. The deployment should show a number 1 for each value, such that the desired number of pods matches the up-to-date and running number. The pod should show zero restarts.

```
student@master:~$ kubectl get deployment,pod
```

```
1 NAME          READY  UP-TO-DATE  AVAILABLE  AGE
2 deployment.apps/firstpod  1/1    1           1           2m42s
3
4 NAME          READY  STATUS    RESTARTS  AGE
5 pod/firstpod-7d88d7b6cf-lrsbk  1/1    Running   0           2m42s
```

3. View the details of the deployment, then the pod. Work through the output slowly. Knowing what a healthy deployment and looks like can be helpful when troubleshooting issues. Again the **Tab** key can be helpful when using long auto-generated object names. You should be able to type firstpod**Tab** and the name will complete when viewing the pod.

```
student@master:~$ kubectl describe deployment firstpod
```



```

1 Name:          firstpod
2 Namespace:     default
3 CreationTimestamp: Wed, 15 Apr 2020 17:17:25 +0000
4 Labels:        app=firstpod
5 Annotations:    deployment.kubernetes.io/revision=1
6 Selector:      app=firstpod
7 Replicas:      1 desired | 1 updated | 1 total | 1 available...
8 StrategyType:   RollingUpdate
9 MinReadySeconds: 0
10 <output_omitted>

```

```
student@master:~$ kubectl describe pod firstpod-6bb4574d94-rqk76
```

```

1 Name:          firstpod-6bb4574d94-rqk76
2 Namespace:     default
3 Priority:       0
4 PriorityClassName: <none>
5 Node:          master/10.128.0.2
6 Start Time:    Wed, 15 Apr 2020 17:17:25 +0000
7 Labels:        pod-template-hash=2660130850
8                app=firstpod
9 Annotations:    cni.projectcalico.org/podIP: 192.168.200.65/32
10 Status:       Running
11 IP:           192.168.200.65
12 Controlled By: ReplicaSet/firstpod-6bb4574d94
13
14 <output_omitted>

```

4. Note that the resources are in the default namespace. Get a list of available namespaces.

```
student@master:~$ kubectl get namespaces
```

NAME	STATUS	AGE
default	Active	20m
kube-node-lease	Active	20m
kube-public	Active	20m
kube-system	Active	20m

5. There are four default namespaces. Look at the pods in the kube-system namespace.

```
student@master:~$ kubectl get pod -n kube-system
```

NAME	READY	STATUS	RESTARTS	AGE
calico-node-5ftrr	2/2	Running	0	24m
calico-node-f7zrw	2/2	Running	0	21m
coredns-fb8b8dccb-cmkds	1/1	Running	0	24m
coredns-fb8b8dccb-grltk	1/1	Running	0	24m
etcd-v141-r24p	1/1	Running	0	23m

6. Now look at the pods in a namespace that does not exist. Note you do not receive an error.

```
student@master:~$ kubectl get pod -n fakenamespace
```

```
1 No resources found in fakenamespace namespace.
```

7. You can also view resources in all namespaces at once. Use the `--all-namespaces` options to select objects in all namespaces at once.

```
student@master:~$ kubectl get pod --all-namespaces
```

```
1 NAMESPACE      NAME                                     READY   STATUS    RESTARTS   AGE
2 default        firstpod-69cfd8d9-kj6ql                1/1     Running   0           44m
3 kube-system    calico-node-5ftrr                      2/2     Running   0           92m
4 kube-system    calico-node-f7zrw                      2/2     Running   0           89m
5 kube-system    coredns-fb8b8dcf-cmkds                 1/1     Running   0           92m
6 <output_omitted>
```

8. View several resources at once. Note that most resources have a short name such as rs for ReplicaSet, po for Pod, svc for Service, and ep for endpoint. Note the endpoint still exists after we deleted the pod.

```
student@master:~$ kubectl get deploy,rs,po,svc,ep
```

```
1 NAME                                READY   UP-TO-DATE   AVAILABLE   AGE
2 deployment.apps/firstpod            1/1     1             1           4m
3
4 NAME                                DESIRED   CURRENT   READY....
5 replicaset.apps/firstpod-6bb4574d94-rqk76 1         1         1 ....
6
7 NAME                                READY   STATUS    RESTARTS   AGE
8 pod/firstpod-6bb4574d94-rqk76        1/1     Running   0           4m
9
10 NAME                                TYPE          CLUSTER-IP   EXTERNAL-IP  PORT(S)    AGE
11 service/basicservice                NodePort      10.108.147.76 <none>       80:31601/TCP 21m
12 service/kubernetes                 ClusterIP      10.96.0.1     <none>       443/TCP    21m
13
14 NAME                                ENDPOINTS      AGE
15 endpoints/basicservice              <none>         21m
16 endpoints/kubernetes                10.128.0.3:6443 21m
```

9. Delete the ReplicaSet and view the resources again. Note that the age on the ReplicaSet and the pod it controls is now less than a minute of age. The deployment operator started a new ReplicaSet operator when we deleted the existing one. The new ReplicaSet started another pod when the desired spec did not match the current status.

```
student@master:~$ kubectl delete rs firstpod-6bb4574d94-rqk76
```

```
1 replicaset.apps "firstpod-6bb4574d94-rqk76" deleted
```

```
student@master:~$ kubectl get deployment,rs,po,svc,ep
```

```
1 NAME                                READY   UP-TO-DATE   AVAILABLE   AGE
2 deployment.apps/firstpod            1/1     1             1           7m
3
4 NAME                                DESIRED   CURRENT....
5 replicaset.apps/firstpod-6bb4574d94-rqk76 1         1         ....
6
7 NAME                                READY   STATUS    RESTARTS   AGE
8 pod/firstpod-7d99ffc75-p9hbw        1/1     Running   0           12s
9
10 NAME                                TYPE          CLUSTER-IP   EXTERNAL-IP  PORT(S)    AGE
11 service/kubernetes                 ClusterIP      10.96.0.1     <none>       443/TCP    24m
12
13 NAME                                ENDPOINTS      AGE
14 endpoints/kubernetes                10.128.0.2:6443 80m
15 endpoints/basicservice              <none>         21m
```

10. This time delete the top-level controller. After about 30 seconds for everything to shut down you should only see the cluster service and endpoint remain for the cluster and the service we created.

```
student@master:~$ kubectl delete deployment firstpod
```

```
1 deployment.apps "firstpod" deleted
```

```
student@master:~$ kubectl get deployment,rs,po,svc,ep
```

```
1 NAME                TYPE          CLUSTER-IP    EXTERNAL-IP  PORT(S)        AGE
2 service/basicservice NodePort      10.108.147.76 <none>       80:31601/TCP   35m
3 kubernetes          ClusterIP    10.96.0.1     <none>       443/TCP        24m
4
5 NAME                ENDPOINTS      AGE
6 endpoints/basicservice <none>        21m
7 kubernetes          10.128.0.3:6443 24m
```

11. As we won't need it for a while, delete the basicservice service as well.

```
student@master:~$ kubectl delete svc basicservice
```

```
1 service "basicservice" deleted
```

Exercise 2.6: Domain Review



Very Important

The source pages and content in this review could change at any time. **IT IS YOUR RESPONSIBILITY TO CHECK THE CURRENT INFORMATION.**

1. Using a browser go to <https://www.cncf.io/certification/ckad/> and read through the program description.
2. In the **Exam Resources** section open the [Curriculum Overview](#) and [Candidate-handbook](#) in new tabs. Both of these should be read and understood prior to sitting for the exam.
3. Navigate to the [Curriculum Overview](#) tab. You should see links for domain information for various versions of the exam. Select the latest version, such as **CKAD_Curriculum_V1.20.1.pdf**. The versions you see may be different. You should see a new page showing a PDF.
4. Read through the document. Be aware that the term Understand, such as Understand Services, is more than just knowing they exist. In this case expect it to also mean create, update, and troubleshoot.
5. Locate the **Core Concepts** section. If you review the lab, you will see we have covered these steps. Again, please note this document will change, distinct from this book. **It remains your responsibility to check for changes in the online document.** They may change on an irregular and unannounced basis.

Certified Kubernetes Application Developer (CKAD) Exam Curriculum v1.15.0

This document provides the curriculum outline of the Knowledge, Skills and Abilities that a Certified Kubernetes Application Developer (CKAD) can be expected to demonstrate.

CKAD Curriculum v1.15.0

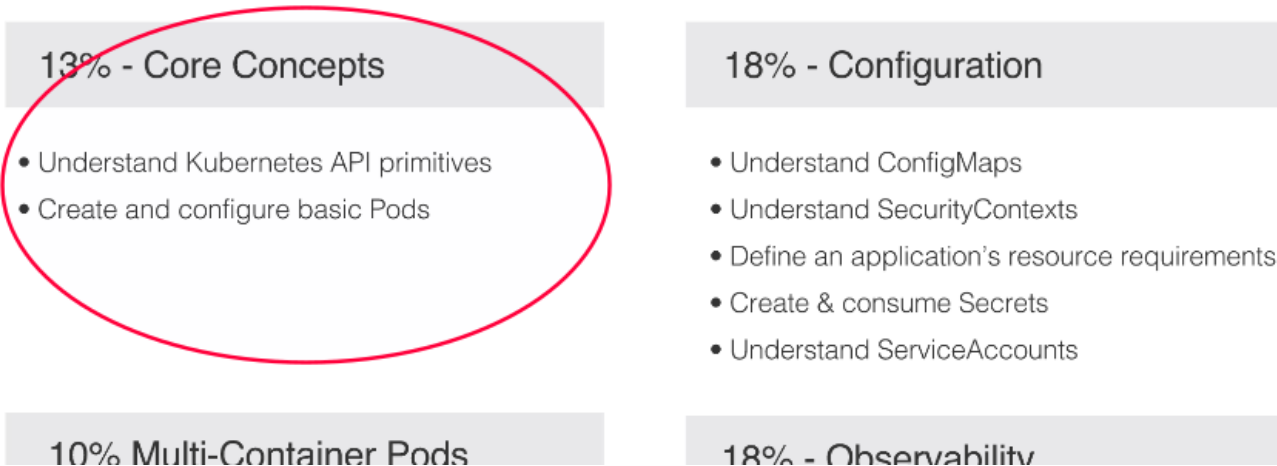


Figure 2.10: Core Concepts Domain

- Navigate to the **Candidate-handbook** tab. You are strongly encourage to read and understand this entire document prior to taking the exam. Again, please note this document will change, distinct from this book. **It remains your responsibility to check for changes in the online document.** They may change on an irregular and unannounced basis.
- Find the **Guidelines and Tips for Use of the Linux server terminal**” section in the document.
- Among other points you will note the current exam version and three (at the time this was written) domains and subdomains you can use, with some stated conditions.

working on the correct cluster.

12. The clusters comprising the exam environment are currently running Kubernetes 1.16

13. You are permitted to use your Chrome or Chromium browser to open one additional tab in order to access assets at <https://kubernetes.io/docs/> and its subdomain, <https://github.com/kubernetes/> and its subdomains, or <https://kubernetes.io/blog/>. No other tabs may be opened and no other sites may be navigated to. *The allowed sites above may contain links that point to external sites. It is the responsibility of the candidate not to click on any links that cause them to navigate to a domain that is not allowed.*

14. Issues with unapproved text within the terminal pane may be resolved by resizing your

Figure 2.11: Exam Handbook Guidelines and Tips

- Using only the allowed browser, URLs, and subdomains search for and bookmark a YAML example to create and configure a basic pod. Ensure it works for the version of the exam you are taking. URLs may change, plan on checking each book mark prior to taking the exam.
- Using a timer and bookmarked YAML files see how long it takes you to create and verify. Try it again and see how much faster you can complete and test each step:
 - A new pod with the **nginx** image. Showing all containers running and a Ready status.
 - A new service exposing the pod as a **nodePort**, which presents a working webserver configured in the previous step.

- Update the pod to run the **nginx:1.11-alpine** image and re-verify you can view the webserver via a nodePort.

11. Find and use the `architecture-review1.yaml` file included in the course tarball. Your path, such as course number, may be different than the one in the example below. Use the **find** output. Determine if the pod is running. Fix any errors you may encounter. The use of **kubectl describe** may be helpful.

```
student@master:~$ find $HOME -name architecture-review1.yaml
```

```
1 /home/student/LFD259/SOLUTIONS/s_02/architecture-review1.yaml
```

```
student@master:~$ cp <copy-paste-from-above> .
```

```
student@master:~$ kubectl create -f architecture-review1.yaml
```

12. Remove any pods or services you may have created as part of the review before moving on to the next section. For example:

```
student@master:~$ kubectl delete -f architecture-review1.yaml
```


Chapter 3

Build



In this chapter, we are going to deploy our new Python application into our Kubernetes cluster. We'll first containerize it, using Docker commands, then create and use our own local registry. After testing that the registry is available to all of our nodes, we will deploy our application in a multi-container pod. In keeping with the ideas of decoupled applications that are transient, we will also use readinessProbes, to make sure that the application is fully running prior to accepting traffic from the cluster, and then, livenessProbes, to ensure that the container and the application continues to run in a healthy state, and startupProbes to allow for a slow-starting application.

3.1	Container Options	39
3.2	Containerizing an Application	41
3.3	Creating the Dockerfile	41
3.4	Hosting a Local Repository	41
3.5	Creating a Deployment	42
3.6	Running Commands in a Container	42
3.7	Multi-Container Pod	42
3.8	readinessProbe	43
3.9	livenessProbe	43
3.10	startupProbe	43
3.11	Testing	43
3.12	Labs	44

3.1 Container Options

- There are many competing organizations working with containers. As an orchestration tool, Kubernetes is being developed to work with many of them, with the overall community moving toward open standards and easy interoperability. The early and strong presence of Docker meant that historically, this was not the focus. As Docker evolved, spreading their vendor-lock characteristics through the container creation and deployment life cycle, new projects and features have become popular. As other container engines become mature Kubernetes continues to become more open and independent.
- A **container runtime** is the component which runs the containerized application upon request. Docker Engine remains the default for Kubernetes, though **cri-o** and others are gaining community support.
- The containerized image is moving from Docker to one that is not bound to higher-level tools and that is more portable across operating systems and environments. The **Open Container Initiative (OCI)** was formed to help with this. Docker

donated their **libcontainer** project to form a new code base called runC to support these goals. More information about runC can be found on GitHub at <https://github.com/opencontainers/runc>.

- Where Docker was once the only real choice for developers, the trend toward open specifications and flexibility indicates that building with vendor-neutral features is a wise choice.

Docker

- Launched in 2013, Docker has become synonymous with running containerized applications. Docker made containerizing, deploying, and consuming applications easy. As a result, it became the default option in production. With an open registry of images, Docker Hub (<https://hub.docker.com/>), you can download and deploy vendor or individual-created images on multiple architectures with a single and easy to use tool set. This ease meant it was the sensible default choice for any developer as well. Kubernetes defaults to using the Docker engine to run containers.
- Over the past few years, Docker has continued to grow and add features, including orchestration, which they call Swarm. These added features addressed some of the pressing needs in production, but also increased the vendor-lock and size of the product. This has led to an increase in open tools and specifications such as **cri-o**. A developer looking toward the future would be wise to work with mostly open tools for containers and Kubernetes, but he or she should understand that Docker is still the production tool of choice outside of a Red Hat environment at the moment.

Container Runtime Interface (CRI)

- The goal of the Container Runtime Interface (CRI) is to allow easy integration of container runtimes with kubelet. By providing a protobuf method for API, specifications and libraries, new runtimes can easily be integrated without needing deep understanding of kubelet internals.
- The project is in early stage, with lots of development in action. Now that Docker-CRI integration is done, new runtimes should be easily added and swapped out. At the moment, **cri-o**, **rktlet** and **frakti** are listed as work-in-progress.

rkt

After Red Hat bought CoreOS work on **rkt** stopped, in favor of **cri-o**. The **rkt** project has become the first archived project of CNCF. While the code is available to be forked and worked on, no one has done so at this point.

- The rkt (<https://github.com/rkt/rkt>) runtime, pronounced rocket, provides a CLI for running containers. Announced by CoreOS in 2014, it is now part of the Cloud Native Computing Foundation family of projects. Learning from early Docker issues, it is focused on being more secure, open and inter-operable. Many of its features have been met by Docker improvements. It is not quite an easy drop-in replacement for Docker, but progress has been made. **rkt** uses the **apoc** specification, and can run Docker, **apoc** and **OCI** images. It deploys immutable pods.
- There has been a lot of attention to the project and it was expected to be the leading replacement for Docker until **cri-o** became part of the official Kubernetes Incubator.

CRI-O

- This project is currently in incubation as part of Kubernetes. It uses the Kubernetes Container Runtime Interface with OCI-compatible runtimes, thus the name CRI-O (<https://github.com/kubernetes-sigs/cri-o>). Currently, there is support for runC (default) and Clear Containers, but a stated goal of the project is to work with any OCI-compliant runtime.
- While newer than Docker or rkt, this project has gained major vendor support due to its flexibility and compatibility.

Containerd

- The intent of the **containerd** project is not to build a user-facing tool; instead, it is focused on exposing highly-decoupled low-level primitives:
 - Defaults to runC to run containers according to the OCI Specifications
 - Intended to be embedded into larger systems
 - Minimal CLI, focused on debugging and development.
- With a focus on supporting the low-level, or backend plumbing of containers, this project is better suited to integration and operation teams building specialized products, instead of typical build, ship, and run application.

3.2 Containerizing an Application

- To containerize an application, you begin by creating your application. Not all applications do well with containerization. The more stateless and transient the application, the better. Also, remove any environmental configuration, as those can be provided using other tools, like `ConfigMaps` and `secrets`. Develop the application until you have a single build artifact which can be deployed to multiple environments without needing to be changed, using decoupled configuration instead. Many legacy applications become a series of objects and artifacts, residing among multiple containers.
- The use of **Docker** has been the industry standard. Large companies such as **Red Hat** are moving to other open source tools. While currently new one may expect them to become the new standard.
 - **buildah** - This tool is focused on creating Open Container Initiative (OCI) images. It allows for creating images with and without a `Dockerfile`. It also does not require super-user privilege. A growing `golang` based API allows for easy integration with other tools. More can be found here: <https://github.com/containers/buildah>
 - **podman** - This "pod manager" tool allows for the life cycle management of a container, including creating, starting, stopping, and updating. One could consider this a replacement for `docker run`. More can be found here: <https://github.com/containers/libpod>

3.3 Creating the Dockerfile

- Create a directory to hold application files. The **docker build** process will pull everything in the directory when it creates the image. Move the scripts and files for the containerized application into the directory.
- Also, in the directory, create a `Dockerfile`. The name is important, as it must be those ten characters, beginning with a capital D; newer versions allow a different filename to be used after `-f <filename>`. This is the expected file for the **docker build** to know how to create the image.
- Each instruction is iterated by the Docker daemon in sequence. The instructions are not case-sensitive, but are often written in uppercase to easily distinguish from arguments. This file must have the `FROM` instruction first. This declares the base image for the container. Following can be a collection of `ADD`, `RUN` and a `CMD` to add resources and run commands to populate the image. More details about the Dockerfile reference at <https://docs.docker.com/engine/reference/builder> can be found in the Docker documentation.
- Test the image by verifying that you can see it listed among other images, and use `docker run <app-name>` to execute it. Once you find the application performs as expected, you can push the image to a local repository in Docker Hub, after creating an account. Here is process summary:
 - Create Dockerfile
 - Build the container


```
sudo docker build -t simpleapp
```
 - Verify the image


```
sudo docker images
sudo docker run simpleapp
```
 - Push to the repository


```
sudo docker push
```

3.4 Hosting a Local Repository

- While Docker has made it easy to upload images to their Hub, these images are then public and accessible to the world. A common alternative is to create a local repository and push images there. While it can add administrative overhead, it can save downloading bandwidth, while maintaining privacy and security.
- Once the local repository has been configured, you can use **docker tag**, then **push** to populate the repository with local images. Consider setting up an insecure repository until you know it works, then configuring TLS access for greater security.

3.5 Creating a Deployment

- Once you can push and pull images using the **docker** command, try to run a new deployment inside Kubernetes using the image. The string passed to the `--image` argument includes the repository to use, the name of the application, then the version.
- Use **kubectl create** to test the image:

```
1 kubectl create deployment <Deploy-Name> --image=<repo>/<app-name>:<version>
```

```
kubectl create deployment time-date --image=10.110.186.162:5000/simpleapp:v2.2
```

- Be aware that the string “latest” is just that: a string, and has no reference to actually being the latest. Only use this string if you have a process in place to name and rename versions of the application as they become available. If this is not understood, you could be using the “latest” release, which is several releases older than the most current.
- Verify the Pod shows a running status and that all included containers are running as well.

```
kubectl get pods
```

- Test that the application is performing as expected, running whichever tempest and QA testing the application has. Terminate a pod and test that the application is as transient as designed.

3.6 Running Commands in a Container

- Part of the testing may be to execute commands within the Pod. What commands are available depend on what was included in the base environment when the image was created. In keeping to a decoupled and lean design, it's possible that there is no shell, or that the Bourne shell is available instead of bash. After testing, you may want to revisit the build and add resources necessary for testing and production.
- Use the `-it` options for an interactive shell instead of the command running without interaction or access.
- If you have more than one container, declare which container:

```
kubectl exec -it <Pod-Name> -- /bin/bash
```

3.7 Multi-Container Pod

- It may not make sense to recreate an entire image to add functionality like a shell or logging agent. Instead, you could add another container to the pod, which would provide the necessary tools.
- Each container in the pod should be transient and decoupled. If adding another container limits the scalability or transient nature of the original application, then a new build may be warranted.

- With most network plugins every container in a pod shares a single IP address and namespace. Each container has equal potential access to storage given to the pod. Kubernetes does not provide any locking, so your configuration or application should be such that containers do not have conflicting writes.

One container could be read only while the other writes, you could configure the containers to write to different directories in the volume, or the application could have built in locking. Without these protections there would be no way to control the order of containers writing to the storage.

- There are three terms often used for multi-container pods, ambassador, adapter, and sidecar. Based off of the documentation you may think there is some setting for each. There is not. Each term is an expression of what a secondary pod is intended to do. All are just multi-container pods.
 - **ambassador** - This type of secondary container would be used to communicate with outside resources, often outside the cluster. Using a proxy, like **Envoy** or other, you can embed a proxy instead of using one provided by the cluster. Helpful if you are unsure of the cluster configuration.
 - **adapter** - This type of secondary container is useful to modify the data the primary container generates. For example the Microsoft version of ASCII is distinct from everyone else. You may need to modify a datastream for proper use.
 - **sidecar** - Similar to a sidecar on a motorcycle, it does not provide the main power, but it does help carry stuff. A sidecar is a secondary container which helps or provides a service not found in the primary application. Logging containers are a common sidecar.

3.8 readinessProbe

- Oftentimes, our application may have to initialize or be configured prior to being ready to accept traffic. As we scale up our application, we may have containers in various states of creation. Rather than communicate with a client prior to being fully ready, we can use a `readinessProbe`. Kubelet will continue to test the container every `periodSeconds` and will not direct a service to send it traffic until the test succeeds.
- With the `exec` statement, the container is not considered ready until a command returns a zero exit code. As long as the return is non-zero, the container is considered not ready and kubelet will continue to test.
- Another type of probe uses an HTTP GET request (`httpGet`). Using a defined header to a particular port and path, the container is not considered healthy until the web server returns a code 200–399. Any other code indicates failure, and the probe will try again.
- The TCP Socket probe (`tcpSocket`) will attempt to open a socket on a predetermined port, and keep trying based on `periodSeconds`. Once the port can be opened, the container is considered healthy.

3.9 livenessProbe

- In addition to deciding if a container is available for traffic you may want a container to be restarted if it fails. In this case you can use a `livenessProbe`. It shares the same syntax as a `readinessProbe`, and only differs in the name.
- Both probes are often configured into a deployment to ensure applications are ready for traffic and are restarted if unhealthy.

3.10 startupProbe

- Becoming beta recently we can also use the `startupProbe`. This probe is useful for testing application which takes a long time to start.
- If kubelet uses a `startupProbe` it will disable liveness and readiness checks until the application passes the test. The duration until the container is considered failed is `failureThreshold` times `periodSeconds`. For example if your `periodSeconds` was set to five seconds, and your `failureThreshold` was set to ten, kubelet would check the application every five seconds until it succeeds or is considered failed after a total of 50 seconds. If you set the `periodSeconds` to 60 seconds, kubelet would keep testing for 300 seconds, or five minutes, before considering the container failed.

3.11 Testing

- With the decoupled and transient nature and great flexibility, there are many possible combinations of deployments. Each deployment would have its own method for testing. No matter which technology is implemented, the goal is the end user getting what is expected. Building a test suite for your newly deployed application will help speed up the development process and limit issues with the Kubernetes integration.
- In addition to overall access, building tools which ensure the distributed application functions properly, especially in a transient environment, is a good idea.
- While custom-built tools may be best at testing a deployment there are some built-in **kubectl** arguments to begin the process. The first is the **describe** and the next would be **logs**

You can see the details, conditions, volumes and events for an object with **describe**. The Events at the end of the output can be helpful during testing as it presents a chronological and node view of cluster actions and associated messages. A simple `nginx` pod may show the following output.

```
kubectl describe pod test1
```

```
1  ....
2  Events:
3      Type      Reason      Age   From              Message
4      ----      -
5      Normal    Scheduled   18s   default-scheduler Successfully assigned default/test1 to master
6      Normal    Pulling     17s   kubelet, master   Pulling image "nginx"
7      Normal    Pulled      16s   kubelet, master   Successfully pulled image "nginx"
8      Normal    Created     16s   kubelet, master   Created container nginx
9      Normal    Started     16s   kubelet, master   Started container nginx
```

- A next step in testing may be to look at the output of containers within a pod. Not all applications will generate logs, so it could be difficult to know if the lack of output is due to error or configuration.

```
kubectl logs test1
```

```
1  
```

A different pod may be configured to show lots of log output, such as the **etcd** pod:

```
kubectl -n kube-system logs etcd-master    ##<-- Use tab to complete for your etcd pod
```

```
1  ....
2  2020-01-15 22:08:31.613174 I | mvcc: store.index: compact 446329
3  2020-01-15 22:08:31.627117 I | mvcc: finished scheduled compaction at 446329 (took 13.540866ms)
4  2020-01-15 22:13:31.622615 I | mvcc: store.index: compact 447019
5  2020-01-15 22:13:31.638184 I | mvcc: finished scheduled compaction at 447019 (took 15.122934ms)
6  2020-01-15 22:18:31.626845 I | mvcc: store.index: compact 447709
7  2020-01-15 22:18:31.640592 I | mvcc: finished scheduled compaction at 447709 (took 13.309975ms)
8  2020-01-15 22:23:31.637814 I | mvcc: store.index: compact 448399
9  2020-01-15 22:23:31.651318 I | mvcc: finished scheduled compaction at 448399 (took 13.170176ms)
```

3.12 Labs

Exercise 3.1: Deploy a New Application

Overview

In this lab we will deploy a very simple **Python** application, test it using Docker, ingest it into Kubernetes and configure probes to ensure it continues to run. This lab requires the completion of the previous lab, the installation and

configuration of a Kubernetes cluster.

Working with Python

1. Install python on your master node. It may already be installed, as is shown in the output below.

```
student@master:~$ sudo apt-get -y install python
1 Reading package lists... Done
2 Building dependency tree
3 Reading state information... Done
4 python is already the newest version (2.7.12-1~16.04).
5 python set to manually installed.
6 0 upgraded, 0 newly installed, 0 to remove and 5 not upgraded.
```

2. Locate the python binary on your system.

```
student@master:~$ which python
```

```
1 /usr/bin/python
```

3. Create and change into a new directory. The Docker build process pulls everything from the current directory into the image file by default. Make sure the chosen directory is empty.

```
student@master:~$ mkdir app1
student@master:~$ cd app1
student@master:~/app1$ ls -l
```

```
1 total 0
```

4. Create a simple python script which prints the time and hostname every 5 seconds. There are six commented parts to this script, which should explain what each part is meant to do. The script is included with others in the course tar file, though you are encouraged to create the file by hand if not already familiar with the process. While the command shows **vim** as an example other text editors such as **nano** work just as well.

```
student@master:~/app1$ vim simple.py
```



simple.py

```
1 #!/usr/bin/python
2 ## Import the necessary modules
3 import time
4 import socket
5
6 ## Use an ongoing while loop to generate output
7 while True :
8
9 ## Set the hostname and the current date
10 host = socket.gethostname()
11 date = time.strftime("%Y-%m-%d %H:%M:%S")
12
13 ## Convert the date output to a string
14 now = str(date)
15
16 ## Open the file named date in append mode
17 ## Append the output of hostname and time
```



```

18 f = open("date.out", "a" )
19 f.write(now + "\n")
20 f.write(host + "\n")
21 f.close()
22
23 ## Sleep for five seconds then continue the loop
24 time.sleep(5)

```

5. Make the file executable and test that it works. Use Ctrl-C to interrupt the while loop after 20 or 30 seconds. The output will be sent to a newly created file in your current directory called `date.out`.

```

student@master:~/app1$ chmod +x simple.py
student@master:~/app1$ ./simple.py

```

```

1 ^CTraceback (most recent call last):
2   File "./simple.py", line 42, in <module>
3     time.sleep(5)
4 KeyboardInterrupt

```

6. and timestamp stamps.

```

student@master:~/app1$ cat date.out

```

```

1 2018-03-22 15:51:38
2 master
3 2018-03-22 15:51:43
4 master
5 2018-03-22 15:51:48
6 master
7 <output_omitted>

```

7. Create a text file named `Dockerfile`.



Very Important

The name is important: it cannot have a suffix.

We will use three statements, `FROM` to declare which version of Python to use, `ADD` to include our script and `CMD` to indicate the action of the container. Should you be including more complex tasks you may need to install extra libraries, shown commented out as `RUN pip install` in the following example.

```

student@master:~/app1$ vim Dockerfile

```



Dockerfile

```

FROM python:2
ADD simple.py /
## RUN pip install pystrich
CMD [ "python", "./simple.py" ]

```

8. Build the container. The output below shows mid-build as necessary software is downloaded. You will need to use **sudo** in order to run this command. After the three step process completes the last line of output should indicate success. Note the dot (.) at the end of the command indicates the current directory.

```

student@master:~/app1$ sudo docker build -t simpleapp .

```

```

1 Sending build context to Docker daemon 3.072 kB
2 Step 1/3 : FROM python:2
3 2: Pulling from library/python
4 4176fe04cefe: Pull complete
5 851356ecf618: Pull complete
6 6115379c7b49: Pull complete
7 aaf7d781d601: Extracting [=====] 54.03 MB/135 MB
8 40cf661a3cc4: Download complete
9 c582f0b73e63: Download complete
10 6c1ea8f72a0d: Download complete
11 7051a41ae6b7: Download complete
12 <output_omitted>
13 Successfully built c4e0679b9c36

```

9. Verify you can see the new image among others downloaded during the build process, installed to support the cluster, or you may have already worked with. The newly created `simpleapp` image should be listed first.

```
student@master:~/app1$ sudo docker images
```

```

1 REPOSITORY          TAG          IMAGE ID      CREATED        SIZE
2 simpleapp           latest       ba54e4910397  6 seconds ago  902MB
3 k8s.gcr.io/kube-proxy v1.20.1     4e68534e24f6  7 days ago    117MB
4 k8s.gcr.io/kube-apiserver v1.20.1    a595af0107f9  7 days ago    173MB
5 <output_omitted>

```

10. Use **sudo docker** to run a container using the new image. While the script is running you won't see any output and the shell will be occupied running the image in the background. After 30 seconds use **ctrl-c** to interrupt. The local `date.out` file will not be updated with new times, instead that output will be a file of the container image.

```
student@master:~$ sudo docker run simpleapp
```

```

1 ^CTraceback (most recent call last):
2   File "./simple.py", line 24, in <module>
3     time.sleep(5)
4 KeyboardInterrupt

```

11. Locate the newly created `date.out` file. The following command should show two files of this name, the one created when we ran `simple.py` and another under `/var/lib/docker` when run via a Docker container.

```
student@master:~/app1$ sudo find / -name date.out
```

```

1 /home/student/app1/date.out
2 /var/lib/docker/overlay2/ee814320c900bd24fad0c5db4a258d3c2b78a19cde
3 629d7de7d27270d6a0c1f5/diff/date.out

```

12. View the contents of the `date.out` file created via Docker. Note the need for **sudo** as Docker created the file this time, and the owner is root. The long name is shown on several lines in the example, but would be a single line when typed or copied.

```

student@master:~/app1$ sudo tail \
/var/lib/docker/overlay2/ee814320c900bd24fad0c5db4a258d3c2b78a19cde
629d7de7d27270d6a0c1f5/diff/date.out

```

```

1 2018-03-22 16:13:46
2 53e1093e5d39
3 2018-03-22 16:13:51
4 53e1093e5d39
5 2018-03-22 16:13:56
6 53e1093e5d39

```

✍ Exercise 3.2: Configure A Local Docker Repo

While we could create an account and upload our application to hub.docker.com, thus sharing it with the world, we will instead create a local repository and make it available to the nodes of our cluster.

1. We'll need to complete a few steps with special permissions, for ease of use we'll become root using **sudo**.

```
student@master:~/app1$ cd
student@master:~$ sudo -i
```

2. Install the **docker-compose** software and utilities to work with the **nginx** server which will be deployed with the registry.

```
root@master:~# apt-get install -y docker-compose apache2-utils
```

```
1 <output_omitted>
```

3. Create a new directory for configuration information. We'll be placing the repository in the root filesystem. A better location may be chosen in a production environment.

```
root@master:~# mkdir -p /localdocker/data
root@master:~# cd /localdocker/
```

4. Create a Docker compose file. Inside is an entry for the **nginx** web server to handle outside traffic and a registry entry listening to loopback port 5000 for running a local Docker registry.

```
root@master:/localdocker# vim docker-compose.yaml
```

YAML

docker-compose.yaml

```
1 nginx:
2   image: "nginx:1.17"
3   ports:
4     - 443:443
5   links:
6     - registry:registry
7   volumes:
8     - /localdocker/nginx:/etc/nginx/conf.d
9 registry:
10  image: registry:2
11  ports:
12    - 127.0.0.1:5000:5000
13  environment:
14    REGISTRY_STORAGE_FILESYSTEM_ROOTDIRECTORY: /data
15  volumes:
16    - /localdocker/data:/data
```

5. Use the **docker-compose up** command to create the containers declared in the previous step YAML file. This will capture the terminal and run until you use **ctrl-c** to interrupt. There should be five `registry_1` entries with info messages about memory and which port is being listened to. Once we're sure the Docker file works we'll convert to a Kubernetes tool. **Let it run. You will use ctrl-c in a few steps.**

```
root@master:/localdocker# docker-compose up
```

```
1 Pulling nginx (nginx:1.17)...
2 1.17: Pulling from library/nginx
3 2a72cbf407d6: Pull complete
4 f37cbdc183b2: Pull complete
```



```

5 78b5ad0b466c: Pull complete
6 Digest: sha256:edad623fc7210111e8803b4359ba4854e101bcca1fe7f46bd1d35781f4034f0c
7 Status: Downloaded newer image for nginx:1.17
8 Creating localdocker_registry_1
9 Creating localdocker_nginx_1
10 Attaching to localdocker_registry_1, localdocker_nginx_1
11 registry_1 | time="2018-03-22T18:32:37Z" level=warning msg="No HTTP secret provided - generated ran
12 <output_omitted>

```

6. Test that you can access the repository. Open a second terminal to the master node. Use the **curl** command to test the repository. It should return `{}`, but does not have a carriage-return so will be on the same line as the following prompt. You should also see the GET request in the first, captured terminal, without error. Don't forget the trailing slash. You'll see a "Moved Permanently" message if the path does not match exactly.

```
student@master:~/localdocker$ curl http://127.0.0.1:5000/v2/
```

```
1 {}student@master:~/localdocker$
```

7. Now that we know **docker-compose** format is working, ingest the file into Kubernetes using **kompose**. Use **ctrl-c** to stop the previous **docker-compose** command.

```

1 ^CGracefully stopping... (press Ctrl+C again to force)
2 Stopping localdocker_nginx_1 ... done
3 Stopping localdocker_registry_1 ... done

```

8. Download the kompose binary and make it executable. The command can run on a single line. Note that the option following the dash is the letter as in output. Also that is a zero, not capital O (ohh) in the short URL. The short URL goes here: <https://github.com/kubernetes/kompose/releases/download/v1.1.0/kompose-linux-amd64>

```
root@master:/localdocker# curl -L https://bit.ly/2tN0bEa -o kompose
```

	% Total	% Received	% Xferd	Average Speed	Time	Time	Time	Current
				Dload Upload	Total	Spent	Left	Speed
3	100	609	0	609	0	0	1963	0 --:--:-- --:--:-- --:--:-- 1970
4	100	45.3M	100	45.3M	0	0	16.3M	0 0:00:02 0:00:02 --:--:-- 25.9M

```
root@master:/localdocker# chmod +x kompose
```

9. Move the binary to a directory in our \$PATH. Then return to your non-root user.

```

root@master:/localdocker# mv ./kompose /usr/local/bin/kompose
root@master:/localdocker# exit

```

10. Create two physical volumes in order to deploy a local registry for Kubernetes. 200Mi for each should be enough for each of the volumes. Use the **hostPath** storageclass for the volumes.

More details on how persistent volumes and persistent volume claims are covered in an upcoming chapter, Deployment Configuration.

```
student@master:~$ vim vol1.yaml
```

YAML

vol1.yaml

```

1 apiVersion: v1
2 kind: PersistentVolume
3 metadata:
4   labels:
5     type: local
6   name: task-pv-volume
7 spec:

```

YAML

```

8  accessModes:
9    - ReadWriteOnce
10 capacity:
11   storage: 200Mi
12 hostPath:
13   path: /tmp/data
14 persistentVolumeReclaimPolicy: Retain

```

```
student@master:~$ vim vol2.yaml
```

YAML**vol2.yaml**

```

1  apiVersion: v1
2  kind: PersistentVolume
3  metadata:
4    labels:
5      type: local
6    name: registryvm
7  spec:
8    accessModes:
9      - ReadWriteOnce
10   capacity:
11     storage: 200Mi
12   hostPath:
13     path: /tmp/nginx
14   persistentVolumeReclaimPolicy: Retain

```

11. Create both volumes.

```
student@master:~$ kubectl create -f vol1.yaml
```

```
1 persistentvolume/task-pv-volume created
```

```
student@master:~$ kubectl create -f vol2.yaml
```

```
1 persistentvolume/registryvm created
```

12. Verify both volumes have been created. They should show an Available status.

```
student@master:~$ kubectl get pv
```

1	NAME	CAPACITY	ACCESS MODES	RECLAIM POLICY	STATUS
2	CLAIM	STORAGECLASS	REASON	AGE	
3	registryvm	200Mi	RWO	Retain	Available
4					27s
5	task-pv-volume	200Mi	RWO	Retain	Available
6					32s

13. Go to the configuration file directory for the local Docker registry.

```
student@master:~$ cd /localdocker/
student@master:~/localdocker$ ls
```

```
1 data  docker-compose.yaml  nginx
```

14. Convert the Docker file into a single YAML file for use with Kubernetes. Not all objects convert exactly from Docker to **kompose**, you may get errors about the mount syntax for the new volumes. They can be safely ignored.

```
student@master:~/localdocker$ sudo kompose convert -f docker-compose.yaml -o localregistry.yaml
```

```
1 WARN Volume mount on the host "/localdocker/nginx/" isn't supported - ignoring path on the host
2 WARN Volume mount on the host "/localdocker/data" isn't supported - ignoring path on the host
```

15. Review the file. You'll find that multiple Kubernetes objects will have been created such as Services, Persistent Volume Claims and Deployments using environmental parameters and volumes to configure the container within.

```
student@master:~/localdocker$ less localregistry.yaml
```

```
1 apiVersion: v1
2 items:
3 - apiVersion: v1
4   kind: Service
5   metadata:
6     annotations:
7       kompose.cmd: kompose convert -f docker-compose.yaml -o localregistry.yaml
8       kompose.version: 1.1.0 (36652f6)
9     creationTimestamp: null
10    labels:
11  <output_omitted>
```

16. View the cluster resources prior to deploying the registry. Only the cluster service and two available persistent volumes should exist in the default namespace.

```
student@master:~/localdocker$ kubectl get pods,svc,pvc,pv,deploy
```

```
1 NAME          TYPE          CLUSTER-IP    EXTERNAL-IP    PORT(S)    AGE
2 kubernetes    ClusterIP     10.96.0.1     <none>         443/TCP    4h
3
4 NAME          CAPACITY    ACCESS MODES    RECLAIM POLICY
5 STATUS        CLAIM       STORAGECLASS    REASON    AGE
6 persistentvolume/registryvm    200Mi     RWO         Retain
7 Available
8 persistentvolume/task-pv-volume 200Mi     RWO         Retain
9 Available
10 Available          17s
```

17. To illustrate the fast changing nature of Kubernetes you will show that the API has changed for Deployments. With each new release of Kubernetes you may want to plan on a YAML review. First determine new object settings and configuration, then compare and contrast to your existing YAML files. Edit and test for the new configurations.

Another more common way to find YAML issues is to attempt to create an object using previous YAML in a new version of Kubernetes and track down errors. Not suggested, but what often happens instead of the following process.

To view the current cluster requirements use the `--dry-run` option for the **kubectl create** command to see what the API now uses. We can compare the current values to our existing (previous version) YAML files. This will help determine what to edit for the local registry in an upcoming step.



Very Important

If you are not familiar with Linux you may want to skip the following couple steps and copy and use the `edited-localregistry.yaml` file included in the course tarball. If so skip to the step which says to: Use kubectl to create the local docker registry.

```
student@master:~/localdocker$ kubectl create deployment drytry --image=nginx --dry-run=client -o yaml
```



drytry

```

1  apiVersion: apps/v1
2  kind: Deployment
3  metadata:
4    creationTimestamp: null
5    labels:
6      app: drytry
7      name: drytry
8  spec:
9    replicas: 1
10   selector:
11     matchLabels:
12       app: drytry
13   strategy: {}
14   template:
15     <output_omitted>

```

18. From the previous command output and comparing line by line to objects in the existing `localregistry.yaml` file output we can see that the `apiVersion` of the Deployment object has changed, and we need to add `selector`, add `matchLabels`, and a `label` line. The three lines to add will be part of the `replicaSet` information, right after the `replicas` line, with `selector` the same indentation as `replicas`.

Following is a **diff** output, a common way to compare two files to each other, before and after an edit. Use the **man** page to decode the output if you are not already familiar with the command.

```
student@master:~/localdocker$ sudo cp localregistry.yaml old-localregistry.yaml
```

```
student@master:~/localdocker$ sudo vim localregistry.yaml
<make edits>
```

```
student@master:~/localdocker$ diff localregistry.yaml old-localregistry.yaml
```

```

41c41
< - apiVersion: apps/v1
---
> - apiVersion: extensions/v1beta1
53,55d52
<   selector:
<     matchLabels:
<       io.kompose.service: nginx
93c90
< - apiVersion: apps/v1
---
> - apiVersion: extensions/v1beta1
105,107d101
<   selector:
<     matchLabels:
<       io.kompose.service: registry

```

19. Use **kubectl** to create the local docker registry.

```
student@master:~/localdocker$ kubectl create -f localregistry.yaml
```

```

1  service/nginx created
2  service/registry created
3  deployment.apps/nginx created
4  persistentvolumeclaim/nginx-claim0 created
5  deployment.apps/registry created
6  persistentvolumeclaim/registry-claim0 created

```

20. View the newly deployed resources. The persistent volumes should now show as Bound. Be aware that due to the manner that volumes are bound it is possible that the registry claim may not to be bound to the registry volume. Find the `service IP` for the registry. It should be sharing port 5000. In the example below the IP address is 10.110.186.162, yours may be different.

```
student@master:~/localdocker$ kubectl get pods,svc,pvc,pv,deploy
```

```

1  NAME                                READY   STATUS    RESTARTS   AGE
2  pod/nginx-6b58d9cdfd-95zxq          1/1     Running   0           1m
3  pod/registry-795c6c8b8f-b8z4k       1/1     Running   0           1m
4
5  NAME                                TYPE                CLUSTER-IP    EXTERNAL-IP   PORT(S)    AGE
6  service/kubernetes                  ClusterIP          10.96.0.1     <none>        443/TCP    1h
7  service/nginx                      ClusterIP          10.106.82.218 <none>        443/TCP    1m
8  service/registry                   ClusterIP          10.110.186.162 <none>        5000/TCP    1m
9
10 NAME                                STATUS    VOLUME
11 CAPACITY  ACCESS MODES  STORAGECLASS  AGE
12 persistentvolumeclaim/nginx-claim0    Bound     registryvm
13 200Mi      RWO
14 persistentvolumeclaim/registry-claim0 Bound     task-pv-volume
15 200Mi      RWO
16
17 NAME                                CAPACITY  ACCESS MODES  RECLAIM POLICY
18 STATUS  CLAIM  STORAGECLASS  REASON  AGE
19 persistentvolume/registryvm          200Mi     RWO
20 Bound
21 default/nginx-claim0
22 persistentvolume/task-pv-volume      200Mi     RWO
23 Bound
24 default/registry-claim0
25
26 NAME                                READY   UP-TO-DATE   AVAILABLE   AGE
27 deployment.apps/nginx                1/1     1             1           12s
28 deployment.apps/registry             1/1     1             1           12s

```

21. Verify you get the same `{}` response using the Kubernetes deployed registry as we did when using **docker-compose**. Note you must use the trailing slash after `v2`. Please also note that if the connection hangs it may be due to a firewall issue. If running your nodes using GCE ensure your instances are using VPC setup and all ports are allowed. If using AWS also make sure all ports are being allowed.

Edit the IP address to that of your registry service.

```
student@master:~/localdocker$ curl http://10.110.186.162:5000/v2/
```

```
1 {}student@master:~/localdocker$
```

22. Edit the Docker configuration file to allow insecure access to the registry. In a production environment steps should be taken to create and use TLS authentication instead. Use the IP and port of the registry you verified in the previous step.

```
student@master:~$ sudo vim /etc/docker/daemon.json
```

```
1 { "insecure-registries":["10.110.186.162:5000"] }
```

23. Restart docker on the local system. It can take up to a minute for the restart to take place. Ensure the service is active. It should report that the service recently became status as well.

```
student@master:~$ sudo systemctl restart docker.service
student@master:~$ sudo systemctl status docker.service | grep Active
```

```
1 Active: active (running) since Tue 2019-09-24 15:24:36 UTC; 40s ago
```

24. Download and tag a typical image from hub.docker.com. Tag the image using the IP and port of the registry. We will also use the latest tag.

```
student@master:~$ sudo docker pull ubuntu
```

```
1 Using default tag: latest
2 latest: Pulling from library/ubuntu
3 <output_omitted>
4 Digest: sha256:9ee3b83bcaa383e5e3b657f042f4034c92cdd50c03f73166c145c9ceaea9ba7c
5 Status: Downloaded newer image for ubuntu:latest
```

```
student@master:~$ sudo docker tag ubuntu:latest 10.110.186.162:5000/tagtest
```

25. Push the newly tagged image to your local registry. If you receive an error about an HTTP request to an HTTPS client check that you edited the `/etc/docker/daemon.json` file correctly and restarted the service.

```
student@master:~$ sudo docker push 10.110.186.162:5000/tagtest
```

```
1 The push refers to a repository [10.110.186.162:5000/tagtest]
2 db584c622b50: Pushed
3 52a7ea2bb533: Pushed
4 52f389ea437e: Pushed
5 88888b9b1b5b: Pushed
6 a94e0d5a7c40: Pushed
7 latest: digest: sha256:0847cc7fed1bfafac713b0aa4ddfb8b9199a99092ae1fc4e718cb28e8528f65f size: 1357
```

26. We will test to make sure we can also pull images from our local repository. Begin by removing the local cached images.

```
student@master:~$ sudo docker image remove ubuntu:latest
```

```
1 Untagged: ubuntu:latest
2 Untagged: ubuntu@sha256:e348fbb0e0a0e73ab0370de151e7800684445c509d46195aef73e090a49bd6
```

```
student@master:~$ sudo docker image remove 10.110.186.162:5000/tagtest
```

```
1 Untagged: 10.110.186.162:5000/tagtest:latest
2 <output_omitted>
```

27. Pull the image from the local registry. It should report the download of a newer image.

```
student@master:~$ sudo docker pull 10.110.186.162:5000/tagtest
```

```
1 Using default tag: latest
2 latest: Pulling from tagtest
3 Digest: sha256:0847cc7fed1bfafac713b0aa4ddfb8b9199a99092ae1fc4e718cb28e8528f65f
4 Status: Downloaded newer image for 10.110.186.162:5000/tagtest:latest
```

28. Use docker tag to assign the `simpleapp` image and then push it to the local registry. The image and dependent images should be pushed to the local repository.

```
student@master:~$ sudo docker tag simpleapp 10.110.186.162:5000/simpleapp
```

```
student@master:~$ sudo docker push 10.110.186.162:5000/simpleapp
```

```
1 The push refers to a repository [10.110.186.162:5000/simpleapp]
2 321938b97e7e: Pushed
3 ca82a2274c57: Pushed
4 de2fbb43bd2a: Pushed
5 4e32c2de91a6: Pushed
6 6e1b48dc2ccc: Pushed
7 ff57bdb79ac8: Pushed
8 6e5e20cbf4a7: Pushed
9 86985c679800: Pushed
10 8fad67424c4e: Pushed
11 latest: digest: sha256:67ea3e11570042e70cdcbad684a1e2986f59aaf53703e51725accdf5c70d475a size: 2218
```

29. Configure the worker (second) node to use the local registry running on the master server. Connect to the worker node. Edit the Docker `daemon.json` file with the same values as the master node and restart the service. Ensure it is active.

```
student@worker:~$ sudo vim /etc/docker/daemon.json
```

```
1 { "insecure-registries":["10.110.186.162:5000"] }
```

```
student@worker:~$ sudo systemctl restart docker.service
```

```
student@worker:~$ sudo systemctl status docker.service
```

30. From the worker node, pull the recently pushed image from the registry running on the master node.

```
student@worker:~$ sudo docker pull 10.110.186.162:5000/simpleapp
```

```
1 Using default tag: latest
2 latest: Pulling from simpleapp
3 f65523718fc5: Pull complete
4 1d2dd88bf649: Pull complete
5 c09558828658: Pull complete
6 0e1d7c9e6c06: Pull complete
7 c6b6fe164861: Pull complete
8 45097146116f: Pull complete
9 f21f8abae4c4: Pull complete
10 1c39556edcd0: Pull complete
11 85c79f0780fa: Pull complete
12 Digest: sha256:67ea3e11570042e70cdcbad684a1e2986f59aaf53703e51725acddf5c70d475a
13 Status: Downloaded newer image for 10.110.186.162:5000/simpleapp:latest
```

31. Return to the master node and deploy the `simpleapp` in Kubernetes with several replicas. We will name the deployment `try1`. Scale to have six replicas. Multiple replicas the scheduler should run some containers on each node.

```
student@master:~$ kubectl create deployment try1 --image=10.110.186.162:5000/simpleapp
```

```
1 deployment.apps/try1 created
```

```
student@master:~$ kubectl scale deployment try1 --replicas=6
```

```
1 deployment.apps/try1 scaled
```

32. View the running pods. You should see six replicas of `simpleapp` as well as two running the locally hosted image repository.

```
student@master:~$ kubectl get pods
```

	NAME	READY	STATUS	RESTARTS	AGE
1	nginx-6b58d9cdfd-j6jm6	1/1	Running	1	13m
2	registry-795c6c8b8f-5jnpn	1/1	Running	1	13m
3	try1-857bdcd888-6klrr	1/1	Running	0	25s
4	try1-857bdcd888-9pwnp	1/1	Running	0	25s
5	try1-857bdcd888-9xkth	1/1	Running	0	25s
6	try1-857bdcd888-tw58z	1/1	Running	0	25s
7	try1-857bdcd888-xj9lk	1/1	Running	0	25s
8	try1-857bdcd888-znprm8	1/1	Running	0	25s

33. On the second node use `sudo docker ps` to verify containers of `simpleapp` are running. The scheduler will usually balance pod count across nodes. As the master already has several pods running the new pods may be on the worker.

```
student@worker:~$ sudo docker ps | grep simple
```

```

1 3ae4668d71d8 \
2 10.110.186.162:5000/simpleapp@sha256:67ea3e11570042e70cdc6ad684a1e2986f59aaf53703e51725accdf5c70d475a \
3 "python ./simple.py" 48 seconds ago Up 48 seconds \
4 k8s_try1_try1-857bdc888-9xkth_default_2e94b97e-322a-11e8-af56-42010a800004_0
5 ef6448764625 \
6 10.110.186.162:5000/simpleapp@sha256:67ea3e11570042e70cdc6ad684a1e2986f59aaf53703e51725accdf5c70d475a \
7 "python ./simple.py" 48 seconds ago Up 48 seconds \
8 k8s_try1_try1-857bdc888-znpm8_default_2e99f356-322a-11e8-af56-42010a800004_0

```

34. Return to the master node. Save the try1 deployment as YAML.

```

student@master:~/app1$ cd ~/app1/
student@master:~/app1$ kubectl get deployment try1 -o yaml > simpleapp.yaml

```

35. Delete and recreate the try1 deployment using the YAML file. Verify the deployment is running with the expected six replicas.

```

student@master:~$ kubectl delete deployment try1

```

```

1 deployment.apps "try1" deleted

```

```

student@master:~/app1$ kubectl create -f simpleapp.yaml

```

```

1 deployment.apps/try1 created

```

```

student@master:~/app1$ kubectl get deployment

```

```

1 NAME      READY   UP-TO-DATE   AVAILABLE   AGE
2 nginx     1/1     1             1           15m
3 registry  1/1     1             1           15m
4 try1      6/6     6             6           5s

```

Exercise 3.3: Configure Probes

When large datasets need to be loaded or a complex application launched prior to client access, a `readinessProbe` can be used. The pod will not become available to the cluster until a test is met and returns a successful exit code. Both `readinessProbes` and `livenessProbes` use the same syntax and are identical other than the name. Where the `readinessProbe` is checked prior to being ready, then not again, the `livenessProbe` continues to be checked.

There are three types of liveness probes: a command returns a zero exit value, meaning success, an HTTP request returns a response code in the 200 to 399 range, and the third probe uses a TCP socket. In this example we'll use a command, `cat`, which will return a zero exit code when the file `/tmp/healthy` has been created and can be accessed.

1. Edit the YAML deployment file and add the stanza for a `readinessprobe`. Remember that when working with YAML whitespace matters. Indentation is used to parse where information should be associated within the stanza and the entire file. Do not use tabs. If you get an error about validating data, check the indentation. It can also be helpful to paste the file to this website to see how indentation affects the JSON value, which is actually what Kubernetes ingests: <https://www.json2yaml.com/> An edited file is also included in the tarball, but requires the image name to be edited to match your registry IP address.

```

student@master:~/app1$ vim simpleapp.yaml

```




simpleapp.yaml

```

1  ....
2      spec:
3          containers:
4              - image: 10.111.235.60:5000/simpleapp:latest
5                imagePullPolicy: Always
6                name: simpleapp
7                readinessProbe:          #<--This line and next five
8                    periodSeconds: 5
9                    exec:
10                       command:
11                         - cat
12                         - /tmp/healthy
13                 resources: {}
14  ....

```

2. Delete and recreate the try1 deployment.

```
student@master:~/app1$ kubectl delete deployment try1
```

```
1 deployment.apps "try1" deleted
```

```
student@master:~/app1$ kubectl create -f simpleapp.yaml
```

```
1 deployment.apps/try1 created
```

3. The new try1 deployment should reference six pods, but show zero available. They are all missing the `/tmp/healthy` file.

```
student@master:~/app1$ kubectl get deployment
```

```

1  NAME          READY   UP-TO-DATE   AVAILABLE   AGE
2  nginx         1/1     1            1           19m
3  registry     1/1     1            1           19m
4  try1          0/6     6            0           15s

```

4. Take a closer look at the pods. Choose one of the try1 pods as a test to create the health check file.

```
student@master:~/app1$ kubectl get pods
```

```

1  NAME                                READY   STATUS    RESTARTS   AGE
2  nginx-6b58d9cdfd-g7lnk              1/1     Running   1           40m
3  registry-795c6c8b8f-7vwdn          1/1     Running   1           40m
4  try1-9869bdb88-2wfnr                0/1     Running   0           26s
5  try1-9869bdb88-6bknl                0/1     Running   0           26s
6  try1-9869bdb88-786v8                0/1     Running   0           26s
7  try1-9869bdb88-gmvs4                0/1     Running   0           26s
8  try1-9869bdb88-lfv1x                0/1     Running   0           26s
9  try1-9869bdb88-rtchc                0/1     Running   0           26s

```

5. Run the bash shell interactively and touch the `/tmp/healthy` file.

```
student@master:~/app1$ kubectl exec -it try1-9869bdb88-rtchc -- /bin/bash
```

```
root@try1-9869bdb88-rtchc:/# touch /tmp/healthy
```

```
root@try1-9869bdb88-rtchc:/# exit
```

```
1 exit
```

6. Wait at least five seconds, then check the pods again. Once the probe runs again the container should show available quickly. The pod with the existing `/tmp/healthy` file should be running and show 1/1 in a READY state. The rest will continue to show 0/1.

```
student@master:~/app1$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-6b58d9cdfd-g7lnk	1/1	Running	1	44m
registry-795c6c8b8f-7vwdn	1/1	Running	1	44m
try1-9869bdb88-2wfnr	0/1	Running	0	4m
try1-9869bdb88-6bknl	0/1	Running	0	4m
try1-9869bdb88-786v8	0/1	Running	0	4m
try1-9869bdb88-gmvs4	0/1	Running	0	4m
try1-9869bdb88-lfv1x	0/1	Running	0	4m
try1-9869bdb88-rtchc	1/1	Running	0	4m

7. Touch the file in the remaining pods. Consider using a **for** loop, as an easy method to update each pod. Note the `>` shown in the output represents the secondary prompt, you would not type in that character

```
student@master:~$ for name in try1-9869bdb88-2wfnr try1-9869bdb88-6bknl \
> try1-9869bdb88-786v8 try1-9869bdb88-gmvs4 try1-9869bdb88-lfv1x
> do
> kubectl exec $name -- touch /tmp/healthy
> done
```

8. It may take a short while for the probes to check for the file and the health checks to succeed.

```
student@master:~/app1$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-6b58d9cdfd-g7lnk	1/1	Running	1	1h
registry-795c6c8b8f-7vwdn	1/1	Running	1	1h
try1-9869bdb88-2wfnr	1/1	Running	0	22m
try1-9869bdb88-6bknl	1/1	Running	0	22m
try1-9869bdb88-786v8	1/1	Running	0	22m
try1-9869bdb88-gmvs4	1/1	Running	0	22m
try1-9869bdb88-lfv1x	1/1	Running	0	22m
try1-9869bdb88-rtchc	1/1	Running	0	22m

9. Now that we know when a pod is healthy, we may want to keep track that it stays healthy, using a `livenessProbe`. You could use one probe to determine when a pod becomes available and a second probe, to a different location, to ensure ongoing health.

Edit the deployment again. Add in a `livenessProbe` section as seen below. This time we will add a `Sidecar` container to the pod running a simple application which will respond to port 8080. Note that the dash (-) in front of the name. Also `goproxy` is indented the same number of spaces as the - in front of the `image:` line for `simpleapp` earlier in the file. In this example that would be seven spaces

```
student@master:~/app1$ vim simpleapp.yaml
```

YAML

simpleapp.yaml

```
1 ....
2     terminationMessagePath: /dev/termination-log
3     terminationMessagePolicy: File
4     - name: goproxy                #<-- Indented 7 spaces, add lines from here...
```



```

5     image: k8s.gcr.io/goproxy:0.1
6     ports:
7       - containerPort: 8080
8     readinessProbe:
9       tcpSocket:
10        port: 8080
11        initialDelaySeconds: 5
12        periodSeconds: 10
13     livenessProbe:           #<-- This line is 9 spaces indented, fyi
14       tcpSocket:
15        port: 8080
16        initialDelaySeconds: 15
17        periodSeconds: 20     #<-- ....to here
18     dnsPolicy: ClusterFirst
19     restartPolicy: Always
20 ....

```

10. Delete and recreate the deployment.

```
student@master:~$ kubectl delete deployment try1
```

```
1 deployment.apps "try1" deleted
```

```
student@master:~$ kubectl create -f simpleapp.yaml
```

```
1 deployment.apps/try1 created
```

11. View the newly created pods. You'll note that there are two containers per pod, and only one is running. The new simpleapp containers will not have the `/tmp/healthy` file, so they will not become available until we touch the `/tmp/healthy` file again. We could include a command which creates the file into the container arguments. The output below shows it can take a bit for the old pods to terminate.

```
student@master:~$ kubectl get pods
```

	NAME	READY	STATUS	RESTARTS	AGE
2	nginx-6b58d9cdfd-g7lnk	1/1	Running	1	13h
3	registry-795c6c8b8f-7vwdn	1/1	Running	1	13h
4	try1-76cc5ffcc6-4rjvh	1/2	Running	0	3s
5	try1-76cc5ffcc6-bk5f5	1/2	Running	0	3s
6	try1-76cc5ffcc6-d8n5q	0/2	ContainerCreating	0	3s
7	try1-76cc5ffcc6-mm6tw	1/2	Running	0	3s
8	try1-76cc5ffcc6-r9q5n	1/2	Running	0	3s
9	try1-76cc5ffcc6-tx4dz	1/2	Running	0	3s
10	try1-9869bdb88-2wfnr	1/1	Terminating	0	12h
11	try1-9869bdb88-6bknl	1/1	Terminating	0	12h
12	try1-9869bdb88-786v8	1/1	Terminating	0	12h
13	try1-9869bdb88-gmvs4	1/1	Terminating	0	12h
14	try1-9869bdb88-lfvlx	1/1	Terminating	0	12h
15	try1-9869bdb88-rtchc	1/1	Terminating	0	12h

12. Create the health check file for the readinessProbe. You can use a **for** loop again for each action, this setup will leverage labels so you don't have to look up the pod names. As there are now two containers in the pod, you should include the container name for which one will execute the command. If no name is given, it will default to the first container. Depending on how you edited the YAML file try1 should be the first pod and goproxy the second. To ensure the correct container is updated, add **-c simpleapp** to the **kubectl** command. Your pod names will be different. Use the names of the newly started containers from the **kubectl get pods** command output. Note the `>` character represents the secondary prompt, you would not type in that character.

```
student@master:~$ for name in $(kubectl get pod -l app=try1 -o name)
> do
> kubectl exec $name -c simpleapp -- touch /tmp/healthy
> done
```

13. In the next minute or so the Sidecar container in each pod, which was not running, will change status to Running. Each should show 2/2 containers running.

```
student@master:~$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-6b58d9cfd-g7lnk	1/1	Running	1	13h
registry-795c6c8b8f-7vwdn	1/1	Running	1	13h
try1-76cc5ffcc6-4rjvh	2/2	Running	0	3s
try1-76cc5ffcc6-bk5f5	2/2	Running	0	3s
try1-76cc5ffcc6-d8n5q	2/2	Running	0	3s
try1-76cc5ffcc6-mm6tw	2/2	Running	0	3s
try1-76cc5ffcc6-r9q5n	2/2	Running	0	3s
try1-76cc5ffcc6-tx4dz	2/2	Running	0	3s

14. View the events for a particular pod. Even though both containers are currently running and the pod is in good shape, note the events section shows the issue.

```
student@master:~/app1$ kubectl describe pod try1-76cc5ffcc6-tx4dz | tail
```

```
1 Normal SuccessfulMountVolume 9m kubelet, master-lab-x6dj
2 MountVolume.SetUp succeeded for volume "default-token-jf69w"
3 Normal Pulling 9m kubelet, master-lab-x6dj
4 pulling image "10.108.143.90:5000/simpleapp"
5 Normal Pulled 9m kubelet, master-lab-x6dj
6 Successfully pulled image "10.108.143.90:5000/simpleapp"
7 Normal Created 9m kubelet, master-lab-x6dj
8 Created container
9 Normal Started 9m kubelet, master-lab-x6dj
10 Started container
11 Normal Pulling 9m kubelet, master-lab-x6dj
12 pulling image "k8s.gcr.io/goproxy:0.1"
13 Normal Pulled 9m kubelet, master-lab-x6dj
14 Successfully pulled image "k8s.gcr.io/goproxy:0.1"
15 Normal Created 9m kubelet, master-lab-x6dj
16 Created container
17 Normal Started 9m kubelet, master-lab-x6dj
18 Started container
19 Warning Unhealthy 4m (x60 over 9m) kubelet, master-lab-x6dj
20 Readiness probe failed: cat: /tmp/healthy: No such file or directory
```

15. If you look for the status of each container in the pod, they should show that both are Running and ready showing True.

```
student@master:~/app1$ kubectl describe pod try1-76cc5ffcc6-tx4dz | grep -E 'State|Ready'
```

```
1 State: Running
2 Ready: True
3 State: Running
4 Ready: True
5 Ready: True
6 ContainersReady: True
```

Exercise 3.4: Domain Review

**Very Important**

The source pages and content in this review could change at any time. **IT IS YOUR RESPONSIBILITY TO CHECK THE CURRENT INFORMATION.**

Revisit the CKAD domain list on [Curriculum Overview](#) and locate some of the topics we have covered in this chapter.

- Understand Multi-Container Pod design patterns (e.g. ambassador, adapter, sidecar)
- Understand LivenessProbes and ReadinessProbes
- Understand container logging

Figure 3.1: **Observability Domain**

Focus on ensuring you have all the necessary files and processes understood first. Repeat the review until you are sure you have bookmarks of YAML samples and can complete each step quickly.

1. Using the three URL locations allowed by the exam, find and bookmark working YAML examples for LivenessProbes, ReadinessProbes, and multi-container pods.
2. Deploy a new nginx webserver. Add a LivenessProbe and a ReadinessProbe on port 80. Test that both probes and the webserver work.
3. Use the `build-review1.yaml` file to create a non-working deployment. Fix the deployment such that both containers are running and in a READY state. The web server listens on port 80, and the proxy listens on port 8080.
4. View the default page of the web server. When successful verify the GET activity logs in the container log. The message should look something like the following. Your time and IP may be different.

```
192.168.124.0 - - [30/Jan/2020:03:30:31 +0000] "GET / HTTP/1.1" 200 612 "-" "curl/7.58.0" "-"
```

5. Remove any resources created in this review.

Chapter 4

APIs and Access



4.1	API Access	64
4.2	Annotations	68
4.3	Working with A Simple Pod	69
4.4	kubectl and API	70
4.5	Swagger and OpenAPI	77
4.6	Labs	79

4.1 API Access

API Access

- API driven architecture
- API groups
- RESTful style
- Standard HTTP verbs
- Deprecation process now being honored

Kubernetes has a powerful **REST**-based API. The entire architecture is API driven. Knowing where to find resource endpoints and understanding how the API changes between versions can be important to ongoing administrative tasks, as there is much ongoing change and growth. Starting with v1.16 deprecated objects are no longer honored by the API server.

As we learned in the Architecture chapter, the main agent for communication between cluster agents and from outside the cluster is the `kube-apiserver`. A **curl** query to the agent will expose the current API groups. Groups may have multiple versions which evolve independently of other groups, and follow a domain-name format with several names reserved, such as single-word domains, the empty group and any name ending in `.k8s.io`.

RESTful

- Responds to typical HTTP verbs (GET, POST, DELETE ...)
- Allows easy interaction with other ecosystems
- Scripting and interaction with deployment tools
- User impersonation headers

kubectl makes API calls on your behalf. You can also make calls externally, using **curl** or other program. With the appropriate certs and keys, you can make requests or pass **json** files to make configuration changes.

```
$ curl --cert userbob.pem --key userBob-key.pem \  
  --cacert /path/to/ca.pem \  
  https://k8sServer:6443/api/v1/pods
```

The ability to impersonate other users or groups, subject to RBAC configuration, allows a manual override authentication. This can be helpful for debugging authorization policies of other users.

Checking Access

- Several ways to authenticate
- `auth can-i` subcommand to query authorization
- Accepts `can-i` and `reconcile` arguments
- More in Security chapter to follow.

While there is more detail on security in a later chapter, it is helpful to check the current authorizations, both as an admin, as well as another user. The following shows what user bob could do in the default namespace and the developer namespace:

```
$ kubectl auth can-i create deployments
```

```
1 yes
```

```
$ kubectl auth can-i create deployments --as bob
```

```
1 no
```

```
$ kubectl auth can-i create deployments --as bob --namespace developer
```

```
1 yes
```

There are currently three APIs which can be applied to set who and what can be queried.

- `SelfSubjectAccessReview`
Access review for any user, helpful for delegating to others.
- `LocalSubjectAccessReview`
Review is restricted to a specific namespace.
- `SelfSubjectRulesReview`
A review which shows allowed actions for a user within a particular namespace.

The use of `reconcile` allows a check of authorization necessary to create an object from a file. No output indicates the creation would be allowed.

Optimistic Concurrency

- Currently leverage JSON
- `resourceVersion`
- Clients must handle 409 CONFLICT Errors

The default serialization for API calls must be JSON. There is an effort to use **Google's** protobuf serialization, but this remains experimental. While we may work with files in YAML format, they are converted to and from JSON.

Kubernetes uses the `resourceVersion` value to determine API updates and implement optimistic concurrency. In other words, an object is not locked from the time it has been read until the object is written.

Instead, upon an updated call to an object, the `resourceVersion` is checked and a 409 CONFLICT is returned should the number have changed. The `resourceVersion` is currently backed via the `modifiedIndex` parameter in the **etcd** database, and is unique to the namespace, kind and server. Operations which do not change an object such as WATCH or GET do not update this value.

4.2 Annotations

Using Annotations

- Distinct from Labels
- Non-identifying metadata
- Key/value maps
- Metadata otherwise held in exterior databases
- Useful for third-party automation

Labels are used to work with objects or collections of objects; annotations are not.

Instead annotations allow for metadata to be included with an object that may be helpful outside of Kubernetes object interaction. Similar to labels, they are key to value maps. They also are able to hold more information, and more human readable information than labels.

Having this kind of metadata can be used to track information such as a timestamp, pointers to related objects from other ecosystems, or even an email from the developer responsible for that object's creation.

The annotation data could otherwise be held in an exterior database, but that would limit the flexibility of the data. The more this metadata is included, the easier to integrate management and deployment tools or shared client libraries.

For example, to annotate only Pods within a namespace, then overwrite the annotation and finally delete it:

```
$ kubectl annotate pods --all description='Production Pods' -n prod
```

```
$ kubectl annotate --overwrite pod webpod description="Old Production Pods" -n prod
```

```
$ kubectl -n prod annotate pod webpod description-
```

4.3 Working with A Simple Pod

Simple Pod

- Lowest compute unit of K8s
- Typically multiple containers grouped together
- Created from PodSpec
- Few required, many optional
 - `apiVersion`
Must match existing API group
 - `kind`
The type of object to create
 - `metadata`
At least a name
 - `spec`
What to create and parameters

As discussed earlier, a Pod is the lowest compute unit and individual object we can work with in Kubernetes. It can be a single container, but often it will consist of a primary application container and one or more supporting containers.

Below is an example of a simple pod manifest in YAML format. You can see the `apiVersion`, the `kind`, the `metadata`, and its `spec`, which define the container that actually runs in this pod:

YAML

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: firstpod
5 spec:
6   containers:
7   - image: nginx
8     name: stan
```

You can use the `kubectl create` command to create this pod in Kubernetes. Once it is created, you can check its status with `kubectl get pods`. Output is omitted to save space:

```
$ kubectl create -f simple.yaml
$ kubectl get pods
$ kubectl get pod firstpod -o yaml
$ kubectl get pod firstpod -o json
```

4.4 kubectl and API

Manage API Resources with kubectl

- API exposed via RESTful interface
- Use **curl** to access and test
- Use verbose mode
- Leverages HTTP verbs

Kubernetes exposes resources via RESTful API calls, which allows all resources to be managed via HTTP, JSON or even XML. the typical protocol being HTTP. The state of the resources can be changed using standard HTTP verbs (e.g. GET, POST, PATCH, DELETE, etc.).

kubectl has a verbose mode argument which shows details from where the command gets and updates information. Other output includes **curl** commands you could use to obtain the same result. While the verbosity accepts levels from zero to any number, there is currently no verbosity value greater than ten. You can check this out for `kubectl get`. The output below has been formatted for clarity:

```
$ kubectl --v=10 get pods firstpod
```

```
1 ....
2 I1215 17:46:47.860958 29909 round_tripper.go:417]
3   curl -k -v -XGET -H "Accept: application/json"
4   -H "User-Agent: kubectl/v1.8.5 (linux/amd64) kubernetes/cce11c6"
5   https://10.128.0.3:6443/api/v1/namespaces/default/pods/firstpod
6 ....
```

If you delete this pod, you will see that the HTTP method changes from XGET to XDELETE.

```
$ kubectl --v=10 delete pods firstpod
```

```
1 ....
2 I1215 17:49:32.166115 30452 round_tripper.go:417]
3   curl -k -v -XDELETE -H "Accept: application/json, */*"
4   -H "User-Agent: kubectl/v1.8.5 (linux/amd64) kubernetes/cce11c6"
5   https://10.128.0.3:6443/api/v1/namespaces/default/pods/firstpod
6 ....
```

Access From Outside The Cluster

- Can use **curl** from outside cluster
- Must use SSL/TLS for secure access
- Information found in `~/.kube/config`
- View server information via **kubectl config view**

The primary tool used from the command line will be **kubectl**, which calls **curl** on your behalf. You can also use the **curl** command from outside the cluster to view or make changes.

The basic server information, with redacted TLS certificate information can be found in the output of

```
$ kubectl config view
```

If you view verbose output from a previous page, you will note that the first line references a config file where this information is pulled from, `~/.kube/config`.

```
1 I1215 17:35:46.725407 27695 loader.go:357]
2 Config loaded from file /home/student/.kube/config
```

Without the certificate authority, key and certificate from this file, only insecure **curl** commands can be used, which will not expose much due to security settings. We will use curl to access our cluster using TLS in an upcoming lab.

~/.kube/config

```

apiVersion: v1
clusters:
- cluster:
    certificate-authority-data: LS0tLS1CRUdF.....
    server: https://10.128.0.3:6443
  name: kubernetes
contexts:
- context:
    cluster: kubernetes
    user: kubernetes-admin
  name: kubernetes-admin@kubernetes
current-context: kubernetes-admin@kubernetes
kind: Config
preferences: {}
users:
- name: kubernetes-admin
  user:
    client-certificate-data: LS0tLS1CRUdJTib.....
    client-key-data: LS0tLS1CRUdJTi.....

```

The output above shows 19 lines of output with each of the keys being heavily truncated. While the keys may look similar close examination shows them to be distinct.

- **apiVersion**
As with other objects, this instructs the kube-apiserver where to assign the data.
- **clusters**
This contains the name of the cluster as well as where to send the API calls. The `certificate-authority-data` is passed to authenticate the **curl** request.
- **contexts**
A setting which allows easy access to multiple clusters, possibly as various users, from one config file. It can be used to set namespace, user, and cluster.
- **current-context**
Shows which cluster and user **kubectl** would use. These settings can also be passed on a per-command basis.
- **kind**
Every object within Kubernetes must have this setting, in this case a declaration of object type `Config`.
- **preferences**
Currently not used optional settings for the **kubectl** command, such as colorizing output.
- **users**
A nickname associated with client credentials which can be client key and certificate, username and password, and a token. Token and username/password are mutually exclusive. These can be configured via the **kubectl config set-credentials** command.

Namespaces

- **Linux** kernel feature
 - Segregates system resources
 - Core functionality of containers
- API Object
 - Four namespaces to begin with:
 - * `default`
 - * `kube-node-lease`
 - * `kube-public`
 - * `kube-system`
 - - **-all-namespaces**

The term **namespace** is used to both reference the Kernel feature as well as the segregation of API objects by Kubernetes. Both are means to keep resources distinct.

Every API call includes a namespace, using `default` if not otherwise declared: <https://10.128.0.3:6443/api/v1/namespaces/default/pods>.

Namespaces are intended to isolate multiple groups and the resources they have access to work with via quotas. Eventually, access control policies will work on namespace boundaries as well. One could use `Labels` to group resources for administrative reasons.

There are four namespaces when a cluster is first created.

- `default`
This is where all resources are assumed unless set otherwise.
- `kube-node-lease` The namespace where worker node lease information is kept.
- `kube-public`
A namespace readable by all, even those not authenticated. General information is often included in this namespaces.
- `kube-system`
Contains infrastructure pods.

Should you want to see all resources on a system you must pass the `--all-namespaces` option to the **kubectl** command.

Working with Namespaces

```
$ kubectl get ns
$ kubectl create ns linuxcon
$ kubectl describe ns linuxcon
$ kubectl get ns/linuxcon -o yaml
$ kubectl delete ns/linuxcon
```

The above commands show how to view, create and delete namespaces. Note that the **describe** subcommand shows several settings such as Labels, Annotations, resource quotas, and resource limits which we will discuss later in the course.

Once a namespace has been created you can reference via YAML when creating resource:

```
$ cat redis.yaml
```

YAML

redis.yaml

```
1 apiVersion: V1
2 kind: Pod
3 metadata:
4   name: redis
5   namespace: linuxcon
6 ...
```

API Resources with kubectl

- All available via **kubectl**
- `kubectl [command] [type] [Name] [flag]`
- `kubectl help` for more information
- Abbreviated names

All API resources exposed are available via **kubectl**. Expect the list to change.

- | | | |
|---|----------------------------------|------------------------------|
| • all | • endpoints (ep) | • pods (po) |
| • certificatesigningrequests (csr) | • events (ev) | • podsecuritypolicies (psp) |
| • clusterrolebindings | • horizontalpodautoscalers (hpa) | • podtemplates |
| • clusterroles | • ingresses (ing) | • replicaset (rs) |
| • clusters (valid only for federation apiservers) | • jobs | • replicationcontrollers rc) |
| • componentstatuses (cs) | • limitranges (limits) | • resourcequotas (quota) |
| • configmaps (cm) | • namespaces (ns) | • rolebindings |
| • controllerrevisions | • networkpolicies (netpol) | • roles |
| • cronjobs | • nodes (no) | • secrets |
| • customresourcedefinition (crd) | • persistentvolumeclaims (pvc) | • serviceaccounts (sa) |
| • daemonsets (ds) | • persistentvolumes (pv) | • services (svc) |
| • deployments (deploy) | • poddisruptionbudgets (pdb) | • statefulsets |
| | • podpreset | • storageclasses |

Additional Resource Methods

- Various Endpoints
- CLI `--help`
- Online documentation

In addition to basic resource management via REST, the API also provides some extremely useful endpoints for certain resources.

For example, you can access the logs of a container, exec into it, and watch changes to it with the following endpoints:

```
$ curl --cert /tmp/client.pem --key /tmp/client-key.pem \  
  --cacert /tmp/ca.pem -v -XGET \  
  https://10.128.0.3:6443/api/v1/namespaces/default/pods/firstpod/log
```

This would be the same as the following. If the container does not have any standard out, there would be no logs.

```
$ kubectl logs firstpod
```

Other calls you could make, following the various API groups on your cluster:

```
GET /api/v1/namespaces/{namespace}/pods/{name}/exec  
GET /api/v1/namespaces/{namespace}/pods/{name}/log  
GET /api/v1/watch/namespaces/{namespace}/pods/{name}
```

4.5 Swagger and OpenAPI

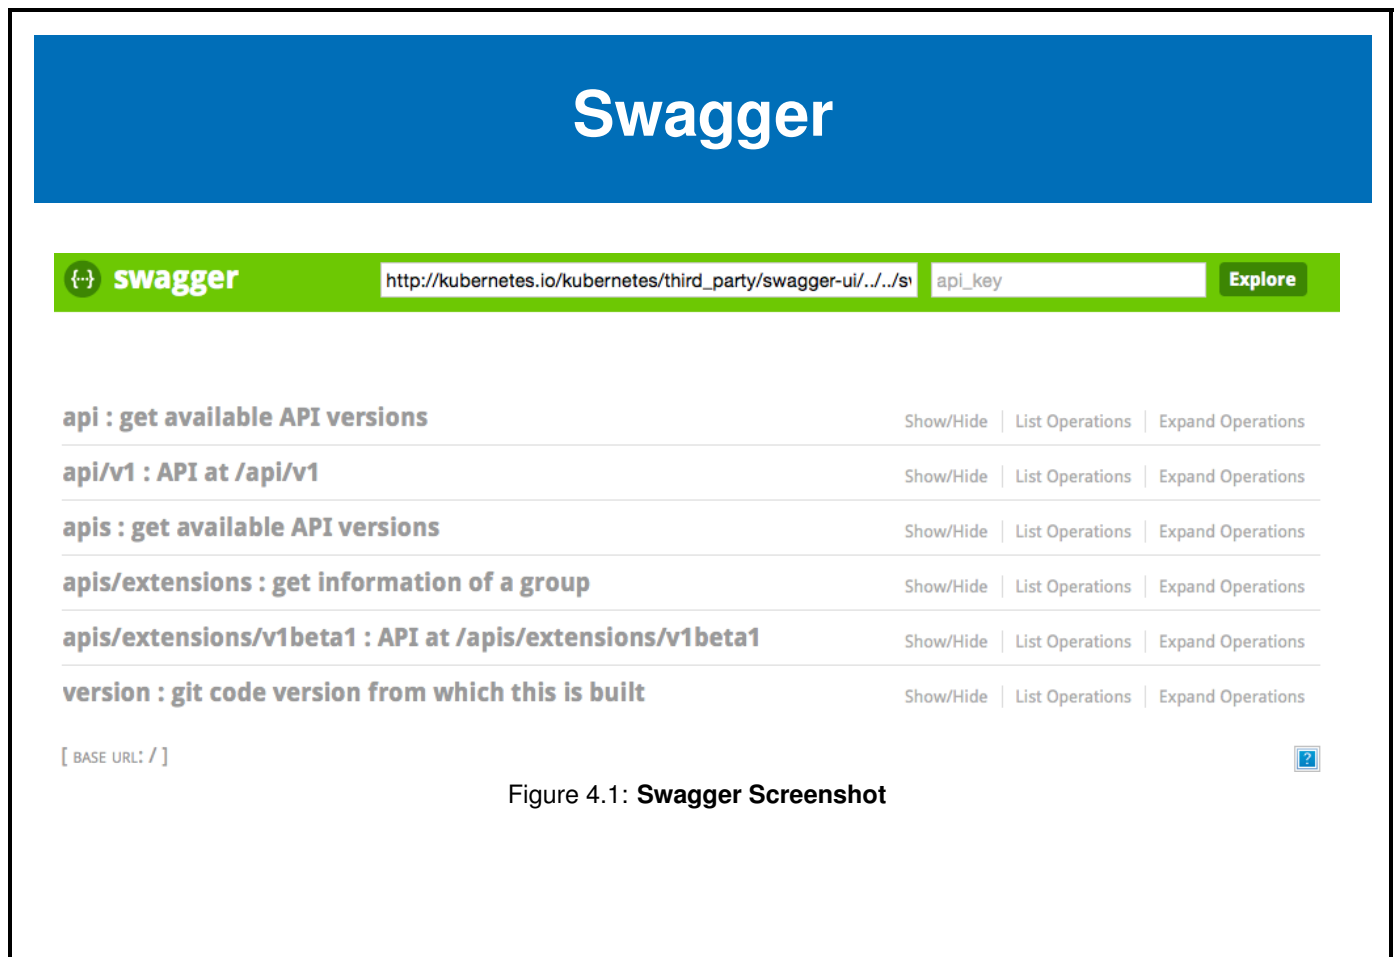


Figure 4.1: Swagger Screenshot

The entire Kubernetes API uses a **Swagger** specification. This is evolving towards the **OpenAPI** initiative. It is extremely useful, as it allows, for example, to auto-generate client code. All the stable resources definitions are available on the documentation site.

You can browse some of the API groups via a Swagger UI at <https://swagger.io/specification/>.

API Maturity

- Versioning of API levels for easier growth
- Not directly tied to software versioning
- Versions imply level of support
 - Alpha
 - Beta
 - Stable

The use of API groups and different versions allows for development to advance without changes to an existing group of APIs. This allows for easier growth and separation of work among separate teams. While there is an attempt to maintain some consistency between API and software versions they are only indirectly linked.

The use of JSON and **Google's** Protobuf serialization scheme will follow the same release guidelines.

An `Alpha` level release, noted with `alpha` in the name, may be buggy and is disabled by default. Features could change or disappear at any time. Only use these features on a test cluster which is often rebuilt.

The `Beta` level, found with `beta` in the name, has more well tested code and is enabled by default. It also ensures that as changes move forward they will be tested for backwards compatibility between versions. It has not been adopted and tested enough to be called stable. Expect some bugs and issues.

Use of the `Stable` version, denoted by only an integer which may be preceded by the letter `v`, is for stable APIs.

4.6 Labs

✎ Exercise 4.1: Configuring TLS Access

Overview

Using the Kubernetes API, **kubectl** makes API calls for you. With the appropriate TLS keys you could run **curl** as well use a **golang** client. Calls to the kube-apiserver get or set a PodSpec, or desired state. If the request represents a new state the **Kubernetes Control Plane** will update the cluster until the current state matches the specified state. Some end states may require multiple requests. For example, to delete a ReplicaSet, you would first set the number of replicas to zero, then delete the ReplicaSet.

An API request must pass information as JSON. **kubectl** converts **.yaml** to JSON when making an API request on your behalf. The API request has many settings, but must include **apiVersion**, **kind** and **metadata**, and **spec** settings to declare what kind of container to deploy. The **spec** fields depend on the object being created.

We will begin by configuring remote access to the kube-apiserver then explore more of the API.

1. Begin by reviewing the **kubectl** configuration file. We will use the three certificates and the API server address.

```
student@master:~$ less $HOME/.kube/config
```

```
1 <output_omitted>
```

2. We will create a variables using certificate information. You may want to double-check each parameter as you set it. Begin with setting the **client-certificate-data** key.

```
student@master:~$ export client=$(grep client-cert $HOME/.kube/config |cut -d " " -f 6)
```

```
student@master:~$ echo $client
```

```
1 LS0tLS1CRUdJTiBDRVJUSUZJQ0FURSB0tLS0tCk1JSUM4akNDQWRxZ0F3SUJ
2 BZ01JRy9wbC9rWEpNdm3RFFZSkvWklodmNOQVFFTEJRQXdGVEVUTUJFR0
3 ExVUUKQXhNS2EzVmlaWEp1WlhSbGN6QWVGdzB4TnpFeU1UTXh0e1EyTXpKY
4 UZ3MHhPREV5TVRNeE56UTJNe1JhTURReApGekFWQmdOVk1JBb1REbk41YzNS
5 <output_omitted>
```

3. Almost the same command, but this time collect the **client-key-data** as the **key** variable.

```
student@master:~$ export key=$(grep client-key-data $HOME/.kube/config |cut -d " " -f 6)
```

```
student@master:~$ echo $key
```

```
1 <output_omitted>
```

4. Finally set the **auth** variable with the **certificate-authority-data** key.

```
student@master:~$ export auth=$(grep certificate-authority-data $HOME/.kube/config |cut -d " " -f 6)
```

```
student@master:~$ echo $auth
```

```
1 <output_omitted>
```

5. Now encode the keys for use with **curl**.

```
student@master:~$ echo $client | base64 -d - > ./client.pem
```

```
student@master:~$ echo $key | base64 -d - > ./client-key.pem
```

```
student@master:~$ echo $auth | base64 -d - > ./ca.pem
```

6. Pull the API server URL from the config file. Your hostname or IP address may be different.

```
student@master:~$ kubectl config view |grep server
```

```
1 server: https://k8smaster:6443
```

7. Use **curl** command and the encoded keys to connect to the API server. Use your hostname, or IP, found in the previous command, which may be different than the example below.

```
student@master:~$ curl --cert ./client.pem \
--key ./client-key.pem \
--cacert ./ca.pem \
https://k8smaster:6443/api/v1/pods
```

```
{
  "kind": "PodList",
  "apiVersion": "v1",
  "metadata": {
    "selfLink": "/api/v1/pods",
    "resourceVersion": "239414"
  },
  <output_omitted>
```

8. If the previous command was successful, create a JSON file to create a new pod. Remember to use **find** and search for this file in the tarball output, it can save you some typing.

```
student@master:~$ vim curlpod.json
```

```
{
  "kind": "Pod",
  "apiVersion": "v1",
  "metadata": {
    "name": "curlpod",
    "namespace": "default",
    "labels": {
      "name": "examplepod"
    }
  },
  "spec": {
    "containers": [{
      "name": "nginx",
      "image": "nginx",
      "ports": [{"containerPort": 80}]
    }]
  }
}
```

9. The previous **curl** command can be used to build a XPOST API call. There will be a lot of output, including the scheduler and taints involved. Read through the output. In the last few lines the phase will probably show Pending, as it's near the beginning of the creation process.

```
student@master:~$ curl --cert ./client.pem \
--key ./client-key.pem --cacert ./ca.pem \
https://k8smaster:6443/api/v1/namespaces/default/pods \
-XPOST -H'Content-Type: application/json' \
-d@curlpod.json
```

```
{
  "kind": "Pod",
  "apiVersion": "v1",
  "metadata": {
    "name": "curlpod",
    <output_omitted>
```


- Verify the new pod exists and shows a Running status.

```
student@master:~$ kubectl get pods
```

	NAME	READY	STATUS	RESTARTS	AGE
1	curlpod	1/1	Running	0	45s

Exercise 4.2: Explore API Calls

- One way to view what a command does on your behalf is to use **strace**. In this case, we will look for the current endpoints, or targets of our API calls. Install the tool, if not present.

```
student@master:~$ sudo apt-get install -y strace
```

```
student@master:~$ kubectl get endpoints
```

	NAME	ENDPOINTS	AGE
1	kubernetes	10.128.0.3:6443	3h

- Run this command again, preceded by **strace**. You will get a lot of output. Near the end you will note several **openat** functions to a local directory, `/home/student/.kube/cache/discovery/k8smaster_6443`. If you cannot find the lines, you may want to redirect all output to a file and `grep` for them. This information is cached, so you may see some differences should you run the command multiple times. As well your IP address may be different.

```
student@master:~$ strace kubectl get endpoints
```

```
1 execve("/usr/bin/kubectl", ["kubectl", "get", "endpoints"], [/*...
2 ....
3 openat(AT_FDCWD, "/home/student/.kube/cache/discovery/k8smaster_6443..
4 <output_omitted>
```

- Change to the parent directory and explore. Your endpoint IP will be different, so replace the following with one suited to your system.

```
student@master:~$ cd /home/student/.kube/cache/discovery/
```

```
student@master:~/kube/cache/discovery$ ls
```

```
1 k8smaster_6443
```

```
student@master:~/kube/cache/discovery$ cd k8smaster_6443/
```

- View the contents. You will find there are directories with various configuration information for kubernetes.

```
student@master:~/kube/cache/discovery/k8smaster_6443$ ls
```

1 admissionregistration.k8s.io	certificates.k8s.io	node.k8s.io
2 apiextensions.k8s.io	coordination.k8s.io	policy
3 apiregistration.k8s.io	crd.projectcalico.org	rbac.authorization.k8s.io
4 apps	discovery.k8s.io	scheduling.k8s.io
5 authentication.k8s.io	events.k8s.io	servergroups.json
6 authorization.k8s.io	extensions	storage.k8s.io
7 autoscaling	flowcontrol.apiserver.k8s.io	v1
8 batch	networking.k8s.io	

- Use the `find` command to list out the subfiles. The prompt has been modified to look better on this page.

```
student@master:~/k8smaster_6443$ find .
```

```

1 .
2 ./storage.k8s.io
3 ./storage.k8s.io/v1beta1
4 ./storage.k8s.io/v1beta1/serverresources.json
5 ./storage.k8s.io/v1
6 ./storage.k8s.io/v1/serverresources.json
7 ./scheduling.k8s.io
8 ./scheduling.k8s.io/v1beta1
9 ./scheduling.k8s.io/v1beta1/serverresources.json
10 <output_omitted>

```

6. View the objects available in version 1 of the API. For each object, or kind:, you can view the verbs or actions for that object, such as create seen in the following example. Note the prompt has been truncated for the command to fit on one line. Some are HTTP verbs, such as GET, others are product specific options, not standard HTTP verbs. The command may be **python**, depending on what version is installed.

```
student@master:~$ python3 -m json.tool v1/serverresources.json
```

JSON **serverresources.json**

```

1 {
2   "apiVersion": "v1",
3   "groupVersion": "v1",
4   "kind": "APIResourceList",
5   "resources": [
6     {
7       "kind": "Binding",
8       "name": "bindings",
9       "namespaced": true,
10      "singularName": "",
11      "verbs": [
12        "create"
13      ]
14    },
15    <output_omitted>

```

7. Some of the objects have shortNames, which makes using them on the command line much easier. Locate the shortName for endpoints.

```
student@master:~$ python3 -m json.tool v1/serverresources.json | less
```

JSON **serverresources.json**

```

1 ....
2 {
3   "kind": "Endpoints",
4   "name": "endpoints",
5   "namespaced": true,
6   "shortNames": [
7     "ep"
8   ],
9   "singularName": "",
10  "verbs": [
11    "create",
12    "delete",
13    ....

```

8. Use the shortName to view the endpoints. It should match the output from the previous command.

```
student@master:~$ kubectl get ep
```

```
1  NAME                ENDPOINTS                AGE
2  kubernetes          10.128.0.3:6443          3h
```

9. We can see there are 37 objects in version 1 file.

```
student@master:~$ python3 -m json.tool v1/serverresources.json | grep kind
```

```
1  "kind": "APIResourceList",
2  "kind": "Binding",
3  "kind": "ComponentStatus",
4  "kind": "ConfigMap",
5  "kind": "Endpoints",
6  "kind": "Event",
7  <output_omitted>
```

10. Looking at another file we find nine more.

```
student@master:~$ python3 -m json.tool \
  apps/v1/serverresources.json | grep kind
```

```
1  "kind": "APIResourceList",
2  "kind": "ControllerRevision",
3  "kind": "DaemonSet",
4  "kind": "DaemonSet",
5  "kind": "Deployment",
6  <output_omitted>
```

11. Delete the curlpod to recoup system resources.a

```
student@master:~$ kubectl delete po curlpod
```

```
1  pod "curlpod" deleted
```

12. Take a look around the other files in this directory as time permits.

Chapter 5

API Objects



5.1	API Objects	86
5.2	The v1 Group	87
5.3	API Resources	89
5.4	RBAC APIs	94
5.5	Labs	95

5.1 API Objects

Overview

- Ongoing growth in API objects
- Track release notes to find new objects
- In this chapter we will introduce common API objects:
 - Deployment the typical object used
 - DaemonSets, ReplicaSets now apps/v1
 - StatefulSets (once called PetSets) part of apps/v1 since v1.9
 - Jobs and CronJob now batch/v1
 - RBAC moved from v1alpha1 all the way to v1 in one release
- Explain which API group contains these new API objects.
- Find additional resources to start using the new API objects.

This chapter is about additional API resources or objects. We will learn about resources in the v1 API group, among others. Stability increases and code becomes more stable as objects move from alpha versions, to beta, then v1 indicating stability.

DaemonSets, which ensure a Pod on every node, and StatefulSets, which stick a container to a node and otherwise act like a deployment, have progressed to apps/v1 stability.

Role-based Access Control (RBAC), essential to security has made the leap from v1alpha1 to the stable v1 status.

As a fast moving project keeping track of changes, and possible changes can be an important part of ongoing system administration. Release notes, as well as discussions to release notes can be found in version-dependent sub-directories at: <https://github.com/kubernetes/enhancements/>. For example, the release feature status can be found here: <https://kubernetes.io/docs/setup/release/notes/>.

Starting with v1.16 deprecated API object versions will respond with an error instead of being accepted. This is an important change from historic behavior.

5.2 The v1 Group

v1 API Group

- Pod
- Node
- Service Account
- Resource Quota
- Endpoint
- More added with each release

The v1 API is no longer a single group, but rather a collection of groups for each main object category. For example there is a v1 group, a storage.k8s.io/v1 group, and rbac.authorization.k8s.io/v1 etc... Currently there are eight v1 groups.

We have touched on several objects in lab exercises. Here are details for some of them:

- **Node**

Represents a machine (physical or virtual) that is part of your Kubernetes cluster. You can get more information about nodes with the `kubectl get nodes` command. You can turn on and off the scheduling to a node with the `kubectl cordon/uncordon` commands.

- **Service Account**

Provides an identifier for processes running in a pod to access the API server and performs actions that it is authorized to do.

- **Resource Quota**

It is an extremely useful tool, allowing you to define quotas per namespace. For example, if you want to limit a specific namespace to only run a given number of pods, you can write a `resourcequota` manifest, create it with **kubectl** and the quota will be enforced.

- **Endpoint**

Generally, you do not manage endpoints. They represent the set of IPs for pods that match a particular service. They are handy when you want to check that a service actually matches some running pods. If an endpoint is empty, then it means that there are no matching pods and something is most likely wrong with your service definition.

Discovering API Groups

```
$ curl https://localhost:6443/apis \
--header "Authorization: Bearer $token" -k

{
  "kind": "APIGroupList",
  "apiVersion": "v1",
  "groups": [
    {
      "name": "apiregistration.k8s.io",
      "versions": [
        {
          "groupVersion": "apiregistration.k8s.io/v1",
          "version": "v1"
        }
      ],
      "preferredVersion": {
        "groupVersion": "apiregistration.k8s.io/v1",
        "version": "v1"
      }
    }
  ]
}
```

We can take a closer look at the output of request for current APIs. Each of the name values can be appended to the URL to see details of that group. For example you could drill down to find included objects at this URL: <https://localhost:6443/apis/apiregistration.k8s.io/v1beta1>

If you follow this URL you will find only one resource, with a name of `apiservices`. If it seems to be listed twice the lower output is for status. You'll note there are different verbs or actions for each. Another entry is if this object is namespaced, or restricted to only one namespace. In this case it is not.

You could **curl** each of these URIs and discover additional API objects, their characteristics and associated verbs.

5.3 API Resources

Deploying an Application

- Deployment
- ReplicaSet
- Pod

Using the **kubectl create** command we can quickly deploy an application. We have looked at the Pods created running the application, like **nginx**. Looking closer you will find that a **Deployment** was created which manages a **ReplicaSet** which then deploys the Pod. Lets take a closer look at each object.

- **Deployment** - A controller which manages the state of ReplicaSets and the pods within. The higher level control allows for more flexibility with upgrades and administration. Unless you have a good reason, use a deployment.
- **ReplicaSet** - Orchestrates individual Pod life cycle and updates. These are newer versions of Replication Controllers which differ only in selector support.
- **Pod** - As we've mentioned the lowest unit we can manage, runs the application container, possibly support containers.

DaemonSets

- Ensures every node runs a single pod
- Similar to ReplicaSet
- Often used for logging, metrics and security pods.
- Can be configured to avoid nodes

Should you want to have a logging application on every node a `DaemonSet` may be a good choice. The controller ensures that a single pod, of the same type, runs on every node in the cluster. When a new node is added to the cluster a Pod, same as deployed on the other nodes, is started. When the node is removed the `DaemonSet` makes sure the local Pod is deleted.

As usual, you get all the CRUD operations via **kubectl**:

```
$ kubectl get daemonsets
$ kubectl get ds
```

StatefulSet

- Similar to Deployment
- Ensures unique pods
- Guarantee ordering
- Stable in v1.9

Pods deployed using a `StatefulSet` use the same Pod specification. How this is different than a `Deployment` is that a `StatefulSet` considers each Pod as unique and provides ordering to Pod deployment.

In order to track each Pod as a unique object the controller uses identity composed of stable storage, stable network identity, and an ordinal. This identity remains with the node regardless to which node the Pod is running on at any one time.

The default deployment scheme is sequential starting with 0, such as `app-0`, `app-1`, `app-2` etc. A following Pod will not launch until the current Pod reaches a running and ready state. They are not deployed in parallel.

Autoscaling

- Agents which add or remove resources from the cluster.
- Horizontal Pod Autoscaling (HPA)
 - Scale based on current CPU usage, or custom metric
 - Must have **Metrics Server** or custom component running
- Vertical Pod Autoscaler (under development)
- Cluster Autoscaler (CA)
 - Add or remove nodes based on utilization
 - Makes request to cloud provider
 - Pods which cannot be evicted prevent scale-down

In the autoscaling group we find the **Horizontal Pod Autoscalers (HPA)**. This is a stable resource. HPAs automatically scale Replication Controllers, ReplicaSets, or Deployments based on a target of 50% CPU usage by default. The usage is checked by kubelet every 30 seconds and retrieved by Metrics Server API call every minute. HPA checks with Metrics Server every 30 seconds. Should a Pod be added or removed HPA waits 180 seconds before further action.

Other metrics can be used and queried via REST. The autoscaler does not collect the metrics, it only makes a request for the aggregated information and increases or decreases the number of replicas to match the configuration.

The Cluster Autoscaler (CA) adds or removes nodes to the cluster based off of inability to deploy a Pod or having nodes with low utilization for at least 10 minutes. This allows dynamic requests of resources from the cloud provider and minimizes expense for unused nodes. If you are using CA nodes should be added and removed through `cluster-autoscaler-` commands. Scale-up and down of nodes is checked every 10 seconds, but decisions are made on a node every 10 minutes. Should a scale-down fail the group will be rechecked in 3 minutes, with the failing node being eligible in five minutes. The total time to allocate a new node is largely dependent on the cloud provider.

Another project still under development is the Vertical Pod Autoscaler. This component will adjust the amount of CPU and memory requested by Pods.

Jobs

- Part of Batch API group
- Jobs run Pod until number of completions reached
 - Batch processing or one-off Pods
 - Ensure specified number of pods successfully terminate
 - Can run multiple Pods in parallel
- Cronjob to run Pod on regular basis
 - Creates a Pod about once per executing time
 - Some issues, job should be idempotent
 - Can run in serial or parallel
 - Same time syntax as Linux cron job

Jobs are part of the batch API group. They are used to run a set number of pods to completion. If a pod fails, it will be restarted until the number of completion is reached.

While they can be seen as a way to do batch processing in Kubernetes, they can also be used to run one-off pods. A Job specification will have a parallelism and a completion key. If omitted, they will be set to one. If they are present, the parallelism number will set the number of pods that can be running concurrently and the completion number will set how many pods need to run successfully for the Job itself to be considered done. Several Job patterns can be implemented, like a traditional work queue.

Cronjobs work in the similar manner to **Linux** jobs with the same time syntax. There are some cases where a job would not be run during a time period or could run twice, as a result the requested Pod should be idempotent.

An option spec field is `.spec.concurrencyPolicy` which determines how to handle existing jobs should the time segment expire. If set to `Allow`, the default, another concurrent job will be run. If set to `Forbid` the current job continues and the new job is skipped. A value of `Replace` cancels the current job and starts a new job in its place.

5.4 RBAC APIs

RBAC

- `rbac.authorization.k8s.io`
- Provide resources
 - `ClusterRole`
 - `ClusterRoleBinding`
 - `RoleBinding`
 - `Role`
- Combined with quotas for typical production deployments

The last API resources that we will look at are in the `rbac.authorization.k8s.io` group. We actually have four resources: `ClusterRole`, `Role`, `ClusterRoleBinding`, and `RoleBinding`. They are used for **Role Based Access Control (RBAC)** to Kubernetes.

```
$ curl localhost:8080/apis/rbac.authorization.k8s.io/v1
```

```
1  ...
2  "groupVersion": "rbac.authorization.k8s.io/v1",
3  "resources": [
4  ...
5      "kind": "ClusterRoleBinding"
6  ...
7      "kind": "ClusterRole"
8  ...
9      "kind": "RoleBinding"
10 ...
11     "kind": "Role"
12 ...
```

These resources allow us to define Roles within a cluster and associate users to these Roles. For example, we can define a Role for someone who can only read pods in a specific namespace, or a Role that can create deployments, but no services.



Please Note

More on RBAC is covered later in the Security chapter.

5.5 Labs

Exercise 5.1: RESTful API Access

Overview

We will continue to explore ways of accessing the control plane of our cluster. In the security chapter we will discuss there are several authentication methods, one of which is use of a Bearer token. We will work with one then deploy a local proxy server for application-level access to the Kubernetes API.

We will use the **curl** command to make API requests to the cluster, in an insecure manner. Once we know the IP address and port, then the token we can retrieve cluster data in a RESTful manner. By default most of the information is restricted, but changes to authentication policy could allow more access.

1. First we need to know the IP and port of a node running a replica of the API server. The master system will typically have one running. Use **kubectl config view** to get overall cluster configuration, and find the server entry. This will give us both the IP and the port.

```
student@master:~$ kubectl config view
```

```
1 apiVersion: v1
2 clusters:
3 - cluster:
4   certificate-authority-data: DATA+OMITTED
5   server: https://k8smaster:6443
6   name: kubernetes
7 <output_omitted>
```

2. Next we need to find the bearer token. This is part of a default token. Look at a list of tokens, first all on the cluster, then just those in the default namespace. There will be a secret for each of the controllers of the cluster.

```
student@master:~$ kubectl get secrets --all-namespaces
```

```
1 NAMESPACE      NAME                                TYPE      ...
2 default         default-token-jdqp7               kubernetes.io/service-account-token...
3 kube-node-lease default-token-j67mt               kubernetes.io/service-account-token...
4 kube-public     default-token-b2prn               kubernetes.io/service-account-token...
5 kube-system     attachdetach-controller-token-ckwvh kubernetes.io/servic...
6 kube-system     bootstrap-signer-token-wpx66      kubernetes.io/service-accou...
7 <output_omitted>
```

```
student@master:~$ kubectl get secrets
```

```
1 NAME                                TYPE                                DATA  AGE
2 default-token-jdqp7                 kubernetes.io/service-account-token 3      23h
```

3. Look at the details of the secret. We will need the token: information from the output.

```
student@master:~$ kubectl describe secret default-token-jdqp7
```

```
1 Name:                 default-token-jdqp7
2 Namespace:            default
3 Labels:               <none>
4 <output_omitted>
5 token:                eyJhbGciOiJSUzI1NiIsInR5cCI6IkpXVCJ9.eyJpc3MiOiJrdWJ1cm5ldGVz
6 L3NlcnZpY2VhY2NvdW50Iiwia3ViZXJuZXRlcy5pby9zZXJ2aWNlYWNjb3VudC9uYW1lc3Bh
7 Y2UiOiJkZWZhdWx0Iiwia3ViZXJuZXRlcy5pby9zZXJ2aWNlYWNjb3VudC9zZWNyZXQubm
8 <output_omitted>
```

4. Using your mouse to cut and paste, or **cut**, or **awk** to save the data, from the first character `eyJh` to the last, to a variable named `token`. Your token data will be different.

```
student@master:~$ export token=$(kubectl describe \
  secret default-token-jdqp7 |grep ^token |cut -f7 -d ' ')
```

5. Test to see if you can get basic API information from your cluster. We will pass it the server name and port, the token and use the **-k** option to avoid using a cert.

```
student@master:~$ curl https://k8smaster:6443/apis \
  --header "Authorization: Bearer $token" -k
```

```
{
  "kind": "APIGroupList",
  "apiVersion": "v1",
  "groups": [
    {
      "name": "apiregistration.k8s.io",
      "versions": [
        {
          "groupVersion": "apiregistration.k8s.io/v1",
          "version": "v1"
        }
      ]
    }
  ]
}
```

<output_omitted>

6. Try the same command, but look at API v1. Note that the path has changed to `api`.

```
student@master:~$ curl https://k8smaster:6443/api/v1 \
  --header "Authorization: Bearer $token" -k
```

```
1 <output_omitted>
```

7. Now try to get a list of namespaces. This should return an error. It shows our request is being seen as `systemserviceaccount`, which does not have the RBAC authorization to list all namespaces in the cluster.

```
student@master:~$ curl \
  https://k8smaster:6443/api/v1/namespaces \
  --header "Authorization: Bearer $token" -k
```

```
1 <output_omitted>
2 "message": "namespaces is forbidden: User \"system:serviceaccount:default...
3 <output_omitted>
```

8. Pods can also make use of included certificates to use the API. The certificates are automatically made available to a pod under the `/var/run/secrets/kubernetes.io/serviceaccount/`. We will deploy a simple Pod and view the resources. If you view the `token` file you will find it is the same value we put into the `$token` variable. The **-i** will request a **-t** terminal session of the busybox container. Once you exit the container will not restart and the pod will show as completed.

```
student@master:~$ kubectl run -i -t busybox --image=busybox --restart=Never
```



Inside container

```
# ls /var/run/secrets/kubernetes.io/serviceaccount/
ca.crt namespace token
# exit
```

9. Clean up by deleting the busybox container.

```
student@master:~$ kubectl delete pod busybox
```



```
1 pod "busybox" deleted
```

Exercise 5.2: Using the Proxy

Another way to interact with the API is via a proxy. The proxy can be run from a node or from within a Pod through the use of a sidecar. In the following steps we will deploy a proxy listening to the loopback address. We will use **curl** to access the API server. If the **curl** request works, but does not from outside the cluster, we have narrowed down the issue to authentication and authorization instead of issues further along the API ingestion process.

1. Begin by starting the proxy. It will start in the foreground by default. There are several options you could pass. Begin by reviewing the help output.

```
student@master:~$ kubectl proxy -h
```

```
1 Creates a proxy server or application-level gateway between localhost
2 and the Kubernetes API Server. It also allows serving static content
3 over specified HTTP path. All incoming data enters through one port
4 and gets forwarded to the remote kubernetes API Server port, except
5 for the path matching the static content path.
6
7 Examples:
8   # To proxy all of the kubernetes api and nothing else, use:
9
10  $ kubectl proxy --api-prefix=/
11  <output_omitted>
```

2. Start the proxy while setting the API prefix, and put it in the background. You may need to use **enter** to view the prompt. Take note of the process ID, 22500 in the example below, we'll use it to kill the process when we are done.

```
student@master:~$ kubectl proxy --api-prefix=/ &
```

```
1 [1] 22500
2 Starting to serve on 127.0.0.1:8001
```

3. Now use the same **curl** command, but point toward the IP and port shown by the proxy. The output should be the same as without the proxy, but may be formatted differently.

```
student@master:~$ curl http://127.0.0.1:8001/api/
```

```
1 <output_omitted>
```

4. Make an API call to retrieve the namespaces. The command did not work in the previous section due to permissions, but should work now as the proxy is making the request on your behalf.

```
student@master:~$ curl http://127.0.0.1:8001/api/v1/namespaces
```

```
{
  "kind": "NamespaceList",
  "apiVersion": "v1",
  "metadata": {
    "selfLink": "/api/v1/namespaces",
    "resourceVersion": "86902"
  }
}
<output_omitted>
```

5. Stop the proxy service as we won't need it any more. Use the process ID from a previous step. Your process ID may be different.

```
student@master:~$ kill 22500
```

Exercise 5.3: Working with Jobs

While most API objects are deployed such that they continue to be available there are some which we may want to run a particular number of times called a Job, and others on a regular basis called a CronJob

Create A Job

1. Create a job which will run a container which sleeps for three seconds then stops.

```
student@master:~$ vim job.yaml
```

YAML

job.yaml

```
1 apiVersion: batch/v1
2 kind: Job
3 metadata:
4   name: sleepy
5 spec:
6   template:
7     spec:
8       containers:
9       - name: resting
10         image: busybox
11         command: ["/bin/sleep"]
12         args: ["3"]
13         restartPolicy: Never
```

2. Create the job, then verify and view the details. The example shows checking the job three seconds in and then again after it has completed. You may see different output depending on how fast you type.

```
student@master:~$ kubectl create -f job.yaml
```

```
1 job.batch/sleepy created
```

```
student@master:~$ kubectl get job
```

```
1 NAME      COMPLETIONS  DURATION  AGE
2 sleepy    0/1          3s        3s
```

```
student@master:~$ kubectl describe jobs.batch sleepy
```

```
1 Name:      sleepy
2 Namespace: default
3 Selector:   controller-uid=24c91245-d0fb-11e8-947a-42010a800002
4 Labels:     controller-uid=24c91245-d0fb-11e8-947a-42010a800002
5             job-name=sleepy
6 Annotations: <none>
7 Parallelism: 1
8 Completions: 1
9 Start Time:  Tue, 16 Oct 2018 04:22:50 +0000
10 Completed At: Tue, 16 Oct 2018 04:22:55 +0000
11 Duration:    5s
12 Pods Statuses: 0 Running / 1 Succeeded / 0 Failed
13 <output_omitted>
```

```
student@master:~$ kubectl get job
```

```
1 NAME      COMPLETIONS  DURATION  AGE
2 sleepy    1/1          5s        17s
```

3. View the configuration information of the job. There are three parameters we can use to affect how the job runs. Use **-o yaml** to see these parameters. We can see that backoffLimit, completions, and the parallelism. We'll add these parameters next.

```
student@master:~$ kubectl get jobs.batch sleepy -o yaml
```

```
1 <output_omitted>
2   uid: c2c3a80d-d0fc-11e8-947a-42010a800002
3 spec:
4   backoffLimit: 6
5   completions: 1
6   parallelism: 1
7   selector:
8     matchLabels:
9 <output_omitted>
```

4. As the job continues to AGE in a completion state, delete the job.

```
student@master:~$ kubectl delete jobs.batch sleepy
```

```
1 job.batch "sleepy" deleted
```

5. Edit the YAML and add the completions: parameter and set it to 5.

```
student@master:~$ vim job.yaml
```

YAML

job.yaml

```
1 <output_omitted>
2 metadata:
3   name: sleepy
4 spec:
5   completions: 5  #<--Add this line
6   template:
7     spec:
8       containers:
9 <output_omitted>
```

6. Create the job again. As you view the job note that COMPLETIONS begins as zero of 5.

```
student@master:~$ kubectl create -f job.yaml
```

```
1 job.batch/sleepy created
```

```
student@master:~$ kubectl get jobs.batch
```

```
1 NAME      COMPLETIONS  DURATION  AGE
2 sleepy    0/5          5s        5s
```

7. View the pods that running. Again the output may be different depending on the speed of typing.

```
student@master:~$ kubectl get pods
```

```
1 NAME                READY  STATUS      RESTARTS  AGE
2 sleepy-z5tnh         0/1    Completed   0          8s
3 sleepy-zd692         1/1    Running     0          3s
4 <output_omitted>
```

8. Eventually all the jobs will have completed. Verify then delete the job.

```
student@master:~$ kubectl get jobs
```

```
1 NAME      COMPLETIONS  DURATION  AGE
2 sleepy    5/5          26s       10m
```

```
student@master:~$ kubectl delete jobs.batch sleepy
```

```
1 job.batch "sleepy" deleted
```

9. Edit the YAML again. This time add in the `parallelism:` parameter. Set it to **2** such that two pods at a time will be deployed.

```
student@master:~$ vim job.yaml
```

YAML

job.yaml

```
1 <output_omitted>
2   name: sleepy
3   spec:
4     completions: 5
5     parallelism: 2    #<-- Add this line
6     template:
7       spec:
8 <output_omitted>
```

10. Create the job again. You should see the pods deployed two at a time until all five have completed.

```
student@master:~$ kubectl create -f job.yaml
```

```
1 job.batch/sleepy created
```

```
student@master:~$ kubectl get pods
```

```
1 NAME                READY   STATUS    RESTARTS   AGE
2 sleepy-8xwpc         1/1     Running   0           5s
3 sleepy-xjqnf         1/1     Running   0           5s
4 <output_omitted>
```

```
student@master:~$ kubectl get jobs
```

```
1 NAME      COMPLETIONS   DURATION   AGE
2 sleepy    3/5            11s        11s
```

11. Add a parameter which will stop the job after a certain number of seconds. Set the `activeDeadlineSeconds:` to 15. The job and all pods will end once it runs for 15 seconds. We will also increase the sleep argument to five, just to be sure does not expire by itself.

```
student@master:~$ vim job.yaml
```

YAML

```
1 <output_omitted>
2   completions: 5
3   parallelism: 2
4   activeDeadlineSeconds: 15    #<-- Add this line
5   template:
6     spec:
7       containers:
8       - name: resting
9         image: busybox
10        command: ["/bin/sleep"]
11        args: ["5"]             #<-- Edit this line
12 <output_omitted>
```

12. Delete and recreate the job again. It should run for 15 seconds, usually 3/5, then continue to age without further completions.

```
student@master:~$ kubectl delete jobs.batch sleepy
```

```
1 job.batch "sleepy" deleted
```

```
student@master:~$ kubectl create -f job.yaml
```

```
1 job.batch/sleepy created
```

```
student@master:~$ kubectl get jobs
```

```
1 NAME      COMPLETIONS  DURATION  AGE
2 sleepy    1/5          6s        6s
```

```
student@master:~$ kubectl get jobs
```

```
1 NAME      COMPLETIONS  DURATION  AGE
2 sleepy    3/5          16s       16s
```

13. View the message: entry in the Status section of the object YAML output.

```
student@master:~$ kubectl get job sleepy -o yaml
```

```
1 <output_omitted>
2 status:
3   conditions:
4   - lastProbeTime: 2018-10-16T05:45:14Z
5     lastTransitionTime: 2018-10-16T05:45:14Z
6     message: Job was active longer than specified deadline
7     reason: DeadlineExceeded
8     status: "True"
9     type: Failed
10  failed: 2
11  startTime: 2018-10-16T05:44:59Z
12  succeeded: 3
```

14. Delete the job.

```
student@master:~$ kubectl delete jobs.batch sleepy
```

```
1 job.batch "sleepy" deleted
```

Create a CronJob

A CronJob creates a watch loop which will create a batch job on your behalf when the time becomes true. We Will use our existing Job file to start.

1. Copy the Job file to a new file.

```
student@master:~$ cp job.yaml cronjob.yaml
```

2. Edit the file to look like the annotated file shown below. Edit the lines mentioned below. The three parameters we added will need to be removed. Other lines will need to be further indented.

```
student@master:~$ vim cronjob.yaml
```



```

1 apiVersion: batch/v1beta1      #<-- Add beta1 to be v1beta1
2 kind: CronJob                  #<-- Update this line to CronJob
3 metadata:
4   name: sleepy
5 spec:
6   schedule: "*/* * * * *"      #<-- Add Linux style cronjob syntax
7   jobTemplate:                  #<-- New jobTemplate and spec move
8     spec:
9       template:                 #<-- This and following lines move
10        spec:                    #<-- four spaces to the right
11          containers:
12            - name: resting
13              image: busybox
14              command: ["/bin/sleep"]
15              args: ["5"]
16              restartPolicy: Never

```

3. Create the new CronJob. View the jobs. It will take two minutes for the CronJob to run and generate a new batch Job.

```
student@master:~$ kubectl create -f cronjob.yaml
```

```
1 cronjob.batch/sleepy created
```

```
student@master:~$ kubectl get cronjobs.batch
```

```

1 NAME      SCHEDULE      SUSPEND   ACTIVE   LAST SCHEDULE   AGE
2 sleepy    */2 * * * *   False    0        <none>          8s

```

```
student@master:~$ kubectl get jobs.batch
```

```
1 No resources found.
```

4. After two minutes you should see jobs start to run.

```
student@master:~$ kubectl get cronjobs.batch
```

```

1 NAME      SCHEDULE      SUSPEND   ACTIVE   LAST SCHEDULE   AGE
2 sleepy    */2 * * * *   False    0        21s             2m1s

```

```
student@master:~$ kubectl get jobs.batch
```

```

1 NAME                COMPLETIONS   DURATION   AGE
2 sleepy-1539722040    1/1           5s         18s

```

```
student@master:~$ kubectl get jobs.batch
```

```

1 NAME                COMPLETIONS   DURATION   AGE
2 sleepy-1539722040    1/1           5s         5m17s
3 sleepy-1539722160    1/1           6s         3m17s
4 sleepy-1539722280    1/1           6s         77s

```

5. Ensure that if the job continues for more than 10 seconds it is terminated. We will first edit the **sleep** command to run for 30 seconds then add the `activeDeadlineSeconds`: entry to the container.

```
student@master:~$ vim cronjob.yaml
```

YAML

```

1 .....
2   jobTemplate:
3     spec:
4       template:
5         spec:
6           activeDeadlineSeconds: 10 #<-- Add this line
7           containers:
8             - name: resting
9             .....
10          command: ["/bin/sleep"]
11          args: ["30"] #<-- Edit this line
12          restartPolicy: Never
13 .....

```

6. Delete and recreate the CronJob. It may take a couple of minutes for the batch Job to be created and terminate due to the timer.

```
student@master:~$ kubectl delete cronjobs.batch sleepy
```

```
1 cronjob.batch "sleepy" deleted
```

```
student@master:~$ kubectl create -f cronjob.yaml
```

```
1 cronjob.batch/sleepy created
```

```
student@master:~$ kubectl get jobs
```

```

1 NAME                COMPLETIONS  DURATION  AGE
2 sleepy-1539723240    0/1          61s       61s

```

```
student@master:~$ kubectl get cronjobs.batch
```

```

1 NAME    SCHEDULE    SUSPEND  ACTIVE  LAST SCHEDULE  AGE
2 sleepy  */2 * * * *  False   1       72s          94s

```

```
student@master:~$ kubectl get jobs
```

```

1 NAME                COMPLETIONS  DURATION  AGE
2 sleepy-1539723240    0/1          75s       75s

```

```
student@master:~$ kubectl get jobs
```

```

1 NAME                COMPLETIONS  DURATION  AGE
2 sleepy-1539723240    0/1          2m19s     2m19s
3 sleepy-1539723360    0/1          19s       19s

```

```
student@master:~$ kubectl get cronjobs.batch
```

```

1 NAME    SCHEDULE    SUSPEND  ACTIVE  LAST SCHEDULE  AGE
2 sleepy  */2 * * * *  False   2       31s          2m53s

```

7. Clean up by deleting the CronJob.

```
student@master:~$ kubectl delete cronjobs.batch sleepy
```

```
1 cronjob.batch "sleepy" deleted
```


Chapter 6

Design



In this chapter, we are going to talk about resource requirements for our applications, and the ability to set limits for CPU, memory and storage in the root filesystem of the node. We're also going to talk about containers that can assist our primary application, such as a sidecar, which is typically something that enhances the primary application. For example, if our application does not include logging, we can add a second container to the pod, to handle logging on our behalf. An adapter container might conform the traffic to match the other applications or needs of our cluster. This allows for full flexibility for the primary application, to remain generating the data that it does, and not have to be modified to fit that particular cluster. There is also an ambassador container, which is our representation to the outside world. It can act as a proxy then, so that perhaps I might want to split inbound traffic: if it's a read, it goes to one container, and for write, it goes to a different container, allowing me to optimize my environment and my application. We'll also consider other application design concepts, that can be helpful in the planning stage for the deployment of your application. We'll then finish by talking about jobs, whether it's a batch-like job that will run a container once and be done, or cron jobs, which will run a container on a regular basis on our behalf.

6.1	Traditional Applications: Considerations	105
6.2	Decoupled Resources	105
6.3	Transience	106
6.4	Flexible Framework	106
6.5	Managing Resource Usage	106
6.6	Using Label Selectors	107
6.7	Multi-Container Pods	108
6.8	Sidecar Container	109
6.9	Adapter Container	109
6.10	Ambassador	109
6.11	Points to Ponder	109
6.12	Jobs	110
6.13	Labs	110

6.1 Traditional Applications: Considerations

- One of the larger hurdles towards implementing Kubernetes in a production environment is the suitability of application design. Optimal Kubernetes deployment design changes are more than just the simple containerization of an application. Traditional applications were built and deployed with the expectation of long-term processes and strong interdependence.

- For example, an Apache web server allows for incredible customization. Often, the server would be configured and tweaked without interruption. As demand grows, the application may be migrated to larger and larger servers. The build and maintenance of the application assumes the instance would run without reset and have persistent and tightly coupled connections to other resources, such as networks and storage.
- In early usage of containers, applications were containerized without redevelopment. This led to issues after resource failure, or upon upgrade, or configuration. The cost and hassle of redesign and re-implementation should be taken into account.

6.2 Decoupled Resources

- The use of decoupled resources is integral to Kubernetes. Instead of an application using a dedicated port and socket, for the life of the instance, the goal is for each component to be decoupled from other resources. The expectation and software development toward separation allows for each component to be removed, replaced, or rebuilt.
- Instead of hard-coding a resource in an application, an intermediary, such as a Service, enables connection and reconnection to other resources, providing flexibility. A single machine is no longer required to meet the application and user needs; any number of systems could be brought together to meet the needs when, and, as long as, necessary.
- As Kubernetes grows, even more resources are being divided out, which allows for an easier deployment of resources. Also, Kubernetes developers can optimize a particular function with fewer considerations of other objects.

6.3 Transience

- Equally important is the expectation of transience. Each object should be developed with the expectation that other components will die and be rebuilt. With any and all resources planned for transient relationships to others, we can update versions or scale usage in an easy manner.
- An upgrade is perhaps not quite the correct term, as the existing application does not survive. Instead, a controller terminates the container and deploys a new one to replace it, using a different version of the application or setting. Typically, traditional applications were not written this way, opting toward long-term relationships for efficiency and ease of use.

6.4 Flexible Framework

- Like a school of fish, or pod of whales, multiple independent resources working together, but decoupled from each other and without expectation of individual permanent relationship, gain flexibility, higher availability and easy scalability. Instead of a monolithic Apache server, we can deploy a flexible number of **nginx** servers, each handling a small part of the workload. The goal is the same, but the framework of the solution is distinct.
- A decoupled, flexible and transient application and framework is not the most efficient. In order for the Kubernetes orchestration to work, we need a series of agents, otherwise known as controllers or watch-loops, to constantly monitor the current cluster state and make changes until that state matches the declared configuration.
- The commoditization of hardware has enabled the use of many smaller servers to handle a larger workload, instead of single, huge systems.

6.5 Managing Resource Usage

- As with any application, an understanding of resource usage can be helpful for a successful deployment. Kubernetes allows us to easily scale clusters, larger or smaller, to meet demand. An understanding of how the Kubernetes clusters view the resources is an important consideration. The **kube-scheduler**, or a custom scheduler, uses the **PodSpec** to determine the best node for deployment.

- In addition to administrative tasks to grow or shrink the cluster or the number of Pods, there are autoscalers which can add or remove nodes or pods, with plans for one which uses cgroup settings to limit CPU and memory usage by individual containers.
- By default, Pods use as much CPU and memory as the workload requires, behaving and coexisting with other Linux processes. Through the use of resource requests, the scheduler will only schedule a Pod to a node if resources exist to meet all requests on that node. The scheduler takes these and several other factors into account when selecting a node for use.
- Monitoring the resource usage cluster-wide is not an included feature of Kubernetes. Other projects, like Prometheus, are used instead. In order to view resource consumption issues locally, use the `kubectl describe pod` command. You may only know of issues after the pod has been terminated.

CPU

- CPU requests are made in CPU units, each unit being a millicore, using *mille* – the Latin word for thousand. Some documentation uses millicore, others use millicpu, but both have the same meaning. Thus, a request for .7 of a CPU would be 700 millicore. Should a container use more resources than allowed, it won't be killed. The exact amount of overuse is not definite. Note the notation often found in documentation. Each dot would represent a new line and indent if converted to YAML.
 - `spec.containers[].resources.limits.cpu`
 - `spec.containers[].resources.requests.cpu`
- The value of CPUs is not relative. It does not matter how many exist, or if other Pods have requirements. One CPU, in Kubernetes, is equivalent to:
 - 1 AWS vCPU
 - 1 GCP Core
 - 1 Azure vCore
 - 1 Hyperthread on a bare-metal Intel processor with Hyperthreading.

Memory (RAM)

- With Docker engine, the `limits.memory` value is converted to an integer value and becomes the value to the `docker run --memory <value> <image>` command. The handling of a container which exceeds its memory limit is not definite. It may be restarted, or, if it asks for more than the memory request setting, the entire Pod may be evicted from the node.
 - `spec.containers[].resources.limits.memory`
 - `spec.containers[].resources.requests.memory`

Ephemeral Storage (Beta Feature in 1.14)

- Container files, logs, and EmptyDir storage, as well as Kubernetes cluster data, reside on the root filesystem of the host node. As storage is a limited resource, you may need to manage it as other resources. The scheduler will only choose a node with enough space to meet the sum of all the container requests. Should a particular container, or the sum of the containers in a Pod, use more than the limit, the Pod will be evicted.
 - `spec.containers[].resources.limits.ephemeral-storage`
 - `spec.containers[].resources.requests.ephemeral-storage`

6.6 Using Label Selectors

- Labels allow for objects to be selected which may not share other characteristics. For example if a developer were to label their pods using their name they could affect all of their pods, regardless of the application or deployment the pods were using.
- Labels are how operators, also known as watch-loops, track and manage objects. As a result if you were to hard-code the same label for two objects they may be managed by different operators and cause issues. For example one deployment may call for ten pods, while another with the same label calls for five. All the pods would be in a constant state of restarting as each operator tries to start or stop pods until the status matches the spec.
- Consider the possible ways you may want to group your pods and other objects in production. Perhaps you may use `development` and `production` labels to differentiate the state of the application. Perhaps add labels for the department, team, and primary application developer.
- Selectors are namespace scoped. Use the `--all-namespaces` argument to select matching objects in all namespaces.
- The labels, annotations, name, and metadata of an object can be found near the top of `kubectl get {obj-name} -o yaml`. For example:

```
ckad1$ kubectl get pod examplepod-1234-vtlzd -o yaml
```

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   annotations:
5     cni.projectcalico.org/podIP: 192.168.113.220/32
6   creationTimestamp: "2020-01-31T15:13:08Z"
7   generateName: examplepod-1234-
8   labels:
9     app: examplepod
10    pod-template-hash: 1234
11    name: examplepod-1234-vtlzd
12 ....
```

- Using this output one way we could select the pod would be to use the `app` pod: label. In the command line the colon(:) would be replaced with an equals(=) sign and the space removed. The use of `-l` or `--selector` options can be used with `kubectl`.

```
ckad1$ kubectl -n test2 get --selector app=examplepod pod
```

NAME	READY	STATUS	RESTARTS	AGE
examplepod-1234-vtlzd	1/1	Running	0	25m

- There are several built-in object labels. For example nodes have labels such as the `arch`, `hostname`, and `os`, which could be used for assigning pods to a particular node, or type of node.

```
ckad1$ kubectl get node worker
```

```
1 ....
2   creationTimestamp: "2020-05-12T14:23:04Z"
3   labels:
4     beta.kubernetes.io/arch: amd64
5     beta.kubernetes.io/os: linux
6     kubernetes.io/arch: amd64
7     kubernetes.io/hostname: worker
8     kubernetes.io/os: linux
9   managedFields:
10  ....
```

- The `nodeSelector`: entry in the podspec could use this label to cause a pod to be deployed on a particular node with an entry such as:

```
spec:
  nodeSelector:
    kubernetes.io/hostname: worker
  containers:
```

6.7 Multi-Container Pods

- The idea of multiple containers in a Pod goes against the architectural idea of decoupling as much as possible. One could run an entire operating system inside a container, but would lose much of the granular scalability Kubernetes is capable of. But there are certain needs in which a second or third co-located container makes sense. By adding a second container, each container can still be optimized and developed independently, and both can scale and be repurposed to best meet the needs of the workload.

6.8 Sidecar Container

- The idea for a sidecar container is to add some functionality not present in the main container. Rather than bloating code, which may not be necessary in other deployments, adding a container to handle a function such as logging solves the issue, while remaining decoupled and scalable. Prometheus monitoring and Fluentd logging leverage sidecar containers to collect data.

6.9 Adapter Container

- The basic purpose of an adapter container is to modify data, either on ingress or egress, to match some other need. Perhaps, an existing enterprise-wide monitoring tools has particular data format needs. An adapter would be an efficient way to standardize the output of the main container to be ingested by the monitoring tool, without having to modify the monitor or the containerized application. An adapter container transforms multiple applications to singular view.

6.10 Ambassador

- An ambassador allows for access to the outside world without having to implement a service or another entry in an ingress controller:
 - Proxy local connection
 - Reverse proxy
 - Limits HTTP requests
 - Re-route from the main container to the outside world.
- **Ambassador** (<https://www.getambassador.io/>) is an "open source, Kubernetes-native API gateway for microservices built on Envoy".

6.11 Points to Ponder

- Is my application as **decoupled** as it could possibly be?
Is there anything that I could take out, or make its own container?

These questions, while essential, often require an examination of the application within the container. Optimal use of Kubernetes is not typically found by containerization of a legacy application and deployment. Rather most applications are rebuilt entirely, with a focus on decoupled micro-services.

When every container can survive any other containers regular termination, without the end-user noticing, you have probably decoupled the application properly.

- Is each container **transient**, does it expect and have code to properly react when other containers are transient?

Could I run **Chaos Monkey** to kill **ANY** and multiple pods, and my end user would not notice?

Most Kubernetes deployments are not as transient as they should be, especially among legacy applications. The developer hold on to a previous approach and does not create the containerized application such that every connection and communication is transient and will be terminated. With code waiting for the `service` to connect to a replacement Pod and the containers within.

Some will note that this is not efficient as it could be. This is correct. We are not optimizing and working against the orchestration tool. Instead we are taking advantage of the decoupled and transient environment to scale and use only the particular micro-service the end user needs.

- Can I scale any particular component to meet workload demand?

The minimization of an application, breaking it down into the smallest part possible often goes against what developers have spent a career doing. Each division requires more code to handle errors and communication, and is less efficient than a tightly coupled application.

Machine code is very efficient, for example, but not portable. Instead code is written in a higher level language, which may not be interpreted to run as efficiently, but will run in varied environments. Perhaps approach code meant for Kubernetes in the sense that you are making the highest, and most non-specific way possible. If you have broken down each component then you can **scale** only the most necessary component. In this way the actual workload is more efficient, as the only software consuming CPU cycles is that which the end-user requires. There would be minimal application overhead and waste.

Ensure you understand how to leverage a **service** to handle networking. Every Pod is given a single IP address, shared by all containers in a pod. The watch-loop for services will manipulate firewall rules to get traffic to pods with matching **labels**, and interact with the network plugin, such as Calico.

- Have I used the most open standard stable enough to meet my needs?

This item can take a while to investigate, but usually a good investment. With tight production schedules some may use what they know best, or what is easiest at the moment. Due to the fast changing nature, or CI/CD, of a production environment this may lead to more work in the long run.

6.12 Jobs

- Just as we may need to redesign our applications to be decoupled we may also consider that microservices may not need to run all the time. The use of `Jobs` and `CronJobs` can further assist with implementing decoupled and transient microservices.
- Jobs are part of the `batch` API group. They are used to run a set number of pods to completion. If a pod fails, it will be restarted until the number of completion is reached.
- While they can be seen as a way to do batch processing in Kubernetes, they can also be used to run one-off pods. A Job specification will have a `parallelism` and a `completion` key. If omitted, they will be set to one. If they are present, the `parallelism` number will set the number of pods that can run concurrently, and the `completion` number will set how many pods need to run successfully for the Job itself to be considered done. Several Job patterns can be implemented, like a traditional work queue.
- `Cronjobs` work in a similar manner to Linux jobs, with the same time syntax. There are some cases where a job would not be run during a time period or could run twice; as a result, the requested Pod should be idempotent.
- An option `spec` field is `.spec.concurrencyPolicy` which determines how to handle existing jobs, should the time segment expire. If set to `Allow`, the default, another concurrent job will be run. If set to `Forbid`, the current job continues and the new job is skipped. A value of `Replace` cancels the current job and starts a new job in its place.

6.13 Labs

Exercise 6.1: Planning the Deployment

Overview

In this exercise we will investigate common network plugins. Each **kubelet** agent uses one plugin at a time. Due to complexity, the entire cluster uses one plugin which is configured prior to application deployment. Some plugins don't honor security configurations such as network policies. Should you design a deployment which and use a network policy there wouldn't be an error; the policy would have no effect. While developers typically wouldn't care much about the mechanics of it can affect the availability of features and troubleshooting of newly decoupled microservices.

While still new, the community is moving towards the **Container Network Interface (CNI)** specification (<https://github.com/containernetworking/cni>). This provides the most flexibility and features in the fast changing space of container networking.

A common alternative is **kubenet**, a basic plugin which relies on the cloud provider to handle routing and cross-node networking. In a previous lab exercise we configured **Project Calico**. Classic and external modes are also possible. Several software defined network projects intended for Kubernetes have been created recently, with new features added regularly. Be aware that **Calico** is a dynamic project with ongoing and frequent changes.

Evaluate Network Plugins

While developers don't need to configure cluster networking, they may need to understand the plugin in use, and it's particular features and quirks. This section to ensure you have made a quick review of the most commonly used plugins, and where to find more information as necessary.

1. Verify your nodes are using a CNI plugin. Read through the startup process of CNI. Each message begins with a time stamp, type of message and what is reporting the message.

```
student@master:~$ less /var/log/calico/cni/cni.log
```

```
1 2020-10-05 16:50:00.960 [INFO][21091] ipam.go 936: Attempting to assign 1 addresses
2 from block block=192.168.219.64/26 handle="k8s-pod-network.d303bf71efef750a3f420486f
3 e7b8b9e945abe841ef192673a4435af1c20428c" host="master"
```

2. There are many CNI providers possible. The following list represents some of the more common choices, but it is not exhaustive. With many new plugins being developed there may be another which better serves your needs. Use these websites to answer questions which follow. While we strive to keep the answers accurate, please be aware that this area has a lot of attention and development and changes often.

- **Project Calico**
<https://docs.projectcalico.org/v3.0/introduction/>
- **Calico with Canal**
<https://docs.projectcalico.org/v3.0/getting-started/kubernetes/installation/hosted/canal>
- **Weave Works**
<https://www.weave.works/docs/net/latest/kubernetes/kube-addon>
- **Flannel**
<https://github.com/coreos/flannel>
- **Romana**
<http://romana.io/how/romana.basics/>
- **Kube Router**
<https://www.kube-router.io>
- **Kopeio**
<https://github.com/kopeio/networking>

3. Which of the plugins allow vxlans?
4. Which are layer 2 plugins?

5. Which are layer 3?
6. Which allow network policies?
7. Which can encrypt all TCP and UDP traffic?

Multi-container Pod Considerations

Using the information learned from this chapter, consider the following questions:

1. Which deployment method would allow the most flexibility, multiple applications per pod or one per pod?
2. Which deployment method allows for the most granular scalability?
3. Which have the best performance?
4. How many IP addresses are assigned per pod?
5. What are some ways containers can communicate within the same pod?
6. What are some reasons you should have multiple containers per pod?

Do you really know?

When and why would you use a multi-container pod?

Have you found a YAML example online?

Go back and review multi-container pod types and content on decoupling if you can't easily answer these questions. We touched on adding a second logging and a readiness container in a previous chapter and will work more with logging a future exercise.

✓ Solution 6.1

Plugin Answers

1. Which of the plugins allow vxlans?
Canal, Project Calico, Flannel, Kopeio-networking, Weave Net
2. Which are layer 2 plugins?
Canal, Flannel, Kopeio-networking, Weave Net
3. Which are layer 3?
Project Calico, Romana, Kube Router
4. Which allow network policies?
Project Calico, Canal, Kube Router, Romana Weave Net
5. Which can encrypt all TCP and UDP traffic?
Project Calico, Kopeio, Weave Net

Multi Pod Answers

1. Which deployment method would allow the most flexibility, multiple applications per pod or one per Pod?
One per pod
2. Which deployment method allows for the most granular scalability?
One per pod
3. Which have the best inter-container performance?
Multiple per pod.
4. How many IP addresses are assigned per pod?
One
5. What are some ways containers can communicate within the same pod?
IPC, loopback or shared filesystem access.
6. What are some reasons you should have multiple containers per pod?
Lean containers may not have functionality like logging. Able to maintain lean execution but add functionality as necessary, like Ambassadors and Sidecar containers.

Exercise 6.2: Designing Applications With Duration: Create a Job

While most applications are deployed such that they continue to be available there are some which we may want to run a particular number of times called a Job, and others on a regular basis called a CronJob

1. Create a job which will run a container which sleeps for three seconds then stops.

```
student@master:~$ vim job.yaml
```

YAML

job.yaml

```
1 apiVersion: batch/v1
2 kind: Job
3 metadata:
4   name: sleepy
5 spec:
6   template:
7     spec:
8       containers:
9       - name: resting
10         image: busybox
11         command: ["/bin/sleep"]
12         args: ["3"]
13         restartPolicy: Never
```

2. Create the job, then verify and view the details. The example shows checking the job three seconds in and then again after it has completed. You may see different output depending on how fast you type.

```
student@master:~$ kubectl create -f job.yaml
```

```
1 job.batch/sleepy created
```

```
student@master:~$ kubectl get job
```

```

1 NAME      COMPLETIONS  DURATION  AGE
2 sleepy    0/1          3s        3s

```

```
student@master:~$ kubectl describe jobs.batch sleepy
```

```

1 Name:      sleepy
2 Namespace: default
3 Selector:   controller-uid=24c91245-d0fb-11e8-947a-42010a800002
4 Labels:     controller-uid=24c91245-d0fb-11e8-947a-42010a800002
5             job-name=sleepy
6 Annotations: <none>
7 Parallelism: 1
8 Completions: 1
9 Start Time:  Sun, 03 Nov 2019  04:22:50 +0000
10 Completed At: Sun, 03 Nov 2019  04:22:55 +0000
11 Duration:    5s
12 Pods Statuses: 0 Running / 1 Succeeded / 0 Failed
13 <output_omitted>

```

```
student@master:~$ kubectl get job
```

```

1 NAME      COMPLETIONS  DURATION  AGE
2 sleepy    1/1          5s        17s

```

3. View the configuration information of the job. There are three parameters we can use to affect how the job runs. Use `-o yaml` to see these parameters. We can see that `backoffLimit`, `completions`, and the `parallelism`. We'll add these parameters next.

```
student@master:~$ kubectl get jobs.batch sleepy -o yaml
```

```

1 <output_omitted>
2   uid: c2c3a80d-d0fc-11e8-947a-42010a800002
3 spec:
4   backoffLimit: 6
5   completions: 1
6   parallelism: 1
7   selector:
8     matchLabels:
9 <output_omitted>

```

4. As the job continues to AGE in a completion state, delete the job.

```
student@master:~$ kubectl delete jobs.batch sleepy
```

```
1 job.batch "sleepy" deleted
```

5. Edit the YAML and add the `completions:` parameter and set it to 5.

```
student@master:~$ vim job.yaml
```

YAML

job.yaml

```

1 <output_omitted>
2 metadata:
3   name: sleepy
4 spec:

```

YAML

```

5   completions: 5   #<--Add this line
6   template:
7     spec:
8       containers:
9 <output_omitted>

```

6. Create the job again. As you view the job note that COMPLETIONS begins as zero of 5.

```
student@master:~$ kubectl create -f job.yaml
```

```
1 job.batch/sleepy created
```

```
student@master:~$ kubectl get jobs.batch
```

```

1 NAME      COMPLETIONS  DURATION  AGE
2 sleepy    0/5          5s        5s

```

7. View the pods that running. Again the output may be different depending on the speed of typing.

```
student@master:~$ kubectl get pods
```

```

1 NAME                                READY  STATUS      RESTARTS  AGE
2 nginx-67f8fb575f-g4468              1/1    Running     2          2d
3 registry-56cffc98d6-xlhhf          1/1    Running     1          2d
4 sleepy-z5tnh                        0/1    Completed   0          8s
5 sleepy-zd692                        1/1    Running     0          3s
6 <output_omitted>

```

8. Eventually all the jobs will have completed. Verify then delete the job.

```
student@master:~$ kubectl get jobs
```

```

1 NAME      COMPLETIONS  DURATION  AGE
2 sleepy    5/5          26s       10m

```

```
student@master:~$ kubectl delete jobs.batch sleepy
```

```
1 job.batch "sleepy" deleted
```

9. Edit the YAML again. This time add in the `parallelism:` parameter. Set it to 2 such that two pods at a time will be deployed.

```
student@master:~$ vim job.yaml
```

YAML**job.yaml**

```

1 <output_omitted>
2   name: sleepy
3   spec:
4     completions: 5
5     parallelism: 2   #<-- Add this line
6     template:
7       spec:
8 <output_omitted>

```

10. Create the job again. You should see the pods deployed two at a time until all five have completed.

```
student@master:~$ kubectl create -f job.yaml
```

```
student@master:~$ kubectl get pods
```

	NAME	READY	STATUS	RESTARTS	AGE
1	nginx-67f8fb575f-g4468	1/1	Running	2	2d
2	registry-56cffc98d6-xlhhf	1/1	Running	1	2d
3	sleepy-8xwpc	1/1	Running	0	5s
4	sleepy-xjqnf	1/1	Running	0	5s
5	try1-c9cb54f5d-b45gl	2/2	Running	0	8h
6	<output_omitted>				

```
student@master:~$ kubectl get jobs
```

	NAME	COMPLETIONS	DURATION	AGE
1	sleepy	3/5	11s	11s
2				

11. Add a parameter which will stop the job after a certain number of seconds. Set the `activeDeadlineSeconds:` to 15. The job and all pods will end once it runs for 15 seconds.

```
student@master:~$ vim job.yaml
```

YAML

job.yaml

```
1 <output_omitted>
2   completions: 5
3   parallelism: 2
4   activeDeadlineSeconds: 15  #<-- Add this line
5   template:
6     spec:
7       containers:
8       - name: resting
9         image: busybox
10        command: ["/bin/sleep"]
11        args: ["3"]
12 <output_omitted>
```

12. Delete and recreate the job again. It should run for four times then continue to age without further completions.

```
student@master:~$ kubectl delete jobs.batch sleepy
```

```
1 job.batch "sleepy" deleted
```

```
student@master:~$ kubectl create -f job.yaml
```

```
1 job.batch/sleepy created
```

```
student@master:~$ kubectl get jobs
```

	NAME	COMPLETIONS	DURATION	AGE
1	sleepy	2/5	6s	6s
2				

```
student@master:~$ kubectl get jobs
```

	NAME	COMPLETIONS	DURATION	AGE
1	sleepy	4/5	16s	16s
2				

- View the message: entry in the Status section of the object YAML output. You may see less status if the job has yet to run. Wait and try again, if so.

```
student@master:~$ kubectl get job sleepy -o yaml
```

```
1 <output_omitted>
2 status:
3   conditions:
4   - lastProbeTime: "2019-11-03T16:06:10Z"
5     lastTransitionTime: "2019-11-03T16:06:10Z"
6     message: Job was active longer than specified deadline
7     reason: DeadlineExceeded
8     status: "True"
9     type: Failed
10  failed: 1
11  startTime: "2019-11-03T16:05:55Z"
12  succeeded: 4
```

- Delete the job.

```
student@master:~$ kubectl delete jobs.batch sleepy
```

```
1 job.batch "sleepy" deleted
```

Exercise 6.3: Designing Applications With Duration: Create a CronJob

A CronJob creates a watch loop which will create a batch job on your behalf when the time becomes true. We will use our existing Job file to start.

- Copy the Job file to a new file.

```
student@master:~$ cp job.yaml cronjob.yaml
```

- Edit the file to look like the annotated file shown below.

```
student@master:~$ vim cronjob.yaml
```

YAML

cronjob.yaml

```
1 apiVersion: batch/v1beta1    #<-- Add beta1 to be v1beta1
2 kind: CronJob    #<-- Change this line
3 metadata:
4   name: sleepy
5 spec:
6   schedule: "*/2 * * * *"    #<-- Remove completions:, parallelism:, and activeDeadlineSeconds:
7   jobTemplate:               #<-- Add Linux style cronjob syntax
8     spec:                   #<-- New jobTemplate and spec
9       template:             #<-- This and following lines space four to right
10        spec:
11          containers:
12            - name: resting
13              image: busybox
14              command: ["/bin/sleep"]
15              args: ["3"]
16          restartPolicy: Never
```

3. Create the new CronJob. View the jobs. It will take two minutes for the CronJob to run and generate a new batch Job.

```
student@master:~$ kubectl create -f cronjob.yaml
```

```
1 cronjob.batch/sleepy created
```

```
student@master:~$ kubectl get cronjobs.batch
```

```
1 NAME          SCHEDULE      SUSPEND   ACTIVE   LAST SCHEDULE   AGE
2 sleepy        */2 * * * *   False    0        <none>          8s
```

```
student@master:~$ kubectl get job
```

```
1 No resources found in default namespace.
```

4. After two minutes you should see jobs start to run.

```
student@master:~$ kubectl get cronjobs.batch
```

```
1 NAME          SCHEDULE      SUSPEND   ACTIVE   LAST SCHEDULE   AGE
2 sleepy        */2 * * * *   False    0        21s             2m1s
```

```
student@master:~$ kubectl get jobs.batch
```

```
1 NAME                COMPLETIONS   DURATION   AGE
2 sleepy-1539722040    1/1           5s         18s
```

```
student@master:~$ kubectl get jobs.batch
```

```
1 NAME                COMPLETIONS   DURATION   AGE
2 sleepy-1539722040    1/1           5s         5m17s
3 sleepy-1539722160    1/1           6s         3m17s
4 sleepy-1539722280    1/1           6s         77s
```

5. Ensure that if the job continues for more than 10 seconds it is terminated. We will first edit the **sleep** command to run for 30 seconds then add the `activeDeadlineSeconds`: entry to the container.

```
student@master:~$ vim cronjob.yaml
```

Y
A
ML

cronjob.yaml

```
1 ....
2   jobTemplate:
3     spec:
4       template:
5         spec:
6           activeDeadlineSeconds: 10 #<-- Add this line
7           containers:
8             - name: resting
9             ....
10          command: ["/bin/sleep"]
11          args: ["30"] #<-- Edit this line
12          restartPolicy: Never
```

6. Delete and recreate the CronJob. It may take a couple of minutes for the batch Job to be created and terminate due to the timer.

```
student@master:~$ kubectl delete cronjobs.batch sleepy
```

```
1 cronjob.batch "sleepy" deleted
```

```
student@master:~$ kubectl create -f cronjob.yaml
```

```
1 cronjob.batch/sleepy created
```

```
student@master:~$ sleep 120 ; kubectl get jobs
```

```
1 NAME                COMPLETIONS  DURATION  AGE
2 sleepy-1539723240    0/1          61s       61s
```

```
student@master:~$ kubectl get cronjobs.batch
```

```
1 NAME    SCHEDULE    SUSPEND    ACTIVE    LAST SCHEDULE    AGE
2 sleepy  */2 * * * *  False     1         72s           94s
```

```
student@master:~$ kubectl get jobs
```

```
1 NAME                COMPLETIONS  DURATION  AGE
2 sleepy-1539723240    0/1          75s       75s
```

```
student@master:~$ kubectl get jobs
```

```
1 NAME                COMPLETIONS  DURATION  AGE
2 sleepy-1539723240    0/1          2m19s     2m19s
3 sleepy-1539723360    0/1          19s       19s
```

```
student@master:~$ kubectl get cronjobs.batch
```

```
1 NAME    SCHEDULE    SUSPEND    ACTIVE    LAST SCHEDULE    AGE
2 sleepy  */2 * * * *  False     2         31s           2m53s
```

7. Clean up by deleting the CronJob.

```
student@master:~$ kubectl delete cronjobs.batch sleepy
```

```
1 cronjob.batch "sleepy" deleted
```

Exercise 6.4: Using Labels

Create and work with labels. We will understand how the deployment, replicaSet, and pod labels interact.

1. Create a new deployment called design2

```
student@master:~$ kubectl create deployment design2 --image=nginx
```

```
1 deployment.apps/design2 created
```

2. View the wide **kubectl get** output for the design2 deployment and make note of the SELECTOR

```
student@master:~$ kubectl get deployments.apps design2 -o wide
```

```
1 NAME    READY  UP-TO-DATE  AVAILABLE  AGE    CONTAINERS  IMAGES  SELECTOR
2 design2 1/1    1           1          2m13s  nginx       nginx   app=design2
```

3. Use the `-l` option to use the selector to list the pods running inside the deployment. There should be only one pod running.

```
student@master:~$ kubectl get -l app=design2 pod
```

	NAME	READY	STATUS	RESTARTS	AGE
1	design2-766d48574f-5w274	1/1	Running	0	3m1s

4. View the pod details in YAML format using the deployment selector. This time use the `--selector` option. Find the pod label in the output. It should match that of the deployment.

```
student@master:~$ kubectl get --selector app=design2 pod -o yaml
```

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   annotations:
5     cnf.projectcalico.org/podIP: 192.168.113.222/32
6   creationTimestamp: "2020-01-31T16:29:37Z"
7   generateName: design2-766d48574f-
8   labels:
9     app: design2
10    pod-template-hash: 766d48574f
11 ....
```

5. Edit the pod label to be your favorite color.

```
student@master:~$ kubectl edit pod design2-766d48574f-5w274
```

YAML

```
1 ....
2   labels:
3     app: orange                                #<<-- Edit this line
4     pod-template-hash: 766d48574f
5     name: design2-766d48574f-5w274
6 ....
```

6. Now view how many pods are in the deployment. Then how many have design2 in their name. Note the AGE of the pods.

```
student@master:~$ kubectl get deployments.apps design2 -o wide
```

	NAME	READY	UP-TO-DATE	AVAILABLE	AGE	CONTAINERS	IMAGES	SELECTOR
1	design2	1/1	1	1	56s	nginx	nginx	app=design2

```
student@master:~$ kubectl get pods | grep design2
```

1	design2-766d48574f-5w274	1/1	Running	0	82s
2	design2-766d48574f-xttgg	1/1	Running	0	2m12s

7. Delete the design2 deployment.

```
student@master:~$ kubectl delete deploy design2
```

```
1 deployment.apps "design2" deleted
```

8. Check again for pods with design2 in their names. You should find one pod, with an AGE of when you first created the deployment. Once the label was edited the deployment created a new pod in order that the status matches the spec and there be a replica running with the intended label.

```
student@master:~$ kubectl get pods | grep design2
```



```
1 design2-766d48574f-5w274 1/1 Running 0 38m
```

9. Delete the pod using the `-l` and the label you edited to be your favorite color in a previous step. The command details have been omitted. Use previous steps to figure out these commands.

Exercise 6.5: Setting Pod Resource Limits and Requirements

1. Create a new pod running the `vish/stress` image. A YAML `stress.yaml` file has been included in the course tarball.
2. Run the `top` command on the master and worker nodes. You should find a `stress` command consuming the majority of the CPU on one node, the worker. Use `ctrl-c` to exit from `top`. Delete the deployment.
3. Edit the `stress.yaml` file add in the following limits and requests.

```
student@master:~$
```

YAML

```
1 ....
2     name: stressmeout
3     resources:                                #<--- Add this and following six lines
4         limits:
5             cpu: "1"
6             memory: "1Gi"
7         requests:
8             cpu: "0.5"
9             memory: "500Mi"
10    args:
11    - -cpus
12    ....
```

4. Create the deployment again. Check the status of the pod. You should see that it shows an `OOMKilled` status and a growing number of restarts. You may see a status of `Running` if you catch the pod in early in a restart. If you wait long enough you may see `CrashLoopBackOff`.

```
student@master:~$ kubectl get pod stressmeout-7fbbbcc887-v9kvb
```

```
1 NAME                                READY   STATUS    RESTARTS   AGE
2 stressmeout-7fbbbcc887-v9kvb        0/1     OOMKilled  2          32s
```

5. Delete then edit the deployment. Change the `limit:` parameters such that pod is able to run, but not too much extra resources. Try setting the memory limit to exactly what the stress command requests.

As we allow the pod to run on the master node, this could cause issues, such as the `kube-apiserver` restarting due to lack of resources. We will also add a `nodeSelector` to use the built in label of `kubernetes.io/hostname`.

```
student@master:~$ kubectl delete -f stress.yaml
```

```
student@master:~$ vim stress.yaml
```

YAML

```
1 ....
2     spec:
3         nodeSelector:                        #<--- Uncomment and edit
4             kubernetes.io/hostname: worker #<--- to by YOUR worker hostname
5         containers:
6
7     ....
8     resources:
```

YAML

```

9      limits:
10         cpu: "2"
11         memory: "2Gi"
12     requests:
13         ....

```

6. Create the deployment and ensure the pod runs without error. Use **top** to verify the stress command is running on one of the nodes and view the pod details to ensure the CPU and memory limits are in use. Also use the **kubectl describe node** command to view the resources your master and worker node are using. The command details have been omitted. Use previous steps to figure out the commands.

Exercise 6.6: Domain Review



Very Important

The source pages and content in this review could change at any time. **IT IS YOUR RESPONSIBILITY TO CHECK THE CURRENT INFORMATION.**

Revisit the CKAD domain list on [Curriculum Overview](#) and locate some of the topics we have covered in this chapter. They may be in multiple sections. The graphic below shows the topics covered in this chapter.

- Define an application's resource requirements
- Understand Jobs and CronJobs
- Understand how to use Labels, Selectors, and Annotations

Figure 6.1: Multiple Domain

Focus on ensuring you have all the necessary files and processes understood first. Repeat the review until you are sure you have bookmarks of necessary YAML samples and can complete each step quickly, and ensure each object is running properly.

1. Find and use the `design-review1.yaml` file to create a pod.
2. Determine the CPU and memory resource requirements of `design-pod1`.
3. Edit the pod resource requirements such that the CPU limit is exactly twice the amount requested by the container. (Hint: subtract .22)
4. Increase the memory resource limit of the pod until the pod shows a `Running` status. This may require multiple edits and attempts. Determine the minimum amount necessary for the `Running` status to persist at least a minute.
5. Use the `design-review2.yaml` file to create several pods with various labels.
6. Using **only** the `--selector` value `tux` to delete only those pods. This should be half of the pods. Hint, you will need to view pod settings to determine the key value as well.
7. Create a new cronjob which runs `busybox` and the `sleep 30` command. Have the cronjob run every three minutes. View the job status to check your work. Change the settings so the pod runs 10 minutes from the current time, every week. For example, if the current time was 2:14PM, I would configure the job to run at 2:24PM, every Monday.
8. Delete any objects created during this review. You may want to delete all but the cronjob if you'd like to see if it runs in 10 minutes. Then delete that object as well.

Chapter 7

Deployment Configuration



In this chapter, we are going to cover some of the typical tasks necessary for full application deployment inside of Kubernetes. We will begin by talking about attaching storage to our containers. By default, storage would only last as long as the container using it. Through the use of Persistent Volumes and Persistent Volume Claims, we can attach storage such that it lives longer than the container and/or the pod the container runs inside of. We'll cover adding dynamic storage, which might be made available from a cloud provider, and we will talk about ConfigMaps and secrets. These are ways of attaching configuration information into our containers in a flexible manner. They are ingested in almost the exact same way. The difference between them, currently, is a secret is encoded. There are plans for it to be encrypted in the future. We will also talk about how do we update our applications. We can do a rolling update, where there's always some containers available for response to client requests, or update all of them at once. We'll also cover how to roll back an application to a previous version, looking at our entire history of updates, and choosing a particular one.

7.1	Volumes Overview	123
7.2	Introducing Volumes	124
7.3	Volume Spec	125
7.4	Volume Types	126
7.5	Shared Volume Example	126
7.6	Persistent Volumes and Claims	127
7.7	Persistent Volume	128
7.8	Persistent Volume Claim	128
7.9	Dynamic Provisioning	129
7.10	Secrets	130
7.11	Using Secrets via Environment Variables	130
7.12	Mounting Secrets as Volumes	131
7.13	Portable Data with ConfigMaps	131
7.14	Using ConfigMaps	132
7.15	Deployment Configuration Status	132
7.16	Scaling and Rolling Updates	133
7.17	Deployment Rollbacks	134
7.18	Labs	134

7.1 Volumes Overview

- Container engines have traditionally not offered storage that outlives the container. As containers are considered transient, this could lead to a loss of data, or complex exterior storage options. A Kubernetes volume shares at least the Pod

lifetime, not the containers within. Should a container terminate, the data would continue to be available to the new container. A volume can persist longer than a pod, and can be accessed by multiple pods, using `PersistentVolumeClaims`. This allows for **state persistence**.

- A volume is a directory, possibly pre-populated, made available to containers in a Pod. The creation of the directory, the backend storage of the data and the contents depend on the volume type. As of v1.14, there are 28 different volume types ranging from `rbd` to `Ceph`, to `NFS`, to dynamic volumes from a cloud provider like Google's `gcePersistentDisk`. Each has particular configuration options and dependencies.
- Adoption of **Container Storage Interface (CSI)** enables the goal of an industry standard interface for container orchestration to allow access to arbitrary storage systems. Currently, volume plugins are "in-tree", meaning they are compiled and built with the core Kubernetes binaries. This "out-of-tree" object will allow storage vendors to develop a single driver and allow the plugin to be containerized. This will replace the existing Flex plugin which requires elevated access to the host node, a large security concern.
- Should you want your storage lifetime to be distinct from a Pod, you can use **Persistent Volumes**. These allow for empty or pre-populated volumes to be claimed by a Pod using a **Persistent Volume Claim**, then outlive the Pod. Data inside the volume could then be used by another Pod, or as a means of retrieving data.
- There are two API Objects which exist to provide data to a Pod already. Encoded data can be passed using a `Secret` and non-encoded data can be passed with a `ConfigMap`. These can be used to pass important data like SSH keys, passwords, or even a configuration file like `/etc/hosts`.

7.2 Introducing Volumes

- A Pod specification can declare one or more volumes and where they are made available. Each requires a name, a type, and a mount point.
- Keeping acquired data or ingesting it into other containers is a common task, typically requiring the use of a **Persistent Volume Claim (PVC)**.
- The same volume can be made available to multiple containers within a Pod, which can be a method of container-to-container communication. A volume can be made available to multiple Pods, with each given an `access mode` to write. There is no concurrency checking, which means data corruption is probable, unless outside locking takes place.
- A particular `access mode` is part of a Pod request. As a request, the user may be granted more, but not less access, though a direct match is attempted first. The cluster groups volumes with the same mode together, then sorts volumes by size, from smallest to largest. The claim is checked against each in that access mode group, until a volume of sufficient size matches. The three access modes are `RWO` (`ReadWriteOnce`), which allows read-write by a single node, `ROX` (`ReadOnlyMany`), which allows read-only by multiple nodes, and `RWX` (`ReadWriteMany`), which allows read-write by many nodes.
- We had seen an example of containers within a pod while learning about networking. Using the same example we see that the same volume can be used by all containers. In this case both are using the `/data/` directory. `MainApp` reads and writes to the volume, and `Logger` uses the volume read-only. It then writes to a different volume `/log/`.




Figure 7.1: Multiple volumes in a pod

- When a volume is requested, the local **kubelet** uses the `kubelet.pods.go` script to map the raw devices, determine and make the mount point for the container, then create the symbolic link on the host node filesystem to associate the storage to the container. The API server makes a request for the storage to the StorageClass plugin, but the specifics of the requests to the backend storage depend on the plugin in use.
- If a request for a particular StorageClass was not made, then the only parameters used will be access mode and size. The volume could come from any of the storage types available, and there is no configuration to determine which of the available ones will be used.

7.3 Volume Spec

- One of the many types of storage available is an `emptyDir`. The kubelet will create the directory in the container, but not mount any storage. Any data created is written to the shared container space. As a result, it would not be persistent storage. When the Pod is destroyed, the directory would be deleted along with the container.

YAML

```

1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: busybox
5   namespace: default
6 spec:
7   containers:
8   - image: busybox
9     name: busy
10    command:
11      - sleep
12      - "3600"
13    volumeMounts:
14      - mountPath: /scratch
15        name: scratch-volume
16  volumes:
17  - name: scratch-volume
18    emptyDir: {}

```

- The YAML file above would create a Pod with a single container with a volume named `scratch-volume` created, which would create the `/scratch` directory inside the container.

7.4 Volume Types

- There are several types that you can use to define volumes, each with their pros and cons. Some are local, and many make use of network-based resources.
- In GCE or AWS, you can use volumes of type `GCEPersistentDisk` or `awsElasticBlockStore`, which allows you to mount GCE and EBS disks in your Pods, assuming you have already set up accounts and privileges.
- `emptyDir` and `hostPath` volumes are easy to use. As mentioned, `emptyDir` is an empty directory that gets erased when the Pod dies, but is recreated when the container restarts. The `hostPath` volume mounts a resource from the host node filesystem. The resource could be a directory, file socket, character, or block device. These resources must already exist on the host to be used. There are two types, `DirectoryOrCreate` and `FileOrCreate`, which create the resources on the host, and use them if they don't already exist.
- NFS (Network File System) and iSCSI (Internet Small Computer System Interface) are straightforward choices for multiple readers scenarios.
- `rbd` for block storage or `CephFS` and `GlusterFS`, if available in your Kubernetes cluster, can be a good choice for multiple writer needs.
- Besides the volume types we just mentioned, there are many other possible, with more being added: `azureDisk`, `azureFile`, `csi`, `downwardAPI`, `fc` (fibre channel), `flocker`, `gitRepo`, `local`, `projected`, `portworxVolume`, `quobyte`, `scaleIO`, `secret`, `storageos`, `vsphereVolume`, `persistentVolumeClaim`, etc.

7.5 Shared Volume Example

- The following YAML file, for pod `exampleA` creates a pod with two containers, one called `alphacont`, the other called `betacont` both with access to a shared volume called `sharevol`:

YAML

```

1 ....
2   containers:
3   - name: alphacont
4     image: busybox

```

YAML

```

5  volumeMounts:
6    - mountPath: /alphadir
7      name: sharevol
8  - name: betacont
9    image: busybox
10   volumeMounts:
11     - mountPath: /betadir
12       name: sharevol
13   volumes:
14     - name: sharevol
15       emptyDir: {}

```

```

$ kubectl exec -ti exampleA -c betacont -- touch /betadir/foobar
$ kubectl exec -ti exampleA -c alphacont -- ls -l /alphadir

```

```

1 total 0
2 -rw-r--r-- 1 root root 0 Nov 19 16:26 foobar

```

- You could use `emptyDir` or `hostPath` easily, since those types do not require any additional setup, and will work in your Kubernetes cluster.
- Note that one container wrote, and the other container had immediate access to the data. There is nothing to keep the containers from overwriting the other's data. Locking or versioning considerations must be part of the application to avoid corruption.

7.6 Persistent Volumes and Claims

- A persistent volume (pv) is a storage abstraction used to retain data longer than the Pod using it. Pods define a volume of type `persistentVolumeClaim` (pvc) with various parameters for size and possibly the type of backend storage known as its `StorageClass`. The cluster then attaches the persistentVolume.
- Kubernetes will dynamically use volumes that are available, irrespective of its storage type, allowing claims to any backend storage.
- There are several phases to persistent storage:
 - **Provisioning** can be from pvs created in advance by the cluster administrator, or requested from a dynamic source, such as the cloud provider.
 - **Binding** occurs when a control loop on the master notices the PVC, containing an amount of storage, access request, and optionally, a particular `StorageClass`. The watcher locates a matching PV or waits for the `StorageClass` provisioner to create one. The pv must match at least the storage amount requested, but may provide more.
 - The **use** phase begins when the bound volume is mounted for the Pod to use, which continues as long as the Pod requires.
 - **Releasing** happens when the Pod is done with the volume and an API request is sent, deleting the PVC. The volume remains in the state from when the claim is deleted until available to a new claim. The resident data remains depending on the `persistentVolumeReclaimPolicy`.
 - The **reclaim** phase has three options:
 - * **Retain**, which keeps the data intact, allowing for an administrator to handle the storage and data.
 - * **Delete** tells the volume plugin to delete the API object, as well as the storage behind it.
 - * The **Recycle** option runs an `rm -rf /mountpoint` and then makes it available to a new claim. With the stability of dynamic provisioning, the `Recycle` option is planned to be deprecated.

```

$ kubectl get pv
$ kubectl get pvc

```

7.7 Persistent Volume

- The following example shows a basic declaration of a PersistentVolume using the hostPath type.

YAML

```
1 kind: PersistentVolume
2 apiVersion: v1
3 metadata:
4   name: 10Gpv01
5   labels:
6     type: local
7 spec:
8   capacity:
9     storage: 10Gi
10  accessModes:
11    - ReadWriteOnce
12  hostPath:
13    path: "/somepath/data01"
```

- Each type will have its own configuration settings. For example, an already created Ceph or GCE Persistent Disk would not need to be configured, but could be claimed from the provider.
- Persistent volumes are cluster-scoped, but persistent volume claims are namespace-scoped. An alpha feature since v1.11 this allows for static provisioning of Raw Block Volumes, which currently support the Fibre Channel plugin. There is a lot of development and change in this area, with plugins adding dynamic provisioning.

7.8 Persistent Volume Claim

- With a persistent volume created in your cluster, you can then write a manifest for a claim and use that claim in your pod definition. In the Pod, the volume uses the persistentVolumeClaim.

YAML

```
1 kind: PersistentVolumeClaim
2 apiVersion: v1
3 metadata:
4   name: myclaim
5 spec:
6   accessModes:
7     - ReadWriteOnce
8   resources:
9     requests:
10       storage: 8Gi
11 (In the Pod)
12 ....
13 spec:
14   containers:
15     ....
16   volumes:
17     - name: test-volume
18       persistentVolumeClaim:
19         claimName: myclaim
```

- The Pod configuration could also be as complex as this:

YAML

```

1 volumeMounts:
2   - name: Cephpd
3     mountPath: /data/rbd
4   volumes:
5     - name: rbdpd
6       rbd:
7         monitors:
8           - '10.19.14.22:6789'
9           - '10.19.14.23:6789'
10          - '10.19.14.24:6789'
11         pool: k8s
12         image: client
13         fsType: ext4
14         readOnly: true
15         user: admin
16         keyring: /etc/ceph/keyring
17         imageformat: "2"
18         imagefeatures: "layering"

```

Point to Ponder

If you had one application ingesting data but also want to archive, ingest into a data lake, and forward the data, how would you use persistent volume claims?

7.9 Dynamic Provisioning

- While handling volumes with a persistent volume definition and abstracting the storage provider using a claim is powerful, a cluster administrator still needs to create those volumes in the first place. Starting with Kubernetes v1.4, **Dynamic Provisioning** allowed for the cluster to request storage from an exterior, pre-configured source. API calls made by the appropriate plugin allow for a wide range of dynamic storage use.
- The StorageClass API resource allows an administrator to define a persistent volume provisioner of a certain type, passing storage-specific parameters.
- With a StorageClass created, a user can request a claim, which the API Server fills via auto-provisioning. The resource will also be reclaimed as configured by the provider. AWS and GCE are common choices for dynamic storage, but other options exist, such as a Ceph cluster or iSCSI. Single, default class is possible via annotation.
- Here is an example of a StorageClass using GCE:

YAML

```

1 apiVersion: storage.k8s.io/v1
2 kind: StorageClass
3 metadata:
4   name: fast # Could be any name
5   provisioner: kubernetes.io/gce-pd
6   parameters:
7     type: pd-ssd

```

- There are also providers you can create inside of your cluster, which use Custom Resource Definitions and local operators to manage storage.
 - **Rook** - From <https://rook.io> is a CNCF project in incubation status which allows easy management of several back-end storage types, including an easy way to deploy **Ceph**
 - **Longhorn** - From <https://longhorn.io> is a CNCF project in sandbox status, developed by **Rancher** which also allows management of local storage.

7.10 Secrets

- Pods can access local data using volumes, but there is some data you don't want readable to the naked eye. Passwords may be an example. Someone reading through a YAML file may read a password and remember it. Using the Secret API resource, the same password could be encoded. A casual reading would not give away the password. You can create, get, or delete secrets:

```
$ kubectl get secrets
```

- Secrets can be manually encoded with `kubectl create secret`:

```
$ kubectl create secret generic --help
```

```
$ kubectl create secret generic mysql --from-literal=password=root
```

- A secret is not encrypted by default, only base64-encoded. You can see the encoded string inside the secret with **kubectl**. The secret will be decoded and be presented as a string saved to a file. The file can be used as an environmental variable or in a new directory, similar to the presentation of a volume.

In order to encrypt secrets one must create a `EncryptionConfiguration` object with a key and proper identity. Then the **kube-apiserver** needs the `--encryption-provider-config` flag set to a previously configured provider such as `aescbc` or `kms`. Once this is enabled you need to recreate every secret as they are encrypted upon write. Multiple keys are possible. Each key for a provider is tried during decryption. The first key of the first provider is used for encryption. To rotate keys first create a new key, restart (all) `kube-apiserver` processes, then recreate every secret.

- A secret can be made manually as well, then inserted into a YAML file:

```
$ echo LFTr@1n | base64
```

```
1 TEZUckAxbgo=
```

```
$ vim secret.yaml
```

YAML

```
1 apiVersion: v1
2 kind: Secret
3 metadata:
4   name: LF-secret
5 data:
6   password: TEZUckAxbgo=
```

- Prior to Kubernetes v1.18 secrets (and configMaps) were automatically updated. This could lead to issues if a configuration were updated and a pod restarted it may be configured differently than other replicas. In newer versions these objects can be made immutable.

7.11 Using Secrets via Environment Variables

- A secret can be used as an environmental variable in a Pod. You can see one being configured in the following example:

YAML

```
1 ...
2 spec:
3   containers:
4     - image: mysql:5.5
5       env:
6         - name: MYSQL_ROOT_PASSWORD
7           valueFrom:
```

YAML

```

8     secretKeyRef:
9         name: mysql
10        key: password
11        name: mysql

```

- There is no limit to the number of Secrets used, but there is a 1MB limit to their size. Each secret occupies memory, along with other API objects, so very large numbers of secrets could deplete memory on a host.
- They are stored in the **tmpfs** storage on the host node, and are only sent to the host running Pod. All volumes requested by a Pod must be mounted before the containers within the Pod are started. So, a secret must exist prior to being requested.

7.12 Mounting Secrets as Volumes

- You can also mount secrets as files using a volume definition in a pod manifest. The mount path will contain a file whose name will be the key of the secret created with the `kubectl create secret` step earlier.

YAML

```

1  ...
2  spec:
3      containers:
4      - image: busybox
5        command:
6            - sleep
7            - "3600"
8        volumeMounts:
9            - mountPath: /mysqlpassword
10              name: mysql
11        name: busy
12      volumes:
13      - name: mysql
14        secret:
15            secretName: mysql

```

- Once the pod is running, you can verify that the secret is indeed accessible in the container:

```
$ kubectl exec -ti busybox -- cat /mysqlpassword/password
```

```
1 LFTTr@in
```

7.13 Portable Data with ConfigMaps

- A similar API resource to Secrets is the ConfigMap, except the data is not encoded. In keeping with the concept of decoupling in Kubernetes, using a ConfigMap decouples a container image from configuration artifacts.
- They store data as sets of key-value pairs or plain configuration files in any format. The data can come from a collection of files or all files in a directory. It can also be populated from a literal value.
- A ConfigMap can be used in several different ways. A Pod can use the data as environmental variables from one or more sources. The values contained inside can be passed to commands inside the pod. A Volume or a file in a Volume can be created, including different names and particular access modes. In addition, cluster components like controllers can use the data.

- Let's say you have a file on your local filesystem called `config.js`. You can create a ConfigMap that contains this file. The configmap object will have a data section containing the content of the file:

YAML

```
1 $ kubectl get configmap foobar -o yaml
2 kind: ConfigMap
3 apiVersion: v1
4 metadata:
5   name: foobar
6 data:
7   config.js: |
8   {
9   ...
```

- ConfigMaps can be consumed in various ways:
 - Pod environmental variables from single or multiple ConfigMaps
 - Use ConfigMap values in Pod commands
 - Populate Volume from ConfigMap
 - Add ConfigMap data to specific path in Volume
 - Set file names and access mode in Volume from ConfigMap data
 - Can be used by system components and controllers.

7.14 Using ConfigMaps

- Like secrets, you can use ConfigMaps as environment variables or using a volume mount. They must exist prior to being used by a Pod, unless marked as `optional`. They also reside in a specific namespace.
- In the case of environment variables, your pod manifest will use the `valueFrom` key and the `configMapKeyRef` value to read the values. For instance:

YAML

```
1 env:
2 - name: SPECIAL_LEVEL_KEY
3   valueFrom:
4     configMapKeyRef:
5       name: special-config
6       key: special.how
```

- With volumes, you define a volume with the `configMap` type in your pod and mount it where it needs to be used.

YAML

```
1 volumes:
2 - name: config-volume
3   configMap:
4     name: special-config
```

7.15 Deployment Configuration Status

- The Status output is generated when the information is requested:

YAML

```

1 status:
2   availableReplicas: 2
3   conditions:
4   - lastTransitionTime: 2017-12-21T13:57:07Z
5     lastUpdateTime: 2017-12-21T13:57:07Z
6     message: Deployment has minimum availability.
7     reason: MinimumReplicasAvailable
8     status: "True"
9     type: Available
10  observedGeneration: 2
11  readyReplicas: 2
12  replicas: 2
13  updatedReplicas: 2

```

- The output above shows what the same deployment were to look like if the number of replicas were increased to two. The times are different than when the deployment was first generated.
 - `availableReplicas`
Indicates how many were configured by the ReplicaSet and to be compared to the later value of `readyReplicas`, which would be used to determine if all replicas have been fully generated and without error.
 - `observedGeneration`
Shows how often the deployment has been updated. This information can be used to understand the rollout and rollback situation of the deployment.

7.16 Scaling and Rolling Updates

- The API server allows for the configurations settings to be updated for most values. There are some immutable values, which may be different depending on the version of Kubernetes you have deployed.
- A common update is to change the number of replicas running. If this number is set to zero, there would be no containers, but there would still be a ReplicaSet and Deployment. This is the backend process when a Deployment is deleted.

```
$ kubectl scale deploy/dev-web --replicas=4
```

```
1 deployment "dev-web" scaled
```

```
$ kubectl get deployments
```

```

1 NAME          DESIRED  CURRENT  UP-TO-DATE  AVAILABLE  AGE
2 dev-web       4        4        4           1          12m

```

- Non-immutable values can be edited via a text editor, as well. Use **edit** to trigger an update. For example, to change the deployed version of the **nginx** web server to an older version:

```
$ kubectl edit deployment nginx
```

YAML

```

1 ....
2   containers:
3   - image: nginx:1.8 #<---Set to an older version
4     imagePullPolicy: IfNotPresent
5     name: dev-web
6   ....

```

- This would trigger a rolling update of the deployment. While the deployment would show an older age, a review of the Pods would show a recent update and older version of the web server application deployed.

7.17 Deployment Rollbacks

- With all the ReplicaSets of a Deployment being kept, you can also roll back to a previous revision by scaling up and down the ReplicaSets the other way. Next, we will have a closer look at rollbacks, using the `--record` option of the `kubectl` command, which allows annotation in the resource definition. The `create` generator does not have a record function.

```
$ kubectl set image deployment ghost --image=ghost:0.9 --record
$ kubectl get deployments ghost -o yaml
```

```
1 metadata:
2   annotations:
3     deployment.kubernetes.io/revision: "1"
4     kubernetes.io/change-cause: kubectl set image deployment ghost --image=ghost:0.9 --record
```

- Should an update fail, due to an improper image version, for example, you can roll back the change to a working version with `kubectl rollout undo`:

```
$ kubectl set image deployment/ghost ghost=ghost:0.9 --all
$ kubectl rollout history deployment/ghost deployments "ghost":
```

```
1 REVISION    CHANGE-CAUSE
2 1           <none>
3 2           kubectl set image deployment/ghost ghost=ghost:0.9 --all
```

```
$ kubectl get pods
```

```
1 NAME                                READY   STATUS    RESTARTS   AGE
2 ghost-2141819201-tcths              0/1     ImagePullBackOff    0           1m
```

```
$ kubectl rollout undo deployment/ghost
$ kubectl get pods
```

```
1 NAME                                READY   STATUS    RESTARTS   AGE
2 ghost-3378155678-eq5i6              1/1     Running    0           7s
```

- You can roll back to a specific revision with the `--to-revision=2` option. You can also edit a Deployment using the `kubectl edit` command. One can also pause a Deployment, and then resume.

```
$ kubectl rollout pause deployment/ghost
$ kubectl rollout resume deployment/ghost
```

- Please note that you can still do a rolling update on ReplicationControllers with the `kubectl rolling-update` command, but this is done on the client side. Hence, if you close your client, the rolling update will stop.

7.18 Labs

Exercise 7.1: Configure the Deployment: Secrets and ConfigMap



Very Important

Save a copy of your `$HOME/app1/simpleapp.yaml` file, in case you would like to repeat portions of the labs, or you find your file difficult to use due to typos and whitespace issues.

```
student@master:~$ cp $HOME/app1/simpleapp.yaml $HOME/beforeLab5.yaml
```

Overview

In this lab we will add resources to our deployment with further configuration you may need for production.

There are three different ways a **ConfigMap** can ingest data, from a literal value, from a file, or from a directory of files.

1. Create a **ConfigMap** containing primary colors. We will create a series of files to ingest into the **ConfigMap**. First create a directory **primary** and populate it with four files. Then we create a file in our home directory with our favorite color.

```
student@master:~/app1$ cd

student@master:~$ mkdir primary
student@master:~$ echo c > primary/cyan
student@master:~$ echo m > primary/magenta
student@master:~$ echo y > primary/yellow
student@master:~$ echo k > primary/black
student@master:~$ echo "known as key" >> primary/black
student@master:~$ echo blue > favorite
```

2. Generate a **configMap** using each of the three methods.

```
student@master:~$ kubectl create configmap colors \
  --from-literal=text=black \
  --from-file=./favorite \
  --from-file=./primary/
```

```
1 configmap/colors created
```

3. View the newly created **configMap**. Note the way the ingested data is presented.

```
student@master:~$ kubectl get configmap colors
```

```
1 NAME      DATA      AGE
2 colors    6          11s
```

```
student@master:~$ kubectl get configmap colors -o yaml
```

```
1 apiVersion: v1
2 data:
3   black: |
4     k
5     known as key
6   cyan: |
7     c
8   favorite: |
9     blue
10  magenta: |
11    m
12  text: black
13  yellow: |
14    y
15 kind: ConfigMap
16 metadata:
17  <output_omitted>
```

4. Update the YAML file of the application to make use of the **configMap** as an environmental parameter. Add the six lines from the `env:` line to `key:favorite`.

```
student@master:~$ vim $HOME/app1/simpleapp.yaml
```

YAML

simpleapp.yaml

```

1  ....
2      spec:
3          containers:
4              - image: 10.105.119.236:5000/simpleapp
5                env: #<-- Add from here
6                  - name: ilike
7                    valueFrom:
8                      configMapKeyRef:
9                        name: colors
10                       key: favorite #<-- to here
11          imagePullPolicy: Always
12  ....

```

5. Delete and re-create the deployment with the new parameters.

```
student@master-lab-7txt:~$ kubectl delete deployment try1
```

```
1 deployment.apps "try1" deleted
```

```
student@master-lab-7txt:~$ kubectl create -f $HOME/app1/simpleapp.yaml
```

```
1 deployment.apps/try1 created
```

6. Even though the try1 pod is not in a fully ready state, it is running and useful. Use **kubectl exec** to view a variable's value. View the pod state then verify you can see the ilike value within the simpleapp container. Note that the use of double dash (- -) tells the shell to pass the following as standard in.

```
student@master:~$ kubectl get pod
```

```
1 <output_omitted>
```

```
student@master:~$ kubectl exec -c simpleapp -it try1-5db9bc6f85-whxbf \
-- /bin/bash -c 'echo $ilike'
```

```
1 blue
```

7. Edit the YAML file again, this time adding the another method of using a **configMap**. Edit the file to add three lines. envFrom should be indented the same amount as env earlier in the file, and configMapRef should be indented the same as configMapKeyRef.

```
student@master:~$ vim $HOME/app1/simpleapp.yaml
```

YAML

simpleapp.yaml

```

1  ....
2          configMapKeyRef:
3              name: colors
4              key: favorite
5          envFrom: #<-- Add this and the following two lines
6          - configMapRef:
7              name: colors
8          imagePullPolicy: Always
9  ....

```

8. Again delete and recreate the deployment. Check the pods restart.


```
student@master:~$ kubectl delete deployment try1
```

```
1 deployment.apps "try1" deleted
```

```
student@master:~$ kubectl create -f $HOME/app1/simpleapp.yaml
```

```
1 deployment.apps/try1 created
```

```
student@master:~$ kubectl get pods
```

	NAME	READY	STATUS	RESTARTS	AGE
2	nginx-6b58d9cdfd-9fnl4	1/1	Running	1	23h
3	registry-795c6c8b8f-hl5w	1/1	Running	2	23h
4	try1-d4fbf76fd-46pkb	1/2	Running	0	40s
5	try1-d4fbf76fd-9kw24	1/2	Running	0	39s
6	try1-d4fbf76fd-bx9j9	1/2	Running	0	39s
7	try1-d4fbf76fd-jw8g7	1/2	Running	0	40s
8	try1-d4fbf76fd-lppl5	1/2	Running	0	39s
9	try1-d4fbf76fd-xtfd4	1/2	Running	0	40s

9. View the settings inside the try1 container of a pod. The following output is truncated in a few places. Omit the container name to observe the behavior. Also execute a command to see all environmental variables instead of logging into the container first.

```
student@master:~$ kubectl exec -it try1-d4fbf76fd-46pkb -- /bin/bash -c 'env'
```

```
1 Defaulting container name to simpleapp.
2 Use 'kubectl describe pod/try1-d4fbf76fd-46pkb -n default' to see all of the containers in this pod.
3 REGISTRY_PORT_5000_TCP_ADDR=10.105.119.236
4 HOSTNAME=try1-d4fbf76fd-46pkb
5 TERM=xterm
6 yellow=y
7 <output_omitted>
8 REGISTRY_SERVICE_HOST=10.105.119.236
9 KUBERNETES_SERVICE_PORT=443
10 REGISTRY_PORT_5000_TCP=tcp://10.105.119.236:5000
11 KUBERNETES_SERVICE_HOST=10.96.0.1
12 text=black
13 REGISTRY_SERVICE_PORT_5000=5000
14 <output_omitted>
15 black=k
16 known as key
17
18 <output_omitted>
19 ilike=blue
20 <output_omitted>
21 magenta=m
22
23 cyan=c
24 <output_omitted>
```

10. For greater flexibility and scalability **ConfigMaps** can be created from a YAML file, then deployed and redeployed as necessary. Once ingested into the cluster the data can be retrieved in the same manner as any other object. Create another **configMap**, this time from a YAML file.

```
student@master:~$ vim car-map.yaml
```



car-map.yaml

```

1  apiVersion: v1
2  kind: ConfigMap
3  metadata:
4    name: fast-car
5    namespace: default
6  data:
7    car.make: Ford
8    car.model: Mustang
9    car.trim: Shelby

```

```
student@master:~$ kubectl create -f car-map.yaml
```

```
1  configmap/fast-car created
```

11. View the ingested data, note that the output is just as in file created.

```
student@master:~$ kubectl get configmap fast-car -o yaml
```

```

1  apiVersion: v1
2  data:
3    car.make: Ford
4    car.model: Mustang
5    car.trim: Shelby
6  kind: ConfigMap
7  metadata:
8  <output_omitted>

```

12. Add the **configMap** settings to the `simpleapp.yaml` file as a volume. Both containers in the `try1` deployment can access to the same volume, using `volumeMounts` statements. Remember that the volume stanza is of equal depth to the containers stanza, and should come after the containers have been declared, the example below has the volume added just before the `status: output..`

```
student@master:~$ vim $HOME/app1/simpleapp.yaml
```



simpleapp.yaml

```

1  ....
2  spec:
3    containers:
4      - image: 10.105.119.236:5000/simpleapp
5        volumeMounts:                                #<-- Add this and following two lines
6          - mountPath: /etc/cars
7            name: car-vol
8        env:
9          - name: ilike
10  ....
11  securityContext: {}
12  terminationGracePeriodSeconds: 30
13  volumes:                                             #<-- Add this and following four lines
14    - name: car-vol
15      configMap:
16        defaultMode: 420
17        name: fast-car
18  status:
19  ....

```

13. Delete and recreate the deployment.

```
student@master:~$ kubectl delete deployment try1
```

```
1 deployment.apps "try1" deleted
```

```
student@master:~$ kubectl create -f $HOME/app1/simpleapp.yaml
```

```
1 deployment.apps/try1 created
```

14. Verify the deployment is running. Note that we still have not automated the creation of the `/tmp/healthy` file inside the container, as a result the AVAILABLE count remains zero until we use the `for` loop to create the file. We will remedy this in the next step.

```
student@master:~$ kubectl get deployment
```

```
1 NAME      READY   UP-TO-DATE   AVAILABLE   AGE
2 nginx     1/1     1            1           1d
3 registry  1/1     1            1           1d
4 try1      0/6     6            0           39s
```

15. Our health check was the successful execution of a command. We will edit the command of the existing readinessProbe to check for the existence of the mounted configMap file and re-create the deployment. After a minute both containers should become available for each pod in the deployment.

```
student@master:~$ kubectl delete deployment try1
```

```
1 deployment.apps "try1" deleted
```

```
student@master:~$ vim $HOME/app1/simpleapp.yaml
```

YAML

simpleapp.yaml

```
1 ....
2     readinessProbe:
3       exec:
4         command:
5           - ls /tmp/healthy #<-- Add/Edit this and following line.
6           - /etc/cars
7       periodSeconds: 5
8 ....
```

```
student@master:~$ kubectl create -f $HOME/app1/simpleapp.yaml
```

```
1 deployment.apps/try1 created
```

16. Wait about a minute and view the deployment and pods. All six replicas should be running and report that 2/2 containers are in a ready state within.

```
student@master:~$ kubectl get deployment
```

```
1 NAME      READY   UP-TO-DATE   AVAILABLE   AGE
2 nginx     1/1     1            1           1d
3 registry  1/1     1            1           1d
4 try1      6/6     6            6           1m
```

```
student@master:~$ kubectl get pods
```

	NAME	READY	STATUS	RESTARTS	AGE
2	nginx-6b58d9cdfd-9fnl4	1/1	Running	1	1d
3	registry-795c6c8b8f-hl5wf	1/1	Running	2	1d
4	try1-7865dcb948-2dzc8	2/2	Running	0	1m
5	try1-7865dcb948-7fkh7	2/2	Running	0	1m
6	try1-7865dcb948-d85bc	2/2	Running	0	1m
7	try1-7865dcb948-djrcj	2/2	Running	0	1m
8	try1-7865dcb948-kwlv8	2/2	Running	0	1m
9	try1-7865dcb948-stb2n	2/2	Running	0	1m

17. View a file within the new volume mounted in a container. It should match the data we created inside the configMap. Because the file did not have a carriage-return it will appear prior to the following prompt.

```
student@master:~$ kubectl exec -c simpleapp -it try1-7865dcb948-stb2n \
-- /bin/bash -c 'cat /etc/cars/car.trim'
```

```
1 Shelby student@master:~$
```

Exercise 7.2: Configure the Deployment: Attaching Storage

There are several types of storage which can be accessed with Kubernetes, with flexibility of storage being essential to scalability. In this exercise we will configure an NFS server. With the NFS server we will create a new **persistent volume (pv)** and a **persistent volume claim (pvc)** to use it.

1. Search for pv and pvc YAML example files on <http://kubernetes.io/docs> and <http://kubernetes.io/blog>.
2. Use the `CreateNFS.sh` script from the tarball to set up NFS on your master node. This script will configure the server, export `/opt/sfw` and create a file `/opt/sfw/hello.txt`. Use the `find` command to locate the file if you don't remember where you extracted the tar file. This example narrows the search to your `$HOME` directory. Change for your environment directory. You may find the same file in more than one sub-directory of the tarfile.

```
student@master:~$ find $HOME -name CreateNFS.sh
```

```
1 /home/student/LFD259/SOLUTIONS/s_05/CreateNFS.sh
```

```
student@master:~$ cp /home/student/LFD259/SOLUTIONS/s_05/CreateNFS.sh $HOME
```

```
student@master:~$ bash $HOME/CreateNFS.sh
```

```
1 Hit:1 http://us-central1.gce.archive.ubuntu.com/ubuntu xenial InRelease
2 Get:2 http://us-central1.gce.archive.ubuntu.com/ubuntu xenial-updates InRelease [102 kB]
3
4 <output_omitted>
5
6 Should be ready. Test here and second node
7
8 Export list for localhost:
9 /opt/sfw *
```

3. Test by mounting the resource from your **second node**. Begin by installing the client software.

```
student@worker:~$ sudo apt-get -y install nfs-common nfs-kernel-server
```

```
1 <output_omitted>
```

4. Test you can see the exported directory using `showmount` from you second node.

```
student@worker:~$ showmount -e master #<-- Edit to be first node's name or IP
```

```
1 Export list for master:
2 /opt/sfw *
```

5. Mount the directory. Be aware that unless you edit `/etc/fstab` this is not a persistent mount. Change out the node name for that of your master node.

```
student@worker:~$ sudo mount master:/opt/sfw /mnt
```

6. Verify the `hello.txt` file created by the script can be viewed.

```
student@worker:~$ ls -l /mnt
1 total 4
2 -rw-r--r-- 1 root root 9 Sep 28 17:55 hello.txt
```

7. Return to the master node and create a YAML file for an object with kind **PersistentVolume**. The included example file needs an edit to the `server:` parameter. Use the hostname of the master server and the directory you created in the previous step. Only syntax is checked, an incorrect name or directory will not generate an error, but a Pod using the incorrect resource will not start. Note that the `accessModes` do not currently affect actual access and are typically used as labels instead.

```
student@master:~$ find $HOME -name PVol.yaml
1 /home/student/LFD259/SOLUTIONS/s_05/PVol.yaml

student@master:~$ cp /home/student/LFD259/SOLUTIONS/s_05/PVol.yaml $HOME

student@master:~$ vim PVol.yaml
```

YAML

PVol.yaml

```
1 apiVersion: v1
2 kind: PersistentVolume
3 metadata:
4   name: pvvol-1
5 spec:
6   capacity:
7     storage: 1Gi
8   accessModes:
9     - ReadWriteMany
10  persistentVolumeReclaimPolicy: Retain
11  nfs:
12    path: /opt/sfw
13    server: master
14    readOnly: false
#<-- Edit to match master node name or IP
```

8. Create and verify you have a new 1Gi volume named **pvvol-1**. Note the status shows as Available. Remember we made two persistent volumes for the image registry earlier.

```
student@master:~$ kubectl create -f PVol.yaml
1 persistentvolume/pvvol-1 created
```

```
student@master:~$ kubectl get pv
```

	NAME	CAPACITY	ACCESS MODES	RECLAIM POLICY	STATUS	CLAIM	STORAGECLASS	REASON	AGE
1									
2	pvvol-1	1Gi	RWX	Retain	Available				4s
3	registryvm	200Mi	RWO	Retain	Bound	default/nginx-claim0			4d
4	task-pv-volume	200Mi	RWO	Retain	Bound	default/registry-claim0			4d

9. Now that we have a new volume we will use a **persistent volume claim (pvc)** to use it in a Pod. We should have two existing claims from our local registry.

```
student@master:~/$ kubectl get pvc
```

	NAME	STATUS	VOLUME	CAPACITY	ACCESS MODES	STORAGECLASS	AGE
1	nginx-claim0	Bound	registryvm	200Mi	RWO		4d
2	registry-claim0	Bound	task-pv-volume	200Mi	RWO		4d

10. Create or copy a yaml file with the kind **PersistentVolumeClaim**.

```
student@master:~$ vim pvc.yaml
```

YAML

pvc.yaml

```
1 apiVersion: v1
2 kind: PersistentVolumeClaim
3 metadata:
4   name: pvc-one
5 spec:
6   accessModes:
7     - ReadWriteMany
8   resources:
9     requests:
10      storage: 200Mi
```

11. Create and verify the new pvc status is bound. Note the size is 1Gi, even though 200Mi was suggested. Only a volume of at least that size could be used, the first volume with found with at least that much space was chosen.

```
student@master:~$ kubectl create -f pvc.yaml
```

```
1 persistentvolumeclaim/pvc-one created
```

```
student@master:~$ kubectl get pvc
```

	NAME	STATUS	VOLUME	CAPACITY	ACCESS MODES	STORAGECLASS	AGE
1	nginx-claim0	Bound	registryvm	200Mi	RWO		4d
2	pvc-one	Bound	pvvol-1	1Gi	RWX		4s
3	registry-claim0	Bound	task-pv-volume	200Mi	RWO		4d

12. Now look at the status of the physical volume. It should also show as bound.

```
student@master:~$ kubectl get pv
```

	NAME	CAPACITY	ACCESS MODES	RECLAIM POLICY	STATUS
1	CLAIM	STORAGECLASS	REASON	AGE	
2	pvvol-1	1Gi	RWX	Retain	Bound
3	default/pvc-one			14m	
4	registryvm	200Mi	RWO	Retain	Bound
5	default/nginx-claim0			4d	
6	task-pv-volume	200Mi	RWO	Retain	Bound
7	default/registry-claim0			4d	

13. Edit the `simpleapp.yaml` file to include two new sections. One section for the container while will use the volume mount point, you should have an existing entry for `car-vol`. The other section adds a volume to the deployment in general, which you can put after the `configMap` volume section.

```
student@master:~$ vim $HOME/app1/simpleapp.yaml
```



simpleapp.yaml

```

1  ....
2      volumeMounts:
3      - name: car-vol
4        mountPath: /etc/cars
5      - name: nfs-vol
6        mountPath: /opt
7  ....
8      volumes:
9      - name: car-vol
10        configMap:
11          defaultMode: 420
12          name: fast-car
13      - name: nfs-vol
14        persistentVolumeClaim:
15          claimName: pvc-one
16  status:
17  ....

```

#<-- Add this and following line

#<-- Add this and following two lines

14. Delete and re-create the deployment.

```
student@master:~$ kubectl delete deployment try1 ; kubectl create -f $HOME/app1/simpleapp.yaml
```

```

1 deployment.apps "try1" deleted
2 deployment.apps/try1 created

```

15. View the details any of the pods in the deployment, you should see `nfs-vol` mounted under `/opt`. The use to command line completion with the `tab` key can be helpful for using a pod name.

```
student@master:~$ kubectl describe pod try1-594fbb5fc7-5k7sj
```

```

1 <output_omitted>
2   Mounts:
3     /etc/cars from car-vol (rw)
4     /opt from nfs-vol (rw)
5     /var/run/secrets/kubernetes.io/serviceaccount from default-token-j7cqd (ro)
6 <output_omitted>

```

✍ Exercise 7.3: Using ConfigMaps Configure Ambassador Containers

In an earlier lab we added a second Ambassador container to handle logging. Now that we have learned about using ConfigMaps and attaching storage we will use configure our basic pod.

1. Review the YAML for our earlier simple pod. Recall that we added an Ambassador style logging container to the pod but had not fully configured the logging.

```
student@master:~$ cat basic.yaml
```

```

1 <output_omitted>
2   containers:
3   - name: webcont
4     image: nginx
5     ports:
6     - containerPort: 80
7   - name: fdlogger
8     image: fluent/fluentd

```

- Let us begin by adding shared storage to each container. We will use the `hostPath` storage class to provide the PV and PVC. First we create the directory.

```
student@master:~$ sudo mkdir /tmp/weblog
```

- Now we create a new PV to use that directory for the `hostPath` storage class. We will use the `storageClassName` of `manual` so that only PVCs which use that name will bind the resource.

```
student@master:~$ vim weblog-pv.yaml
```

YAML
weblog-pv.yaml

```
1 kind: PersistentVolume
2 apiVersion: v1
3 metadata:
4   name: weblog-pv-volume
5   labels:
6     type: local
7 spec:
8   storageClassName: manual
9   capacity:
10    storage: 100Mi
11  accessModes:
12    - ReadWriteOnce
13  hostPath:
14    path: "/tmp/weblog"
```

- Create and verify the new PV exists and shows an `Available` status.

```
student@master:~$ kubectl create -f weblog-pv.yaml
```

```
1 persistentvolume/weblog-pv-volume created
```

```
student@master:~$ kubectl get pv weblog-pv-volume
```

```
1 NAME          CAPACITY ACCESS MODES RECLAIM POLICY
2 STATUS        CLAIM    STORAGECLASS REASON  AGE
3
4 weblog-pv-volume 100Mi    RWO          Retain
5 Available              manual                21s
```

- Next we will create a PVC to use the PV we just created.

```
student@master:~$ vim weblog-pvc.yaml
```

YAML
weblog-pvc.yaml

```
1 kind: PersistentVolumeClaim
2 apiVersion: v1
3 metadata:
4   name: weblog-pv-claim
5 spec:
6   storageClassName: manual
7   accessModes:
8     - ReadWriteOnce
9   resources:
10    requests:
11      storage: 100Mi
```


6. Create the PVC and verify it shows as Bound to the the PV we previously created.

```
student@master:~$ kubectl create -f weblog-pvc.yaml
```

```
1 persistentvolumeclaim/weblog-pv-claim created
```

```
student@master:~$ kubectl get pvc weblog-pv-claim
```

```
1 NAME              STATUS  VOLUME          CAPACITY  ACCESS MODES
2 STORAGECLASS      AGE
3 weblog-pv-claim   Bound  weblog-pv-volume 100Mi    RWO
4 manual            79s
```

7. We are ready to add the storage to our pod. We will edit three sections. The first will declare the storage to the pod in general, then two more sections which tell each container where to make the volume available.

```
student@master:~$ vim basic.yaml
```

YA
ML

basic.yaml

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: basicpod
5   labels:
6     type: webserver
7 spec:
8   volumes:                                     #<-- Add three lines, same depth as containers
9     - name: weblog-pv-storage
10       persistentVolumeClaim:
11         claimName: weblog-pv-claim
12   containers:
13     - name: webcont
14       image: nginx
15       ports:
16         - containerPort: 80
17       volumeMounts:                             #<-- Add three lines, same depth as ports
18         - mountPath: "/var/log/nginx/"
19           name: weblog-pv-storage                # Must match volume name above
20     - name: fdlogger
21       image: fluent/fluentd
22       volumeMounts:                             #<-- Add three lines, same depth as image:
23         - mountPath: "/var/log"
24           name: weblog-pv-storage                # Must match volume name above
```

8. At this point we can create the pod again. When we create a shell we will find that the `access.log` for **nginx** is no longer a symbolic link pointing to `stdout` it is a writable, zero length file. Leave a **tailf** of the log file running.

```
student@master:~$ kubectl create -f basic.yaml
```

```
1 pod/basicpod created
```

```
student@master:~$ kubectl exec -c webcont -it basicpod -- /bin/bash
```



On Container

```
root@basicpod:/# ls -l /var/log/nginx/access.log
```

```
1 -rw-r--r-- 1 root root 0 Oct 18 16:12 /var/log/nginx/access.log
```



```
root@basicpod:/# tail -f /var/log/nginx/access.log
```

9. Open a second connection to your master node. We will use the pod IP as we have not yet configured a service to expose the pod.

```
student@master:~$ kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE
NOMINATED NODE						
basicpod 2/2		Running	0	3m26s	192.168.213.181	master
<none>						

10. Use **curl** to view the welcome page of the webserver. When the command completes you should see a new entry added to the log. Right after the GET we see a 200 response indicating success. You can use **ctrl-c** and **exit** to return to the host shell prompt.

```
student@master:~$ curl http://192.168.213.181
```

```
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<output_omitted>
```



On Container

```
192.168.32.128 - - [18/Oct/2018:16:16:21 +0000] "GET / HTTP/1.1" 200 612 "-" "curl/7.47.0" "-"
```

11. Now that we know the webcont container is writing to the PV we will configure the logger to use that directory as a source. For greater flexibility we will configure **fluentd** using a **configMap**.

Fluentd has many options for input and output of data. We will read from a file of the webcont container and write to standard out of the fdlogger container. The details of the data settings can be found in **fluentd** documentation here: <https://docs.fluentd.org/v1.0/categories/config-file>

```
student@master:~$ vim weblog-configmap.yaml
```



weblog-configmap.yaml

```
1 apiVersion: v1
2 kind: ConfigMap
3 metadata:
4   name: fluentd-config
5   namespace: default
6 data:
7   fluentd.conf: |
8     <source>
9       @type tail
10      format none
11      path /var/log/access.log
12      tag count.format1
13    </source>
14
15    <match *.*>
16      @type stdout
17      id stdout_output
```



```
18     </match>
```

12. Create the new configMap.

```
student@master:~$ kubectl create -f weblog-configmap.yaml
```

```
1 configmap/fluentd-config created
```

13. View the logs for both containers in the basicpod. You should see some startup information, but not the HTTP traffic.

```
student@master:~$ kubectl logs basicpod webcont
```

```
1 /docker-entrypoint.sh: /docker-entrypoint.d/ is not empty, will attempt to perform configuration
2 /docker-entrypoint.sh: Looking for shell scripts in /docker-entrypoint.d/
3 /docker-entrypoint.sh: Launching /docker-entrypoint.d/10-listen-on-ipv6-by-default.sh
4 10-listen-on-ipv6-by-default.sh: Getting the checksum of /etc/nginx/conf.d/default.conf
5 10-listen-on-ipv6-by-default.sh: Enabled listen on IPv6 in /etc/nginx/conf.d/default.conf
6 /docker-entrypoint.sh: Launching /docker-entrypoint.d/20-envsubst-on-templates.sh
7 /docker-entrypoint.sh: Configuration complete; ready for start up
```

```
student@master:~$ kubectl logs basicpod fdlogger
```

```
1 2020-09-02 19:32:59 +0000 [info]: reading config file path="/etc/fluentd-config/fluentd.conf"
2 2020-09-02 19:32:59 +0000 [info]: starting fluentd-0.12.29
3 2020-09-02 19:32:59 +0000 [info]: gem 'fluent-mixin-config-placeholders' version '0.4.0'
4 2020-09-02 19:32:59 +0000 [info]: gem 'fluent-mixin-plaintextformatter' version '0.2.6'
5
6 <output_omitted>
7
8 <source>
9   @type tail
10  format none
11  path /var/log/access.log
12 <output_omitted>
```

14. Now we will edit the pod yaml file so that the **fluentd** container will mount the configmap as a volume and reference the variables inside the config file. You will add three areas, the volume declaration to the pod, the env parameter and the mounting of the volume to the fluentd container

```
student@master:~$ vim basic.yaml
```



basic.yaml

```
1 ....
2 volumes:
3   - name: weblog-pv-storage
4     persistentVolumeClaim:
5       claimName: weblog-pv-claim
6   - name: log-config                                #<-- This and two lines following
7     configMap:
8       name: fluentd-config                          # Must match existing configMap
9   ....
10  image: fluent/fluentd
11  env:                                                #<-- This and two lines following
12    - name: FLUENTD_OPT
13      value: -c /etc/fluentd-config/fluentd.conf
14  ....
```



```

15     volumeMounts:
16       - mountPath: "/var/log"
17         name: weblog-pv-storage
18       - name: log-config                #<-- This and next line
19         mountPath: "/etc/fluentd-config"

```

15. At this point we can delete and re-create the pod, which would cause the configmap to be used by the new pod, among other changes.

```
student@master:~$ kubectl delete pod basicpod
```

```
1 pod "basicpod" deleted
```

```
student@master:~$ kubectl create -f basic.yaml
```

```
1 pod/basicpod created
```

```
student@master:~$ kubectl get pod basicpod -o wide
```

```

1 NAME          READY   STATUS    RESTARTS   AGE   IP              NODE          NOMINATED....
2 basicpod      2/2     Running   0           8s    192.168.171.122 worker        <none>      ....

```

16. Use **curl** a few times to look at the default page served by basicpod

```
student@master:~$ curl http://192.168.171.122
```

```

1 <!DOCTYPE html>
2 <html>
3 <head>
4 <title>Welcome to nginx!</title>
5 <style>
6   body {
7 <output_omitted>

```

17. Look at the logs for both containers. In addition to the standard startup information, you should also see the HTTP requests from the curl commands you just used at the end of the fdlogger output.

```
student@master:~$ kubectl logs basicpod webcont
```

```

1 /docker-entrypoint.sh: /docker-entrypoint.d/ is not empty, will attempt to perform configuration
2 /docker-entrypoint.sh: Looking for shell scripts in /docker-entrypoint.d/
3 /docker-entrypoint.sh: Launching /docker-entrypoint.d/10-listen-on-ipv6-by-default.sh
4 10-listen-on-ipv6-by-default.sh: Getting the checksum of /etc/nginx/conf.d/default.conf
5 10-listen-on-ipv6-by-default.sh: Enabled listen on IPv6 in /etc/nginx/conf.d/default.conf
6 /docker-entrypoint.sh: Launching /docker-entrypoint.d/20-envsubst-on-templates.sh
7 /docker-entrypoint.sh: Configuration complete; ready for start up

```

```
student@master:~$ kubectl logs basicpod fdlogger
```

```

1 2020-09-02 19:32:59 +0000 [info]: reading config file path="/etc/fluentd-config/fluentd.conf"
2 2020-09-02 19:32:59 +0000 [info]: starting fluentd-0.12.29
3 2020-09-02 19:32:59 +0000 [info]: gem 'fluent-mixin-config-placeholders' version '0.4.0'
4 2020-09-02 19:32:59 +0000 [info]: gem 'fluent-mixin-plaintextformatter' version '0.2.6'
5
6 <output_omitted>
7
8 <source>
9   @type tail

```

```

10     format none
11     path /var/log/access.log
12
13 <output_omitted>
14
15 2020-09-02 19:47:38 +0000 count.format1: {"message":"192.168.219.64 - - [02/Sep/2020:19:47:38 +0000] \"GET / HTTP/1.1\" 200 12345"}
16 2020-09-02 19:47:41 +0000 count.format1: {"message":"192.168.219.64 - - [02/Sep/2020:19:47:41 +0000] \"GET / HTTP/1.1\" 200 12345"}
17 2020-09-02 19:47:47 +0000 count.format1: {"message":"192.168.219.64 - - [02/Sep/2020:19:47:47 +0000] \"GET / HTTP/1.1\" 200 12345"}

```

✎ Exercise 7.4: Rolling Updates and Rollbacks

When we started working with `simpleapp` we used a **Docker** tag called `latest`. While this is the default tag when pulling an image, and commonly used, it remains just a string, it may not be the actual latest version of the image.

1. Make a slight change to our source and create a new image. We will use updates and rollbacks with our application. Adding a comment to the last line should be enough for a new image to be generated.

```
student@master:~$ cd ~/app1
```

```
student@master:~/app1$ vim simple.py
```

```
<output_omitted>
```

```
## Sleep for five seconds then continue the loop
```

```
time.sleep(5)
```

```
## Adding a new comment so image is different.
```

2. Build the image again. A new container and image will be created. Verify when successful. There should be a different image ID and a recent creation time.

```
student@master:~/app1$ sudo docker build -t simpleapp .
```

```

1 Sending build context to Docker daemon 7.168 kB
2 Step 1/3 : FROM python:2
3 ---> 2863c80c418c
4 Step 2/3 : ADD simple.py /
5 ---> cde8ecf8492b
6 Removing intermediate container 3e908b76b5b4
7 Step 3/3 : CMD python ./simple.py
8 ---> Running in 354620c97bf5
9 ---> cc6bba0ea213
10 Removing intermediate container 354620c97bf5
11 Successfully built cc6bba0ea213

```

```
student@master:~/app1$ sudo docker images
```

```

1 REPOSITORY          TAG
2 IMAGE ID            CREATED             SIZE
3 simpleapp           latest
4 cc6bba0ea213        8 seconds ago      886 MB
5 10.105.119.236:5000/simpleapp latest
6 15b5ad19d313        4 days ago         886 MB
7 <output_omitted>

```

3. Tag and push the updated image to your locally hosted registry. A reminder your IP address will be different than the example below. Use the tag `v2` this time instead of `latest`.

```
student@master:~/app1$ sudo docker tag simpleapp \
10.105.119.236:5000/simpleapp:v2
```

```
student@master:~/app1$ sudo docker push 10.105.119.236:5000/simpleapp:v2
```

```
1 The push refers to a repository [10.105.119.236:5000/simpleapp]
2 d6153c8cc7c3: Pushed
3 ca82a2274c57: Layer already exists
4 de2fbb43bd2a: Layer already exists
5 4e32c2de91a6: Layer already exists
6 6e1b48dc2ccc: Layer already exists
7 ff57bdb79ac8: Layer already exists
8 6e5e20cbf4a7: Layer already exists
9 86985c679800: Layer already exists
10 8fad67424c4e: Layer already exists
11 v2: digest: sha256:6cf74051d09463d89f1531fceb9c44cbf99006f8d9b407
12 dd91d8f07baeee7e9c size: 2218
```

4. Connect to a terminal running on your second node. Pull the latest image, then pull v2. Note the latest did not pull the new version of the image. Again, remember to use the IP for your locally hosted registry. You'll note the digest is different.

```
student@worker:~$ sudo docker pull 10.105.119.236:5000/simpleapp
```

```
1 Using default tag: latest
2 latest: Pulling from simpleapp
3 Digest: sha256:cefa3305c36101d32399baf0919d3482ae8a53c926688be33
4 86f9bbc04e490a5
5 Status: Image is up to date for 10.105.119.236:5000/simpleapp:latest
```

```
student@worker:~$ sudo docker pull 10.105.119.236:5000/simpleapp:v2
```

```
1 v2: Pulling from simpleapp
2 f65523718fc5: Already exists
3 1d2dd88bf649: Already exists
4 c09558828658: Already exists
5 0e1d7c9e6c06: Already exists
6 c6b6fe164861: Already exists
7 45097146116f: Already exists
8 f21f8abae4c4: Already exists
9 1c39556edcd0: Already exists
10 fa67749bf47d: Pull complete
11 Digest: sha256:6cf74051d09463d89f1531fceb9c44cbf99006f8d9b407dd91d8
12 f07baeee7e9c
13 Status: Downloaded newer image for 10.105.119.236:5000/simpleapp:v2
```

5. Use **kubectl edit** to update the image for the `try1` deployment to use v2. As we are only changing one parameter we could also use the **kubectl set** command. Note that the configuration file has not been updated, so a delete or a replace command would not include the new version. It can take the pods up to a minute to delete and to recreate each pod in sequence.

```
student@master:~/app1$ kubectl edit deployment try1
```

```
....
containers:
- image: 10.105.119.236:5000/simpleapp:v2  #<-- Edit tag
  imagePullPolicy: Always
....
```

6. Verify each of the pods has been recreated and is using the new version of the image. Note some messages will show the scaling down of the old **replicaset**, others should show the scaling up using the new image.

```
student@master:~/app1$ kubectl get events
```

```

1 42m      Normal    ScalingReplicaSet   Deployment   Scaled up replica set try1-7fdbb5d557 to 6
2 32s      Normal    ScalingReplicaSet   Deployment   Scaled up replica set try1-7fd7459fc6 to 2
3 32s      Normal    ScalingReplicaSet   Deployment   Scaled down replica set try1-7fdbb5d557 to 5
4 32s      Normal    ScalingReplicaSet   Deployment   Scaled up replica set try1-7fd7459fc6 to 3
5 23s      Normal    ScalingReplicaSet   Deployment   Scaled down replica set try1-7fdbb5d557 to 4
6 23s      Normal    ScalingReplicaSet   Deployment   Scaled up replica set try1-7fd7459fc6 to 4
7 22s      Normal    ScalingReplicaSet   Deployment   Scaled down replica set try1-7fdbb5d557 to 3
8 22s      Normal    ScalingReplicaSet   Deployment   Scaled up replica set try1-7fd7459fc6 to 5
9 18s      Normal    ScalingReplicaSet   Deployment   Scaled down replica set try1-7fdbb5d557 to 2
10 18s      Normal    ScalingReplicaSet   Deployment   Scaled up replica set try1-7fd7459fc6 to 6
11 8s       Normal    ScalingReplicaSet   Deployment   (combined from similar events):
12 Scaled down replica set try1-7fdbb5d557 to 0
13
```

7. View the images of a Pod in the deployment. Narrow the output to just view the images. The goproxy remains unchanged, but the simpleapp should now be v2.

```
student@master:~/app1$ kubectl describe pod try1-895fccfb-ttqdn |grep Image
```

```

1      Image:          10.105.119.236:5000/simpleapp:v2
2      Image ID:\
3          docker-pullable://10.105.119.236:5000/simpleapp@sha256:6cf74051d09
4 463d89f1531fceb9c44cbf99006f8d9b407dd91d8f07baeee7e9c
5      Image:          k8s.gcr.io/goproxy:0.1
6      Image ID:\
7          docker-pullable://k8s.gcr.io/goproxy@sha256:5334c7ad43048e3538775c
8 b09aaf184f5e8acf4b0ea60e3bc8f1d93c209865a5
```

8. View the update history of the deployment.

```
student@master:~/app1$ kubectl rollout history deployment try1
```

```

1 deployments "try1"
2 REVISION  CHANGE-CAUSE
3 1          <none>
4 2          <none>
```

9. Compare the output of the **rollout history** for the two revisions. Images and labels should be different, with the image v2 being the change we made.

```
student@master:~/app1$ kubectl rollout history deployment try1 --revision=1 > one.out
```

```
student@master:~/app1$ kubectl rollout history deployment try1 --revision=2 > two.out
```

```
student@ckad-~/app1:~$ diff one.out two.out
```

```

1 1c1
2 < deployments "try1" with revision #1
3 ---
4 > deployments "try1" with revision #2
5 3c3
6 <   Labels:          pod-template-hash=1509661973
7 ---
8 >   Labels:          pod-template-hash=45197796
9 7c7
10 <   Image:           10.105.119.236:5000/simpleapp
11 ---
12 >   Image:           10.105.119.236:5000/simpleapp:v2
```

10. View what would be undone using the **dry-run** option while undoing the rollout. This allows us to see the new template prior to using it.

```
student@master:~/app1$ kubectl rollout undo --dry-run=client deployment/try1
```

```
1 deployment.apps/try1
2 Pod Template:
3   Labels:      pod-template-hash=1509661973
4     run=try1
5   Containers:
6     try1:
7       Image:      10.105.119.236:5000/simpleapp:latest
8       Port:      <none>
9 <output_omitted>
```

11. View the pods. Depending on how fast you type the try1 pods should be about 2 minutes old.

```
student@master:~/app1$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-6b58d9cdfd-9fnl4	1/1	Running	1	5d
registry-795c6c8b8f-hl5wf	1/1	Running	2	5d
try1-594fbb5fc7-7dl7c	2/2	Running	0	2m
try1-594fbb5fc7-8mxlb	2/2	Running	0	2m
try1-594fbb5fc7-jr7h7	2/2	Running	0	2m
try1-594fbb5fc7-s24wt	2/2	Running	0	2m
try1-594fbb5fc7-xfffg	2/2	Running	0	2m
try1-594fbb5fc7-zfmz8	2/2	Running	0	2m

12. In our case there are only two revisions, which is also the default number kept. Were there more we could choose a particular version. The following command would have the same effect as the previous, without the **-dry-run** option.

```
student@master:~/app1$ kubectl rollout undo deployment try1 --to-revision=1
```

```
1 deployment.apps/try1 rolled back
```

13. Again, it can take a bit for the pods to be terminated and re-created. Keep checking back until they are all running again.

```
student@master:~/app1$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-6b58d9cdfd-9fnl4	1/1	Running	1	5d
registry-795c6c8b8f-hl5wf	1/1	Running	2	5d
try1-594fbb5fc7-7dl7c	2/2	Terminating	0	3m
try1-594fbb5fc7-8mxlb	0/2	Terminating	0	2m
try1-594fbb5fc7-jr7h7	2/2	Terminating	0	3m
try1-594fbb5fc7-s24wt	2/2	Terminating	0	2m
try1-594fbb5fc7-xfffg	2/2	Terminating	0	3m
try1-594fbb5fc7-zfmz8	1/2	Terminating	0	2m
try1-895fccfb-8dn4b	2/2	Running	0	22s
try1-895fccfb-kz72j	2/2	Running	0	10s
try1-895fccfb-rxxtw	2/2	Running	0	24s
try1-895fccfb-srwq4	1/2	Running	0	11s
try1-895fccfb-vkymb	2/2	Running	0	31s
try1-895fccfb-z46qr	2/2	Running	0	31s

✍ Exercise 7.5: Domain Review



Very Important

The source pages and content in this review could change at any time. **IT IS YOUR RESPONSIBILITY TO CHECK THE CURRENT INFORMATION.**

Revisit the CKAD domain list on [Curriculum Overview](#) and locate some of the topics we have covered in this chapter. The graphic shows bullet points covered in this chapter.




Figure 7.2: **Deployment Related Domain Topics**

Focus on ensuring you have all the necessary files and processes understood first. Repeat the review until you are sure you have bookmarks of necessary YAML samples and can complete each step quickly, and ensure each object is running properly.

Using only the allowed browser, URLs, and subdomains search for and bookmark a YAML example to create and configure the resources called for in this review.

1. Create a new secret called `specialofday` using the key `entree` and the value `meatloaf`.
2. Create a new deployment called `foodie` running the `nginx` image.
3. Add the `specialofday` secret to pod mounted as a volume under the `/food/` directory.
4. Execute a bash shell inside a `foodie` pod and verify the secret has been properly mounted.
5. Update the deployment to use the `nginx:1.12.1-alpine` image and verify the new image is in use.
6. Roll back the deployment and verify the typical, current stable version of `nginx` is in use again.
7. Create a new 200M NFS volume called `reviewvol` using the NFS server configured earlier in the lab.
8. Create a new PVC called `reviewpvc` which will uses the `reviewvol` volume.
9. Edit the deployment to use the PVC and mount the volume under `/newvol`
10. Execute a bash shell into the `nginx` container and verify the volume has been mounted.
11. Delete any resources created during this review.

Chapter 8

Custom Resource Definition



8.1	Overview	156
8.2	Custom Resource Definitions	157
8.3	Aggregated APIs	161
8.4	Labs	162

8.1 Overview

Custom Resources

- Flexible to meet changing needs
- Create your own API objects
- Manage your API objects via **kubectl**
- Custom Resource Definitions (CRD)
- Aggregated APIs (AA)

We have been working with built-in resources, or API endpoints. The flexibility of Kubernetes allows for dynamic addition of new resources as well. Once these Custom Resources have been added the objects can be created and accessed using standard calls and commands like **kubectl**. The creation of a new object stores new structured data in the **etcd** database and allows access via **kube-apiserver**.

To make a new, custom resource part of a declarative API there needs to be a controller to retrieve the structured data continually and act to meet and maintain the declared state. This controller, or operator, is an agent to create and manage one or more instances of a specific stateful application. We have worked with built-in controllers such for Deployments, DaemonSets and other resources.

The functions encoded into a custom operator should be all the tasks a human would need to perform if deploying the application outside of Kubernetes. The details of building a custom controller are outside the scope of this course and not included.

There are two ways to add custom resources to your Kubernetes cluster. The easiest, but less flexible, way is by adding a Custom Resource Definition to the cluster. The second which is more flexible is the use of Aggregated APIs which requires a new API server to be written and added to the cluster.

Either way of a new object to the cluster, as distinct from a built-in resource, is called a Custom Resource.

If you are using RBAC for authorization you probably will need to grant access to the new CRD resource and controller. If using an Aggregated API, you can use the same or different authentication process.

8.2 Custom Resource Definitions

Custom Resource Definitions

- Easy to deploy
- Typically does not require programming
- Does not require another API server
- Namespaced or cluster-scoped

As we have learned decoupled nature of Kubernetes depends on a collection of watcher loops, or controllers, interrogating the **kube-apiserver** to determine if a particular configuration is true. If the current state does not match the declared state the controller makes API calls to modify the state until they do match. If you add a new API object and controller you can use the existing **kube-apiserver** to monitor and control the object. The addition of a Custom Resource Definition will be added to the cluster API path, currently under apiextensions.k8s.io/v1.

While this is the easiest way to add a new object to the cluster it may not be flexible enough for your needs. Only the existing API functionality can be used. Objects must respond to REST requests and have their configuration state validated and stored in the same manner as built-in objects. They would also need to exist with the protection rules of built-in objects.

A CRD allows the resource to be deployed in a namespace or available in the entire cluster. The YAML file sets this with the `scope` parameter which can be set to `Namespaced` or `Cluster`.

Prior to v1.8 there was a resource type called `ThirdPartyResource` (TPR). This has been deprecated and is no longer available. All resources will need to be rebuilt as CRD. After upgrade existing TPRs will need to be removed and replaced by CRDs such that the API URL points to functional object.

Configuration Example

```
apiVersion: apiextensions.k8s.io/v1
kind: CustomResourceDefinition
metadata:
  name: backups.stable.linear.com
spec:
  group: stable.linear.com
  version: v1
  scope: Namespaced
  names:
    plural: backups
    singular: backup
    shortNames:
      - bks
    kind: BackUp
```

apiVersion:

Should match the current level of stability, currently `apiextensions.k8s.io/v1`

kind: CustomResourceDefinition

The object type being inserted by the **kube-apiserver**.

name: backups.stable.linear.com

The name must match the `spec` field declared later. The syntax must be `<plural name>.<group>`.

group: stable.linear.com

The group name will become part of the REST API under `/apis/<group>/<version>` or `/apis/stable/v1` in this case with the version set to **v1**.

scope

Determines if the object exists in a single namespace or is cluster-wide.

plural

Defines the last part of the API URL such as `apis/stable/v1/backups`.

singular and shortNames

represent the name with displayed and make CLI usage easier.

kind

A CamelCased singular type used in resource manifests.

New Object Configuration

```
apiVersion: "stable.linux.com/v1"
kind: BackUp
metadata:
  name: a-backup-object
spec:
  timeSpec: "* * * * */5"
  image: linux-backup-image
  replicas: 5
```

Note that the `apiVersion` and `kind` match the CRD we created in a previous step. The `spec` parameters depend on the controller.

The object will be evaluated by the controller. If the syntax, such as `timeSpec` does not match the expected value you will receive an error, should validation be configured. Without validation only the existence of the variable is checked, not its details.

Optional Hooks

- Finalizer

```
metadata:
  finalizers:
    - finalizer.stable.linux.com
```

- Validation

```
validation:
  openAPIV3Schema:
    properties:
      spec:
        properties:
          timeSpec:
            type: string
            pattern: '^(\\d+|\\*)(\\/\\d+)?(\\s+(\\d+|\\*)(\\/\\d+)?){4}$'
          replicas:
            type: integer
            minimum: 1
            maximum: 10
```

Just as with built-in objects you can use an asynchronous pre-delete hook known as a Finalizer. If an API delete request is received the object metadata field `metadata.deletionTimestamp` is updated. The controller then triggers whichever finalizer has been configured. When the finalizer completes it is removed from the list. The controller continues to compete and remove finalizers until the string is empty. Then the object itself is deleted.

A feature in beta starting with v1.9 allows for validation of custom objects via the OpenAPI v3 schema. This will check various properties of the object configuration being passed the API server. In the example above the `timeSpec` must be a string matching a particular pattern and the number of allowed replicas is between one and 10. If the validation does not match the error returned is the failed line of validation.

8.3 Aggregated APIs

Understanding Aggregated APIs (AA)

- Usually requires non-trivial programming
- More control over API behavior
- Subordinate API server behind primary
- Primary acts as proxy
- Leverages extension resource
- Mutual TLS auth between API servers
- RBAC rule to allow addition of API service objects

The use of Aggregated APIs allows adding additional Kubernetes-type API servers to the cluster. The added server acts as a subordinate to **kube-apiserver** which as of v1.7 runs the aggregation layer in-process. When an extension resource is registered the aggregation layer watches a passed URL path and proxy any requests to the newly registered API service.

The aggregation layer is easy to enable. Edit the flags passed during startup of the **kube-apiserver** to include `--enable-aggregator-routing=true`. Some vendors enable this feature by default.

The creation of the exterior can be done via YAML configuration files or APIs. Configuring TLS auth between components and RBAC rules for various new objects is also required. A sample API server is available on github here: <https://github.com/kubernetes/sample-apiserver>. A project currently in incubation stage is an API server builder which should handle much of the security and connection configuration. It can be found here: <https://github.com/kubernetes-incubator/apiserver-builder>

8.4 Labs

✍ Exercise 8.1: Create a Custom Resource Definition

Overview

The use of CustomResourceDefinitions (CRD), has become a common manner to deploy new objects and operators. Creation of a new operator is beyond the scope of this course, basically it is a watch-loop comparing a spec to the current status, and making changes until the states match. A good discussion of creating a controller can be found here: <https://coreos.com/blog/introducing-operators.html>.

First we will examine an existing CRD, then make a simple CRD, but without any particular action. It will be enough to find the object ingested into the API and responding to commands.

1. View the existing CRDs.

```
student@master:~$ kubectl get crd --all-namespaces
```

```
1 NAME                                CREATED AT
2 bgpconfigurations.crd.projectcalico.org 2020-04-19T17:29:02Z
3 bgppeers.crd.projectcalico.org          2020-04-19T17:29:02Z
4 blockaffinities.crd.projectcalico.org    2020-04-19T17:29:02Z
5 <output_omitted>
```

2. We can see from the names that these CRDs are all working on Calico, our network plugin. View the `calico.yaml` file we used when we initialized the cluster to see how these objects were created, and some CRD templates to review.

```
student@master:~$ less calico.yaml
```

```
1 <output_omitted>
2 ---
3 # Source: calico/templates/kdd-crds.yaml
4
5 apiVersion: apiextensions.k8s.io/v1beta1
6 kind: CustomResourceDefinition
7 metadata:
8   name: bgpconfigurations.crd.projectcalico.org
9 <output_omitted>
```

3. Now that we have seen some examples, we will create a new YAML file.

```
student@master:~$ vim crd.yaml
```

YAML

crd.yaml

```
1 apiVersion: apiextensions.k8s.io/v1beta1
2 kind: CustomResourceDefinition
3 metadata:
4   name: crontabs.training.lfs458.com
5     # This name must match names below.
6     # <plural>.<group> syntax
7 spec:
8   scope: Cluster      #Could also be Namespaced
9   group: training.lfs458.com
10  version: v1
11  names:
12    kind: CronTab      #Typically CamelCased for resource manifest
13    plural: crontabs   #Shown in URL
```

**YA
ML**

```

14     singular: crontab #Short name for CLI alias
15     shortNames:
16     - ct              #CLI short name

```

4. Add the new resource to the cluster.

```
student@master:~$ kubectl create -f crd.yaml
```

```

1 customresourcedefinition.apiextensions.k8s.io/crontabs.training.lfs458.com
2 created

```

5. View and describe the resource. The new line may be in the middle of the output. You'll note the **describe** output is unlike other objects we have seen so far.

```
student@master:~$ kubectl get crd
```

```

1 NAME                                CREATED AT
2 <output_omitted>
3 crontabs.training.lfs458.com        2018-08-03T05:25:20Z
4 <output_omitted>

```

```
student@master:~$ kubectl describe crd crontab<Tab>
```

```

1 Name:                crontabs.training.lfs458.com
2 Namespace:
3 Labels:              <none>
4 Annotations:         <none>
5 API Version:         apiextensions.k8s.io/v1
6 Kind:                CustomResourceDefinition
7 <output_omitted>

```

6. Now that we have a new API resource we can create a new object of that type. In this case it will be a crontab-like image, which does not actually exist, but is being used for demonstration.

```
student@master:~$ vim new-crontab.yaml
```

**YA
ML****new-crontab.yaml**

```

1 apiVersion: "training.lfs458.com/v1"
2   # This is from the group and version of new CRD
3 kind: CronTab
4   # The kind from the new CRD
5 metadata:
6   name: new-cron-object
7 spec:
8   cronSpec: "*/5 * * * *"
9   image: some-cron-image
10   #Does not exist

```

7. Create the new object and view the resource using short and long name.

```
student@master:~$ kubectl create -f new-crontab.yaml
```

```
1 crontab.training.lfs458.com/new-cron-object created
```

```
student@master:~$ kubectl get CronTab
```

```

1 NAME          AGE
2 new-cron-object 22s

```

```
student@master:~$ kubectl get ct
```

```

1 NAME          AGE
2 new-cron-object 29s

```

```
student@master:~$ kubectl describe ct
```

```

1 Name:          new-cron-object
2 Namespace:
3 Labels:        <none>
4
5 <output_omitted>
6
7 Spec:
8   Cron Spec:   */5 * * * *
9   Image:       some-cron-image
10  Events:      <none>

```

8. To clean up the resources we will delete the CRD. This should delete all of the endpoints and objects using it as well.

```
student@master:~$ kubectl delete -f crd.yaml
```

```

1 customresourcedefinition.apiextensions.k8s.io
2 "crontabs.training.lfs458.com" deleted

```

```
student@master:~$ kubectl get ct
```

```

1 Error from server (NotFound): Unable to list "training.lfs458.com/v1, Resource=crontabs":
2 the server could not find the requested resource (get crontabs.training.lfs458.com)

```

Chapter 9

Scheduling



9.1	Overview	166
9.2	Scheduler Settings	167
9.3	Policies	169
9.4	Affinity Rules	172
9.5	Taints and Tolerations	177
9.6	Labs	180

9.1 Overview

kube-scheduler

- Topology-aware algorithm to determine Pod placement
- Pod priority and preemption
- Labels used for each method
- **podAffinity** label encourage Pod on specific nodes
- Avoid via **podAntiAffinity**
- Node **taints** to repel specific Pods
- Pod **tolerations** of node **taints**
- **nodeSelector** to force specific Pods
- Require, prefer and evict

The larger and more diverse a Kubernetes deployment becomes the more administration of scheduling can be important. The **kube-scheduler** determines which nodes will run a Pod.

Users can set the priority of a pod, which will allow preemption of lower priority pods. The eviction of lower priority pods would then allow the higher priority pod to be scheduled.

The scheduler tracks the set of nodes in your cluster, filters them based on a set of predicates, then uses priority functions to determine on which node each Pod should be scheduled. The Pod spec as part of a request is sent to the **kubelet** on the node for creation.

The default scheduling decision can be affected through the use of Labels on nodes or Pods. Labels of **podAffinity**, **taints**, and **pod bindings** allow for configuration from the Pod or the node perspective. Some like **tolerations** allow a Pod to work with a Node, even when the Node has a **taint** that would otherwise preclude a Pod being scheduled.

Not all labels are drastic. Affinity settings may encourage a Pod to be deployed on a node, but would deploy the Pod elsewhere if the node was not available. Sometimes documentation may use the term require, but practice shows the setting to be more of a request. As beta features expect the specifics to change. Some settings will evict Pods from a node should the required condition no longer be true such as `requiredDuringSchedulingRequiredDuringExecution`.

Others options, like a custom scheduler, need to be programmed and deployed into your Kubernetes cluster.

9.2 Scheduler Settings

Predicates

- **Predicates** check used to filter out inappropriate nodes
- Particular order from [predicates.go](#)

```
predicatesOrdering = []string{CheckNodeConditionPred,  
GeneralPred, HostNamePred, PodFitsHostPortsPred,  
MatchNodeSelectorPred, PodFitsResourcesPred, NoDiskConflictPred,  
PodToleratesNodeTaintsPred, PodToleratesNodeNoExecuteTaintsPred,  
CheckNodeLabelPresencePred, checkServiceAffinityPred,  
MaxEBSVolumeCountPred, MaxGCEPDVolumeCountPred,  
MaxAzureDiskVolumeCountPred, CheckVolumeBindingPred,  
NoVolumeZoneConflictPred, CheckNodeMemoryPressurePred,  
CheckNodeDiskPressurePred, MatchInterPodAffinityPred}
```

The scheduler goes through a set of filters, or predicates to find available nodes, then ranks each node using priority functions. The node with the highest rank is selected to run the Pod.

The **predicates**, such as PodFitsHost or NoDiskConflict are evaluated in a particular and configurable order. In this way a node has the least amount of checks for new Pod deployment, which can be useful to exclude a node from unnecessary checks if node is not in the proper condition.

For example, there is a filter called HostNamePred, which is also known as HostName, which filters out nodes that do not match the node name specified in the pod specification. Another predicate for PodFitsResources to make sure that the available CPU and memory can fit the resources required by the Pod.

The scheduler can be updated by passing a configuration of kind: Policy which can order predicates, give special weights to priorities and even hardPodAffinitySymmetricWeight which deploys Pods such that if we set Pod A to run with Pod B, then Pod B should automatically be run with Pod A.

Priorities

- Set of **priority functions** used to weight resources
- Iteration of current status each adding or removing 1 to priority
- Highest value node chosen by scheduler
- SelectorSpreadPriority assigned to node with least pods
- View list at [master/pkg/scheduler/algorithm/priorities](#)
- **PriorityClasses** can be used to allocate and preempt other pods
- Use a Pod Disruption Budget to limit disruption by preemption

Unless Pod and node affinity has been configured the SelectorSpreadPriority setting, which ranks nodes based on the number of existing running pods, will select the node with the least amount of Pods. This is a basic way to spreading Pods across the cluster.

Other priorities can be used for particular cluster needs. The ImageLocalityPriorityMap favors nodes which already have downloaded container images. The total sum of image size is compared with the largest having highest priority, but does not check the image about to be used.

There currently are more than ten included priorities which range from checking the existence of a label to choosing a node with the most requested CPU and memory usage.

A feature stable as of v1.14 allows the setting of a PriorityClass and assigning pods via the use of PriorityClassName settings. This allows user to preempt, or evict, lower priority pods so that their higher priority pods can be scheduled. The **kube-scheduler** determines a node where the pending pod could run if one or more existing pods were evicted. If a node is found the low priority pod(s) are evicted and the higher priority pod scheduled. The use of a Pod Disruption Budget (PDB) is a way to limit the number of pods preemption evicts to ensure enough pods remain running. The scheduler will remove pods even if the PDB is violated if no other options are available.

9.3 Policies

Scheduling Policies

```
"kind" : "Policy",
"apiVersion" : "v1",
"predicates" : [
  {"name" : "MatchNodeSelector", "order": 6},
  {"name" : "PodFitsHostPorts", "order": 2},
  {"name" : "PodFitsResources", "order": 3},
  {"name" : "NoDiskConflict", "order": 4},
  {"name" : "PodToleratesNodeTaints", "order": 5},
  {"name" : "PodFitsHost", "order": 1}
],
"priorities" : [
  {"name" : "LeastRequestedPriority", "weight" : 1},
  {"name" : "BalancedResourceAllocation", "weight" : 1},
  {"name" : "ServiceSpreadingPriority", "weight" : 2},
  {"name" : "EqualPriority", "weight" : 1}
],
"hardPodAffinitySymmetricWeight" : 10
}
```

The default scheduler contains a number of predicates and priorities; however, these can be changed via a scheduler policy file.

Typically, you will configure a scheduler with this policy using the `--policy-config-file` parameter and define a name for this scheduler using the `--scheduler-name` parameter. You will then have two schedulers running and will be able to specify which scheduler to use in the Pod specification.

With multiple schedulers there could be conflict in Pod allocation. Each Pod should declare which scheduler should be used. But if separate schedulers determine a node is eligible because of available resources and both attempt to deploy, causing the resource to no longer be available, a conflict would occur. The current solution is for the local `kubelet` to return the Pods to the scheduler for reassignment. Eventually one Pod will succeed and the other be scheduled elsewhere.

Pod Specification

- Pod Specification fields

- nodeName
- nodeSelector
- affinity
- tolerations
- schedulerName

Most scheduling decisions can be made as part of the the Pod spec. The `nodeName` and `nodeSelector` options allow a Pod to be assigned to a single node or a group of nodes with particular labels.

Affinity and anti-affinity can be used to require or prefer which node is used by the scheduler. If using a preference instead a matching node is chosen first, but other nodes would be used if no match is present.

The use of `taints` allows a node to be labeled such that Pods would not be scheduled for some reason, such as the `master` node after initialization. A `toleration` allows a Pod to ignore the taint and be scheduled assuming other requirements are met.

Should none of these options meet the needs of the cluster there is also the ability to deploy a custom scheduler. Each Pod could then include a `schedulerName` to choose which schedule to use.

Specifying Node Label

- Match a label with `nodeSelector`
- Pod remains in `Pending` state until a suitable Node is found

```
spec:
  containers:
  - name: redis
    image: redis
  nodeSelector:
    net: fast
```

The `nodeSelector` field in a pod specification provides a straightforward way to target a node or set of nodes, using one or more key-value pairs.

Setting the `nodeSelector` tells the scheduler to place the pod on a node that matches the labels. All listed selectors must be met, but the node could have more labels. In the above example any node with a key of `net` set to `fast` would be a candidate for scheduling. Remember that labels are admin created tags, with no tie to actual resources. This node could have a slow network.

The pod would remain `Pending` until a node is found with the matching labels.

The use of affinity/anti-affinity should be able to express every feature as `nodeSelector`.

9.4 Affinity Rules

Pod Affinity Rules

- Uses In, NotIn, Exists, and DoesNotExist operators
- `requiredDuringSchedulingIgnoredDuringExecution`
- `preferredDuringSchedulingIgnoredDuringExecution`
- `affinity`
 - `podAffinity`
 - `podAntiAffinity`
- `topologyKey`

Pods which may communicate a lot or share data may operate best if co-located, which would be a form of affinity. For greater fault tolerance you may want Pods to be as separate as possible, which would be anti-affinity. These settings are used by the scheduler based on labels of Pods already running. As a result the scheduler must interrogate each node and track the labels of running Pods. Clusters larger than several hundred nodes may see significant performance loss.

The use of `requiredDuringSchedulingIgnoredDuringExecution` means that the Pod will not be scheduled on a node unless the following operator is true. If the operator changes to become false in the future the Pod will continue to run. This could be seen as a hard rule.

Similar is `preferredDuringSchedulingIgnoredDuringExecution` which will choose a node with the desired setting before those without. Should no properly labeled nodes be available the Pod will execute anyway. This is more of a soft settings which declares a preference instead of a requirement.

With the use of `podAffinity` the scheduler will try to schedule Pods together. The use of `podAntiAffinity` would cause the scheduler to keep Pods on different nodes.

The `topologyKey` allows a general grouping of Pod deployment. Affinity (or the inverse anti-affinity) will try to run on nodes with the declared topology key and running Pods with a particular label.

The `topologyKey` could be any legal key, but there some important considerations:

- If using `requiredDuringScheduling` and the admission controller `LimitPodHardAntiAffinityTopology` setting the `topologyKey` must be set to `kubernetes.io/hostname`.
- If using `PreferredDuringScheduling` an empty `topologyKey` is assumed to be all, or the combination of `kubernetes.io/hostname`, `topology.kubernetes.io/zone` and `topology.kubernetes.io/region`.

podAffinity Example

```
spec:
  affinity:
    podAffinity:
      requiredDuringSchedulingIgnoredDuringExecution:
        - labelSelector:
            matchExpressions:
              - key: security
                operator: In
                values:
                  - S1
          topologyKey: topology.kubernetes.io/zone
```

An example of affinity and podAffinity settings. This also requires a particular label to be matched when the Pod starts, but not required if the label is later removed.

Inside the declared topology zone the Pod can be scheduled on a node running a Pod with a key label of `security` and a value of `S1`. If this requirement is not met the Pod will remain in a `Pending` state.

podAntiAffinity Example

```
podAntiAffinity:
  preferredDuringSchedulingIgnoredDuringExecution:
  - weight: 100
    podAffinityTerm:
      labelSelector:
        matchExpressions:
        - key: security
          operator: In
          values:
          - S2
      topologyKey: kubernetes.io/hostname
```

With `podAntiAffinity` we can prefer to avoid nodes with a particular label. In this case the scheduler will prefer to avoid a node with a key set to `security` and value of `S2`.

In a large, varied, environment there may be multiple situations to be avoided. As a preference this settings tries to avoid certain labels, but will still schedule the Pod on some node. As the Pod will still run we can provide a `weight` to a particular rule. The weights can be declared in a value from 1 to 100. The scheduler then tries to choose, or avoid, the node with the greatest combined value.

Node Affinity Rules

- Similar to Pod affinity rules, declared with affinity
- Uses In, NotIn, Exists, and DoesNotExist operators
- `requiredDuringSchedulingIgnoredDuringExecution`
- `preferredDuringSchedulingIgnoredDuringExecution`
- Planned for future
 - `requiredDuringSchedulingRequiredDuringExecution`

Where Pod affinity/anti-affinity has to do with other Pods, the use of `nodeAffinity` allows Pod scheduling based of node labels. This is similar, and will some day replace, use of the `nodeSelector` setting. The scheduler will not look at other Pods on the system, but the labels of the nodes. This should have much less performance impact on the cluster, even with a large number of nodes.

Until `nodeSelector` has been fully deprecated both the selector and required labels must be met for a Pod to be scheduled.

Node Affinity Example

```
spec:
  affinity:
    nodeAffinity:
      requiredDuringSchedulingIgnoredDuringExecution:
        nodeSelectorTerms:
          - matchExpressions:
              - key: kubernetes.io/colo-tx-name
                operator: In
                values:
                  - tx-aus
                  - tx-dal
            preferredDuringSchedulingIgnoredDuringExecution:
              - weight: 1
                preference:
                  matchExpressions:
                    - key: disk-speed
                      operator: In
                      values:
                        - fast
                        - quick
```

The first `nodeAffinity` rule requires a node with a key of `kubernetes.io/colo-tx-name` which has one of two possible values `tx-aus` or `tx-dal`.

The second rule gives extra weight to nodes with a key of `disk-speed` with a value of `fast` or `quick`. The Pod will be scheduled on some node in any case this just prefers a particular label.

9.5 Taints and Tolerations

Taint

- Expressed as `key=value:effect`
- `key` and `value` created by admin
- Effect must be **NoSchedule**, **PreferNoSchedule**, or **NoExecute**
- Only nodes can be tainted currently
- Multiple taints possible on a node

A node with a particular taint will repel Pods without tolerations for that taint.

The `key` and `value` used can be any legal string, while allows flexibility to prevent Pods from running on nodes based off of any need. If a Pod does not have an existing `toleration` the scheduler will not consider the tainted node.

There are three effects, or ways to handle Pod scheduling.

- **NoSchedule** The scheduler will not schedule a Pod on this node unless the Pod has this `toleration`. Existing Pods continue to run, regardless of `toleration`.
- **PreferNoSchedule** The scheduler will avoid using this node unless there are no untainted nodes for the Pods `toleration`. Existing Pods are unaffected.
- **NoExecute** This taint will cause existing Pods to be evacuated and no future Pods scheduled. Should an existing Pod have a `toleration` it will continue to run. If the Pod `tolerationSeconds` value is set the Pod will remain for that many seconds then be evicted. Certain node issues will cause **kubelet** to add 300 second `tolerations` to avoid unnecessary evictions.

If a node has multiple taints the scheduler ignores those with matching tolerations. The remaining un-ignored taints have their typical effect.

The use of `TaintBasedEvictions` is still an alpha feature. The **kubelet** uses taints to rate-limit evictions when the node has problems.

Tolerations

- Pod setting to run on tainted nodes
- Same **effects** as node taints
- Two operators
 - Exists
 - Equal which requires value

```
tolerations:  
- key: "server"  
  operator: "Equal"  
  value: "ap-east"  
  effect: "NoExecute"  
  tolerationSeconds: 3600
```

Setting tolerations on a node are used to schedule on tainted nodes. This provides an easy way to avoid Pods using the node. Only those with a particular toleration would be scheduled.

An operator can be included in a Pod spec, defaulting to `Equal` if not declared. The use of operator `Equal` requires a value to match. The `Exists` operator should not be specified. If an empty key uses the `Exists` operator it will tolerate every taint. If there is no effect, but a key and operator are declared all effects are matched with the declared key.

In the above example the Pod will remain on the server with a key of `server` and value of `ap-east` for 3600 seconds after the node has been tainted with `NoExecute`. When the time runs out the Pod will be evicted.

Custom Scheduler

- Create and deploy custom scheduler as a container
- Run multiple schedulers simultaneously
- Pods can declare which scheduler to use
- Scheduler must be running or Pod remains `Pending`
- View scheduler and other information with **kubectl get events**

If the default scheduling mechanisms are not flexible enough for your needs you can write your own scheduler. The programming of a custom scheduler is outside the scope of this course, but you may want to start with the existing scheduler code which can be found here: <https://github.com/kubernetes/kubernetes/tree/master/pkg/scheduler>

If a Pod spec does not declare which scheduler to use the standard scheduler is used by default. If the Pod declares a scheduler and that container is not running the Pod would remain in `Pending` state forever.

The end result of the scheduling process is that a pod gets a **binding** that specifies which node it should run on. A binding is a Kubernetes API primitive in the **api/v1** group. Technically without any scheduler running, you could still schedule a pod on a node, by specifying a binding for that pod.

9.6 Labs

Exercise 9.1: Assign Pods Using Labels

Overview

While allowing the system to distribute Pods on your behalf is typically the best route, you may want to determine which nodes a Pod will use. For example you may have particular hardware requirements to meet for the workload. You may want to assign VIP Pods to new, faster hardware and everyone else to older hardware.

In this exercise we will use `labels` to schedule Pods to a particular node. Then we will explore `taints` to have more flexible deployment in a large environment.

1. Begin by getting a list of the nodes. They should be in the ready state and without added labels or taints.

```
student@master:~$ kubectl get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
master	Ready	master	44h	v1.19.0
worker	Ready	<none>	43h	v1.19.0

2. View the current labels and taints for the nodes.

```
student@master:~$ kubectl describe nodes |grep -A5 -i label
```

```
1 Labels:          beta.kubernetes.io/arch=amd64
2                  beta.kubernetes.io/os=linux
3                  kubernetes.io/arch=amd64
4                  kubernetes.io/hostname=c-jg5n
5                  kubernetes.io/os=linux
6                  node-role.kubernetes.io/master=
7 --
8 Labels:          beta.kubernetes.io/arch=amd64
9                  beta.kubernetes.io/os=linux
10                 kubernetes.io/arch=amd64
11                 kubernetes.io/hostname=c-v2t3
12                 kubernetes.io/os=linux
13 --
```

```
student@master:~$ kubectl describe nodes |grep -i taint
```

```
1 Taints:          <none>
2 Taints:          <none>
```

3. Get a count of how many containers are running on both the master and worker nodes. There are about 24 containers running on the master in the following example, and eight running on the worker. There are status lines which increase the `wc` count. You may have more or less, depending on previous labs and cleaning up of resources. If you are using **cri-o** you can view containers using **crictrl ps**.

```
student@master:~$ kubectl get deployments --all-namespaces
```

NAMESPACE	NAME	READY	UP-TO-DATE	AVAILABLE	AGE
default	secondapp	1	1	1	37m
default	thirdpage	1	1	1	24m
kube-system	calico-kube-controllers	0	0	0	44h
kube-system	coredns	2/2	2	2	44h

```
student@master:~$ sudo docker ps |wc -l
```

```
1 21
```

```
student@worker:~$ sudo docker ps |wc -l
```

```
1 11
```

4. For the purpose of the exercise we will assign the master node to be VIP hardware and the secondary node to be for others.

```
student@master:~$ kubectl label nodes master status=vip
```

```
1 node/master labeled
```

```
student@master:~$ kubectl label nodes worker status=other
```

```
1 node/worker labeled
```

5. Verify your settings. You will also find there are some built in labels such as hostname, os and architecture type. The output below appears on multiple lines for readability.

```
student@master:~$ kubectl get nodes --show-labels
```

```
1 NAME          STATUS    ROLES    AGE   VERSION   LABELS
2 master        Ready     master   44h   v1.19.0   beta.kubernetes.io/arch=
3 amd64,beta.kubernetes.io/os=linux,kubernetes.io/hostname=master,
4 node-role.kubernetes.io/master=,status=vip
5 worker        Ready     <none>    44h   v1.19.0   beta.kubernetes.io/arch=
6 amd64,beta.kubernetes.io/os=linux,kubernetes.io/hostname=worker,status=other
```

6. Create `vip.yaml` to spawn four busybox containers which sleep the whole time. Include the `nodeSelector` entry.

```
student@master:~$ vim vip.yaml
```

YAML

`vip.yaml`

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: vip
5 spec:
6   containers:
7   - name: vip1
8     image: busybox
9     args:
10    - sleep
11    - "1000000"
12   - name: vip2
13     image: busybox
14     args:
15    - sleep
16    - "1000000"
17   - name: vip3
18     image: busybox
19     args:
20    - sleep
21    - "1000000"
22   - name: vip4
23     image: busybox
```

YAML

```

24     args:
25     - sleep
26     - "1000000"
27     nodeSelector:
28     status: vip

```

7. Deploy the new pod. Verify the containers have been created on the master node. It may take a few seconds for all the containers to spawn. Check both the master and the secondary nodes.

```
student@master:~$ kubectl create -f vip.yaml
```

```
1 pod/vip created
```

```
student@master:~$ sudo docker ps |wc -l
```

```
1 26
```

```
student@worker:~$ sudo docker ps |wc -l
```

```
1 11
```

8. Delete the pod then edit the file, commenting out the `nodeSelector` lines. It may take a while for the containers to fully terminate.

```
student@master:~$ kubectl delete pod vip
```

```
1 pod "vip" deleted
```

```
student@master:~$ vim vip.yaml
```

```

1 ....
2 # nodeSelector:
3 #   status: vip

```

9. Create the pod again. Containers should now be spawning on either node. You may see pods for the daemonsets as well.

```
student@master:~$ kubectl get pods
```

```

1 NAME                                READY   STATUS    RESTARTS   AGE
2 secondapp-85765cd95c-2q9sx          1/1     Running   0           43m
3 thirdpage-7c9b56bfd-2q5pr           1/1     Running   0           30m

```

```
student@master:~$ kubectl create -f vip.yaml
```

```
1 pod/vip created
```

10. Determine where the new containers have been deployed. They should be more evenly spread this time.

```
student@master:~$ sudo docker ps |wc -l
```

```
1 21
```

```
student@worker:~$ sudo docker ps |wc -l
```

```
1 16
```

11. Create another file for other users. Change the names from vip to others, and uncomment the nodeSelector lines.

```
student@master:~$ cp vip.yaml other.yaml

student@master:~$ sed -i s/vip/other/g other.yaml

student@master:~$ vim other.yaml
```

YAML

other.yaml

```
1  ....
2  nodeSelector:
3  status: other
```

12. Create the other containers. Determine where they deploy.

```
student@master:~$ kubectl create -f other.yaml
```

```
1 pod/other created
```

```
student@master:~$ sudo docker ps |wc -l
```

```
1 21
```

```
student@worker:~$ sudo docker ps |wc -l
```

```
1 21
```

13. Shut down both pods and verify they terminated. Only our previous pods should be found.

```
student@master:~$ kubectl delete pods vip other
```

```
1 pod "vip" deleted
2 pod "other" deleted
```

```
student@master:~$ kubectl get pods
```

	NAME	READY	STATUS	RESTARTS	AGE
2	secondapp-85765cd95c-2q9sx	1/1	Running	0	51m
3	thirdpage-7c9b56bfd-2q5pr	1/1	Running	0	40m

Exercise 9.2: Using Taints to Control Pod Deployment

Use taints to manage where Pods are deployed or allowed to run. In addition to assigning a Pod to a group of nodes, you may also want to limit usage on a node or fully evacuate Pods. Using taints is one way to achieve this. You may remember that the master node begins with a NoSchedule taint. We will work with three taints to limit or remove running pods.

1. Verify that the master and secondary node have the minimal number of containers running.

```
student@master:~$ kubectl delete deployment secondapp \
thirdpage
```

```
1 deployment.apps "secondapp" deleted
2 deployment.apps "thirdpage" deleted
```

2. Create a deployment which will deploy eight **nginx** containers. Begin by creating a YAML file.

```
student@master:~$ vim taint.yaml
```



taint.yaml

```

1  apiVersion: apps/v1
2  kind: Deployment
3  metadata:
4    name: taint-deployment
5  spec:
6    replicas: 8
7    selector:
8      matchLabels:
9        app: nginx
10   template:
11     metadata:
12       labels:
13         app: nginx
14     spec:
15       containers:
16       - name: nginx
17         image: nginx:1.16.1
18         ports:
19         - containerPort: 80

```

3. Apply the file to create the deployment.

```
student@master:~$ kubectl apply -f taint.yaml
```

```
1 deployment.apps/taint-deployment created
```

4. Determine where the containers are running. In the following example three have been deployed on the master node and five on the secondary node. Remember there will be other housekeeping containers created as well. Your numbers may be slightly different.

```
student@master:~$ sudo docker ps |grep nginx
```

```

1  00c1be5df1e7      nginx@sha256:e3456c851a152494c3e.....
2  <output_omitted>

```

```
student@master:~$ sudo docker ps |wc -l
```

```
1 27
```

```
student@worker:~$ sudo docker ps |wc -l
```

```
1 17
```

5. Delete the deployment. Verify the containers are gone.

```
student@master:~$ kubectl delete deployment taint-deployment
```

```
1 deployment.apps "taint-deployment" deleted
```

```
student@master:~$ sudo docker ps |wc -l
```

```
1 21
```

6. Now we will use a taint to affect the deployment of new containers. There are three taints, `NoSchedule`, `PreferNoSchedule` and `NoExecute`. The taints having to do with schedules will be used to determine newly deployed containers, but will not affect running containers. The use of `NoExecute` will cause running containers to move.

Taint the secondary node, verify it has the taint then create the deployment again. We will use the key of `bubba` to illustrate the key name is just some string an admin can use to track Pods.


```
student@master:~$ kubectl taint nodes worker \
    bubba=value:PreferNoSchedule
```

```
1 node/worker tainted
```

```
student@master:~$ kubectl describe node |grep Taint
```

```
1 Taints:                bubba=value:PreferNoSchedule
2 Taints:                <none>
```

```
student@master:~$ kubectl apply -f taint.yaml
```

```
1 deployment.apps/taint-deployment created
```

7. Locate where the containers are running. We can see that more containers are on the master, but there still were some created on the secondary. Delete the deployment when you have gathered the numbers.

```
student@master:~$ sudo docker ps |wc -l
```

```
1 21
```

```
student@worker:~$ sudo docker ps |wc -l
```

```
1 23
```

```
student@master:~$ kubectl delete deployment taint-deployment
```

```
1 deployment.apps "taint-deployment" deleted
```

8. Remove the taint, verify it has been removed. Note that the key is used with a minus sign appended to the end.

```
student@master:~$ kubectl taint nodes worker bubba-
```

```
1 node/worker untainted
```

```
student@master:~$ kubectl describe node |grep Taint
```

```
1 Taints:                <none>
2 Taints:                <none>
```

9. This time use the NoSchedule taint, then create the deployment again. The secondary node should not have any new containers, with only daemonsets and other essential pods running.

```
student@master:~$ kubectl taint nodes worker \
    bubba=value:NoSchedule
```

```
1 node/worker tainted
```

```
student@master:~$ kubectl apply -f taint.yaml
```

```
1 deployment.apps/taint-deployment created
```

```
student@master:~$ sudo docker ps |wc -l
```

```
1 21
```

```
student@worker:~$ sudo docker ps |wc -l
```

```
1 23
```

10. Remove the taint and delete the deployment. When you have determined that all the containers are terminated create the deployment again. Without any taint the containers should be spread across both nodes.

```
student@master:~$ kubectl delete deployment taint-deployment
```

```
1 deployment.apps "taint-deployment" deleted
```

```
student@master:~$ kubectl taint nodes worker bubba-
```

```
1 node/worker untainted
```

```
student@master:~$ kubectl apply -f taint.yaml
```

```
1 deployment.apps/taint-deployment created
```

```
student@master:~$ sudo docker ps |wc -l
```

```
1 27
```

```
student@worker:~$ sudo docker ps |wc -l
```

```
1 17
```

11. Now use the NoExecute to taint the secondary (**worker**) node. Wait a minute then determine if the containers have moved. The DNS containers can take a while to shutdown. Some containers will remain on the worker node to continue communication from the cluster.

```
student@master:~$ kubectl taint nodes worker \
    bubba=value:NoExecute
```

```
1 node "worker" tainted
```

```
student@master:~$ sudo docker ps |wc -l
```

```
1 37
```

```
student@worker:~$ sudo docker ps |wc -l
```

```
1 5
```

12. Remove the taint. Wait a minute. Note that all of the containers did not return to their previous placement.

```
student@master:~$ kubectl taint nodes worker bubba-
```

```
1 node/worker untainted
```

```
student@master:~$ sudo docker ps |wc -l
```

```
1 32
```

```
student@worker:~$ sudo docker ps |wc -l
```

```
1 6
```

13. Remove the deployment a final time to free up resources.

```
student@master:~$ kubectl delete deployment taint-deployment
```

```
1 deployment.apps "taint-deployment" deleted
```

Chapter 10

Security



In this chapter, we are going to talk about how an API call is ingested into the cluster. We will go through the details of the three different phases each call goes through. We'll look at the different types of authentication that are available to us, and work with RBAC, which handles the authorization on our behalf. We'll also configure pod policies. The use of a pod policy allows us to configure the details of what a container can do, and how processes run within that container.

We'll finish by understanding a network policy. If the network plugin honors a network policy, this allows you to isolate a pod from other pods in the environment. The default Kubernetes architecture says that all pods should be able to see all pods. So, this is a change from that historic approach. But, as we use it in a production environment, you may want to limit ingress

10.1	Security Overview	187
10.2	Accessing the API	187
10.3	Authentication	188
10.4	Authorization	189
10.5	ABAC	189
10.6	RBAC	190
10.7	RBAC Process Overview	190
10.8	Admission Controller	190
10.9	Security Contexts	191
10.10	Pod Security Policies	191
10.11	Network Security Policies	192
10.12	Network Security Policy Example	192
10.13	Default Policy Example	193
10.14	Labs	194

10.1 Security Overview

- Security is a big and complex topic, especially in a distributed system like Kubernetes. Thus, we are just going to cover some of the concepts that deal with security in the context of Kubernetes.
- Then, we are going to focus on the authentication aspect of the API server and we will dive into authorization, looking at things like ABAC and RBAC, which is now the default configuration when you bootstrap a Kubernetes cluster with **kubeadm**.

- We are going to look at the admission control system, which lets you look at and possibly modify the requests that are coming in, and do a final deny or accept on those requests.
- Following that, we're going to look at a few other concepts, including how you can secure your Pods more tightly using security contexts and pod security policies, which are full-fledged API objects in Kubernetes.
- Finally, we will look at network policies. By default, we tend not to turn on network policies, which let any traffic flow through all of our pods, in all the different namespaces. Using network policies, we can actually define Ingress rules so that we can restrict the Ingress traffic between the different namespaces. The network tool in use, such as Flannel or Calico will determine if a network policy can be implemented. As Kubernetes becomes more mature, this will become a strongly suggested configuration.

10.2 Accessing the API

- To perform any action in a Kubernetes cluster, you need to access the API and go through three main steps:
 - Authentication
 - Authorization (ABAC or RBAC)
 - Admission Control.
- These steps are described in more detail in the official documentation at <https://kubernetes.io/docs/reference/access-authn-authz/controlling-access/> about controlling access to the API and illustrated by the picture below, from <https://kubernetes.io/docs/reference/command-line-tools-reference/kube-apiserver/>:

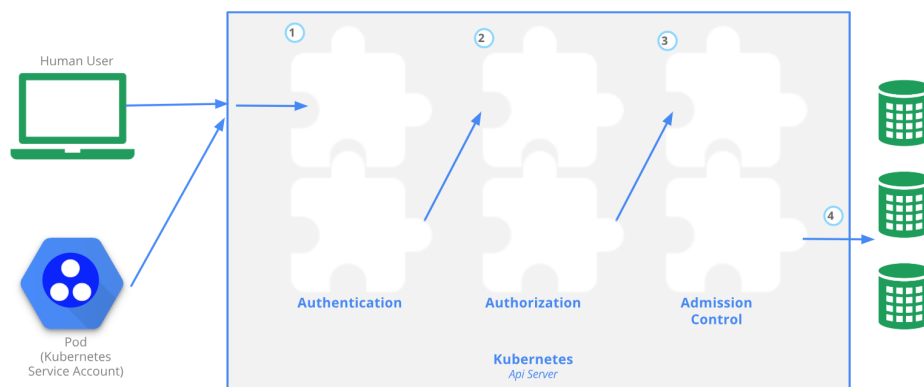


Figure 10.1: Accessing the API

- Once a request reaches the API server securely, it will first go through any authentication module that has been configured. The request can be rejected if authentication fails or it gets authenticated and passed to the authorization step.
- At the authorization step, the request will be checked against existing policies. It will be authorized if the user has the permissions to perform the requested actions. Then, the requests will go through the last step of admission. In general, admission controllers will check the actual content of the objects being created and validate them before admitting the request.
- In addition to these steps, the requests reaching the API server over the network are encrypted using TLS. This needs to be properly configured using SSL certificates. If you use kubeadm, this configuration is done for you; otherwise, follow Kelsey Hightower's guide, *Kubernetes The Hard Way* at <https://github.com/kelseyhightower/kubernetes-the-hard-way>, or the API server configuration options at <https://kubernetes.io/docs/reference/command-line-tools-reference/kube-apiserver/>.

10.3 Authentication

- There are three main points to remember with authentication in Kubernetes:
 1. In its straightforward form, authentication is done with certificates, tokens or basic authentication (i.e. username and password).
 2. Users are not created by the API, but should be managed by an external system.
 3. System accounts are used by processes to access the API.
(<https://kubernetes.io/docs/tasks/configure-pod-container/configure-service-account/>)
- There are two more advanced authentication mechanisms: **Webhooks** which can be used to verify bearer tokens, and connection with an external **OpenID** provider.
- The type of authentication used is defined in the kube-apiserver startup options. Below are four examples of a subset of configuration options that would need to be set depending on what choice of authentication mechanism you choose:
 - `--basic-auth-file`
 - `--oidc-issuer-url`
 - `--token-auth-file`
 - `--authorization-webhook-config-file`
- One or more Authenticator Modules are used: x509 Client Certs; static token, bearer or bootstrap token; static password file; service account and OpenID connect tokens. Each is tried until successful, and the order is not guaranteed. Anonymous access can also be enabled, otherwise you will get a 401 response. Users are not created by the API, and should be managed by an external system.
- To learn more about authentication, see the official Kubernetes documentation at <https://kubernetes.io/docs/reference/access-authn-authz/authentication/>.

10.4 Authorization

- Once a request is authenticated, it needs to be authorized to be able to proceed through the Kubernetes system and perform its intended action.
- There are three main authorization modes and two global Deny/Allow settings. The three main modes are:
 1. ABAC
 2. RBAC
 3. WebHook
- They can be configured as kube-apiserver startup options:
 - `--authorization-mode=ABAC` `authorization-mode=RBAC`
 - `--authorization-mode=Webhook`
 - `--authorization-mode=AlwaysDeny`
 - `--authorization-mode=AlwaysAllow`
- The authorization modes implement policies to allow requests. Attributes of the requests are checked against the policies (e.g. user, group, namespace, verb).

10.5 ABAC

- **ABAC** stands for **Attribute Based Access Control**. It was the first authorization model in Kubernetes that allowed administrators to implement the right policies. Today, **RBAC** is becoming the default authorization mode.
- Policies are defined in a JSON file and referenced to by a **kube-apiserver** startup option:

```
--authorization-policy-file=my_policy.json
```

- For example, the policy file shown below authorizes user Bob to read pods in the namespace foobar:

```
{ "apiVersion":
  "abac.authorization.kubernetes.io/v1beta1", "kind":
  "Policy", "spec": { "user": "bob", "namespace":
    "foobar", "resource": "pods", "readonly": true } }
```

You can check other policy examples in the Kubernetes documentation.

10.6 RBAC

- **RBAC** (<https://kubernetes.io/docs/reference/access-authn-authz/rbac/>) stands for **Role Based Access Control**.
- All resources are modeled API objects in Kubernetes, from Pods to Namespaces. They also belong to **API Groups**, such as `core` and `apps`. These resources allow operations such as Create, Read, Update, and Delete (CRUD), which we have been working with so far. Operations are called verbs inside YAML files. Adding to these basic components, we will add more elements of the API, which can then be managed via RBAC.
- **Rules** are operations which can act upon an API group. **Roles** are a group of rules which affect, or scope, a single namespace, whereas `ClusterRoles` have a scope of the entire cluster.
- Each operation can act upon one of three subjects, which are **User Accounts** which don't exist as API objects, `Service Accounts`, and `Groups` which are known as `clusterrolebinding` when using **kubectl**.
- RBAC is then writing rules to allow or deny operations by users, roles or groups upon resources.

10.7 RBAC Process Overview

- While RBAC can be complex, the basic flow is to create a certificate for a user. As a user is not an API object of Kubernetes, we are requiring outside authentication, such as OpenSSL certificates. After generating the certificate against the cluster certificate authority, we can set that credential for the user using a context.
- Roles can then be used to configure an association of `apiGroups`, `resources`, and the verbs allowed to them. The user can then be bound to a role limiting what and where they can work in the cluster.
- Here is a summary of the RBAC process:
 - Determine or create namespace
 - Create certificate credentials for user
 - Set the credentials for the user to the namespace using a context
 - Create a role for the expected task set
 - Bind the user to the role
 - Verify the user has limited access.

10.8 Admission Controller

- The last step in letting an API request into Kubernetes is admission control.
- Admission controllers are pieces of software that can access the content of the objects being created by the requests. They can modify the content or validate it, and potentially deny the request.
- Admission controllers are needed for certain features to work properly. Controllers have been added as Kubernetes matured. Starting with 1.13.1 release of the **kube-apiserver** the admission controllers are now compiled into the binary, instead of a list passed during execution. To enable or disable you can pass the following options, changing out the plugins you want to enable or disable.

```
--enable-admission-plugins=NamespaceLifecycle,LimitRanger
--disable-admission-plugins=PodNodeSelector
```

- The first controller is `Initializers` which will allow the dynamic modification of the API request, providing great flexibility. Each admission controller functionality is explained in the documentation. For example, the `ResourceQuota` controller will ensure that the object created does not violate any of the existing quotas.

10.9 Security Contexts

- Pods and containers within pods can be given specific security constraints to limit what processes running in containers can do. For example, the UID of the process, the Linux capabilities, and the filesystem group can be limited.
- This security limitation is called a **security context**. It can be defined for the entire pod or per container, and is represented as additional sections in the resources manifests. The notable difference is that Linux capabilities are set at the container level.
- For example, if you want to enforce a policy that containers cannot run their process as the root user, you can add a pod security context like the one below:

YAML

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: nginx
5 spec:
6   securityContext:
7     runAsNonRoot: true
8   containers:
9   - image: nginx
10     name: nginx
```

- Then, when you create this pod, you will see a warning that the container is trying to run as root and that it is not allowed. Hence, the Pod will never run:

```
$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx	0/1	container has runAsNonRoot and image will run as root	0	10s

- You can read more in the Kubernetes documentation about configuring security contexts to give proper constraints to your pods or containers at <https://kubernetes.io/docs/tasks/configure-pod-container/security-context/>.

10.10 Pod Security Policies

- To automate the enforcement of security contexts, you can define Pod Security Policies (**PSP**). (See <https://kubernetes.io/docs/concepts/policy/pod-security-policy/>). A PSP is defined via a standard Kubernetes manifest following the PSP API schema. An example is presented on the next page.
- These policies are cluster-level rules that govern what a pod can do, what they can access, what user they run as, etc.
- For instance, if you do not want any of the containers in your cluster to run as the root user, you can define a PSP to that effect. You can also prevent containers from being privileged or use the host network namespace, or the host PID namespace.
- While **PSP** has been helpful there are other methods gaining popularity. The **Open Policy Agent (OPA)**, often pronounced as "oh-pa", provides a unified set of tools and policy framework. This allows a single point of configuration for all of your cloud deployments.
- **OPA** can be deployed as an admission controller inside of Kubernetes, which allows OPA to enforce or mutate requests as they are received. Using the **OPA Gatekeeper** it can be deployed using Custom Resource Definitions. More can be found here: <https://www.openpolicyagent.org/>
- You can see an example of a PSP below:

YAML

```
1 apiVersion: policy/v1beta1
2 kind: PodSecurityPolicy
3 metadata:
4   name: restricted
5 spec:
6   seLinux:
7     rule: RunAsAny
8   supplementalGroups:
9     rule: RunAsAny
10  runAsUser:
11    rule: MustRunAsNonRoot
12  fsGroup:
13    rule: RunAsAny
```

- For Pod Security Policies to be enabled, you need to configure the admission controller of the controller-manager to contain PodSecurityPolicy. These policies make even more sense when coupled with the **RBAC** configuration in your cluster. This will allow you to finely tune what your users are allowed to run and what capabilities and low level privileges their containers will have.
- See the PSP RBAC example at <https://github.com/kubernetes/examples/blob/master/staging/podsecuritypolicy/rbac/README.md> on **GitHub** for more details.

10.11 Network Security Policies

- By default, all pods can reach each other; all ingress and egress traffic is allowed. This has been a high-level networking requirement in Kubernetes. However, network isolation can be configured and traffic to pods can be blocked. In newer versions of Kubernetes, egress traffic can also be blocked. This is done by configuring a **NetworkPolicy**. As all traffic is allowed, you may want to implement a policy that drops all traffic, then, other policies which allow desired ingress and egress traffic.
- The `spec` of the policy can narrow down the effect to a particular namespace, which can be handy. Further settings include a `podSelector`, or label, to narrow down which Pods are affected. Further ingress and egress settings declare traffic to and from IP addresses and ports.

- Not all network providers support the NetworkPolicies kind. A non-exhaustive list of providers with support includes Calico, Romana, Cilium, Kube-router, and WeaveNet.
- In previous versions of Kubernetes, there was a requirement to annotate a namespace as part of network isolation, specifically the `net.beta.kubernetes.io/network-policy=` value. Some network plugins may still require this setting.
- On the next page, you can find an example of a NetworkPolicy recipe. More network policy recipes can be found on GitHub at <https://github.com/ahmetb/kubernetes-network-policy-recipes>.

10.12 Network Security Policy Example

- The use of policies has become stable, noted with the `v1` `apiVersion`. The example below narrows down the policy to affect the default namespace.
- Only Pods with the label of `role: db` will be affected by this policy, and the policy has both Ingress and Egress settings.
- The ingress setting includes a `172.17` network, with a smaller range of `172.17.1.0` IPs being excluded from this traffic.

YAML

```
apiVersion: networking.k8s.io/v1
2 kind: NetworkPolicy
3 metadata:
4   name: ingress-egress-policy
5   namespace: default
6 spec:
7   podSelector:
8     matchLabels:
9       role: db
10  policyTypes:
11    - Ingress
12    - Egress
13  ingress:
14    - from:
15      - ipBlock:
16          cidr: 172.17.0.0/16
17          except:
18            - 172.17.1.0/24
19      - namespaceSelector:
20          matchLabels:
21            project: myproject
22      - podSelector:
23          matchLabels:
24            role: frontend
25    ports:
26      - protocol: TCP
27        port: 6379
28  egress:
29    - to:
30      - ipBlock:
31          cidr: 10.0.0.0/24
32      ports:
33        - protocol: TCP
34          port: 5978
```

- These rules change the namespace for the following settings to be labeled `project: myproject`. The affected Pods also would need to match the label `role: frontend`. Finally, TCP traffic on port 6379 would be allowed from these Pods.
- The egress rules have the `to` settings, in this case the `10.0.0.0/24` range TCP traffic to port 5978.

- The use of empty ingress or egress rules denies all type of traffic for the included Pods, though this is not suggested. Use another dedicated NetworkPolicy instead.
- Note that there can also be complex matchExpressions statements in the spec, but this may change as NetworkPolicy matures.

```
podSelector:
  matchExpressions:
    - {key: inns, operator: In, values: ["yes"]}
```

10.13 Default Policy Example

- The empty braces will match all Pods not selected by other NetworkPolicy and will not allow ingress traffic. Egress traffic would be unaffected by this policy.

YAML

```
1 apiVersion: networking.k8s.io/v1
2 kind: NetworkPolicy
3 metadata:
4   name: default-deny
5 spec:
6   podSelector: {}
7   policyTypes:
8     - Ingress
```

- With the potential for complex ingress and egress rules, it may be helpful to create multiple objects which include simple isolation rules and use easy to understand names and labels.
- Some network plugins, such as WeaveNet, may require annotation of the Namespace. The following shows the setting of a DefaultDeny for the myns namespace:

YAML

```
1 kind: Namespace
2 apiVersion: v1
3 metadata:
4   name: myns
5   annotations:
6     net.beta.kubernetes.io/network-policy: |
7       {
8         "ingress": {
9           "isolation": "DefaultDeny"
10        }
11      }
```

10.14 Labs

Exercise 10.1: Set SecurityContext for a Pod and Container

Working with Security: Overview

In this lab we will implement security features for new applications, as the simpleapp YAML file is getting long and more difficult to read. Kubernetes architecture favors smaller, decoupled, and transient applications working together. We'll continue to emulate that in our exercises.

In this exercise we will create two new applications. One will be limited in its access to the host node, but have access to encoded data. The second will use a network security policy to move from the default all-access Kubernetes policies to a mostly closed network. First we will set security contexts for pods and containers, then create and consume secrets, then finish with configuring a network security policy.

1. Begin by making a new directory for our second application. Change into that directory.

```
student@master:~$ mkdir $HOME/app2
```

```
student@master:~$ cd $HOME/app2/
```

2. Create a YAML file for the second application. In the example below we are using a simple image, busybox, which allows access to a shell, but not much more. We will add a `runAsUser` to both the pod as well as the container.

```
student@master:~/app2$ vim second.yaml
```

YAML

second.yaml

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: secondapp
5 spec:
6   securityContext:
7     runAsUser: 1000
8   containers:
9     - name: busy
10      image: busybox
11      command:
12        - sleep
13        - "3600"
14      securityContext:
15        runAsUser: 2000
16        allowPrivilegeEscalation: false
```

3. Create the secondapp pod and verify it's running. Unlike the previous deployment this application is running as a pod. Look at the YAML output, to compare and contrast with what a deployment looks like. The status section probably has the largest contrast.

```
student@master:~/app2$ kubectl create -f second.yaml
```

```
1 pod/secondapp created
```

```
student@master:~/app2$ kubectl get pod secondapp
```

```
1 NAME          READY    STATUS    RESTARTS   AGE
2 secondapp    1/1     Running   0           21s
```

```
student@master:~/app2$ kubectl get pod secondapp -o yaml
```

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   annotations:
5     cni.projectcalico.org/podIP: 192.168.158.97/32
6     creationTimestamp: "2019-11-03T21:23:12Z"
7     name: secondapp
8 <output_omitted>
```

4. Execute a Bourne shell within the Pod. Check the user ID of the shell and other processes. It should show the container setting, not the pod. This allows for multiple containers within a pod to customize their UID if desired. As there is only one container in the pod we do not need to use the **-c busy** option.

```
student@master:~/app2$ kubectl exec -it secondapp -- sh
```

On Container

```
/ $ ps aux
```

```
1  PID    USER     TIME   COMMAND
2      1      2000      0:00   sleep 3600
3      8      2000      0:00    sh
4     12      2000      0:00   ps aux
```

5. While here check the capabilities of the kernel. In upcoming steps we will modify these values.

On Container

```
/ $ grep Cap /proc/1/status
```

```
1  CapInh:      00000000a80425fb
2  CapPrm:      0000000000000000
3  CapEff:      0000000000000000
4  CapBnd:      00000000a80425fb
5  CapAmb:      0000000000000000
```

```
/ $ exit
```

6. Use the capability shell wrapper tool, the **capsh** command, to decode the output. We will view and compare the output in a few steps. Note that there are 14 comma separated capabilities listed.

```
student@master:~/app2$ capsh --decode=00000000a80425fb
```

```
1  0x00000000a80425fb=cap_chown,cap_dac_override,cap_fowner,cap_fsetid,
2  cap_kill,cap_setgid,cap_setuid,cap_setpcap,cap_net_bind_service,
3  cap_net_raw,cap_sys_chroot,cap_mknod,cap_audit_write,cap_setfcap
```

7. Edit the YAML file to include new capabilities for the container. A capability allows granting of specific, elevated privileges without granting full root access. We will be setting **NET_ADMIN** to allow interface, routing, and other network configuration. We'll also set **SYS_TIME**, which allows system clock configuration. More on kernel capabilities can be read here: <https://github.com/torvalds/linux/blob/master/include/uapi/linux/capability.h>

It can take up to a minute for the pod to fully terminate, allowing the future pod to be created.

```
student@master:~/app2$ kubectl delete pod secondapp
```

```
1  pod "secondapp" deleted
```

```
student@master:~/app2$ vim second.yaml
```

YAML

second.yaml

```
1  <output_omitted>
2      - sleep
3      - "3600"
```



```

4     securityContext:
5       runAsUser: 2000
6       allowPrivilegeEscalation: false
7       capabilities:                                #<-- Add this and following line
8         add: ["NET_ADMIN", "SYS_TIME"]

```

8. Create the pod again. Execute a shell within the container and review the Cap settings under `/proc/1/status`. They should be different from the previous instance.

```
student@master:~/app2$ kubectl create -f second.yaml
```

```
1 pod/secondapp created
```

```
student@master:~/app2$ kubectl exec -it secondapp -- sh
```



On Container

```
/ $ grep Cap /proc/1/status
```

```

1 CapInh:      00000000aa0435fb
2 CapPrm:      0000000000000000
3 CapEff:      0000000000000000
4 CapBnd:      00000000aa0435fb
5 CapAmb:      0000000000000000

```

```
/ $ exit
```

9. Decode the output again. Note that the instance now has 16 comma delimited capabilities listed. **cap_net_admin** is listed as well as **cap_sys_time**.

```
student@master:~/app2$ capsh --decode=00000000aa0435fb
```

```

1 0x00000000aa0435fb=cap_chown,cap_dac_override,cap_fowner,
2 cap_fsetid,cap_kill,cap_setgid,cap_setuid,cap_setpcap,
3 cap_net_bind_service,cap_net_admin,cap_net_raw,cap_sys_chroot,
4 cap_sys_time,cap_mknod,cap_audit_write,cap_setfcap

```

✍ Exercise 10.2: Create and consume Secrets

Secrets are consumed in a manner similar to ConfigMaps, covered in an earlier lab. While at-rest encryption is just now enabled, historically a secret was just base64 encoded. There are three types of encryption which can be configured.

1. Begin by generating an encoded password.

```
student@master:~/app2$ echo LFTr@1n | base64
```

```
1 TEZUckAxbgo=
```

2. Create a YAML file for the object with an API object kind set to Secret. Use the encoded key as a password parameter.

```
student@master:~/app2$ vim secret.yaml
```

YAML**secret.yaml**

```

1  apiVersion: v1
2  kind: Secret
3  metadata:
4    name: lfsecret
5  data:
6    password: TEZUckAxbgo=

```

3. Ingest the new object into the cluster.

```
student@master:~/app2$ kubectl create -f secret.yaml
```

```
1  secret/lfsecret created
```

4. Edit secondapp YAML file to use the secret as a volume mounted under `/mysqlpassword`. `volumeMounts`: lines up with the container `name`: and `volumes`: lines up with `containers`: Note the pod will restart when the sleep command finishes every 3600 seconds, or every hour.

```
student@master:~/app2$ vim second.yaml
```

YAML**second.yaml**

```

1  ....
2      runAsUser: 2000
3      allowPrivilegeEscalation: false
4      capabilities:
5        add: ["NET_ADMIN", "SYS_TIME"]
6      volumeMounts:                                #<-- Add this and six following lines
7        - name: mysql
8          mountPath: /mysqlpassword
9      volumes:
10     - name: mysql
11       secret:
12         secretName: lfsecret

```

```
student@master:~/app2$ kubectl delete pod secondapp
```

```
1  pod "secondapp" deleted
```

```
student@master:~/app2$ kubectl create -f second.yaml
```

```
1  pod/secondapp created
```

5. Verify the pod is running, then check if the password is mounted where expected. We will find that the password is available in its clear-text, decoded state.

```
student@master:~/app2$ kubectl get pod secondapp
```

```

1  NAME          READY   STATUS    RESTARTS   AGE
2  secondapp     1/1     Running   0           34s

```

```
student@master:~/app2$ kubectl exec -ti secondapp -- /bin/sh
```



On Container

```
/ $ cat /mysqlpassword/password
```

```
1 LFTr@1n
```

- View the location of the directory. Note it is a symbolic link to `..data` which is also a symbolic link to another directory. After taking a look at the filesystem within the container, exit back to the node.



On Container

```
/ $ cd /mysqlpassword/
```

```
/mysqlpassword $ ls
```

```
1 password
```

```
/mysqlpassword $ ls -al
```

```
1 total 4
2 drwxrwxrwt  3 root    root    100 Apr 11 07:24 .
3 drwxr-xr-x 21 root    root    4096 Apr 11 22:30 ..
4 drwxr-xr-x  2 root    root     60 Apr 11 07:24 ..4984_11_04_07_24_47.831222818
5 lrwxrwxrwx  1 root    root     31 Apr 11 07:24 ..data -> ..4984_11_04_07_24_47.831222818
6 lrwxrwxrwx  1 root    root     15 Apr 11 07:24 password -> ..data/password
```

```
/mysqlpassword $ exit
```

Exercise 10.3: Working with ServiceAccounts

We can use ServiceAccounts to assign cluster roles, or the ability to use particular HTTP verbs. In this section we will create a new ServiceAccount and grant it access to view secrets.

- Begin by viewing secrets, both in the default namespace as well as all.

```
student@master:~/app2$ cd
```

```
student@master:~$ kubectl get secrets
```

NAME	TYPE	DATA	AGE
default-token-c4rdg	kubernetes.io/service-account-token	3	4d16h
lfsecret	Opaque	1	6m5s

```
student@master:~$ kubectl get secrets --all-namespaces
```

NAMESPACE	NAME	TYPE	DATA	AGE
default	default-token-c4rdg	kubernetes.io/service-account-token	3	4d16h
kube-public	default-token-zqzbg	kubernetes.io/service-account-token	3	4d16h
kube-system	attachdetach-controller-token-wxzvc	kubernetes.io/service-account-token	3	4d16h

<output_omitted>

- We can see that each agent uses a secret in order to interact with the API server. We will create a new ServiceAccount which will have access.

```
student@master:~$ vim serviceaccount.yaml
```

YAML
serviceaccount.yaml

```
1 apiVersion: v1
2 kind: ServiceAccount
3 metadata:
4   name: secret-access-sa
```

```
student@master:~$ kubectl create -f serviceaccount.yaml
```

```
1 serviceaccount/secret-access-sa created
```

```
student@master:~$ kubectl get serviceaccounts
```

```
1 NAME           SECRETS  AGE
2 default         1        1d17h
3 secret-access-sa 1        34s
```

- Now we will create a ClusterRole which will list the actual actions allowed cluster-wide. We will look at an existing role to see the syntax.

```
student@master:~$ kubectl get clusterroles
```

```
1 NAME                               AGE
2 admin                             1d17h
3 calico-cni-plugin                 1d17h
4 calico-kube-controllers            1d17h
5 cluster-admin                     1d17h
6 <output_omitted>
```

- View the details for the admin and compare it to the cluster-admin. The admin has particular actions allowed, but cluster-admin has the meta-character '*' allowing all actions.

```
student@master:~$ kubectl get clusterroles admin -o yaml
```

```
1 <output_omitted>
```

```
student@master:~$ kubectl get clusterroles cluster-admin -o yaml
```

```
1 <output_omitted>
```

- Using some of the output above, we will create our own file.

```
student@master:~$ vim clusterrole.yaml
```

YAML
clusterrole.yaml

```
1 apiVersion: rbac.authorization.k8s.io/v1
2 kind: ClusterRole
3 metadata:
4   name: secret-access-cr
5 rules:
6 - apiGroups:
```


YAML

```

7   - ""
8   resources:
9   - secrets
10  verbs:
11  - get
12  - list

```

6. Create and verify the new ClusterRole.

```
student@master:~$ kubectl create -f clusterrole.yaml
```

```
1 clusterrole.rbac.authorization.k8s.io/secret-access-cr created
```

```
student@master:~$ kubectl get clusterrole secret-access-cr -o yaml
```

```

1 apiVersion: rbac.authorization.k8s.io/v1
2 kind: ClusterRole
3 metadata:
4   creationTimestamp: 2018-10-18T19:27:24Z
5   name: secret-access-cr
6 <output_omitted>

```

7. Now we bind the role to the account. Create another YAML file which uses roleRef::

```
student@master:~$ vim rolebinding.yaml
```

YAML**rolebinding.yaml**

```

1 apiVersion: rbac.authorization.k8s.io/v1
2 kind: RoleBinding
3 metadata:
4   name: secret-rb
5 subjects:
6 - kind: ServiceAccount
7   name: secret-access-sa
8 roleRef:
9   kind: ClusterRole
10  name: secret-access-cr
11 apiGroup: rbac.authorization.k8s.io

```

8. Create the new RoleBinding and verify.

```
student@master:~$ kubectl create -f rolebinding.yaml
```

```
1 rolebinding.rbac.authorization.k8s.io/secret-rb created
```

```
student@master:~$ kubectl get rolebindings
```

```

1 NAME      AGE
2 secret-rb 17s

```

9. View the secondapp pod and **grep** for secret settings. Note that it uses the default settings.

```
student@master:~$ kubectl describe pod secondapp |grep -i secret
```

```

1  /var/run/secrets/kubernetes.io/serviceaccount from
2  default-token-c4rdg (ro)
3  Type:          Secret (a volume populated by a Secret)
4  SecretName:    lfsecret
5  Type:          Secret (a volume populated by a Secret)
6  SecretName:    default-token-c4rdg

```

10. Edit the `second.yaml` file and add the use of the `serviceAccount`.

```
student@master:~$ vim $HOME/app2/second.yaml
```

YA
ML

second.yaml

```

1  ....
2  name: secondapp
3  spec:
4  serviceAccountName: secret-access-sa #<-- Add this line
5  securityContext:
6  runAsUser: 1000
7  ....

```

11. We will delete the `secondapp` pod if still running, then create it again. View what the secret is by default.

```
student@master:~$ kubectl delete pod secondapp ; kubectl create -f $HOME/app2/second.yaml
```

```

1  pod "secondapp" deleted
2  pod/secondapp created

```

```
student@master:~$ kubectl describe pod secondapp | grep -i secret
```

```

1  /var/run/secrets/kubernetes.io/serviceaccount from
2  secret-access-sa-token-wd7vm (ro)
3  secret-access-sa-token-wd7vm:
4  Type:          Secret (a volume populated by a Secret)
5  SecretName:    secret-access-sa-token-wd7vm
6

```

✍ Exercise 10.4: Implement a NetworkPolicy

An early architecture decision with Kubernetes was non-isolation, that all pods were able to connect to all other pods and nodes by design. In more recent releases the use of a `NetworkPolicy` allows for pod isolation. The policy only has effect when the network plugin, like **Project Calico**, are capable of honoring them. If used with a plugin like **flannel** they will have no effect. The use of `matchLabels` allows for more granular selection within the namespace which can be selected using a `namespaceSelector`. Using multiple labels can allow for complex application of rules. More information can be found here: <https://kubernetes.io/docs/concepts/services-networking/network-policies>

1. Begin by creating a default policy which denies all traffic. Once ingested into the cluster this will affect every pod not selected by another policy, creating a mostly-closed environment. If you want to only deny ingress or egress traffic you can remove the other policyType.

```
student@master:~$ cd $HOME/app2/
```

```
student@master:~/app2$ vim allclosed.yaml
```

YAML**allclosed.yaml**

```

1 apiVersion: networking.k8s.io/v1
2 kind: NetworkPolicy
3 metadata:
4   name: deny-default
5 spec:
6   podSelector: {}
7   policyTypes:
8   - Ingress
9   - Egress

```

2. Before we can test the new network policy we need to make sure network access works without it applied. Update **secondapp** to include a new container running **nginx**, then test access. Begin by adding two lines for the **nginx** image and name **webserver**, as found below. It takes a bit for the pod to terminate, so we'll delete then edit the file.

```
student@master:~/app2$ kubectl delete pod secondapp
```

```
1 pod "secondapp" deleted
```

```
student@master:~/app2$ vim second.yaml
```

YAML**second.yaml**

```

1   runAsUser: 1000
2 containers:
3   - name: webserver                                #<-- Add this and following line
4     image: nginx
5   - name: busy
6     image: busybox
7     command:

```

3. Create the new pod. Be aware the pod will move from ContainerCreating to Error to CrashLoopBackOff, as only one of the containers will start. We will troubleshoot the error in following steps.

```
student@master:~/app2$ kubectl create -f second.yaml
```

```
1 pod/secondapp created
```

```
student@master:~/app2$ kubectl get pods
```

```

1 NAME                                READY   STATUS             RESTARTS   AGE
2 nginx-6b58d9cdfd-9fnl4              1/1     Running            1          2d
3 Registry-795c6c8b8f-hl5wf          1/1     Running            2          2d
4 secondapp                           1/2     CrashLoopBackOff   1          13s
5 <output_omitted>

```

4. Take a closer look at the events leading up to the failure. The images were pulled and the container was started. It was the full execution of the container which failed.

```
student@master:~/app2$ kubectl get event
```

```

1 <output_omitted>
2 25s      Normal    Scheduled    Pod    Successfully assigned default/secondapp to master
3 4s       Normal    Pulling     Pod    pulling image "nginx"
4 2s       Normal    Pulled      Pod    Successfully pulled image "nginx"
5 2s       Normal    Created     Pod    Created container
6 2s       Normal    Started     Pod    Started container

```

```

7 23s      Normal    Pulling    Pod    pulling image "busybox"
8 21s      Normal    Pulled     Pod    Successfully pulled image "busybox"
9 21s      Normal    Created    Pod    Created container
10 21s      Normal    Started    Pod    Started container
11 1s       Warning   BackOff    Pod    Back-off restarting failed container

```

5. View the logs of the **webserver** container mentioned in the previous output. Note there are errors about the user directive and not having permission to make directories.

```
student@master:~/app2$ kubectl logs secondapp webserver
```

```

1 2018/04/13 19:51:13 [warn] 1#1: the "user" directive makes sense
2 only if the master process runs with super-user privileges,
3 ignored in /etc/nginx/nginx.conf:2
4 nginx: [warn] the "user" directive makes sense only if the master
5 process runs with super-user privileges,
6 ignored in /etc/nginx/nginx.conf:2
7 2018/04/13 19:51:13 [emerg] 1#1: mkdir() "/var/cache/nginx/client_temp"
8 failed (13: Permission denied)
9 nginx: [emerg] mkdir() "/var/cache/nginx/client_temp" failed
10 (13: Permission denied)

```

6. Delete the pods. Edit the YAML file to comment out the setting of a UID for the entire pod.

```
student@master:~/app2$ kubectl delete -f second.yaml
```

```
1 pod "secondapp" deleted
```

```
student@master:~/app2$ vim second.yaml
```

YAML

second.yaml

```

1 spec:
2   serviceAccountName: secret-access-sa
3   # securityContext:                                #<-- Comment this and following line
4   #   runAsUser: 1000
5   containers:
6   - name: webserver

```

7. Create the pod again. This time both containers should run. You may have to wait for the previous pod to fully terminate, depending on how fast you type.

```
student@master:~/app2$ kubectl create -f second.yaml
```

```
1 pod/secondapp created
```

```
student@master:~/app2$ kubectl get pods
```

```

1 NAME                                READY    STATUS    RESTARTS   AGE
2 <output_omitted>
3 secondapp                           2/2      Running   0           5s

```

8. Expose the **webserver** using a NodePort service. Expect an error due to lack of labels.

```
student@master:~/app2$ kubectl expose pod secondapp --type=NodePort --port=80
```

```

1 error: couldn't retrieve selectors via --selector flag or
2 introspection: the pod has no labels and cannot be exposed
3 See 'kubectl expose -h' for help and examples.

```

9. Edit the YAML file to add a label in the metadata, adding the example: second label right after the pod name. Note you can delete several resources at once by passing the YAML file to the delete command. Delete and recreate the pod. It may take up to a minute for the pod to shut down.

```
student@master:~/app2$ kubectl delete -f second.yaml
```

```
1 pod "secondapp" deleted
```

```
student@master:~/app2$ vim second.yaml
```

YAML

second.yaml

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: secondapp
5   labels:                                #<-- This and following line
6     example: second
7 spec:
8   # securityContext:
9   #   runAsUser: 1000
10  <output_omitted>
```

```
student@master:~/app2$ kubectl create -f second.yaml
```

```
1 pod/secondapp created
```

```
student@master:~/app2$ kubectl get pods
```

```
1 NAME                                READY    STATUS    RESTARTS   AGE
2 <output_omitted>
3 secondapp                           2/2      Running   0           15s
```

10. This time we will expose a NodePort again, and create the service separately, then add a label to illustrate how labels are essential for tying resources together inside of Kubernetes.

```
student@master:~/app2$ kubectl create service nodeport secondapp --tcp=80
```

```
1 service/secondapp created
```

11. Look at the details of the service. Note the selector is set to app: secondapp. Also take note of the nodePort, which is 31655 in the example below, yours may be different.

```
student@master:~/app2$ kubectl get svc secondapp -o yaml
```

```
1 apiVersion: v1
2 kind: Service
3 metadata:
4   creationTimestamp: "2020-04-16T04:40:07Z"
5   labels:
6     example: second
7   managedFields:
8
9   ....
10
11 spec:
12   clusterIP: 10.97.96.75
13   externalTrafficPolicy: Cluster
14   ports:
15   - nodePort: 31655
```

```

16     port: 80
17     protocol: TCP
18     targetPort: 80
19     selector:
20       example: second
21     sessionAffinity: None
22     type: NodePort
23 status:
24   loadBalancer: {}

```

12. Test access to the service using **curl** and the ClusterIP shown in the previous output. As the label does not match any other resources, the **curl** command should fail. If it hangs **control-c** to exit back to the shell.

```
student@master:~/app2$ curl http://10.97.96.75
```

13. Edit the service. We will change the label to match **secondapp**, and set the nodePort to a new port, one that may have been specifically opened by our firewall team, port 32000.

```
student@master:~/app2$ kubectl edit svc secondapp
```

YAML

secondapp service

```

1  <output_omitted>
2  ports:
3  - name: "80"
4    nodePort: 32000      #<-- Edit this line
5    port: 80
6    protocol: TCP
7    targetPort: 80
8  selector:
9    example: second      #<-- Edit this line
10 sessionAffinity: None
11 <output_omitted>

```

14. Verify the updated port number is showing properly, and take note of the ClusterIP. The example below shows a ClusterIP of 10.97.96.75 and a port of 32000 as expected.

```
student@master:~/app2$ kubectl get svc
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
<output_omitted>					
secondapp	NodePort	10.97.96.75	<none>	80:32000/TCP	5m

15. Test access to the high port. You should get the default nginx welcome page both if you test from the node to the ClusterIP:<low-port-number> and from the exterior hostIP:<high-port-number>. As the high port is randomly generated make sure it's available. Both of your nodes should be exposing the web server on port 32000. The example shows the use of the **curl** command, you could also use a web browser.

```
student@master:~/app2$ curl http://10.97.96.75
```

```

1  <!DOCTYPE html>
2  <html>
3  <head>
4  <title>Welcome to nginx!</title>
5  <output_omitted>

```

```
[user@laptop ~]$ curl http://35.184.219.5:32000
```

```

1 <!DOCTYPE html>
2 <html>
3 <head>
4 <title>Welcome to nginx!</title>
5 <output_omitted>

```

16. Now test egress from a container to the outside world. We'll use the **netcat** command to verify access to a running web server on port 80. First test local access to nginx, then a remote server.

```
student@master:~/app2$ kubectl exec -it -c busy secondapp -- sh
```



On Container

```

/ $ nc -vz 127.0.0.1 80
1 127.0.0.1 (127.0.0.1:80) open

/ $ nc -vz www.linux.com 80
1 www.linux.com (151.101.185.5:80) open

/ $ exit

```

Exercise 10.5: Testing the Policy

1. Now that we have tested both ingress and egress we can implement the network policy.

```
student@master:~/app2$ kubectl create -f $HOME/app2/allclosed.yaml
```

```
1 networkpolicy.networking.k8s.io/deny-default created
```

2. Use the ingress and egress tests again. Three of the four should eventually timeout. Start by testing from outside the cluster, and interrupt if you get tired of waiting.

```
[user@laptop ~]$ curl http://35.184.219.5:32000
```

```

1 curl: (7) Failed to connect to 35.184.219.5 port
2 32000: Connection timed out

```

3. Then test from the host to the container.

```
student@master:~/app2$ curl http://10.97.96.75:80
```

```
1 curl: (7) Failed to connect to 10.97.96.75 port 80: Connection timed out
```

4. Now test egress. From container to container should work, as the filter is outside of the pod. Then test egress to an external web page. It should eventually timeout.

```
student@master:~/app2$ kubectl exec -it -c busy secondapp -- sh
```



On Container

```

/ $ nc -vz 127.0.0.1 80
1 127.0.0.1 (127.0.0.1:80) open

```



```

/ $ nc -vz www.linux.com 80
1 nc: bad address 'www.linux.com'

/ $ exit

```

5. Update the NetworkPolicy and comment out the Egress line. Then replace the policy.

```
student@master:~/app2$ vim $HOME/app2/allclosed.yaml
```



allclosed.yaml

```

1 ....
2 spec:
3   podSelector: {}
4   policyTypes:
5     - Ingress
6   # - Egress          #<-- Comment out this line

```

```
student@master:~/app2$ kubectl replace -f $HOME/app2/allclosed.yaml
```

```
1 networkpolicy.networking.k8s.io/deny-default replaced
```

6. Test egress access to an outside site. Get the IP address of the **eth0** inside the container while logged in. The IP is 192.168.55.91 in the example below, yours may be different.

```
student@master:~/app2$ kubectl exec -it -c busy secondapp -- sh
```



On Container

```

/ $ nc -vz www.linux.com 80
1 www.linux.com (151.101.185.5:80) open

/ $ ip a
1 1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue qlen 1000
2   link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
3   inet 127.0.0.1/8 scope host lo
4     valid_lft forever preferred_lft forever
5   inet6 ::1/128 scope host
6     valid_lft forever preferred_lft forever
7 2: tunl0@NONE: <NOARP> mtu 1480 qdisc noop qlen 1000
8   link/ipip 0.0.0.0 brd 0.0.0.0
9 4: eth0@if59: <BROADCAST,MULTICAST,UP,LOWER_UP,M-DOWN> mtu 1500 qdisc noqueue
10  link/ether 1e:c8:7d:6a:96:c3 brd ff:ff:ff:ff:ff:ff
11  inet 192.168.55.91/32 scope global eth0
12    valid_lft forever preferred_lft forever
13  inet6 fe80::1cc8:7dff:fe6a:96c3/64 scope link
14    valid_lft forever preferred_lft forever

/ $ exit

```


7. Now add a selector to allow ingress to only the nginx container. Use the IP from the **eth0** range.

```
student@master:~/app2$ vim $HOME/app2/allclosed.yaml
```

YAML

allclosed.yaml

```
1 <output_omitted>
2 policyTypes:
3   - Ingress
4   ingress:                                #<-- Add this and following three lines
5     - from:
6       - ipBlock:
7         cidr: 192.168.0.0/16
8   # - Egress
```

8. Recreate the policy, and verify its configuration.

```
student@master:~/app2$ kubectl replace -f $HOME/app2/allclosed.yaml
```

```
1 networkpolicy.networking.k8s.io/deny-default replaced
```

```
student@master:~/app2$ kubectl get networkpolicy
```

```
1 NAME          POD-SELECTOR  AGE
2 deny-default  <none>       3m2s
```

```
student@master:~/app2$ kubectl get networkpolicy -o yaml
```

```
1 apiVersion: v1
2 items:
3   - apiVersion: networking.k8s.io/v1
4     kind: NetworkPolicy
5     metadata:
6     <output_omitted>
```

9. Test access to the container both using **curl** as well as **ping**, the IP address to use was found from the IP inside the container. You may need to install **iputils-ping** or other software to use **ping**.

```
student@master:~/app2$ curl http://192.168.55.91
```

```
1 <!DOCTYPE html>
2 <html>
3 <head>
4 <title>Welcome to nginx!</title>
5 <output_omitted>
```

```
student@master:~/app2$ ping -c5 192.168.55.91
```

```
1 PING 192.168.55.91 (192.168.55.91) 56(84) bytes of data.
2 64 bytes from 192.168.55.91: icmp_seq=1 ttl=63 time=1.11 ms
3 64 bytes from 192.168.55.91: icmp_seq=2 ttl=63 time=0.352 ms
4 64 bytes from 192.168.55.91: icmp_seq=3 ttl=63 time=0.350 ms
5 64 bytes from 192.168.55.91: icmp_seq=4 ttl=63 time=0.359 ms
6 64 bytes from 192.168.55.91: icmp_seq=5 ttl=63 time=0.295 ms
7
8 --- 192.168.55.91 ping statistics ---
9 5 packets transmitted, 5 received, 0% packet loss, time 4054ms
10 rtt min/avg/max/mdev = 0.295/0.495/1.119/0.312 ms
```

10. Update the policy to only allow ingress for TCP traffic on port 80, then test with **curl**, which should work. The ports entry should line up with the from entry a few lines above.

```
student@master:~/app2$ vim $HOME/app2/allclosed.yaml
```

YAML

allclosed.yaml

```
1 <output_omitted>
2   - Ingress
3   ingress:
4     - from:
5       - ipBlock:
6         cidr: 192.168.0.0/16
7       ports:                                #<-- Add this and two following lines
8         - port: 80
9         protocol: TCP
10  # - Egress
```

```
student@master:~/app2$ kubectl replace -f $HOME/app2/allclosed.yaml
```

```
1 networkpolicy.networking.k8s.io/deny-default replaced
```

```
student@master:~/app2$ curl http://192.168.55.91
```

```
1 <!DOCTYPE html>
2 <html>
3 <head>
4 <title>Welcome to nginx!</title>
5 <output_omitted>
```

11. All five pings should fail, with zero received.

```
student@master:~/app2$ ping -c5 192.168.55.91
```

```
1 PING 192.168.55.91 (192.168.55.91) 56(84) bytes of data.
2
3 --- 192.168.55.91 ping statistics ---
4 5 packets transmitted, 0 received, 100% packet loss, time 4098ms
```

12. You may want to remove the default-deny policy, in case you want to get to your registry or other pods.

```
student@master:~/app2$ kubectl delete networkpolicies deny-default
```

```
1 networkpolicy.networking.k8s.io "deny-default" deleted
```

✍ Exercise 10.6: Domain Review



Very Important

The source pages and content in this review could change at any time. **IT IS YOUR RESPONSIBILITY TO CHECK THE CURRENT INFORMATION.**

Revisit the CKAD domain list on [Curriculum Overview](#) and locate some of the topics we have covered in this chapter.

Figure 10.2: **Domain**

Focus on ensuring you have all the necessary files and processes understood first. Repeat the review until you are sure you have bookmarks of necessary YAML samples and can complete each step quickly, and ensure each object is running properly.

1. Create a new deployment which uses the `nginx` image.
2. Create a new `LoadBalancer` service to expose the newly created deployment. Test that it works.
3. Create a new `NetworkPolicy` called `netblock` which blocks all traffic to pods in this deployment only. Test that all traffic is blocked to deployment.
4. Update the `netblock` policy to allow traffic to the pod on port 80 only. Test that you can access the default `nginx` web page.

5. Find and use the `security-review1.yaml` file to create a pod.

```
student@master:~$ kubectl create -f security-review1.yaml
```

6. View the status of the pod.
7. Use the following commands to figure out why the pod has issues.

```
student@master:~$ kubectl get pod securityreview
```

```
student@master:~$ kubectl describe pod securityreview
```

```
student@master:~$ kubectl logs securityreview
```

8. After finding the errors, log into the container and find the proper id of the `nginx` user.
9. Edit the `yaml` and re-create the pod such that the pod runs without error.
10. Create a new `serviceAccount` called `securityaccount`.
11. Create a `ClusterRole` named `secrole` which only allows `create`, `delete`, and `list` of pods in all `apiGroups`.
12. Bind the `clusterRole` to the `serviceAccount`.
13. Locate the token of the `securityaccount`. Create a file called `/tmp/securitytoken`. Put only the value of `token:` is equal to, a long string that may start with `eyJh` and be several lines long. Careful that only that string exists in the file.
14. Remove any resources you have added during this review

Chapter 11

Exposing Applications



In this chapter, we are going to talk about exposing our applications using services. We'll start by talking about a ClusterIP service. This allows us to expose our application inside of the cluster only. A NodePort will create a ClusterIP and then also, get a high number port from the node. This allows us then to expose our applications internally on a low number port, and externally, but only on a high number port. The use of a LoadBalancer first creates a NodePort on our behalf. The difference is it also sends an asynchronous request to a cloud provider to make use of a LoadBalancer. If we are not using a cloud provider, or if a load balancer is not available, it continues to run as a NodePort. The fourth service that we are going to discuss is called ExternalName. This is useful for handling requests to an outside DNS resource, instead of one inside of the cluster. This is helpful if you were in the process of moving your application from outside of the cluster inside. We'll also discuss the configuration of an Ingress Controller. The two reasons that you might want to use an Ingress Controller is, for example, if you have thousands of services, having them independent can become difficult to manage, and wasteful of resources. You can then consolidate it to a single Ingress Controller. Or multiple Ingress Controllers if you want more flexibility. The second reason is an Ingress Controller allows you to expose a low numbered port to your application. Without that, you could get there, but you'd have to use iptables, which could become difficult and not terribly flexible to manage.

11.1	Service Types	213
11.2	Services Diagram	214
11.3	Service Update Pattern	215
11.4	Accessing an Application with a Service	215
11.5	Service without a Selector	215
11.6	ClusterIP	216
11.7	NodePort	216
11.8	LoadBalancer	216
11.9	ExternalName	216
11.10	Ingress Resource	217
11.11	Ingress Controller	217
11.12	Service Mesh	217
11.13	Labs	218

11.1 Service Types

- Services can be of the following types:
 - ClusterIP

- NodePort
 - LoadBalancer
 - ExternalName
- The ClusterIP service type is the default, and only provides access internally (except if manually creating an external endpoint). The range of ClusterIP used is defined via an API server startup option.
 - The NodePort type is great for debugging, or when a static IP address is necessary, such as opening a particular address through a firewall. The NodePort range is defined in the cluster configuration.
 - The LoadBalancer service was created to pass requests to a cloud provider like GKE or AWS. Private cloud solutions also may implement this service type if there is a cloud provider plugin, such as with CloudStack and OpenStack. Even without a cloud provider, the address is made available to public traffic, and packets are spread among the Pods in the deployment automatically.
 - A newer service is ExternalName, which is a bit different. It has no selectors, nor does it define ports or endpoints. It allows the return of an alias to an external service. The redirection happens at the DNS level, not via a proxy or forward. This object can be useful for services not yet brought into the Kubernetes cluster. A simple change of the type in the future would redirect traffic to the internal objects. As **CoreDNS** has become more stable this service is not used as much.
 - The `kubectl proxy` command creates a local service to access a ClusterIP. This can be useful for troubleshooting or development work.

11.2 Services Diagram

- The **kube-proxy** running on cluster nodes watches the API server service resources. It presents a type of virtual IP address for services other than ExternalName. The mode for this process has changed over versions of Kubernetes.
- In v1.0, services ran in userspace mode as TCP/UDP over IP or Layer 4. In the v1.1 release, the iptables proxy was added and became the default mode starting with v1.2.

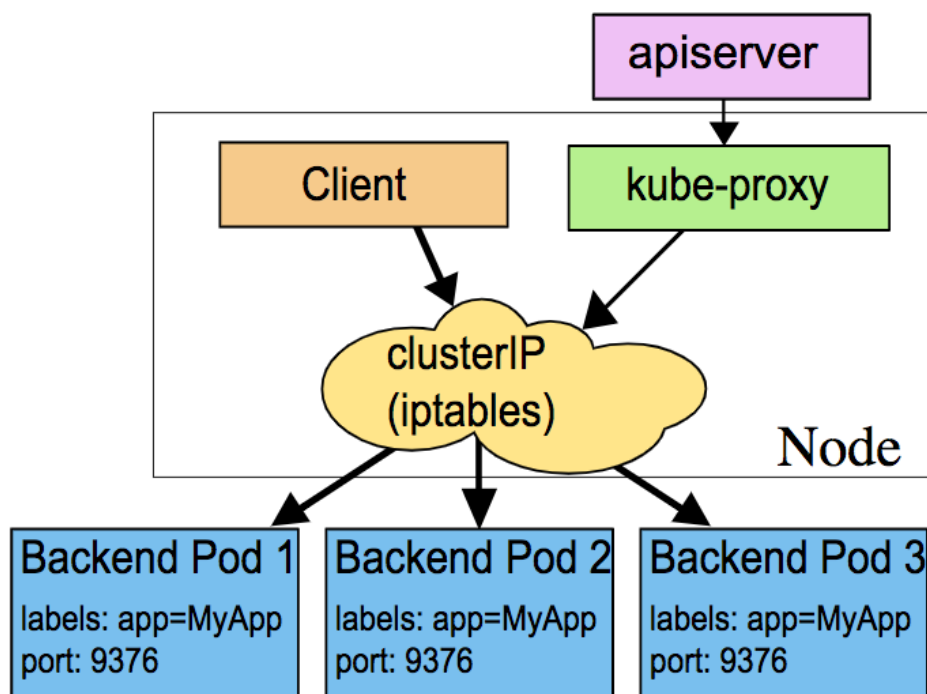


Figure 11.1: Traffic from Cluster IP to Pod

- In the `iptables` proxy mode, **kube-proxy** continues to monitor the API server for changes in Service and Endpoint objects, and updates rules for each object when created or removed. One limitation to the new mode is an inability to connect to a Pod should the original request fail, so it uses a Readiness Probe to ensure all containers are functional prior to connection. This mode allows for up to approximately 5000 nodes. Assuming multiple Services and Pods per node, this leads to a bottleneck in the kernel.
- Another mode introduced in v1.9 and GA in v1.11 is `ipvs`. This mode works in the kernel space for greater speed, and allows for a configurable load-balancing algorithm, such as round-robin, shortest expected delay, least connection and several others. This can be helpful for large clusters, much past the previous 5000 node limitation. This mode assumes IPVS kernel modules are installed and running prior to **kube-proxy**. Clusters built with **kubeadm** do not enable `ipvs` by default, but this can be passed during cluster initialization.
- The **kube-proxy** mode is configured via a flag sent during initialization, such as `mode=iptables` and could also be `IPVS` or `userspace`.

11.3 Service Update Pattern

- Labels are used to determine which Pods should receive traffic from a service. Labels can be dynamically updated for an object, which may affect which Pods continue to connect to a service.
- The default update pattern is for a rolling deployment, where new Pods are added, with different versions of an application, and due to automatic load balancing, receive traffic along with previous versions of the application.
- Should there be a difference in applications deployed, such that clients would have issues communicating with different versions, you may consider a more specific label for the deployment, which includes a version number. When the deployment creates a new replication controller for the update, the label would not match. Once the new Pods have been created, and perhaps allowed to fully initialize, we would edit the labels for which the Service connects. Traffic would shift to the new and ready version, minimizing client version confusion.

11.4 Accessing an Application with a Service

- The basic step to access a new service is to use **kubectl**:

```
$ kubectl expose deployment/nginx --port=80 --type=NodePort
$ kubectl get svc
```

	NAME	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
1	kubernetes	10.0.0.1	<none>	443/TCP	18h
2	nginx	10.0.0.112	<nodes>	80:31230/TCP	5s

```
$ kubectl get svc nginx -o yaml
```

```
1 apiVersion: v1
2 kind: Service
3 ...
4 spec:
5   clusterIP: 10.0.0.112
6   ports:
7     - nodePort: 31230
8   ...
9 Open browser http://<Public IP>:31230
```

- The `kubectl expose` command created a service for the **nginx** deployment. This service used port 80 and generated a random port on all the nodes. A particular port and `targetPort` can also be passed during object creation to avoid random values. The `targetPort` defaults to the `port`, but could be set to any value, including a string referring to a port on a backend Pod. Each Pod could have a different port, but traffic is still passed via the name. Switching traffic to a different port would maintain a client connection, while changing versions of software, for example.

- The `kubectl get svc` command gave you a list of all the existing services, and we saw the **nginx** service, which was created with an internal cluster IP.
- The range of cluster IPs and the range of ports used for the random `NodePort` are configurable in the API server startup options.
- Services can also be used to point to a service in a different namespace, or even a resource outside the cluster, such as a legacy application not yet in Kubernetes.

11.5 Service without a Selector

- Typically, a service creates a new endpoint for connectivity. Should you want to create the service, but later add the endpoint, such as connecting to a remote database, you can use a service without selectors. This can also be used to direct the service to another service, in a different namespace or cluster.

11.6 ClusterIP

- For inter-cluster communication, frontends talking to back-ends can use `ClusterIP`s. These addresses and endpoints only work within the cluster.
- `ClusterIP` is the default type of service created.

```
spec: clusterIP: 10.108.95.67 ports: - name: "443" port:
443 protocol: TCP targetPort: 443
```

11.7 NodePort

- `NodePort` is a simple connection from a high-port routed to a `ClusterIP` using `iptables`, or `ipvs` in newer versions. The creation of a `NodePort` generates a `ClusterIP` by default. Traffic is routed from the `NodePort` to the `ClusterIP`. Only high ports can be used, as declared in the source code. The `NodePort` is accessible via calls to `<NodeIP>:<NodePort>`.

```
spec:
  clusterIP: 10.97.191.46
  externalTrafficPolicy: Cluster
  ports:
  - nodePort: 31070
    port: 80
    protocol: TCP
    targetPort: 800a0
  selector:
    io.kompose.service: nginx
  sessionAffinity: None
  type: NodePort
```

11.8 LoadBalancer

- Creating a `LoadBalancer` service generates a `NodePort`, which then creates a `ClusterIP`. It also sends an asynchronous call to an external load balancer, typically supplied by a cloud provider. The `External-IP` value will remain in a `<Pending>` state until the load balancer returns. Should it not return, the `NodePort` created acts as it would otherwise.


```
Type: LoadBalancer
loadBalancerIP: 12.45.105.12
clusterIP: 10.5.31.33
ports:
- protocol: TCP
  Port: 80
```

- The routing of traffic to a particular backend pod depends on the cloud provider in use.

11.9 ExternalName

- The use of an `ExternalName` service, which is a special type of service without selectors, is to point to an external DNS server. Use of the service returns a `CNAME` record. Working with the `ExternalName` service is handy when using a resource external to the cluster, perhaps prior to full integration.

```
spec:
  Type: ExternalName
  externalName: ext.db.example.com
```

- DNS services are now handled by the `kube-dns` service which uses the **CoreDNS** image on kubeadm created clusters. The **CoreDNS** service has many features in addition to being a DNS server, and is extensible using plugins. More can be found here: <https://coredns.io/manual/toc/>

11.10 Ingress Resource

- An ingress resource is an API object containing a list of rules matched against all incoming requests. Only HTTP rules are currently supported. In order for the controller to direct traffic to the backend, the HTTP request must match both the host and the path declared in the ingress.

11.11 Ingress Controller

- Handling a few services can be easily done. However, managing thousands or tens of thousands of services can create inefficiencies. The use of an **Ingress Controller** manages ingress rules to route traffic to existing services. Ingress can be used for fan out to services, name-based hosting, TLS, or load balancing. Another feature is the ability to expose low-numbered ports. Services have been hard-coded not to expose ports lower than 1024.




Figure 11.2: Ingress Controller

- There are a few Ingress Controllers with **nginx** and **GCE** that are "officially supported" by the community. **Traefik** (pronounced "traffic") <https://traefik.io/> and **HAProxy** <https://haproxy.org> are in common use, as well. More controllers are planned, as is support for more HTTPS/TLS modes, combining L4 and L7 ingress and requesting name or IP via claims.

11.12 Service Mesh

- For more complex connections or resources such as service discovery, rate limiting, traffic management and advanced metrics you may want to implement a `service mesh`.
- A `service mesh` consists of edge and embedded proxies communicating with each other and handing traffic based on rules from a control plane. Various options are available including **Envoy**, **Istio**, and **linkerd**.
 - **Envoy** - a modular and extensible proxy favored due to modular construction, open architecture and dedication to remaining unmonitized. Often used as a data plane under other tools of a service mesh. <https://www.envoyproxy.io>
 - **Istio** - a powerful tool set which leverages Envoy proxies via a multi-component control plane. Built to be platform independent it can be used to make the service mesh flexible and feature filled. <https://istio.io>
 - **linkerd** - Another `service mesh` purpose built to be easy to deploy, fast, and ultralight. <https://linkerd.io/>

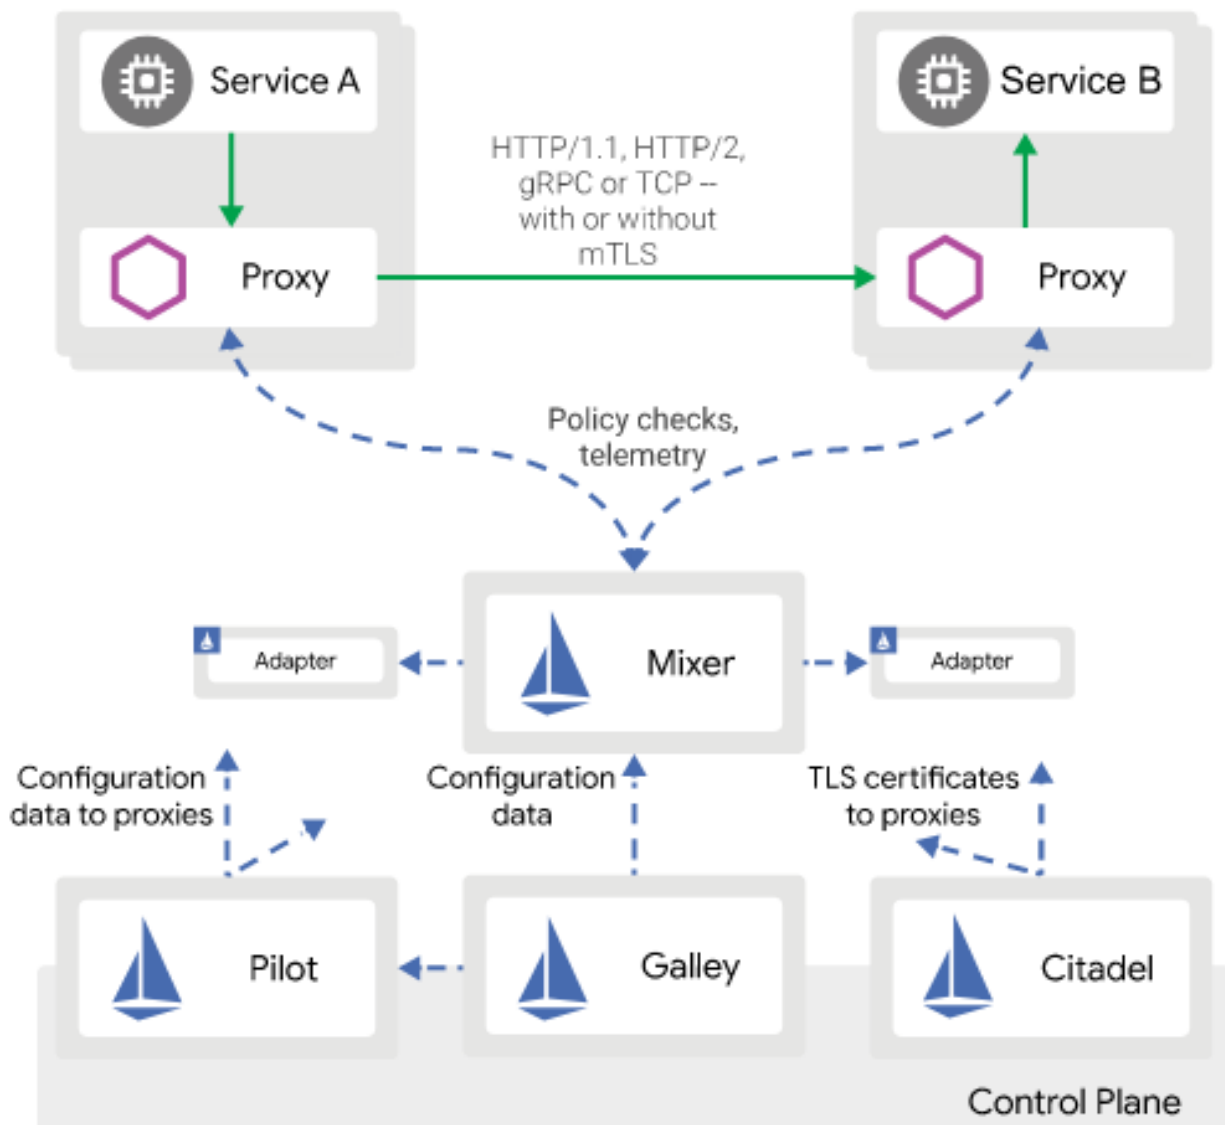


Figure 11.3: Istio Service Mesh

1

11.13 Labs

Exercise 11.1: Exposing Applications: Expose a Service

Overview

In this lab we will explore various ways to expose an application to other pods and outside the cluster. We will add to the NodePort used in previous labs other service options.

1. We will begin by using the default service type ClusterIP. This is a cluster internal IP, only reachable from within the cluster. Begin by viewing the existing services.

¹Image downloaded from <https://istio.io/docs/concepts/what-is-istio/> 23-sep-2019

```
student@master:~$ kubectl get svc
```

	NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
1	kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	8d
2	nginx	ClusterIP	10.108.95.67	<none>	443/TCP	8d
3	registry	ClusterIP	10.105.119.236	<none>	5000/TCP	8d
4	secondapp	NodePort	10.111.26.8	<none>	80:32000/TCP	7h

2. Delete the existing service for secondapp.

```
student@master:~$ cd ~/app2
```

```
student@master:~/app2$ kubectl delete svc secondapp
```

```
1 service "secondapp" deleted
```

3. Create a YAML file for a replacement service, which would be persistent. Use the label to select the secondapp. Expose the same port and protocol of the previous service.

```
student@master:~/app2$ vim service.yaml
```

YAML

service.yaml

```
1 apiVersion: v1
2 kind: Service
3 metadata:
4   name: secondapp
5   labels:
6     run: my-nginx
7 spec:
8   ports:
9     - port: 80
10     protocol: TCP
11   type: NodePort
12   selector:
13     example: second
```

4. Create the service, find the new IP and port. Note there is no high number port as this is internal access only.

```
student@master:~/app2$ kubectl create -f service.yaml
```

```
1 service/secondapp created
```

```
student@master:~/app2$ kubectl get svc
```

	NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
1	kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	8d
2	nginx	ClusterIP	10.108.95.67	<none>	443/TCP	8d
3	registry	ClusterIP	10.105.119.236	<none>	5000/TCP	8d
4	secondapp	NodePort	10.98.148.52	<none>	80:32212/TCP	14s

5. Test access. You should see the default welcome page again.

```
student@master:~/app2$ curl http://10.98.148.52
```

```
1 <!DOCTYPE html>
2 <html>
3 <head>
4 <title>Welcome to nginx!</title>
5 <output_omitted>
```

6. To expose a port to outside the cluster we will create a NodePort. We had done this in a previous step from the command line. When we create a NodePort it will create a new ClusterIP automatically. Edit the YAML file again. Add type: NodePort. Also add the high-port to match an open port in the firewall as mentioned in the previous chapter. You'll have to delete and re-create as the existing IP is immutable, but not able to be reused. The NodePort will try to create a new ClusterIP instead.

```
student@master:~/app2$ vim service.yaml
```

YAML

service.yaml

```
1 ....
2     protocol: TCP
3     nodePort: 32000           #<-- Add this and following line
4     type: NodePort
5     selector:
6     example: second
```

```
student@master:~/app2$ kubectl delete svc secondapp ; kubectl create -f service.yaml
```

```
1 service "secondapp" deleted
2 service/secondapp created
```

7. Find the new ClusterIP and ports for the service.

```
student@master:~/app2$ kubectl get svc
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	8d
nginx	ClusterIP	10.108.95.67	<none>	443/TCP	8d
registry	ClusterIP	10.105.119.236	<none>	5000/TCP	8d
secondapp	NodePort	10.109.134.221	<none>	80:32000/TCP	4s

8. Test the low port number using the new ClusterIP for the secondapp service.

```
student@master:~/app2$ curl 10.109.134.221
```

```
1 <!DOCTYPE html>
2 <html>
3 <head>
4 <title>Welcome to nginx!</title>
5 <output_omitted>
```

9. Test access from an external node to the host IP and the high container port. Your IP and port will be different. It should work, even with the network policy in place, as the traffic is arriving via a 192.168.0.0 port.

```
user@laptop:~/Desktop$ curl http://35.184.219.5:32000
```

```
1 <!DOCTYPE html>
2 <html>
3 <head>
4 <title>Welcome to nginx!</title>
5 <output_omitted>
```

10. The use of a LoadBalancer makes an asynchronous request to an external provider for a load balancer if one is available. It then creates a NodePort and waits for a response including the external IP. The local NodePort will work even before the load balancer replies. Edit the YAML file and change the type to be LoadBalancer.

```
student@master:~/app2$ vim service.yaml
```



service.yaml

```

1 ....
2 - port: 80
3   protocol: TCP
4   nodePort: 32000
5   type: LoadBalancer    #<-- Edit this line
6   selector:
7     example: second

```

```
student@master:~/app2$ kubectl delete svc secondapp ; kubectl create -f service.yaml
```

```

1 service "secondapp" deleted
2 service/secondapp created

```

11. As mentioned the cloud provider is not configured to provide a load balancer; the External-IP will remain in pending state. Some issues have been found using this with VirtualBox.

```
student@master:~/app2$ kubectl get svc
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	8d
nginx	ClusterIP	10.108.95.67	<none>	443/TCP	8d
registry	ClusterIP	10.105.119.236	<none>	5000/TCP	8d
secondapp	LoadBalancer	10.109.26.21	<pending>	80:32000/TCP	4s

12. Test again local and from a remote node. The IP addresses and ports will be different on your node.

```
serevic@laptop:~/Desktop$ curl http://35.184.219.5:32000
```

```

1 <!DOCTYPE html>
2 <html>
3 <head>
4 <title>Welcome to nginx!</title>
5 <output_omitted>

```

13. You can also use DNS names provided by **CoreDNS** which dynamically are added when the service is created. Start by logging into the busy container of secondapp.

```
student@master:~/app2$ kubectl exec -it secondapp -c busy -- sh
```



On Container

- (a) Use the **nslookup** command to find the secondapp service. Then find the registry service we configured to provide container images.

```
/ $ nslookup secondapp
```

```

1 Server:      10.96.0.10
2 Address:     10.96.0.10:53
3
4 Name:        secondapp.default.svc.cluster.local
5 Address: 10.96.214.133
6
7 *** Can't find secondapp.svc.cluster.local: No answer
8 *** Can't find secondapp.cluster.local: No answer
9 *** Can't find secondapp.c.endless-station-188822.internal: No answer
10 <output_omitted>

```



```
/ $ nslookup registry
```

```
1 Server:      10.96.0.10
2 Address:    10.96.0.10:53
3
4 Name:       registry.default.svc.cluster.local
5 Address: 10.110.95.21
6 <output_omitted>
```

- (b) Lookup the FQDN associated with the DNS server IP displayed by the commands. Your IP may be different.

```
/ $ nslookup 10.96.0.10
```

```
1 Server:      10.96.0.10
2 Address:    10.96.0.10:53
3
4 10.0.96.10.in-addr.arpa      name = kube-dns.kube-system.svc.cluster.local
```

- (c) Attempt to resolve the service name, which should not bring back any records. Then try with the FQDN. Read through the errors. You'll note that only the default namespaces is checked. You may have to check the FQDN a few times as it doesn't always reply with an answer.

```
/ $ nslookup kube-dns
```

```
1 Server:      10.96.0.10
2 Address:    10.96.0.10:53
3
4 ** server can't find kube-dns.default.svc.cluster.local: NXDOMAIN
5
6 *** Can't find kube-dns.svc.cluster.local: No answer
7 *** Can't find kube-dns.cluster.local: No answer
8 *** Can't find kube-dns.c.endless-station-188822.internal: No answer
```

```
/ $ nslookup kube-dns.kube-system.svc.cluster.local
```

```
1 Server:      10.96.0.10
2 Address:    10.96.0.10:53
3
4 Name:       kube-dns.kube-system.svc.cluster.local
5 Address: 10.96.0.10
6
7 *** Can't find kube-dns.kube-system.svc.cluster.local: No answer
```

- (d) Exit out of the container

```
/ $ exit
```

14. Create a new namespace named `multitenant` and a new deployment named `mainapp`. Expose the deployment port 80 using the name `shopping`

```
student@master:~/app2$ kubectl create ns multitenant
student@master:~/app2$ kubectl -n multitenant create deployment mainapp --image=nginx
student@master:~/app2$ kubectl -n multitenant expose deployment mainapp --name=shopping \
--type=NodePort --port=80
```

15. Log back into the `secondapp` busy container and test access to `mainapp`.

```
student@master:~/app2$ kubectl exec -it secondapp -c busy -- sh
```



On Container

- (a) Use **nslookup** to determine the address of the new service. Start with using just the service name. Then add the service name and the namespaces. As we have not configured coreDNS at all it may not know about all the possible domains to serve.

```
/ $ nslookup shopping
```

```
1 Server:      10.96.0.10
2 Address:    10.96.0.10:53
3
4 ** server can't find shopping.default.svc.cluster.local: NXDOMAIN
5
6 *** Can't find shopping.svc.cluster.local: No answer
7 <output_omitted>
```

```
/ $ nslookup shopping.multitenant
```

```
1 Server:      10.96.0.10
2 Address:    10.96.0.10:53
3
4 ** server can't find shopping.multitenant: NXDOMAIN
5
6 *** Can't find shopping.multitenant: No answer
```

```
/ $ nslookup shopping.multitenant.svc.cluster.local
```

```
1 Server:      10.96.0.10
2 Address:    10.96.0.10:53
3
4 *** Can't find shopping.multitenant.svc.cluster.local: No answer
5
6 *** Can't find shopping.multitenant.svc.cluster.local: No answer
```

- (b) Now try to use the service name and then the name with namespace, to see if it works. The DNS using the namespace should work, even if you don't have access to the default page. RBAC could be used to grant access.

```
/ $ wget shopping
```

```
1 wget: bad address 'shopping'
```

```
/ $ wget shopping.multitenant
```

```
1 Connecting to shopping.multitenant (10.101.4.142:80)
2 wget: can't open 'index.html': Permission denied
```

Exercise 11.2: Ingress Controller

If you have a large number of services to expose outside of the cluster, or to expose a low-number port on the host node you can deploy an ingress controller. While nginx and GCE have controllers officially supported by Kubernetes.io, the Traefik Ingress Controller is easier to install. At the moment.

1. As we have RBAC configured we need to make sure the controller will run and be able to work with all necessary ports, endpoints and resources. Create a YAML file to declare a clusterrole and a clusterrolebinding


```
student@master:~/app2$ vim ingress.rbac.yaml
```



ingress.rbac.yaml

```
1 kind: ClusterRole
2 apiVersion: rbac.authorization.k8s.io/v1
3 metadata:
4   name: traefik-ingress-controller
5 rules:
6   - apiGroups:
7     - ""
8     resources:
9       - services
10      - endpoints
11      - secrets
12    verbs:
13      - get
14      - list
15      - watch
16   - apiGroups:
17     - extensions
18    resources:
19      - ingresses
20    verbs:
21      - get
22      - list
23      - watch
24 ---
25 kind: ClusterRoleBinding
26 apiVersion: rbac.authorization.k8s.io/v1
27 metadata:
28   name: traefik-ingress-controller
29 roleRef:
30   apiGroup: rbac.authorization.k8s.io
31   kind: ClusterRole
32   name: traefik-ingress-controller
33 subjects:
34 - kind: ServiceAccount
35   name: traefik-ingress-controller
36   namespace: kube-system
```

2. Create the new role and binding.

```
student@master:~/app2$ kubectl create -f ingress.rbac.yaml
```

```
1 clusterrole.rbac.authorization.k8s.io/traefik-ingress-controller created
2 clusterrolebinding.rbac.authorization.k8s.io/traefik-ingress-controller created
```

3. Create the Traefik controller. The source web page changes on a regular basis. You can find a recent release by going here <https://github.com/containous/traefik/releases>. The recent 2.X release has many changes and some "undocumented features" being worked on. Find a copy of the file in the course tarball using the **find** command.

```
student@master:~/app2$ find $HOME -name traefik-ds.yaml
```

4. The output below represents the changes in a **diff** output, from a downloaded version to the edited file in the tarball. You do not have to do this step, it is so you understand the kind of changes one may need for the dashboard to work. One line was added, six lines removed. Also with version 2.0 the dashboard does not appear to work, so we are declaring the use of version 1.7.13.

```
student@master:~/app2$ diff download.yaml traefik-ds.yaml
```

YA
ML
traefik-ds.yaml

```
1 23a24          ## Add the following line 24
2 >      hostNetwork: true
3 34,39d34      ## Remove these lines around line 34
4 <      securityContext:
5 <          capabilities:
6 <              drop:
7 <                  - ALL
8 <          add:
9 <              - NET_BIND_SERVICE
```

The included file looks like this:

YA
ML
traefik-ds.rule.yaml

```
1 ....
2     terminationGracePeriodSeconds: 60
3     hostNetwork: True
4     containers:
5     - image: traefik
6       name: traefik-ingress-lb
7       ports:
8       - name: http
9         containerPort: 80
10        hostPort: 80
11       - name: admin
12         containerPort: 8080
13        hostPort: 8080
14       args:
15       - --api
16 .....
```

5. Create the objects using the edited file.

```
student@master:~/app2$ kubectl apply -f traefik-ds.yaml
```

```
1 serviceaccount/traefik-ingress-controller created
2 daemonset.apps/traefik-ingress-controller created
3 service/traefik-ingress-service created
```

6. Now that there is a new controller we need to pass some rules, so it knows how to handle requests. Note that the host mentioned is `www.example.com`, which is probably not your node name. We will pass a false header when testing. Also the service name needs to match the secondapp we've been working with.

```
student@master:~/app2$ vim ingress.rule.yaml
```

YA
ML
ingress.rule.yaml

```
1 apiVersion: networking.k8s.io/v1beta1
2 kind: Ingress
3 metadata:
4   name: ingress-test
5   annotations:
```



```

6     kubernetes.io/ingress.class: traefik
7 spec:
8   rules:
9     - host: www.example.com
10     http:
11       paths:
12         - backend:
13             serviceName: secondapp
14             servicePort: 80
15       path: /

```

7. Now ingest the rule into the cluster.

```
student@master:~/app2$ kubectl create -f ingress.rule.yaml
```

```
1 ingress.networking.k8s.io/ingress-test2 created
```

8. We should be able to test the internal and external IP addresses, and see the nginx welcome page. The loadbalancer would present the traffic, a curl request in this case, to the externally facing interface. Use **ip** a to find the IP address of the interface which would face the load balancer. In this example the interface would be `ens4`, and the IP would be `10.128.0.7`.

```
student@master:~$ ip a
```

```

1 1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen 1000
2   link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
3   inet 127.0.0.1/8 scope host lo
4       valid_lft forever preferred_lft forever
5   inet6 ::1/128 scope host
6       valid_lft forever preferred_lft forever
7 2: ens4: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1460 qdisc mq state UP group default qlen 1000
8   link/ether 42:01:0a:80:00:03 brd ff:ff:ff:ff:ff:ff
9   inet 10.128.0.7/32 brd 10.128.0.3 scope global ens4
10      valid_lft forever preferred_lft forever
11 <output_omitted>

```

```
student@master:~/app2$ curl -H "Host: www.example.com" http://10.128.0.7/
```

```

1 <!DOCTYPE html>
2 <html>
3 <head>
4 <title>Welcome to nginx!</title>
5 <style>

```

```
user@laptop:~$ curl -H "Host: www.example.com" http://35.193.3.179
```

```

1 <!DOCTYPE html>
2 <html>
3 <head>
4 <title>Welcome to nginx!</title>
5 <style>
6 <output_omitted>

```

9. At this point we would keep adding more and more web servers. We'll configure one more, which would then be a process continued as many times as desired. Begin by deploying another **nginx** server. Give it a label and expose port 80.

```
student@master:~/app2$ kubectl create deployment thirdpage --image=nginx
```

```
1 deployment.apps "thirdpage" created
```

10. Assign a label for the ingress controller to match against. Your pod name is unique, you can use the **Tab** key to complete the name.

```
student@master:~/app2$ kubectl label pod thirdpage-<tab> example=third
```

11. Expose the new server as a NodePort.

```
student@master:~/app2$ kubectl expose deployment thirdpage --port=80 --type=NodePort
```

```
1 service/thirdpage exposed
```

12. Now we will customize the installation. Run a bash shell inside the new pod. Your pod name will end differently. Install **vim** or an editor inside the container then edit the `index.html` file of nginx so that the title of the web page will be Third Page. Much of the command output is not shown below.

```
student@master:~/app2$ kubectl exec -it thirdpage-<Tab> -- /bin/bash
```

On Container

```
root@thirdpage-:/# apt-get update

root@thirdpage-:/# apt-get install vim -y

root@thirdpage-:/# vim /usr/share/nginx/html/index.html
1 <!DOCTYPE html>
2 <html>
3 <head>
4 <title>Third Page</title>      #<-- Edit this line
5 <style>
6 <output_omitted>

root@thirdpage-:/# exit
```

Edit the ingress rules to point the thirdpage service.

13. `student@master:~/app2$ kubectl edit ingress ingress-test`

YAML

ingress test

```
1 ....
2 spec:
3   rules:
4     - host: www.example.com
5       http:
6         paths:
7           - backend:
8               serviceName: secondapp
9               servicePort: 80
10          path: /
11          pathType: ImplementationSpecific
12     - host: thirdpage.org      #<-- Add this and the following 7 lines
13       http:
14         paths:
15           - backend:
```

YAML

```

16         serviceName: thirdpage
17         servicePort: 80
18         path: /
19         pathType: ImplementationSpecific
20     status:
21     ....

```

14. Test the second Host: setting using **curl** locally as well as from a remote system, be sure the <title> shows the non-default page. Use the main IP of either node.

```
student@master:~/app2$ curl -H "Host: thirdpage.org" http://10.128.0.7/
```

```

1 <!DOCTYPE html>
2 <html>
3 <head>
4 <title>Third Page</title>
5 <style>
6 <output_omitted>

```

15. The **Traefik.io** ingress controller also presents a dashboard which allows you to monitor basic traffic. From your local system open a browser and navigate to the public IP of your master node with a like this **<YOURPUBLICIP>:8080/dashboard/**. The trailing slash makes a difference.

Follow the HEALTH and PROVIDERS links at the top, as well as the the node IP links and you can view traffic when you reference the pages, from inside or outside the node. Typo the domain names inside the **curl** command and you can also see 404 error traffic. Explore as time permits.

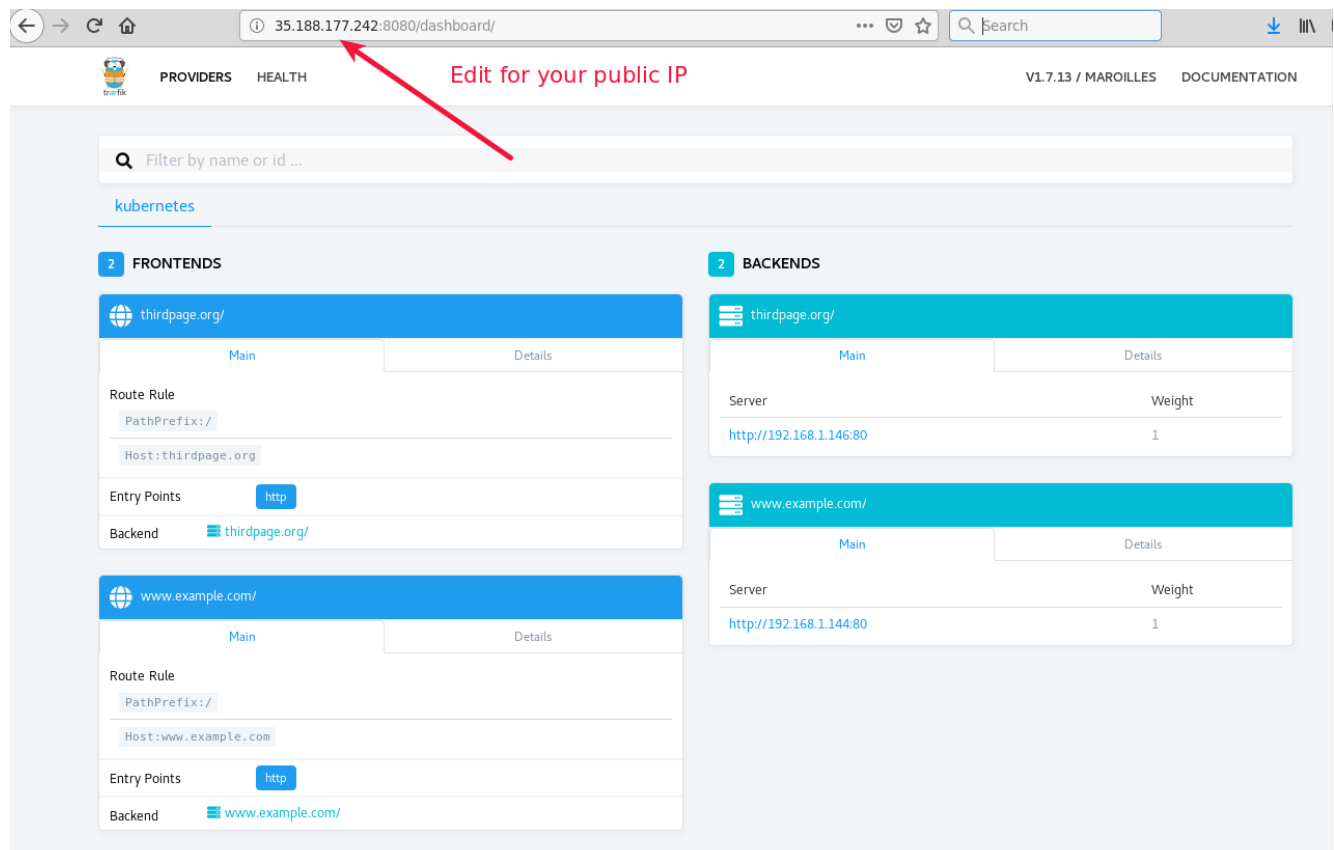


Figure 11.4: Accessing the API

```
student@master:~/app2$ curl -H "Host: thirdpage.org" http://10.128.0.7/
```

```
student@master:~/app2$ curl -H "Host: nopage.net" http://10.128.0.7/
```

✍ Exercise 11.3: Domain Review



Very Important

The source pages and content in this review could change at any time. **IT IS YOUR RESPONSIBILITY TO CHECK THE CURRENT INFORMATION.**

Revisit the CKAD domain list on [Curriculum Overview](#) and locate some of the topics we have covered in this chapter.

• Understand Services

Figure 11.5: **Service Domain Topic**

Focus on ensuring you have all the necessary files and processes understood first. Repeat the review until you are sure you have bookmarks of necessary YAML samples and can complete each step quickly, and ensure each object is running properly.

Using the browser and the three URL locations allowed by the exam, find and bookmark working YAML examples to do the following:

1. Create a new pod called `webone`, running the `nginx` service. Expose port 80.
2. Create a new service named `webone-svc`. The service should be accessible from outside the cluster.
3. Update both the pod and the service with selectors so that traffic for to the service IP shows the web server content.
4. Change the type of the service such that it is only accessible from within the cluster. Test that exterior access no longer works, but access from within the node works.
5. Deploy another pod, called `webtwo`, this time running the `wlniao/website` image. Create another service, called `webtwo-svc` such that only requests from within the cluster work. Note the default page for each server is distinct.
6. Install and configure an ingress controller such that requests for `webone.com` see the `nginx` default page, and requests for `webtwo.org` see the `wlniao/website` default page.
7. Remove any resources created in this review.

Chapter 12

Troubleshooting



Troubleshooting, which we'll cover in this chapter, can be difficult in a distributed and decoupled environment like Kubernetes. We will cover the basic flow that can be used to diagnose and identify issues you may have. We do not have cluster-wide logging built-in into Kubernetes. So, we will talk about other projects you may want to implement, like Prometheus and Fluentd, that offer monitoring and logging in a cluster-wide sense. We will talk about agent logs locally, and node logs locally, as well. With conformance testing, we can ensure that the deployment of Kubernetes that we have chosen conforms to standards used in the community. We'll also talk about cluster certification, which is a way of knowing that the Kubernetes distribution you are using is up to the standards of the community, in general.

12.1	Troubleshooting Overview	231
12.2	Basic Troubleshooting Steps	231
12.3	Ongoing (Constant) Change	232
12.4	Basic Troubleshooting Flow: Pods	232
12.5	Basic Troubleshooting Flow: Node and Security	234
12.6	Basic Troubleshooting Flow: Agents	234
12.7	Monitoring	234
12.8	Logging Tools	235
12.9	Monitoring Applications	235
12.10	System and Agent Logs	236
12.11	Conformance Testing	237
12.12	More Resource	237
12.13	Labs	238

12.1 Troubleshooting Overview

- Kubernetes relies on API calls and is sensitive to network issues. Standard Linux tools and processes are the best method for troubleshooting your cluster. If a shell, such as **bash**, is not available in an affected Pod, consider deploying another similar pod with a shell, like **busybox**. DNS configuration files and tools like **dig** are a good place to start. For more difficult challenges, you may need to install other tools, like **tcpdump**.
- Large and diverse workloads can be difficult to track, so monitoring of usage is essential. Monitoring is about collecting key metrics, such as CPU, memory, disk usage, and network bandwidth on your nodes. It can also be important to monitor well as key metrics in your applications. These features have not been ingested into Kubernetes, so exterior tools will be necessary. Exterior tools such as **Prometheus**. It combines logging, monitoring, and alerting. You can

learn more from the Prometheus website, <https://www.prometheus.io/>. It provides a time-series database, as well as integration with **Grafana** for visualization and dashboards.

- Logging activity across all the nodes is another feature not part of Kubernetes. Using **Fluentd** can be a useful data collector for a unified logging layer. Having aggregated logs can help visualize the issues, and provides the ability to search all logs. It is a good place to start when local network troubleshooting does not expose the root cause. It can be downloaded from the Fluentd website, <https://www.fluentd.org/>.
- We are going to review some of the basic **kubectl** commands that you can use to debug what is happening, and we will walk you through the basic steps to be able to debug your containers, your pending containers, and also the systems in Kubernetes.

12.2 Basic Troubleshooting Steps

- The troubleshooting flow should start with the obvious. If there are errors from the command line, investigate them first. The symptoms of the issue will probably determine the next step to check. Working from the application running inside a container to the cluster as a whole may be a good idea. The application may have a shell you can use, for example:

```
$ kubectl create deployment troubleshoot --image=nginx
$ kubectl exec -ti troubleshoot-<tab> -- /bin/sh
```

- If the Pod is running, use `kubectl logs pod-name` to view the standard out of the container. Without logs, you may consider deploying a **sidecar** container in the Pod to generate and handle logging. The next place to check is networking, including DNS, firewalls and general connectivity, using standard Linux commands and tools.
- Security settings can also be a challenge. **RBAC**, covered in the security chapter, provides mandatory or discretionary access control in a granular manner. **SELinux** and **AppArmor** are also common issues, especially with network-centric applications.
- The issues found with a decoupled system like Kubernetes are similar to those of a traditional datacenter, plus the added layers of Kubernetes controllers:
 - Errors from the command line
 - Pod logs and state of Pods
 - Use shell to troubleshoot Pod DNS and network
 - Check node logs for errors, make sure there are enough resources allocated
 - RBAC, SELinux or AppArmor for security settings
 - API calls to and from controllers to kube-apiserver
 - Inter-node network issues, DNS and firewall
 - Master server controllers (control Pods in pending or error state, errors in log files, sufficient resources, etc).

12.3 Ongoing (Constant) Change

- Unlike typical single vendor programs, there is no one organization responsible for updates to Kubernetes. As a high-velocity open source software (OSS) project, there are thousands of people working all the time. As each one of them works to solve problems or add features, it becomes possible for new problems to be caused, or features to be duplicated. Project teams and regular community meetings try to minimize issues, but they do happen.
- With single vendor products, slow to change by OSS standards, one can get a list of known issues and troubleshoot to discover the particular problem they encounter. With open access, issues are typically fixed quickly, meaning that a traditional approach may not be optimal. Instead, troubleshooting is an ongoing process, where even code or a feature which worked yesterday may have an issue today. Regular revisiting of assumptions and applying an iterative flow to discover and fix issues becomes essential.
- Be aware that the growth of agents follows project-based needs. When people working on cloud interaction needed to wait and integrate each change with overall agent management, a new agent was created, as well as a new project to develop it, called cloud controller manager (CCM). As a result, finding known bug information may require regular review of various websites as the project changes.

12.4 Basic Troubleshooting Flow: Pods

- Should the error not be seen while starting the Pod, investigate from within the Pod. A flow working from the application running inside a container to the larger cluster as a whole may be a good idea. Some applications will have a shell available to run commands, for example:

```
$ kubectl create deployment tester --image=nginx
$ kubectl exec -ti tester-<tab> -- /bin/sh
```

- A feature which is still in alpha state, and may change in anyway or disappear is the ability to attach an image to a running process. This is very helpful for testing the existing pod rather than a new pod, with different configuration. Begin by reading through some of the included examples, then try to attach to an existing pod. You will get an error as `--feature-gates=EphemeralContainers=true` is not set for the kube-apiserver and kube-scheduler by default.

```
$ kubectl alpha debug -h
```

```
1 <output_omitted>
```

```
$ kubectl alpha debug registry-6b5bb79c4-jd8fj --image=busybox
```

```
1 error: ephemeral containers are disabled for this cluster
2 (error from server: "the server could not find the requested resource").
```

- Now test a different feature. Create a container to understand the view a container has of the node it is running on. Change out the node name, `master` for your master node. Explore the commands you have available.

```
$ kubectl alpha debug node master -it --image=busybox
```



On Container

```
/ # ps
```

```
1 ....
2 31585 root      2:00 /coredns -conf /etc/coredns/Corefile
3 31591 root      0:21 containerd-shim -namespace moby -workdir /var/lib/contain
4 31633 root      0:12 /usr/bin/kube-controllers
5 31670 root      0:05 bird -R -s /var/run/calico/bird.ctl -d -c /etc/calico/conf
6 31671 root      0:05 bird6 -R -s /var/run/calico/bird6.ctl -d -c /etc/calico/co
7 31676 root      0:01 containerd-shim -namespace moby -workdir /var/lib/contain
8 31704 root      0:07 /usr/local/bin/kube-proxy --config=/var/lib/kube-proxy/con
```

```
/ # df
```

```
1 <output_omitted>
```

- If the Pod is running, use `kubectl logs pod-name` to view the standard out of the container. Without logs, you may consider deploying a sidecar container in the Pod to generate and handle logging.
- The next place to check is networking, including DNS, firewalls, and general connectivity, using standard Linux commands and tools.
- For pods without the ability to log on their own, you may have to attach a sidecar container. These can be configured to either stream application logs to their own standard out, or run a logging agent.
- Troubleshooting an application begins with the application itself. Is the application behaving as expected? Transient issues are difficult to troubleshoot; difficulties in troubleshooting are also encountered if the issue is intermittent or if it concerns occasional slow performance.

- Assuming the application is not the issue, begin by viewing the pods with the `kubectl get` command. Ensure the pods report a status of Running. A status of Pending typically means a resource is not available from the cluster, such as a properly tainted node, expected storage, or enough resources. Other error codes typically require looking at the logs and events of the containers for further troubleshooting. Also, look for an unusual number of restarts. A container is expected to restart due to several reasons, such as a command argument finishing. If the restarts are not due to that, it may indicate that the deployed application is having an issue and failing due to panic or probe failure.
- View the details of the pod and the status of the containers using the `kubectl describe pod` command. This will report overall pod status, container configurations and container events. Work through each section looking for a reported error.
- Should the reported information not indicate the issue, the next step would be to view any logs of the container, in case there is a misconfiguration or missing resource unknown to the cluster, but required by the application. These logs can be seen with the `kubectl logs <pod-name> <container-name>` command.

12.5 Basic Troubleshooting Flow: Node and Security

- Security settings can also be a challenge. RBAC, covered in the Security chapter, provides mandatory or discretionary access control in a granular manner. SELinux and AppArmor are also common issues, especially with network-centric applications.
- Disabling security for testing would be a common response to an issue. Each tool has its own logs and indications of rule violation. There could be multiple issues to troubleshoot, so be sure to re-enable security and test the workload again.
- Internode networking can also be an issue. Changes to switches, routes, or other network settings can adversely affect Kubernetes. Historically, the primary causes were DNS and firewalls. As Kubernetes integrations have become more common and DNS integration is maturing, this has become less of an issue. Still, check connectivity and for recent infrastructure changes as part of your troubleshooting process. Every so often, an update which was said shouldn't cause an issue may, in fact, be causing an issue.

12.6 Basic Troubleshooting Flow: Agents

- The issues found with a decoupled system like Kubernetes are similar to those of a traditional datacenter, plus the added layers of Kubernetes controllers:
 - Control pods in pending or error state
 - Look for errors in log files
 - Are there enough resources?
 - etc.

12.7 Monitoring

- Monitoring is about collecting metrics from the infrastructure, as well as applications.
- **Prometheus** is part of the Cloud Native Computing Foundation (CNCF). As a Kubernetes plugin, it allows one to scrape resource usage metrics from Kubernetes objects across the entire cluster. It also has several client libraries which allow you to instrument your application code in order to collect application level metrics.
- Collecting the metrics is the first step, making use of the data the next. As open source software we have the ability to expose lots of data points. Graphing the data with a tool like **Grafana** can allow for visual understanding of the cluster and application status. Some facilities fill walls with large TVs offering the entire company a real time view of demand and utilization.



Figure 12.1: Grafana Dashboard Example

1

12.8 Logging Tools

- Logging, like monitoring, is a vast subject in IT. It has many tools that you can use as part of your arsenal.
- Typically, logs are collected locally and aggregated before being ingested by a search engine and displayed via a dashboard which can use the search syntax. While there are many software stacks that you can use for logging, the Elasticsearch, Logstash, and Kibana Stack (ELK) at <https://www.elastic.co/videos/introduction-to-the-elk-stack> has become quite common.
- In Kubernetes, the **kubelet** writes container logs to local files (via the Docker logging driver). The `kubectl logs` command allows you to retrieve these logs.
- Cluster-wide, you can use **Fluentd** (<https://www.fluentd.org/>) to aggregate logs. Check the cluster administration logging concepts at <https://www.fluentd.org/> for a detailed description.
- Fluentd is part of the Cloud Native Computing Foundation and, together with Prometheus, they make a nice combination for monitoring and logging. You can find a detailed walk-through of running Fluentd on Kubernetes in the Kubernetes documentation at <https://kubernetes.io/docs/tasks/debug-application-cluster/logging-elasticsearch-kibana/>.
- Setting up Fluentd for Kubernetes logging is a good exercise in understanding `DaemonSets`. Fluentd agents run on each node via a `DaemonSet`, they aggregate the logs, and feed them to an Elasticsearch instance prior to visualization in a Kibana dashboard.

12.9 Monitoring Applications

- As a distributed system, Kubernetes lacks monitoring and tracing tools which are cluster-aware. Other CNCF projects have started to handle various areas. As they are independent projects, you may find they have some overlap in capability:

¹ Image downloaded from <https://grafana.com/grafana/dashboards/8588> 17-August-2019

– Prometheus

Focused on metrics and alerting. Provides a time-series database, queries and alerts to gather information and alerts for the entire cluster. With integration with Grafana, as well as simple graphs and console templates, you can easily visualize the cluster condition graphically

– Fluentd

A unified logging layer which can collect logs from over 500 plugins and deliver to a multitude of outputs. A single, enterprise-wide tool to filter, buffer, and route messages just about anywhere you may want.

– OpenTracing

While we just learned about logging and metrics, neither of the previously mentioned projects allow for a granular understanding of a transaction across a distributed architecture. This project seems to provide worthwhile instrumentation to propagate tracing among all services, code and packages, so the trace can be complete. Its goal is a “single, standard mechanism” for everything.

– Jaeger

A tracing system developed by Uber, focused on distributed context propagation, transaction monitoring, and root cause analysis, among other features. This has become a common tracing implementation of OpenTracing.

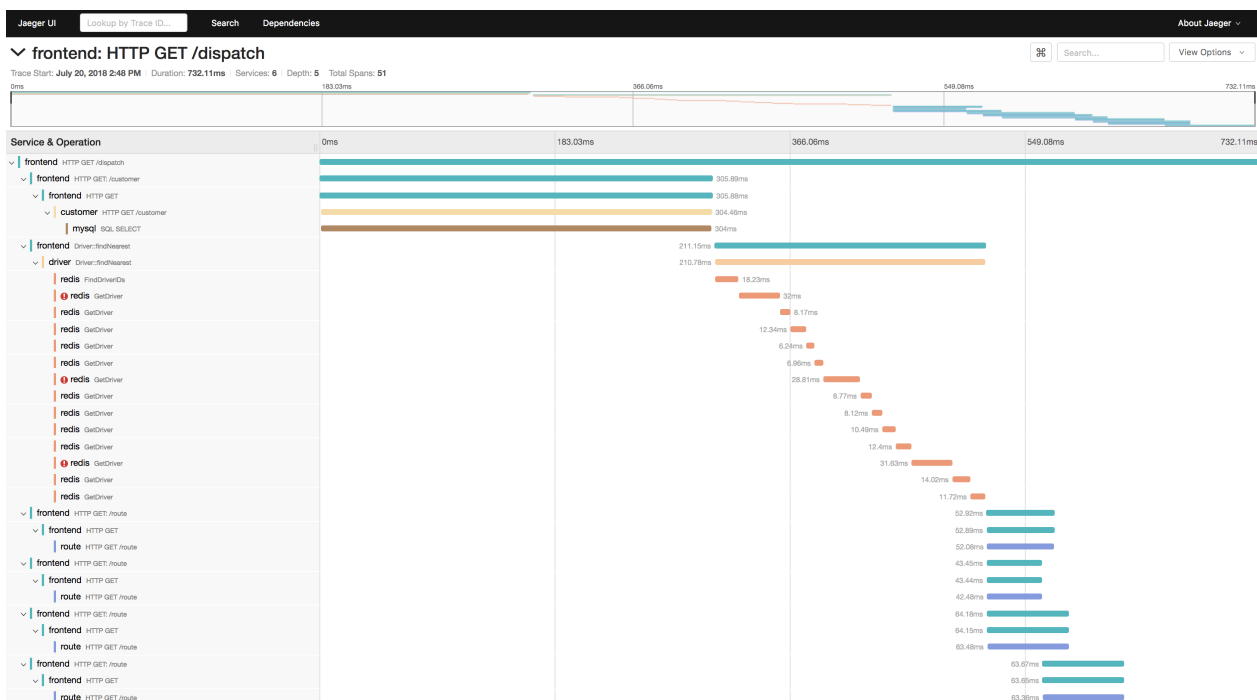


Figure 12.2: Jaeger Tracing Example

2

12.10 System and Agent Logs

- Where system and agent files are found depends on the existence of **systemd**. Those with systemd will log to **journalctl**, which can be viewed with `journalctl -a`. Unless the `/var/log/journal` directory exists, the journal is volatile. As Kubernetes can generate a lot of logs, it would be advisable to configure log rotation, should this directory be created.
- Without **systemd**, the logs will be created under `/var/log/<agent>.log`, in which case it would also be advisable to configure log rotation.
- Container components:

²Image downloaded from <https://www.jaegertracing.io/img/trace-detail-ss.png> 17-August-2019

- kube-scheduler
- kube-proxy
- Non-container components:
 - kubelet
 - Docker
 - Many others.

12.11 Conformance Testing

- The flexibility of Kubernetes can lead to the development of a non-conforming cluster.
 - Meet the demands of your environment
 - Several vendor-provided tools for conformance testing
 - For ease of use, Sonobuoy by Heptio can be used to understand the state of the cluster.

Certified Kubernetes Conformant Program

- With more than 60 known distributions, there is a challenge to consistency and portability. In late 2017, a new program was started by CNCF to certify distributions that meet essential requirements and adhere to complete API functionality:
 - Confidence that workloads from one distribution work on others
 - Guarantees complete API functions
 - Testing and Architecture Special Interest Groups.
- You can read more about submitting conformance results on GitHub at <https://github.com/cncf/k8s-conformance/blob/master/instructions.md>.

12.12 More Resource

- There are several things that you can do to quickly diagnose potential issues with your application and/or cluster. The official documentation offers additional materials to help you get familiar with troubleshooting:
 - General guidelines and instructions (Troubleshooting)
 - <https://kubernetes.io/docs/tasks/debug-application-cluster/troubleshooting/>
 - Troubleshooting applications
 - <https://kubernetes.io/docs/tasks/debug-application-cluster/debug-application/>
 - Troubleshooting clusters
 - <https://kubernetes.io/docs/tasks/debug-application-cluster/debug-cluster/>
 - Debugging Pods
 - <https://kubernetes.io/docs/tasks/debug-application-cluster/debug-pod-replication-controller/>
 - Debugging Services
 - <https://kubernetes.io/docs/tasks/debug-application-cluster/debug-service/>
 - GitHub website for issues and bug tracking
 - <https://github.com/kubernetes/kubernetes/issues>
 - Kubernetes Slack channel
 - <https://kubernetes.slack.com/>

12.13 Labs

Exercise 12.1: Troubleshooting: Monitor Applications

Overview

Troubleshooting can be difficult in a multi-node, decoupled and transient environment. Add in the rapid pace of change and it becomes more difficult. Instead of focusing and remembering a particular error and the fix it may be more useful to learn a flow of troubleshooting and revisit assumptions until the pace of change slows and various areas further mature.

1. View the `secondapp` pod, it should show as `Running`. This may not mean the application within is working properly, but that the pod is running. The restarts are due to the command we have written to run. The pod exists when done, and the controller restarts another container inside. The count depends on how long the labs have been running.

```
student@master/app2:~$ cd
student@master:~$ kubectl get pods secondapp
```

	NAME	READY	STATUS	RESTARTS	AGE
1					
2	secondapp	2/2	Running	49	2d

2. Look closer at the pod. Working slowly through the output check each line. If you have issues, are other pods having issues on the same node or volume? Check the state of each container. Both `busy` and `webserver` should report as `Running`. Note `webserver` has a restart count of zero while `busy` has a restart count of 49. We expect this as, in our case, the pod has been running for 49 hours.

```
student@master:~$ kubectl describe pod secondapp
```

```
1 Name:          secondapp
2 Namespace:     default
3 Node:          worker-wdrq/10.128.0.2
4 Start Time:    Fri, 13 Apr 2018 20:34:56 +0000
5 Labels:        example=second
6 Annotations:   <none>
7 Status:        Running
8 IP:            192.168.55.91
9 Containers:
10  webserver:
11  <output_omitted>
12    State:      Running
13    Started:    Fri, 13 Apr 2018 20:34:58 +0000
14    Ready:      True
15    Restart Count: 0
16  <output_omitted>
17
18  busy:
19  <output_omitted>
20
21    State:      Running
22    Started:    Sun, 15 Apr 2018 21:36:20 +0000
23    Last State: Terminated
24    Reason:     Completed
25    Exit Code:  0
26    Started:    Sun, 15 Apr 2018 20:36:18 +0000
27    Finished:   Sun, 15 Apr 2018 21:36:18 +0000
28    Ready:      True
29    Restart Count: 49
30    Environment: <none>
```

3. There are three values for conditions. Check that the pod reports `Initialized`, `Ready` and `scheduled`.

```

1 <output_omitted>
2 Conditions:
3   Type           Status
4   Initialized     True
5   Ready           True
6   PodScheduled    True
7 <output_omitted>

```

4. Check if there are any events with errors or warnings which may indicate what is causing any problems.

```

1 Events:
2   Type      Reason      Age           From           Message
3   ----      -
4   Normal    Pulling     34m (x50 over 2d) kubelet, worker-wdrq pulling
5   image "busybox"
6   Normal    Pulled      34m (x50 over 2d) kubelet, worker-wdrq Successfully
7   pulled image "busybox"
8   Normal    Created     34m (x50 over 2d) kubelet, worker-wdrq Created
9   container
10  Normal    Started     34m (x50 over 2d) kubelet, worker-wdrq Started
11  container

```

5. View each container log. You may have to sift errors from expected output. Some containers may have no output at all, as is found with busy.

```
student@master:~$ kubectl logs secondapp webserver
```

```

1 192.168.55.0 - - [13/Apr/2018:21:18:13 +0000] "GET / HTTP/1.1" 200
2 612 "-" "curl/7.47.0" "-"
3 192.168.55.0 - - [13/Apr/2018:21:20:35 +0000] "GET / HTTP/1.1" 200
4 612 "-" "curl/7.53.1" "-"
5 127.0.0.1 - - [13/Apr/2018:21:25:29 +0000] "GET" 400 174 "-" "-" "-"
6 127.0.0.1 - - [13/Apr/2018:21:26:19 +0000] "GET index.html" 400 174
7 "-" "-" "-"
8 <output_omitted>

```

```
student@master:~$ kubectl logs secondapp busy
```

```
1 student@master:~$
```

6. Check to make sure the container is able to use DNS and communicate with the outside world. Remember we still have limited the UID for secondapp to be UID **2000**, which may prevent some commands from running. It can also prevent an application from completing expected tasks, and other errors.

```
student@master:~$ kubectl exec -it secondapp -c busy -- sh
```



On Container

```
/ $ nslookup www.linuxfoundation.org
```

```

1 / $ nslookup www.linuxfoundation.org
2 Server:          10.96.0.10
3 Address:         10.96.0.10:53
4
5 Non-authoritative answer:
6 Name:            www.linuxfoundation.org
7 Address: 23.185.0.2
8
9 *** Can't find www.linuxfoundation.org: No answer

```




```
/ $ cat /etc/resolv.conf
1 nameserver 10.96.0.10
2 search default.svc.cluster.local svc.cluster.local
3 cluster.local c.endless-station-188822.internal
4 google.internal
5 options ndots:5
```

Test access to a remote node using **nc (NetCat)**. There are several options to **nc** which can help troubleshoot if the problem is the local node, something between nodes or in the target. In the example below the connect never completes and a **control-c** was used to interrupt.

7. / \$ nc www.linux.com 25

```
1 ^Cpunt!
```

8. Test using an IP address in order to narrow the issue to name resolution. In this case the IP in use is a well known IP for Google's DNS servers. The following example shows that Internet name resolution is working, but our UID issue prevents access to the `index.html` file.

```
/ $ wget http://www.linux.com/
```

```
1 Connecting to www.linux.com (151.101.45.5:80)
2 Connecting to www.linux.com (151.101.45.5:443)
3 wget: can't open 'index.html': Permission denied
4
5 / $ exit
```

9. Make sure traffic is being sent to the correct Pod. Check the details of both the service and endpoint. Pay close attention to ports in use as a simple typo can prevent traffic from reaching the proper pod. Make sure labels and selectors don't have any typos as well.

```
student@master:~$ kubectl get svc
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	10d
nginx	ClusterIP	10.108.95.67	<none>	443/TCP	10d
registry	ClusterIP	10.105.119.236	<none>	5000/TCP	10d
secondapp	LoadBalancer	10.109.26.21	<pending>	80:32000/TCP	1d
thirdpage	NodePort	10.109.250.78	<none>	80:31230/TCP	1h

```
student@master:~$ kubectl get svc secondapp -o yaml
```

```
1 <output_omitted>
2   clusterIP: 10.109.26.21
3   externalTrafficPolicy: Cluster
4   ports:
5   - nodePort: 32000
6     port: 80
7     protocol: TCP
8     targetPort: 80
9   selector:
10     example: second
11 <output_omitted>
```

10. Verify an endpoint for the service exists and has expected values, including namespaces, ports and protocols.

```
student@master:~$ kubectl get ep
```



```

1 NAME           ENDPOINTS           AGE
2 kubernetes     10.128.0.3:6443      10d
3 nginx          192.168.55.68:443    10d
4 registry       192.168.55.69:5000   10d
5 secondapp      192.168.55.91:80     1d
6 thirdpage      192.168.241.57:80    1h

```

```
student@master:~$ kubectl get ep secondapp -o yaml
```

```

1 apiVersion: v1
2 kind: Endpoints
3 metadata:
4   creationTimestamp: 2018-04-14T05:37:32Z
5 <output_omitted>

```

11. If the containers, services and endpoints are working the issue may be with an infrastructure service like **kube-proxy**. Ensure it's running, then look for errors in the logs. As we have two nodes we will have two proxies to look at. As we built our cluster with **kubeadm** the proxy runs as a container. On other systems you may need to use **journalctl** or look under `/var/log/kube-proxy.log`.

```
student@master:~$ ps -elf |grep kube-proxy
```

```

1 4 S root      2864  2847  0 80   0 - 14178 -      15:45 ?
2 00:00:56 /usr/local/bin/kube-proxy --config=/var/lib/kube-proxy/config.conf
3 0 S student  23513 18282  0 80   0 - 3236 pipe_w 22:49 pts/0
4 00:00:00 grep --color=auto kube-proxy

```

```
student@master:~$ journalctl -a | grep proxy
```

```

1 Apr 15 15:44:43 worker-nzjr audit[742]: AVC apparmor="STATUS"
2 operation="profile_load" profile="unconfined" \
3 name="/usr/lib/lxd/lxd-bridge-proxy" pid=742 comm="apparmor_parser"
4 Apr 15 15:44:43 worker-nzjr kernel: audit: type=1400
5 audit(1523807083.011:11): apparmor="STATUS" \
6 operation="profile_load" profile="unconfined" \
7 name="/usr/lib/lxd/lxd-bridge-proxy" pid=742 comm="apparmor_parser"
8 Apr 15 15:45:17 worker-nzjr kubelet[1248]: I0415 15:45:17.153670
9 1248 reconciler.go:217] operationExecutor.VerifyControllerAttachedVolume\
10 started for volume "xtables-lock" \
11 (UniqueName: "kubernetes.io/host-path/e701fc01-38f3-11e8-a142-\
12 42010a800003-xtables-lock") \
13 pod "kube-proxy-t8k4w" (UID: "e701fc01-38f3-11e8-a142-42010a800003")

```

12. Look at both of the proxy logs. Lines which begin with the character **I** are info, **E** are errors. In this example the last message says access to listing an endpoint was denied by RBAC. It was because a default installation via Helm wasn't RBAC aware. If not using command line completion, view the possible pod names first.

```
student@master:~$ kubectl -n kube-system get pod
```

```
student@master:~$ kubectl -n kube-system logs kube-proxy-fsdf
```

```

1 I0405 17:28:37.091224      1 feature_gate.go:190] feature gates: map[]
2 W0405 17:28:37.100565      1 server_others.go:289] Flag proxy-mode=""
3 unknown, assuming iptables proxy
4 I0405 17:28:37.101846      1 server_others.go:138] Using iptables Proxier.
5 I0405 17:28:37.121601      1 server_others.go:171] Tearing down
6 inactive rules.
7 <output_omitted>
8 E0415 15:45:17.086081      1 reflector.go:205] \
9 k8s.io/kubernetes/pkg/client/informers/informers_generated/
10 internalversion/factory.go:85: \

```

```

11 Failed to list *core.Endpoints: endpoints is forbidden: \
12   User "system:serviceaccount:kube-system:kube-proxy" cannot \
13   list endpoints at the cluster scope:\
14 [clusterrole.rbac.authorization.k8s.io "system:node-proxier" not found, \
15   clusterrole.rbac.authorization.k8s.io "system:basic-user" not found,
16   clusterrole.rbac.authorization.k8s.io \
17   "system:discovery" not found]

```

13. Check that the proxy is creating the expected rules for the problem service. Find the destination port being used for the service, **32000** in this case.

```
student@master:~$ sudo iptables-save |grep secondapp
```

```

1 -A KUBE-NODEPORTS -p tcp -m comment --comment "default/secondapp:" \
2 -m tcp --dport 32000 -j KUBE-MARK-MASQ
3 -A KUBE-NODEPORTS -p tcp -m comment --comment "default/secondapp:" \
4 -m tcp --dport 32000 -j KUBE-SVC-DAASHM5XQZF5XI3E
5 -A KUBE-SERVICES !-s 192.168.0.0/16 -d 10.109.26.21/32 -p tcp \
6 -m comment --comment "default/secondapp: \
7   cluster IP" -m tcp --dport 80 -j KUBE-MARK-MASQ
8 -A KUBE-SERVICES -d 10.109.26.21/32 -p tcp -m comment --comment \
9 "default/secondapp: cluster IP" -m tcp \
10   --dport 80 -j KUBE-SVC-DAASHM5XQZF5XI3E
11 <output_omitted>

```

14. Ensure the proxy is working by checking the port targeted by **iptables**. If it fails open a second terminal and view the proxy logs when making a request as it happens.

```
student@master:~$ curl localhost:32000
```

```

1 <!DOCTYPE html>
2 <html>
3 <head>
4 <title>Welcome to nginx!</title>
5 <output_omitted>

```

Exercise 12.2: OPTIONAL LAB: Conformance Testing

The **cncf.io** group is in the process of formalizing what is considered to be a conforming Kubernetes cluster. While that project matures there is an existing tool provided by **Heptio** which can be useful. We will need to make sure a newer version of **Golang** is installed for it to work. You can download the code from github and look around with git or with go, depending on which tool you are most familiar. **Things change quickly these steps may not work....today**

1. Download a compiled binary. A shorter URL is shown first, then the longer, just in case the link changes and you need to navigate. They should download the same file.

```
student@master:~$ curl -sLO https://tinyurl.com/yyu5bs28
```

```
student@master:~$ mv yyu5bs28 sonobuoy.tar.gz
```

```
student@master:~$ tar -xvf sonobuoy.tar.gz
```

```

1 LICENSE
2 sonobuoy

```

```

student@master:~$ curl -sLO \
https://github.com/heptio/sonobuoy/releases/download/v0.15.4/sonobuoy_0.15.4_linux_amd64.tar.gz

```

- Run the test. We will not use the `--wait` option, which will capture the screen until the test finishes. This could take a while to finish. You should get some output indicating testing objects being created.

```
student@master:~$ sudo mv sonobuoy /usr/local/bin/
```

```
student@master:~$ sonobuoy run
```

```
1 WARN[0000] The maximum supported Kubernetes version is 1.15.99, but
2 the server version is v1.16.1. Sonobuoy will continue but unexpected results may occur.
3 INFO[0000] created object                name=sonobuoy namespace= resource=namespaces
4 INFO[0000] created object                name=sonobuoy-serviceaccount namespace=sonobuoy ....
5 INFO[0000] created object                name=sonobuoy-serviceaccount-sonobuoy namespace=....
6 INFO[0000] created object                name=sonobuoy-serviceaccount namespace= resource....
7 INFO[0000] created object                name=sonobuoy-config-cm namespace=sonobuoy resou....
8 INFO[0000] created object                name=sonobuoy-plugins-cm namespace=sonobuoy reso....
9 INFO[0000] created object                name=sonobuoy namespace=sonobuoy resource=pods
10 INFO[0000] created object                name=sonobuoy-master namespace=sonobuoy resource....
```

- View the results inside the sonobuoy pod.

```
student@master:~$ kubectl get pods --all-namespaces
```

```
1 <output_omitted>
2 sonobuoy      sonobuoy                        1/1
3   Running    0          90s
4 sonobuoy      sonobuoy-e2e-job-b3bcb52b4fd54367  2/2
5   Running    0          85s
6 sonobuoy      sonobuoy-systemd-logs-daemon-set-f7ca2bb9a7174908-h47kb  2/2   Running    0          85s
7 sonobuoy      sonobuoy-systemd-logs-daemon-set-f7ca2bb9a7174908-s22d6  2/2   Running    0          85s
```

```
student@master:~$ kubectl -n sonobuoy exec -it sonobuoy -- /bin/bash
```



On Container

- View the files inside the container.

```
root@sonobuoy:/# ls
```

```
1 bin   home   mnt      root      sbin      tmp
2 boot  lib     opt      run        sonobuoy  usr
3 dev   lib64   plugins.d run_master.sh  srv        var
4 etc   media   proc     run_single_node_worker.sh sys
```

- View the `run_master.sh` script. Note that it mentions both the **sonobuoy** command and where to find the results.

```
root@sonobuoy:/# cat run_master.sh
```

```
1 #!/bin/bash
2 #####
3 # Copyright 2017 Heptio Inc.
4 #
5
6 <output_omitted>
7 RESULTS_DIR="${RESULTS_DIR:-/tmp/sonobuoy}"
8 # It's ok for these env vars to be unbound
9 RESULTS_DIR="${RESULTS_DIR}" SONOBUOY_CONFIG="${SONOBUOY_CONFIG}"
10 SONOBUOY_ADVERTISE_IP="${SONOBUOY_ADVERTISE_IP}" /sonobuoy master -v 3 --logtostderr
11
12 echo -n "${RESULTS_DIR}/${ls -t "${RESULTS_DIR}" | grep -v done | head -n 1}" > "${RESULTS_DIR}/done"
```

- View the contents of the `/tmp/sonobuoy` directory. Note the subdirectory is a generated number, yours will be different. The **Tab** key can be used to complete the path.



```

root@sonobuoy:/# ls /tmp/sonobuoy/
1 d39f2629-fa3c-4a0b-9b33-53080e78b57b

root@sonobuoy:/# cd /tmp/sonobuoy/d39f2629-fa3c-4a0b-9b33-53080e78b57b ; ls
1 meta plugins

root@sonobuoy:...57b# find .
1 .
2 ./plugins
3 ./plugins/systemd-logs
4 ./plugins/systemd-logs/results
5 ./plugins/systemd-logs/results/e-6clr
6 ./plugins/systemd-logs/results/e-6clr/systemd_logs
7 ./plugins/systemd-logs/results/e-5c7t
8 ./plugins/systemd-logs/results/e-5c7t/systemd_logs
9 ./meta
10 ./meta/run.log
11 ./meta/config.json

```

7. The **sonobuoy** command has several options. We will use two to explore the test output.

```

root@sonobuoy:...57b# cd /

root@sonobuoy:/# ./sonobuoy status
1
2      PLUGIN      STATUS  RESULT  COUNT
3      e2e         running      1
4      systemd-logs complete      2
5 Sonobuoy is still running. Runs can take up to 60 minutes.

root@sonobuoy:/# ./sonobuoy logs
1 <output_omitted>

```

8. Continue to look through tests and results as time permits. Connect to the other pods in the **sonobuoy** namespace and look for log and result files.

There is also an online, graphical scanner. In testing, inside GCE, the results were blocked and never returned. You may have different outcome in other environments.

Exercise 12.3: Domain Review



Very Important

The source pages and content in this review could change at any time. **IT IS YOUR RESPONSIBILITY TO CHECK THE CURRENT INFORMATION.**

Revisit the CKAD domain list on [Curriculum Overview](#) and locate some of the topics we have covered in this chapter.

- Understand debugging in Kubernetes

Figure 12.3: Troubleshooting Domain Topic

Focus on ensuring you have all the necessary files and processes understood first. Repeat the review until you are sure you have bookmarks of necessary YAML samples and can complete each step quickly, and ensure each object is running properly.

1. Find and use the `troubleshoot-review1.yaml` file to create a deployment. The **create** command will fail. Edit the file to fix issues such that a single pod runs for at least a minute without issue. There are several things to fix.

```
student@master:~$ kubectl create -f troubleshoot-review1.yaml
```

```
1 <Fix any errors found here>
```

When fixed it should look like this:

```
student@master:~$ kubectl get deploy igotttrouble
```

```
1 NAME          READY    UP-TO-DATE    AVAILABLE    AGE
2 igotttrouble  1/1      1             1            5m13s
```

2. Remove any resources created during this review.

Chapter 13

Logging and Troubleshooting - II



13.1	Overview	248
13.2	Troubleshooting Flow	249
13.3	Basic Start Sequence	251
13.4	Monitoring	252
13.5	Plugins	253
13.6	Logging	256
13.7	Troubleshooting Resources	257
13.8	Labs	258

13.1 Overview

Overview

- **Linux** troubleshooting via shell
- Turn on basic monitoring
- Set up cluster-wide logging
- External products **Fluentd**, **Prometheus** helpful.
- Internal **Metrics Server** and API

Kubernetes relies on API calls and is sensitive to network issues. Standard **Linux** tools and processes are the best method for troubleshooting your cluster. If a shell, such as **bash** is not available in an affected Pod, consider deploying another similar pod with a shell, like **busybox**. DNS configuration files and tools like **dig** are a good place to start. For more difficult challenges you may need to install other tools like **tcpdump**.

Large and diverse workloads can be difficult to track, so monitoring of usage is essential. Monitoring is about collecting key metrics such as CPU, memory, and disk usage, and network bandwidth on your nodes, as well as monitoring key metrics in your applications. These features are being ingested into Kubernetes with the **Metric Server**, which is a cut-down version of the now deprecated **Heapster**. Once installed the **Metrics Server** exposes a standard API which can be consumed by other agents such as autoscalers. Once installed this endpoint can be found here on the master server: [/apis/metrics/k8s.io/](https://apis/metrics/k8s.io/)

Logging activity across all the nodes is another feature not part of Kubernetes. Using **Fluentd** can be a useful data collector for a unified logging layer. Having aggregated logs can help visualize the issue and provides the ability to search all logs. A good place to start when local network troubleshooting does not expose the root cause. It can be downloaded from <http://www.fluentd.org>

Another project from cncf.io combines logging, monitoring, and alerting called **Prometheus** can be found here <https://prometheus.io>. It provides a time-series database as well as integration with **Grafana** for visualization and dashboards.

We are going to review some of the basic **kubectl** commands that you can use to debug what is happening, and we will walk you through the basic steps to be able to debug your containers, your pending containers, and also the systems in Kubernetes.

13.2 Troubleshooting Flow

Basic Steps

- Errors from command line
- Pod logs and state of the Pod
- Use shell to troubleshoot Pod DNS and network
- Check node logs for errors. Enough resources
- RBAC, SELinux or AppArmor security settings
- API calls to and from controllers to **kube-apiserver**
- Enable auditing
- Inter-node network issues, DNS and firewall
- Master server controllers.
 - Control Pods in pending or error state
 - Errors in log files
 - Enough resources

The flow of troubleshooting should start with the obvious. If there are errors from the command line investigate them first. The symptoms of the issue will probably determine the next step to check. Working from the application running inside a container to the cluster as a whole may be a good idea. The application may have a shell you can use, for example:

```
$ kubectl create deploy busybox --image=busybox --command sleep 3600
```

```
$ kubectl exec -ti <busybox_pod> -- /bin/sh
```

If the Pod is running use **kubectl logs pod-name** to view standard out of the container. Without logs you may consider deploying a **sidecar** container in the Pod to generate and handle logging. The next place to check is networking including DNS, firewalls and general connectivity using standard **Linux** commands and tools.

Security settings can also be a challenge. RBAC, covered in the security chapter, provides mandatory or discretionary access control in a granular manner. SELinux and AppArmor are also common issues, especially with network-centric applications.

A newer feature of Kubernetes is the ability to enable auditing for the **kube-apiserver**, which can allow a view into actions after the API call has been accepted.

The issues found with a decoupled system like Kubernetes are similar to those of a traditional datacenter, plus the added layers of Kubernetes controllers.

Ephemeral Containers

- Add a tool-filled container to a running pod
- Alpha with v1.16 release

```
kubectl debug buggypod --image debian --attach
```

A feature new to the 1.16 version is the ability to add a container to a running pod. This would allow a feature-filled container to be added to an existing pod without having to terminate and re-create. Intermittent and difficult to determine problems may take a while to reproduce, or not exist with the addition of another container.

As an Alpha stability feature, it may change or be removed at any time. As well they will not be restarted automatically, and several resources such as `ports` or `resources` are not allowed.

These containers are added via the `ephemeralcontainers` handler via an API call, not via the `podSpec`. As a result the use of **`kubectl edit`** is not possible.

You may be able to use the **`kubectl attach`** command to join an existing process within the container. This can be helpful instead of **`kubectl exec`** which executes a new process. The functionality of the attached process depends entirely on what you are attaching to.

13.3 Basic Start Sequence

Cluster Start Sequence

- **systemd** starts `kubelet.service`
 - Uses `/etc/systemd/system/kubelet.service.d/10-kubeadm.conf`
 - Uses `/var/lib/kubelet/config.yaml` config file
 - **staticPodPath** set to `/etc/kubernetes/manifests/`
- **kubelet** creates all pods from `*.yaml` in directory
 - kube-apiserver
 - etcd
 - kube-controller-manager
 - kube-scheduler
- kube-controller-manager control loops use etcd data to start rest

The cluster startup sequence begins with **systemd** if you built the cluster using **kubeadm**. Other tools may leverage a different method. Use **systemctl status kubelet.service** to see the current state and configuration files used to run the **kubelet** binary.

Inside of the `config.yaml` file you will find several settings for the binary including the **staticPodPath** which indicates the directory where **kubelet** will read every yaml file and start every pod. If you put a yaml file in this directory it is a way to troubleshoot the scheduler, as the pod is created with any requests to the scheduler.

The four default yaml files will start the base pods necessary to run the cluster. Once the watch loops and controllers from **kube-controller-manager** run using **etcd** data the rest of the configured objects will be created.

13.4 Monitoring

Monitoring

- Enable the add-ons
- **Metrics Server** and API
- **Prometheus**

Monitoring is about collecting metrics from the infrastructure, as well as applications.

The long used and now deprecated **Heapster** has been replaced with an integrated **Metrics Server**. Once installed and configured the server exposes a standard API which other agents can use to determine usage. It can also be configured to expose custom metrics, which then could also be used by autoscalers to determine if an action should take place.

Prometheus

Prometheus is part of the **Cloud Native Computing Foundation (CNCF)**. As a Kubernetes plugin, it allows one to scrape resource usage metrics from Kubernetes objects across the entire cluster. It also has several client libraries which allow you to instrument your application code in order to collect application level metrics.

13.5 Plugins

Using Krew

- Install the software using steps at <https://krew.dev>

```
$ kubectl krew help
krew is the kubectl plugin manager.
You can invoke krew through kubectl: "kubectl krew [command]..."
```

Usage:

```
krew [command]
```

Available Commands:

help	Help about any command
info	Show information about a kubectl plugin
install	Install kubectl plugins
list	List installed kubectl plugins
search	Discover kubectl plugins
uninstall	Uninstall plugins
update	Update the local copy of the plugin index
upgrade	Upgrade installed plugins to newer versions
version	Show krew version and diagnostics

We have been using the **kubectl** command throughout the course. The basic commands can be used together in a more complex manner extending what can be done. There are over seventy and growing plugins available to interact with Kubernetes objects and components.

At the moment plugins cannot overwrite existing **kubectl** commands. Nor can it add sub-commands to existing commands. Writing new plugins should take into account the command line runtime package and a Go library for plugin authors.

As a plugin the declaration of options such as namespace or container to use must come after the command.

```
$ kubectl sniff bigpod-abcd-123 -c mainapp -n accounting
```

Plugins can be distributed in many ways. The use of **krew** allows for cross-platform packaging and a helpful plugin index, which makes finding new plugins easy.

More information can be found here: <https://kubernetes.io/docs/tasks/extend-kubectl/kubectl-plugins/>

Managing Plugins

- Add paths to plugins to \$PATH variable
- View current plugins with **kubectrl plugin list**
- Find new plugins with **kubectrl krew search**
- Installing using **kubectrl krew install new-plugin**
- Once installed use as **kubectrl** sub-command
- **upgrade** and **uninstall** also available

The help option explains basic operation. After installation ensure the \$PATH includes the plugins. **krew** should allow easy installation and use after that.

```
$ export PATH="${KREW_ROOT:-$HOME/.krew}/bin:$PATH"
$ kubectrl krew search
```

NAME	DESCRIPTION	INSTALLED
access-matrix	Show an RBAC access matrix for server resources	no
advise-ppsp	Suggests PodSecurityPolicies for cluster.	no
....		

```
$ kubectrl krew install tail
```

```
1 Updated the local copy of plugin index.
2 Installing plugin: tail
3 Installed plugin: tail
4 \
5 | Use this plugin:
6
7 ....
8 | | Usage:
9 | |
10 | | # match all pods
11 | | $ kubectrl tail
12 | |
13 | | # match pods in the 'frontend' namespace
14 | | $ kubectrl tail --ns staging
15 ....
```

Sniffing Traffic With Wireshark

- Read installation output for basic usage
- Note some require extra packages and configuration.
- **sniff** requires **Wireshark** and ability to export graphical display
- Pass the command the pod and container to use

```
$ kubectl krew install sniff nginx-123456-abcd -c webcont
```

Cluster network traffic is encrypted making troubleshooting of possible network issues more complex. Using the **sniff** plugin you can view the traffic from within.

The **sniff** command will use the first found container unless you pass the **-c** option to declare which container in the pod to use for traffic monitoring.

13.6 Logging

Logging Tools

- No cluster-wide logging in Kubernetes
- Often aggregated and digested by outside tools like **ElasticSearch**
- **Fluentd**
- **Kibana**

Logging, like monitoring, is a vast subject in IT. It has many tools that you can use as part of your arsenal.

Typically, logs are collected locally and aggregated before being ingested by a search engine and displayed via a dashboard which can use the search syntax. While there are many software stacks that you can use for logging, the **Elasticsearch**, **Logstash**, and **Kibana Stack (ELK)** has become quite common.

In Kubernetes, the **kubelet** writes container logs to local files (via the **Docker** logging driver). The command `kubectl logs` allows you to retrieve these logs.

Cluster-wide, you can use **Fluentd** to aggregate logs: <http://www.fluentd.org/>. Check the cluster administration logging concepts for a detailed description: <https://kubernetes.io/docs/concepts/cluster-administration/logging/>.

Fluentd is part of the **Cloud Native Computing Foundation (CNCF)** and, together with **Prometheus**, makes a nice combination for monitoring and logging. You can find a detailed walk-through of running **Fluentd** on Kubernetes here: <https://kubernetes.io/docs/tasks/debug-application-cluster/logging-elasticsearch-kibana>.

Setting up **Fluentd** for Kubernetes logging can be a good exercise in understanding **DaemonSets**. **Fluentd** agents run on each node via **DaemonSet**, they aggregate the logs, and feed them to an **Elasticsearch** instance prior to visualization in a **Kibana** dashboard.

13.7 Troubleshooting Resources

More Resources

- Official documentation
- Major vendor pages
- Github pages
- Kubernetes **Slack** channel

Other URLs to view for more troubleshooting information:

- General guidelines and instructions:
<https://kubernetes.io/docs/tasks/debug-application-cluster/troubleshooting/>
- Troubleshooting applications:
<https://kubernetes.io/docs/tasks/debug-application-cluster/debug-application>
- Troubleshooting clusters:
<https://kubernetes.io/docs/tasks/debug-application-cluster/debug-cluster>
- Debugging pods:
<https://kubernetes.io/docs/tasks/debug-application-cluster/debug-pod-replication-controller>
- Debugging services:
<https://kubernetes.io/docs/tasks/debug-application-cluster/debug-service/>
- Github Site for issue and bug tracking
<https://github.com/kubernetes/kubernetes/issues>
- Kubernetes Slack channel
kubernetes.slack.com

13.8 Labs

Exercise 13.1: Review Log File Locations

Overview

In addition to various logs files and command output, you can use **journalctl** to view logs from the node perspective. We will view common locations of log files, then a command to view container logs. There are other logging options, such as the use of a **sidecar** container dedicated to loading the logs of another container in a pod.

Whole cluster logging is not yet available with Kubernetes. Outside software is typically used, such as **Fluentd**, part of <http://fluentd.org/>, which is another member project of **CNCF.io**, like Kubernetes.

Take a quick look at the following log files and web sites. As server processes move from node level to running in containers the logging also moves.

1. If using a **systemd**-based Kubernetes cluster, view the node level logs for **kubelet**, the local Kubernetes agent. Each node will have different contents as this is node specific.

```
student@master:~$ journalctl -u kubelet |less
```

```
1 <output_omitted>
```

2. Major Kubernetes processes now run in containers. You can view them from the container or the pod perspective. Use the **find** command to locate the **kube-apiserver** log. Your output will be different, but will be very long.

```
student@master:~$ sudo find / -name "*apiserver*log"
```

```
1 /var/log/containers/kube-apiserver-master_kube-system_kube-apiserver-423
2 d25701998f68b503e64d41dd786e657fc09504f13278044934d79a4019e3c.log
```

3. Take a look at the log file.

```
student@master:~$ sudo less /var/log/containers/kube-apiserver-master_kube-system_kube-
apiserver-423d25701998f68b503e64d41dd786e657fc09504f13278044934d79a4019e3c.log
```

```
1 <output_omitted>
```

4. Search for and review other log files for **coredns**, **kube-proxy**, and other cluster agents.
5. If **not** on a Kubernetes cluster using **systemd** which collects logs via **journalctl** you can view the text files on the master node.
 - (a) **/var/log/kube-apiserver.log**
Responsible for serving the API
 - (b) **/var/log/kube-scheduler.log**
Responsible for making scheduling decisions
 - (c) **/var/log/kube-controller-manager.log**
Controller that manages replication controllers
6. **/var/log/containers**
Various container logs
7. **/var/log/pods/**
More log files for current Pods.
8. Worker Nodes Files (on non-**systemd** systems)

- (a) `/var/log/kubelet.log`
Responsible for running containers on the node
- (b) `/var/log/kube-proxy.log`
Responsible for service load balancing

9. More reading: <https://kubernetes.io/docs/tasks/debug-application-cluster/debug-service/> and <https://kubernetes.io/docs/tasks/debug-application-cluster/determine-reason-pod-failure/>

Exercise 13.2: Viewing Logs Output

Container standard out can be seen via the **kubectl logs** command. If there is no standard out, you would not see any output. In addition, the logs would be destroyed if the container is destroyed.

- View the current Pods in the cluster. Be sure to view Pods in all namespaces.

```
student@master:~$ kubectl get po --all-namespaces
```

	NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
1	kube-system	calico-kube-controllers-7b9dcfcc5-qg6zd	1/1	Running	0	13m
2	kube-system	calico-node-dr279	1/1	Running	0	6d1h
3						
4						
5					
6	kube-system	etcd-master	1/1	Running	2	44h
7	kube-system	kube-apiserver-master	1/1	Running	2	44h
8	kube-system	kube-controller-manager-master	1/1	Running	2	44h
9	kube-system	kube-scheduler-master	1/1	Running	2	44h
10					

- View the logs associated with various infrastructure pods. Using the **Tab** key you can get a list and choose a container. Then you can start typing the name of a pod and use **Tab** to complete the name.

```
student@master:~$ kubectl -n kube-system logs <Tab><Tab>
```

```
1 calico-kube-controllers-7b9dcfcc5-qg6zd
2 calico-node-dr279
3 calico-node-xtvfd
4 coredns-5644d7b6d9-k7kts
5 coredns-5644d7b6d9-rnr2v
6 etcd-master
7 kube-apiserver-master
8 kube-controller-manager-master
9 kube-proxy-qhc4f
10 kube-proxy-s56hl
11 kube-scheduler-f-master
12 traefik-ingress-controller-hw5tv
13 traefik-ingress-controller-mcn47
```

```
student@master:~$ kubectl -n kube-system logs \
    kube-apiserver-master
```

```
1 Flag --insecure-port has been deprecated, This flag will be removed in a future version.
2 I1119 02:31:14.933023      1 server.go:623] external host was not specified, using 10.128.0.3
3 I1119 02:31:14.933356      1 server.go:149] Version: v1.19.0
4 I1119 02:31:15.595131      1 plugins.go:158] Loaded 11 mutating admission controller(s)
5 successfully in the following order: NamespaceLifecycle,LimitRanger,ServiceAccount,
6 NodeRestriction,TaintNodesByCondition,Priority,DefaultTolerationSeconds,DefaultStorageClass,
7 StorageObjectInUseProtection,MutatingAdmissionWebhook,RuntimeClass.
8 I1119 02:31:15.595357      1 plugins.go:161] Loaded 7 validating admission controller(s)
9 successfully in the following order: LimitRanger,ServiceAccount,Priority,
```

```

10 PersistentVolumeClaimResize,ValidatingAdmissionWebhook,RuntimeClass,
11 ResourceQuota.
12 <output_omitted>

```

3. View the logs of other Pods in your cluster.

✍ Exercise 13.3: Adding tools for monitoring and metrics

With the deprecation of **Heapster** the new, integrated **Metrics Server** has been further developed and deployed. The **Prometheus** project of **CNCF.io** has matured from incubation to graduation, is commonly used for collecting metrics, and should be considered as well.

Configure Metrics



Very Important

The metrics-server is written to interact with Docker. If you chose to use crio the logs will show errors and inability to collect metrics.

1. Begin by cloning the software. The **git** command should be installed already. Install it if not found.

```

student@master:~$ git clone \
    https://github.com/kubernetes-incubator/metrics-server.git

```

```
1 <output_omitted>
```

2. As the software may have changed it is a good idea to read the **README.md** file for updated information.

```
student@master:~$ cd metrics-server/ ; less README.md
```

```
1 <output_omitted>
```

3. Create the necessary objects. Be aware as new versions are released there may be some changes to the process and the created objects. Use the components.yaml to create the objects. The backslash is not necessary if you type it all on one line.

```

student@master:~$ kubectl create -f \
    https://github.com/kubernetes-sigs/metrics-server/releases/download/v0.3.7/components.yaml

```

```

1 clusterrole.rbac.authorization.k8s.io/system:aggregated-metrics-reader created
2 clusterrolebinding.rbac.authorization.k8s.io/metrics-server:system:auth-delegator created
3 rolebinding.rbac.authorization.k8s.io/metrics-server-auth-reader created
4 apiservice.apiregistration.k8s.io/v1beta1.metrics.k8s.io created
5 serviceaccount/metrics-server created
6 deployment.apps/metrics-server created
7 service/metrics-server created
8 clusterrole.rbac.authorization.k8s.io/system:metrics-server created
9 clusterrolebinding.rbac.authorization.k8s.io/system:metrics-server created

```

4. View the current objects, which are created in the kube-system namespace. All should show a Running status.

```
student@master:~$ kubectl -n kube-system get pods
```

```

1 <output_omitted>
2 kube-proxy-ld2hb                1/1      Running   0          2d21h
3 kube-scheduler-u16-1-13-1-2f8c  1/1      Running   0          2d21h
4 metrics-server-fc6d4999b-b9rjj  1/1      Running   0          42s

```

5. Edit the metrics-server deployment to allow insecure TLS. The default certificate is x509 self-signed and not trusted by default. In production you may want to configure and replace the certificate. You may encounter other issues as this software is fast-changing. The need for the kubelet-preferred-address-types line has been reported on some platforms.

```
student@master:~$ kubectl -n kube-system edit deployment metrics-server
```

YAML

```
1 .....
2 spec:
3   containers:
4   - args:
5     - --cert-dir=/tmp
6     - --secure-port=4443
7     - --kubelet-insecure-tls #<-- Add this line
8     - --kubelet-preferred-address-types=InternalIP,ExternalIP,Hostname #<--May be needed
9     image: k8s.gcr.io/metrics-server/metrics-server:v0.3.7
10  ....
```

6. Test that the metrics server pod is running and does not show errors. At first you should see a few lines showing the container is listening. As the software changes these messages may be slightly different.

```
student@master:~$ kubectl -n kube-system logs metrics-server<TAB>
```

```
1 I0207 14:08:13.383209      1 serving.go:312] Generated self-signed cert
2 (/tmp/apiserver.crt, /tmp/apiserver.key)
3 I0207 14:08:14.078360      1 secure_serving.go:116] Serving securely on
4 [::]:4443
```

7. Test that the metrics working by viewing pod and node metrics. Your output may have different pods. It can take an minute or so for the metrics to populate and not return an error.

```
student@master:~$ sleep 120 ; kubectl top pod --all-namespaces
```

1	NAMESPACE	NAME	CPU(cores)	MEMORY(bytes)
2	kube-system	calico-kube-controllers-7b9dc5-qg6zd	2m	6Mi
3	kube-system	calico-node-dr279	23m	22Mi
4	kube-system	calico-node-xtvfd	21m	22Mi
5	kube-system	coredns-5644d7b6d9-k7kts	2m	6Mi
6	kube-system	coredns-5644d7b6d9-rnr2v	3m	6Mi
7	<output_omitted>			

```
student@master:~$ kubectl top nodes
```

1	NAME	CPU(cores)	CPU%	MEMORY(bytes)	MEMORY%
2	master	228m	11%	2357Mi	31%
3	worker	76m	3%	1385Mi	18%

8. Using keys we generated in an earlier lab we can also interrogate the API server. Your server IP address will be different.

```
student@master:~$ curl --cert ./client.pem \
  --key ./client-key.pem --cacert ./ca.pem \
  https://k8smaster:6443/apis/metrics.k8s.io/v1beta1/nodes
```

```
{
  "kind": "NodeMetricsList",
  "apiVersion": "metrics.k8s.io/v1beta1",
  "metadata": {
    "selfLink": "/apis/metrics.k8s.io/v1beta1/nodes"
  },
  "items": [
```

```

{
  "metadata": {
    "name": "u16-1-13-1-2f8c",
    "selfLink": "/apis/metrics.k8s.io/v1beta1/nodes/u16-1-13-1-2f8c",
    "creationTimestamp": "2019-01-10T20:27:00Z"
  },
  "timestamp": "2019-01-10T20:26:18Z",
  "window": "30s",
  "usage": {
    "cpu": "215675721n",
    "memory": "2414744Ki"
  }
},
<output_omitted>

```

Configure the Dashboard

While the dashboard looks nice it has not been a common tool in use. Those that could best develop the tool tend to only use the CLI, so it may lack wanted functionality.

Compatibility With Metric Server

The dashboard has not been updated to work with the **Metrics Server** now that **Heapster** has been deprecated. While there is some interest in getting the metrics to show in the dashboard there has been difficulty finding developers to work on the issue. <https://github.com/kubernetes/dashboard/issues/2986>

1. Create the dashboard. The short URL in the step below, which has an capital "oh", not number zero, is for this longer URL: <https://raw.githubusercontent.com/kubernetes/dashboard/v2.0.0-beta6/aio/deploy/recommended.yaml>.

```
student@master:~$ kubectl create -f https://bit.ly/2OFQRMy
```

```

1 secret/kubernetes-dashboard-certs created
2 serviceaccount/kubernetes-dashboard created
3 role.rbac.authorization.k8s.io/kubernetes-dashboard-minimal created
4 rolebinding.rbac.authorization.k8s.io/kubernetes-dashboard-minimal created
5 deployment.apps/kubernetes-dashboard created
6 service/kubernetes-dashboard created

```

2. View the current services in all namespaces. Note that the `kubernetes-dashboard` is a `ClusterIP` and part of the `kube-system` namespace.

```
student@master:~$ kubectl get svc --all-namespaces
```

```

1 <output_omitted>
2 default          thirdpage          NodePort    10.101.190.109    <none>        80:30588/TCP
3 kube-system      coredns            ClusterIP   10.96.0.10        <none>        53/UDP,53/TCP,9153/TCP
4 kube-system      metrics-server     ClusterIP   10.109.37.204     <none>        443/TCP
5 kube-system      traefik-ingress-service ClusterIP   10.103.146.169    <none>        80/TCP,8080/TCP
6 kubernetes-dashboard dashboard-metrics-scraper ClusterIP   10.111.17.237     <none>        8000/TCP
7 kubernetes-dashboard kubernetes-dashboard ClusterIP   10.105.169.5      <none>        443/TCP
8

```

3. Edit the `kubernetes-dashboard` and change the type to a `NodePort`.

```
student@master:~$ kubectl -n kubernetes-dashboard edit svc kubernetes-dashboard
```

YAML

```

1 .....
2   selector:
3     k8s-app: kubernetes-dashboard
4   sessionAffinity: None
5   type: NodePort                                #<-- Edit this line
6   status:
7     loadBalancer: {}

```

4. Check the `kubernetes-dashboard` service again. The Type should show as `NodePort`. Take note of the high-numbered port, which is 30968 in the example below. Yours will be different.

```
student@master:~$ kubectl -n kubernetes-dashboard get svc kubernetes-dashboard
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes-dashboard	NodePort	10.107.224.246	<none>	443:30968/TCP	6m39s

5. There has been some issues with RBAC and the dashboard permissions to see objects. In order to ensure access to view various resources give the dashboard admin access.

```

student@master:~$ kubectl create clusterrolebinding dashaccess \
  --clusterrole=cluster-admin \
  --serviceaccount=kubernetes-dashboard:kubernetes-dashboard

```

```
1 clusterrolebinding.rbac.authorization.k8s.io/dashaccess created
```

6. On your local node open a browser and navigate to an HTTPS URL made of the Public IP and the high-numbered port. You will get a message about an insecure connection. Select the **Advanced** button, then **Add Exception...**, then **Confirm Security Exception**. The page should then show the Kubernetes Dashboard. You may be able to find the public IP address using `curl`.

```
student@master:~$ curl ifconfig.io
```

```
1 35.231.8.178
```

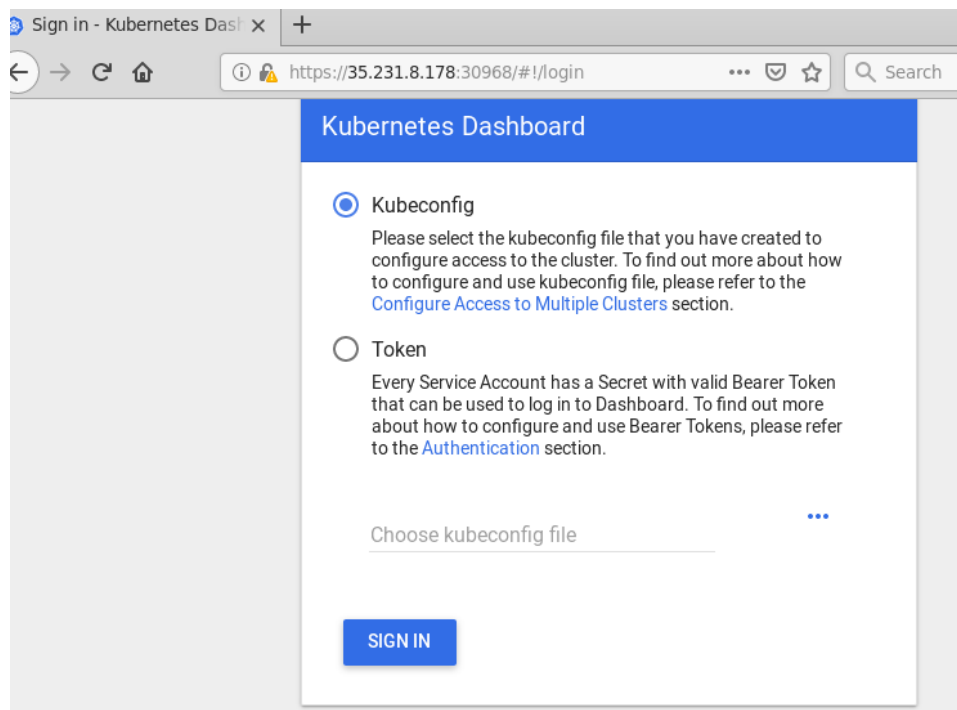


Figure 13.1: External Access via Browser

```

1  ....
2  Data
3  ====
4  ca.crt:          1025 bytes
5  namespace:      11 bytes
6  token:          eyJhbGciOiJSUzI1NiIsImtpZCI6IiJ9.eyJpc3MiOiJrdWJlcm5ldGVzL3N1cnZpY2VhY2NvdW50Iiwia3ViZX
7  JuZXRlc3Y5pby9zZXJ2aWN1YWNjb3VudC9uYW1lc3BhY2UiOiJrdWJlLXN5c3RlbnSIsImt1YmVybmV0ZXMuaW8vc2VydmJlZW
8  291bnQvc2VjcmV0Lm5hbnWUoiOiJrdWJlcm5ldGVzLWRhc2hib2FyZCI6IiJ9.eyJpc3MiOiJrdWJlcm5ldGVzLWRhc2hib2FyZC
9  ZWJjY291bnQvc2VydmJlZW51hY2NvdW50Lm5hbnWUoiOiJrdWJlcm5ldGVzLWRhc2hib2FyZCI6IiJ9.eyJpc3MiOiJrdWJlcm
10 jZWJjY291bnQvc2VydmJlZW51hY2NvdW50Lm5hbnWUoiOiJrdWJlcm5ldGVzLWRhc2hib2FyZCI6IiJ9.eyJpc3MiOiJrdWJlcm
11 I6InN5c3RlbnQvc2VjcmV0Lm5hbnWUoiOiJrdWJlcm5ldGVzLWRhc2hib2FyZCI6IiJ9.eyJpc3MiOiJrdWJlcm5ldGVzLWRhc2
12 qXpq4onn3hLhvz6yLSYxegRd6NysyGUYqnrSFE1trg9i1ftNXKJdzkY5kQzN3AcpUTvyj_BvJgzNh3JM9p7QMjI8LHTz4TrRz
13 rvwJVWitrEn4VnTQuFVcADFD_rKB9FyI_gvT_QiW5fQm24ygtIgfoYd44263oakG8sL64q7UfQNW2wt5S0orMUTybOmX4CXNUYM8
14 G44ejEtv9GW50sVjEmLIGaoEMX7fctwUN_XCyPdzcG2W0xRHahBjmbCuLz2SSWL52q4nXqmhTq_L8VDDpt6LjEqXW6LtDJZGjVC
15 s2MnBLerQz-ZagsVaubbQ

```



8. Navigate around the various sections and use the menu to the left as time allows. As the pod view is of the default namespace, you may want to switch over to the `kube-system` namespace or create a new deployment to view the resources via the GUI. Scale the deployment up and down and watch the responsiveness of the GUI.

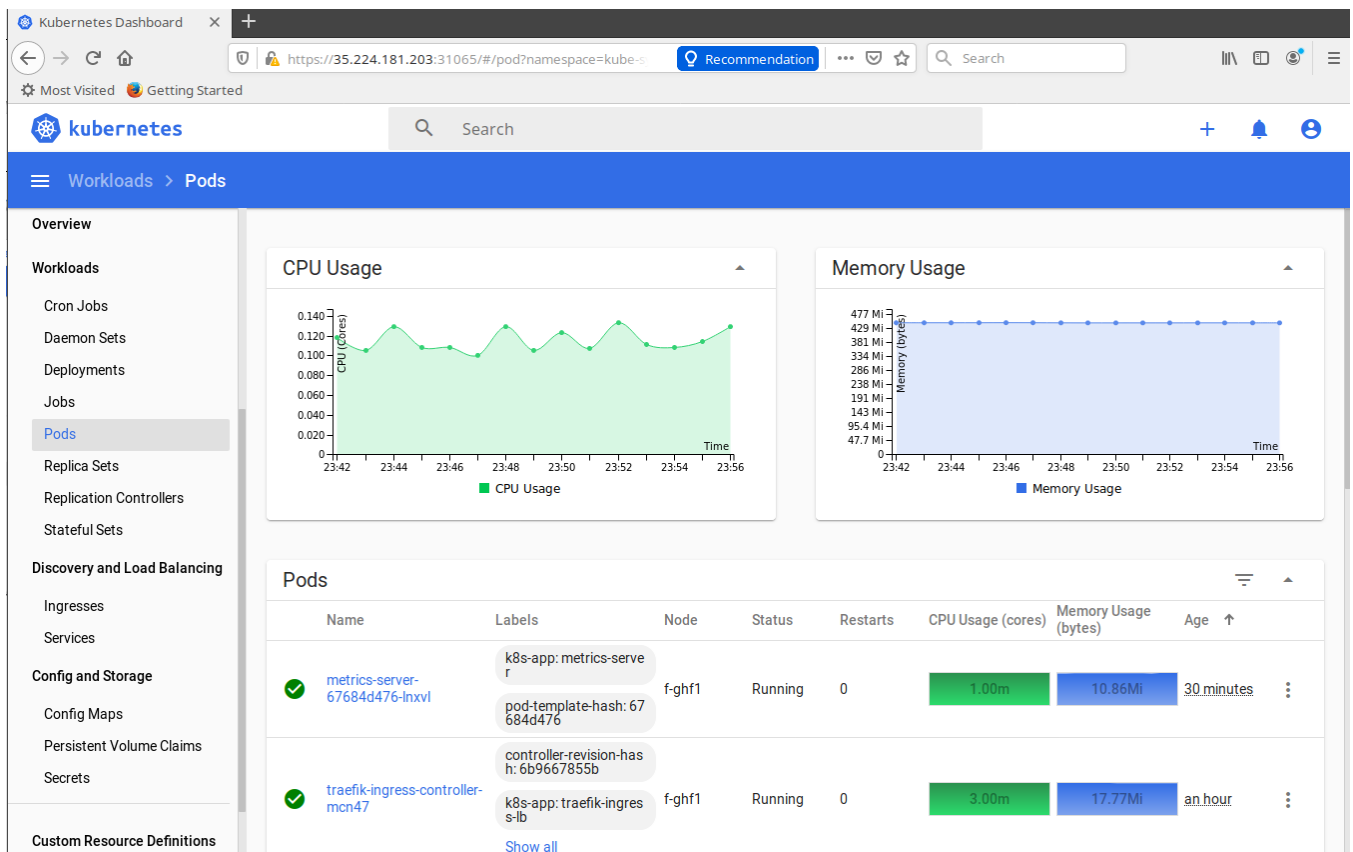


Figure 13.3: External Access via Browser

Chapter 14

Helm



14.1	Overview	268
14.2	Helm	269
14.3	Using Helm	273
14.4	Labs	276

14.1 Overview

Deploying Complex Applications

- Package your Kubernetes application via **chart** template
- **Helm** client to request install of **chart**
- **Tiller** creates cluster resources according to **chart**
- New version 3 changes quite a bit

We have used Kubernetes tools to deploy simple **Docker** applications. Starting with the v1.4 release, the goal was to have a canonical location for software. **Helm** is similar to a package manager like **yum** or **apt**, with a **chart** being similar to a package.

A typical containerized application will have several manifests. Manifests for deployments, services, and ConfigMaps. You will probably also create some secrets, Ingress, and other objects. Each of these will need a manifest.

With **Helm**, you can package all those manifests and make them available as a single tarball. You can put the tarball in a repository, search that repository, discover an application, and then, with a single command, deploy and start the entire application.

The server runs in your Kubernetes cluster, and your client can run anywhere, even a local laptop. With your client, you can connect to multiple repositories of applications, including those provided by vendors.

You will also be able to upgrade, or roll-back, an application easily from the command line.

14.2 Helm

Helm v2 and Tiller




Figure 14.1: **Basic Helm and Tiller Flow**

The **helm** version 2 uses a **Tiller** pod to deploy objects in the cluster. This has led to a lot of issues with security and cluster permissions. The new version 3 does not deploy a pod.

The Helm tool packages a Kubernetes application using a series of YAML files into a **chart**, or package. This allows for simple sharing between users, tuning using a templating scheme, as well as provenance tracking, among other things. **Helm v2** is made of two components:

- A server called **Tiller**, which runs inside your Kubernetes cluster
- A client called **Helm**, which runs on your local machine.

With the **Hel**m client you can browse package repositories containing published **charts**, and deploy those **charts** on your Kubernetes cluster. The **helm** will download the **chart** and pass a request to **Tiller** to create a release, otherwise known as an instance of a Chart. The release will be made of various resources running in the Kubernetes cluster.

Helm v3

- Many changes compared to version 2
- No more Tiller pod
- 3-way strategic merge patches
- Name or generated name now required on install

With the near complete overhaul of Helm the processes and commands have changed quite a bit. Expect to spend some time updating and integrating these changes if you are currently using Helm v2.

One of the most noticeable changes is the removal of the `Tiller` pod. This was an ongoing security issues as the pod needed elevated permissions to deploy charts. The functionality is in the command alone, and no longer requires initialization to use.

In version 2 an update to a chart and deployment used a 2-way strategic merge for patching. This compared the previous manifest to the intended manifest, but not the possible edits done outside of helm commands. The third way now checked is the live state of objects.

Among other changes software installation no longer generates a name automatically. One must be provided, or the `--generated-name` option must be passed.

Chart Contents

```
|-- Chart.yaml
|-- README.md
|-- templates
|   |-- NOTES.txt
|   |-- _helpers.tpl
|   |-- configmap.yaml
|   |-- deployment.yaml
|   |-- pvc.yaml
|   |-- secrets.yaml
|   |-- svc.yaml
|-- values.yaml
```

A chart is an archive set of Kubernetes resource manifests that make up a distributed application. You can check out the GitHub repository where the Kubernetes community is curating charts. Others exist and can be easily created, for example by a vendor providing software. Similar to the use of independent **YUM** repositories.

`Chart.yaml` contains some metadata about the chart, like its name, version, keywords, and so on, in this case for **MariaDB**. `values.yaml` contains keys and values that are used to generate the release in your Cluster. These values are replaced in the resource manifests using the **Go** templating syntax. And finally, the `templates` directory contains the resource manifests that make up this MariaDB application.

More about creating charts can found at:

<https://helm.sh/docs/topics/charts/>.

Templates

```
apiVersion: v1
kind: Secret
metadata:
  name: {{ template "fullname" . }}
  labels:
    app: {{ template "fullname" . }}
    chart: "{{ .Chart.Name }}-{{ .Chart.Version }}"
    release: "{{ .Release.Name }}"
    heritage: "{{ .Release.Service }}"
type: Opaque
data:
  mariadb-root-password: {{ default "" .Values.mariadbRootPassword | b64enc\
                                                                    | quote }}
  mariadb-password: {{ default "" .Values.mariadbPassword | b64enc | quote }}
```

The `template` are resource manifests which use the **Go** templating syntax. Variables defined in the values file, for example, get injected in the template when a release is created. In the **MariadbDB** example provided, the database passwords are stored in a Kubernetes `secret`, and the database configuration is stored in a Kubernetes `ConfigMap`.

We can see that a set of labels are defined in the `secret` metadata using the chart name, Release name, etc. The actual values of the passwords are read from `values.yaml`.

14.3 Using Helm

Initializing Helm v2

```
$ helm init
```

```
...  
Tiller (the helm server side component) has been installed into your \
```

Kubernetes Cluster.

```
Happy Helming!
```

```
$ kubectl get deployments --namespace=kube-system
```

NAMESPACE	NAME	READY	UP-TO-DATE	AVAILABLE	AGE
kube-system	tiller-deploy	1/1	1	1	15s

Helm v3 does not need to be initialized.

As always, you can build **Helm** from source or download a tarball. We expect to see **Linux** packages for the stable release soon. The current RBAC security requirements to deploy **helm** require the creation of a new `serviceaccount` and assigning of permissions and roles. There are several optional settings which can be passed to the **helm init** command, typically for particular security concerns, storage options and also a dry-run option.

The **helm v2** initialization should have created a new `tiller-deploy` pod in your cluster. Please note that this will create a deployment in the `kube-system` namespace.

The client will be able to communicate with the **tiller** pod using port forwarding. Hence, you will not see any service exposing **tiller**.

Chart Repositories

- Search for charts
- Add a new repository
- Simple HTTP servers with index file and tarball of charts
- **helm repo**

A default repository is included when initializing `helm`, but it's common to add other repositories. Repositories are currently simple HTTP servers that contain an index file and a tarball of all the Charts present.

You can interact with a repository using the `helm repo` commands.

```
$ helm repo add testing \
http://storage.googleapis.com/kubernetes-charts-testing
```

```
$ helm repo list
```

```
1 NAME      URL
2 stable    http://storage.googleapis.com/kubernetes-charts
3 local     http://localhost:8879/charts
4 testing   http://storage.googleapis.com/kubernetes-charts...
```

Once you have a repository available, you can search for Charts based on keywords. Below, we search for a redis Chart:

```
$ helm search redis
```

```
1 WARNING: Deprecated index file format. Try 'helm repo update'
2 NAME                VERSION  DESCRIPTION
3 testing/redis-cluster 0.0.5    Highly available Redis cluster with ...
4 testing/redis-standalone 0.0.1    Standalone Redis Master testing/...
```

Once you can find the chart within a repository you can deploy it on your cluster.

Deploying a Chart

```
$ helm install testing/redis-standalone
```

```

Fetched testing/redis-standalone to redis-standalone-0.0.1.tgz
amber-eel
Last Deployed: Fri Oct 21 12:24:01 2016
Namespace: default
Status: DEPLOYED

```

```
Resources:
```

```
==> v1/ReplicationController
```

NAME	DESIRED	CURRENT	READY	AGE
redis-standalone	1	1	0	1s

```
==> v1/Service
```

NAME	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
redis	10.0.81.67	<none>	6379/TCP	0s

To deploy a Chart, you can just use the **helm install** command. There may be several required resources for the installation to be successful, such as available PVs to match chart PVC. Currently the only way to discover which resources need to exist is by reading the READMEs for each chart

You will be able to list the release, delete it, even upgrade it and roll back.

```
$ helm list
```

	NAME	REVISION	UPDATED	STATUS	CHART
1	amber-eel	1	Fri Oct 21 12:24:01 2016	DEPLOYED	redis-stan...

A unique, colorful name will be created for each **helm** instance deployed. You can also use **kubectl** to view new resources **Helm** created in your cluster.

The output of deployment should be carefully reviewed. It often includes information on access to the applications within. If your cluster did not have a required cluster resource, the output is often the first place to begin troubleshooting.

14.4 Labs

Exercise 14.1: Working with Helm and Charts

Overview

helm allows for easy deployment of complex configurations. This could be handy for a vendor to deploy a multi-part application in a single step. Through the use of a Chart, or template file, the required components and their relationships are declared. Local agents like **Tiller** use the API to create objects on your behalf. Effectively its orchestration for orchestration.

There are a few ways to install **Helm**. The newest version may require building from source code. We will download a recent, stable version. Once installed we will deploy a Chart, which will configure **MariaDB** on our cluster.

Install Helm

1. On the master node use **wget** to download the compressed tar file. Various versions can be found here: <https://github.com/helm/helm/releases/>

```
student@master:~$ wget https://get.helm.sh/helm-v3.0.0-linux-amd64.tar.gz
```

```
1 <output_omitted>
2 helm-v3.0.0-linux-a 100%[=====] 11.52M --.-KB/s in 0.1s
3
4 2019-11-25 06:28:39 (112 MB/s) - 'helm-v3.0.0-linux-amd64.tar.gz' saved [12082866/12082866]
```

2. Uncompress and expand the file.

```
student@master:~$ tar -xvf helm-v3.0.0-linux-amd64.tar.gz
```

```
1 linux-amd64/
2 linux-amd64/helm
3 linux-amd64/README.md
4 linux-amd64/LICENSE
```

3. Copy the **helm** binary to the `/usr/local/bin/` directory, so it is usable via the shell search path.

```
student@master:~$ sudo cp linux-amd64/helm /usr/local/bin/helm3
```

4. A Chart is a collection of files to deploy an application. There is a good starting repo available on <https://github.com/kubernetes/charts/tree/master/stable>, provided by vendors, or you can make your own. Search the current Charts in the Helm Hub or an instance of Monocular for available stable databases. Repos change often, so the following output may be different from what you see.

```
student@master:~$ helm3 search hub database
```

URL	DESCRIPTION	CHART VERSION	APP VERSION
https://hub.helm.sh/charts/choerodon/postgresql	Chart for PostgreSQL, an object-relational data...	3.18.4	10.7.0
https://hub.helm.sh/charts/inspur/redis-cluster	Highly available Kubernetes implementation of R...	0.0.1	5.0.6
https://hub.helm.sh/charts/lohmag/clickhouse	ClickHouse is an open source column-oriented da...	0.2.0	19.17
https://hub.helm.sh/charts/halkeye/turtl	The secure, collaborative notebook - Totally pr...	0.1.7	0.7

5. We will begin by adding a common repository and call it `stable`.

```
student@master:~$ helm3 repo add stable https://kubernetes-charts.storage.googleapis.com
```

```
1 "stable" has been added to your repositories
```

```
student@master:~$ helm3 repo update
```

```
1 Hang tight while we grab the latest from your chart repositories...
2 ...Successfully got an update from the "stable" chart repository
3 Update Complete. Happy Helming!
```

6. We will install the **mariadb**. The **- --debug** option will create a lot of output. The output will typically suggest ways to access the software. As well we will indicate that we do not want persistent storage, which would require us to create an available PV.

```
student@master:~$ helm3 --debug install firstdb stable/mariadb \
  --set master.persistence.enabled=false \
  --set slave.persistence.enabled=false
```

```
1 install.go:148: [debug] Original chart version: ""
2 install.go:165: [debug] CHART PATH: /home/student/.cache/helm/repository/mariadb-7.0.1.tgz
3
4 client.go:87: [debug] creating 8 resource(s)
5 NAME: firstdb
6 LAST DEPLOYED: Mon Nov 25 08:03:30 2019
7 <output_omitted>
```

7. Using some of the information at the end of the previous command output we will deploy another container and access the database. We begin by getting the root password for `illmannered-salamander`. Be aware the output lacks a carriage return, so the next prompt will appear on the same line. We will need the password to access the running MariaDB database.

```
student@master:~$ kubectl get secret --namespace default firstdb-mariadb \
  -o jsonpath="{.data.mariadb-root-password}" | base64 --decode
```

```
1 TQ93fCLP13
```

8. Now we will install another container to act as a client for the database. We will use **apt-get** to install client software. Using copy and paste from the installation output may be helpful.

```
student@master:~$ kubectl run firstdb-mariadb-client \
  --rm --tty -i --restart='Never' \
  --image docker.io/bitnami/mariadb:10.3.22-debian-10-r27 \
  --namespace default --command -- bash
```

9. Use the client software to access the database. The following command uses the server name and the root password we found in a previous step. Both of yours will be different.



Inside container

```
I have no name!@firstdb-mariadb-client:/$ mysql -h firstdb-mariadb.default.svc.cluster.local \
  -uroot -p my_database
Enter password: TQ93fCLP13
```

```
1 Welcome to the MariaDB monitor.  Commands end with ; or \ g.
2 Your MariaDB connection id is 153
3 Server version: 10.1.38-MariaDB Source distribution
4
```



```

5 Copyright (c) 2000, 2018, Oracle, MariaDB Corporation Ab and others.
6
7 Type 'help;' or '\ h' for help. Type '\ c' to clear the current input statement.
8
9 MariaDB [(none)]> SHOW DATABASES;
10 +-----+
11 | Database          |
12 +-----+
13 | information_schema |
14 | my_database        |
15 | mysql              |
16 | performance_schema |
17 | test               |
18 +-----+
19 5 rows in set (0.00 sec)
20
21 MariaDB [(none)]>
22 MariaDB [(none)]> quit
23 Bye
24 I have no name!@firstdb-mariadb-client:/$ exit
25

```

10. View the Chart history on the system. The use of the **-a** option will show all Charts including deleted and failed attempts.

```
student@master:~$ helm3 list
```

NAME	NAMESPACE	REVISION	UPDATED
STATUS	CHART	APP VERSION	
firstdb	default	1	2019-11-25 08:03:30.693410615
+0000 UTC	deployed	mariadb-7.0.1	10.3.20

11. Delete the **mariadb** Chart. No releases of mariadb should be found.

```
student@master:~$ helm3 uninstall firstdb
```

```
1 release "firstdb" uninstalled
```

```
student@master:~$ helm3 list
```

NAME	NAMESPACE	REVISION	UPDATED	STATUS	CHART	APP VERSION
------	-----------	----------	---------	--------	-------	-------------

12. Find the downloaded chart. It should be a compressed tarball under the user's home directory. Your **mariadb** version may be slightly different.

```
student@master:~$ find $HOME -name *mariadb*
```

```
1 /home/student/.cache/helm/repository/mariadb-7.3.14.tgz
2
```

13. Move to the archive directory and extract the tarball. Take a look at the files within.

```
student@master:~$ cd $HOME/.cache/helm/repository ; tar -xvf mariadb-*
```

```

1 mariadb/Chart.yaml
2 mariadb/values.yaml
3 mariadb/templates/NOTES.txt
4 mariadb/templates/_helpers.tpl
5 mariadb/templates/initialization-configmap.yaml
6 <output_omitted>

```

14. Copy and rename the `values.yaml` file back to the home directory and return there as well.

```
student@master:~/.cache/helm/repository$ cp mariadb/values.yaml $HOME/custom.yaml ; cd
```

15. Review the `custom.yaml` file, note there are many possible configurations. While most are commented out, take a moment to work slowly through the file.

```
student@master:~$ less custom.yaml
```

```
1 <output_omitted>
```

16. Edit the file. We will change or add two items, the `rootUser` password and storage persistence parameters.

```
student@master:~$ vim custom.yaml
```

```
1 ....
2 rootUser:
3   ## MariaDB admin password
4   ## ref: https://github.com/bitnami/bitnami-docker-mariadb#setting-the-root-password-on-first-run
5   ##
6   password: LFTr@1n                                #<-- Add a password, such as LFTr@1n
7   ##
8   ....
9   persistence:
10    ## If true, use a Persistent Volume Claim, If false, use emptyDir
11    ##
12    enabled: false                                #<--- Change this to false
13    # Enable persistence using an existing PVC
14    # existingClaim:
15    ....
```

17. We will now deploy another **MariaDB** instance using our custom settings. Note we no longer have to pass statements to avoid using persistent storage. Near the end of the output will be commands on using the database.

```
student@master:~$ helm3 install -f custom.yaml seconddb stable/mariadb
```

```
1 NAME: seconddb
2 LAST DEPLOYED: Mon Nov 25 08:54:08 2019
3 NAMESPACE: default
4 STATUS: deployed
5
6
7 <output_omitted>
8
9 To connect to your database:
10
11 1. Run a pod that you can use as a client:
12
13     kubectl run seconddb-mariadb-client --rm --tty -i --restart='Never' --image
14     docker.io/bitnami/mariadb:10.3.20-debian-9-r0 --namespace default --command -- bash
15
16 2. To connect to master service (read/write):
17
18 <output_omitted>
```

18. Copy and paste the command from the output to run a client pod. If you are re-typing, don't include the backslashes. The command could be on typed on a single line. Once in the container log into the database with the password we set and check the default databases.

```
student@master:~$ kubectl run seconddb-mariadb-client \
  --rm --tty -i --restart='Never' \
  --image docker.io/bitnami/mariadb:10.3.22-debian-10-r27 \
  --namespace default --command -- bash
```



On Container

```
I have no name!@iron-pika-mariadb-client:/$ mysql -h seconddb-mariadb.default.svc.cluster.local \
-uroot -p my_database
Enter password: LFTr@1n #<-- The password we set in the config file

1 Welcome to the MariaDB monitor.  Commands end with ; or \g.
2 Your MariaDB connection id is 25
3 Server version: 10.3.17-MariaDB-log Source distribution
4
5 Copyright (c) 2000, 2018, Oracle, MariaDB Corporation Ab and others.
6
7 Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.

MariaDB [my_database]> SHOW DATABASES;

1 +-----+
2 | Database |
3 +-----+
4 | information_schema |
5 | my_database |
6 | mysql |
7 | performance_schema |
8 | test |
9 +-----+
10 5 rows in set (0.001 sec)

MariaDB [my_database]> quit

1 Bye

I have no name!@seconddb-mariadb-client:/$ exit
```

19. Remove anything you have installed using **helm**. Reference earlier in the chapter if you don't remember the command.

Chapter 15

Closing and Evaluation Survey



15.1	Evaluation Survey	281
------	-------------------	-----

15.1 Evaluation Survey

Thank you for taking this **Linux Foundation** course.

Your comments are important to us and we take them seriously, both to measure how well we fulfilled your needs and to help us improve future sessions.

- Please Evaluate your training using the link your instructor will provide.
- Please be sure to check the spelling of your name and use correct capitalization as this is how your name will appear on the certificate.
- This information will be used to generate your **Certificate of Completion**, which will be sent to the email address you have supplied, so make sure it is correct.




Figure 15.1: **Course Survey**